

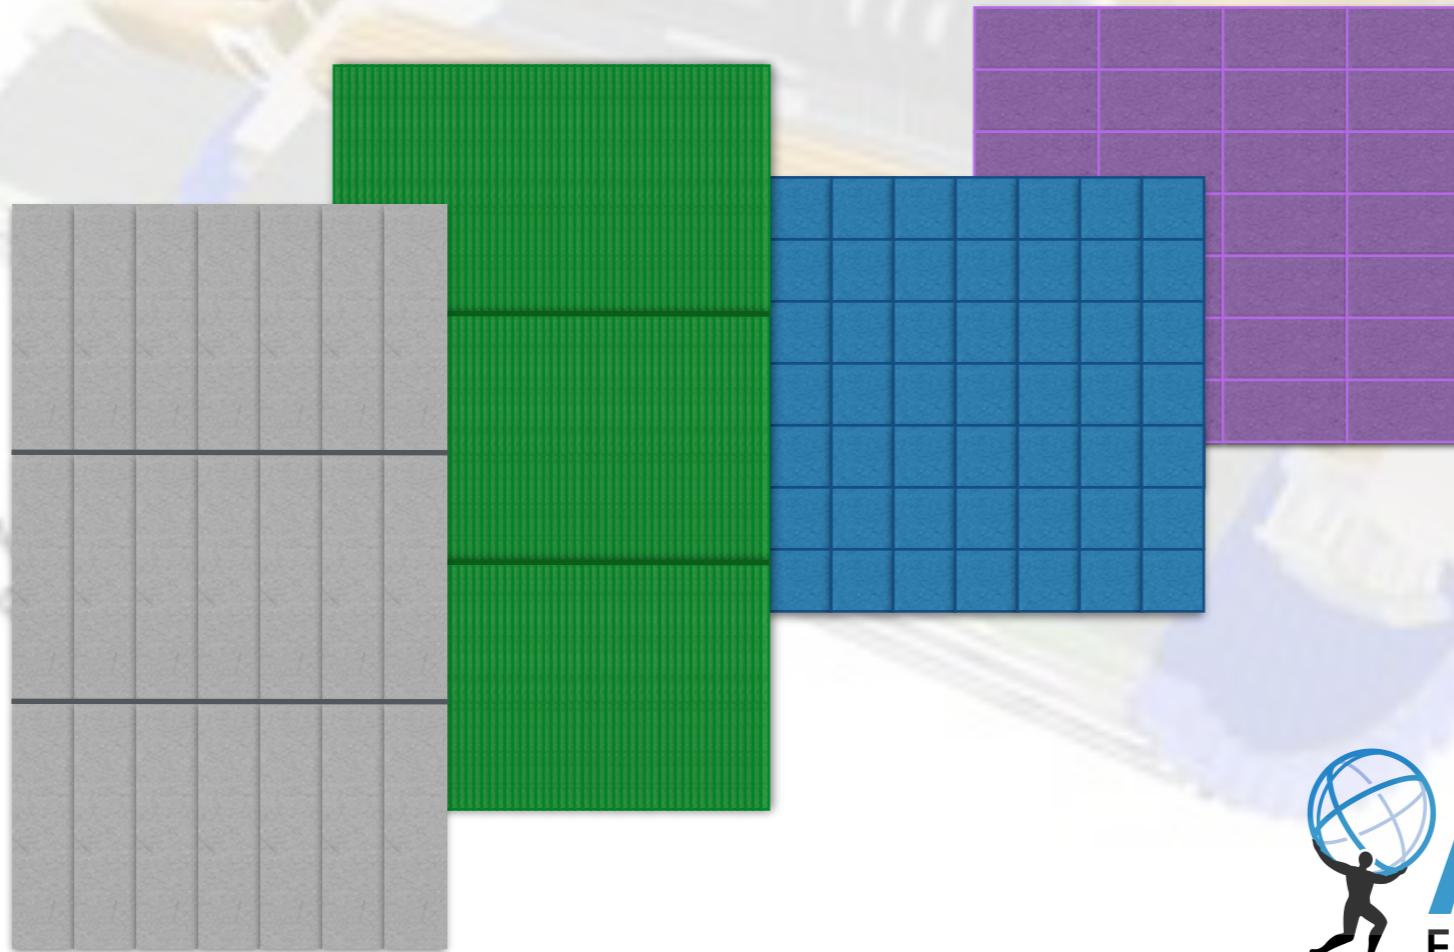
Simulating electromagnetic showers in ATLAS calorimeter with Generative Adversarial Network

Aishik Ghosh

Supervisor: David Rousseau

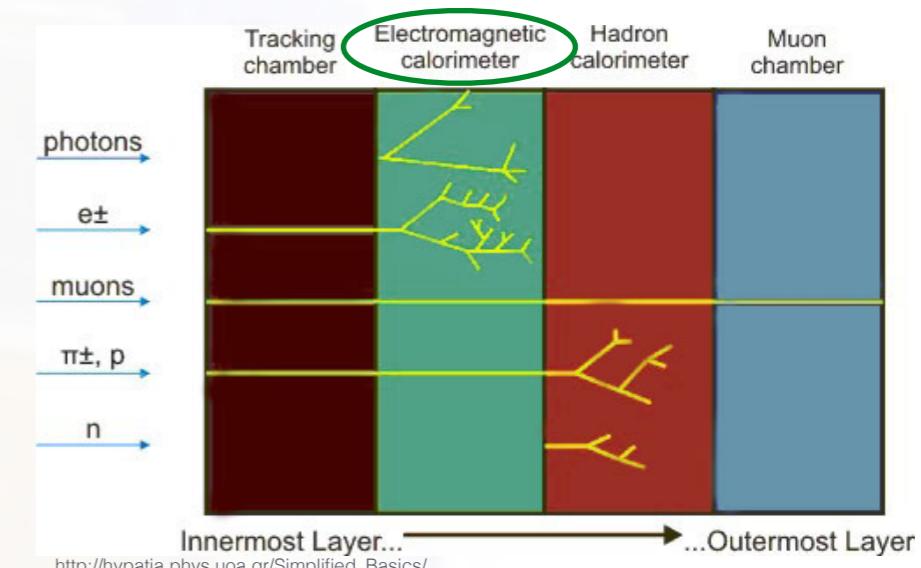
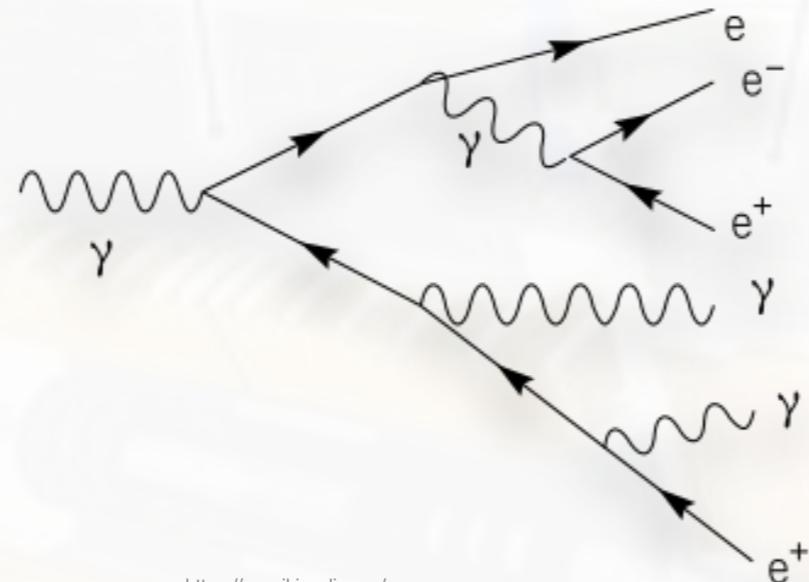
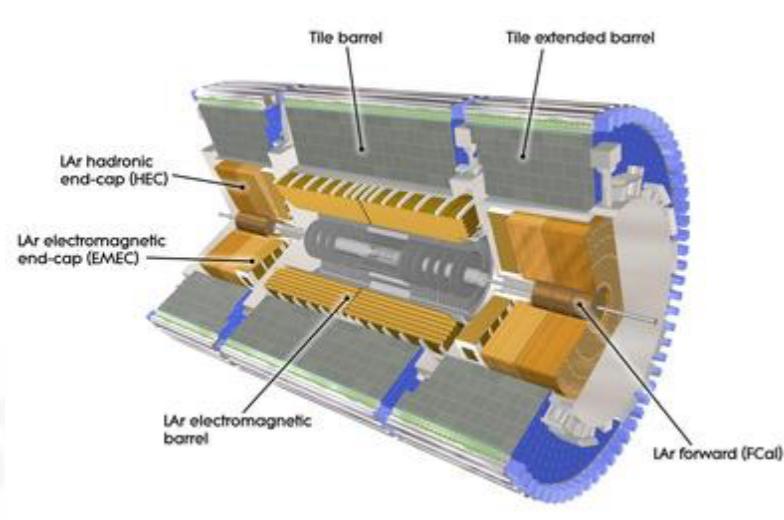
IN2P3 ML Workshop

29 March 2018



Physics Motivation

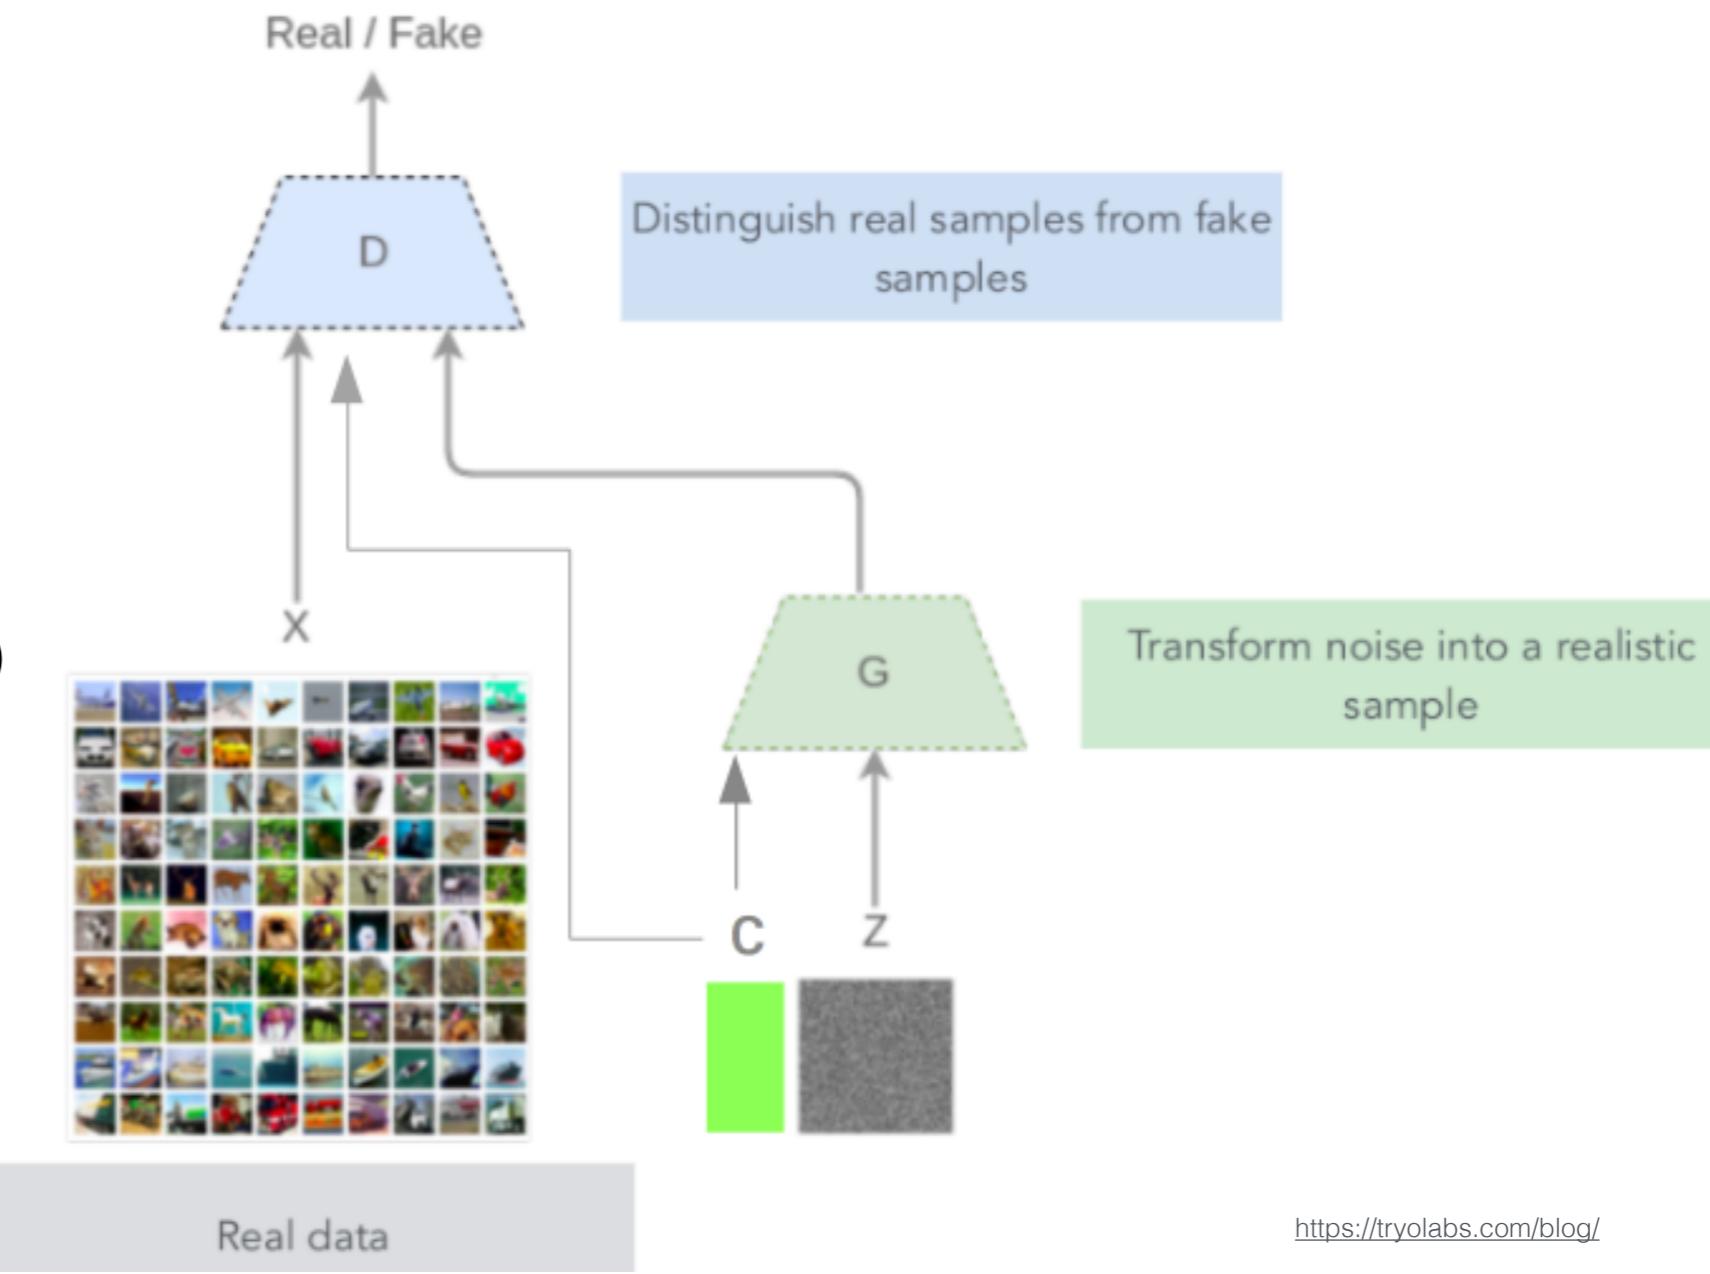
Photons deposit energy into the calorimeter cells as a shower



- Currently >50 % of ATLAS computing time spent on shower simulation
- DNNs based generators could significantly speed up the process
 - An alternative to FastCaloSim [Parameterisation](#), [Frozen Showers](#) approaches
- Training outside [Athena](#) (full ATLAS software framework), apply inside Athena
 - Use state-of-the-art DNN packages, GPU farms for training
 - Use [Lightweight Trained Neural Network](#) to integrate into Athena

What are GANs?

- Generative neural network based model to output images
- G: Generative network takes noise (Z) as input, outputs an image
- D: Discriminator Network learns to differentiate real (X) and fake images (from G)
- D helps G learn the distribution of real data
- Can produce conditional output based on C (e.g. Particle Energy)



<https://tryolabs.com/blog/>

The Calorimeter

<http://inspirehep.net/record/1467455/plots>

2-D Axis: ϕ vs η

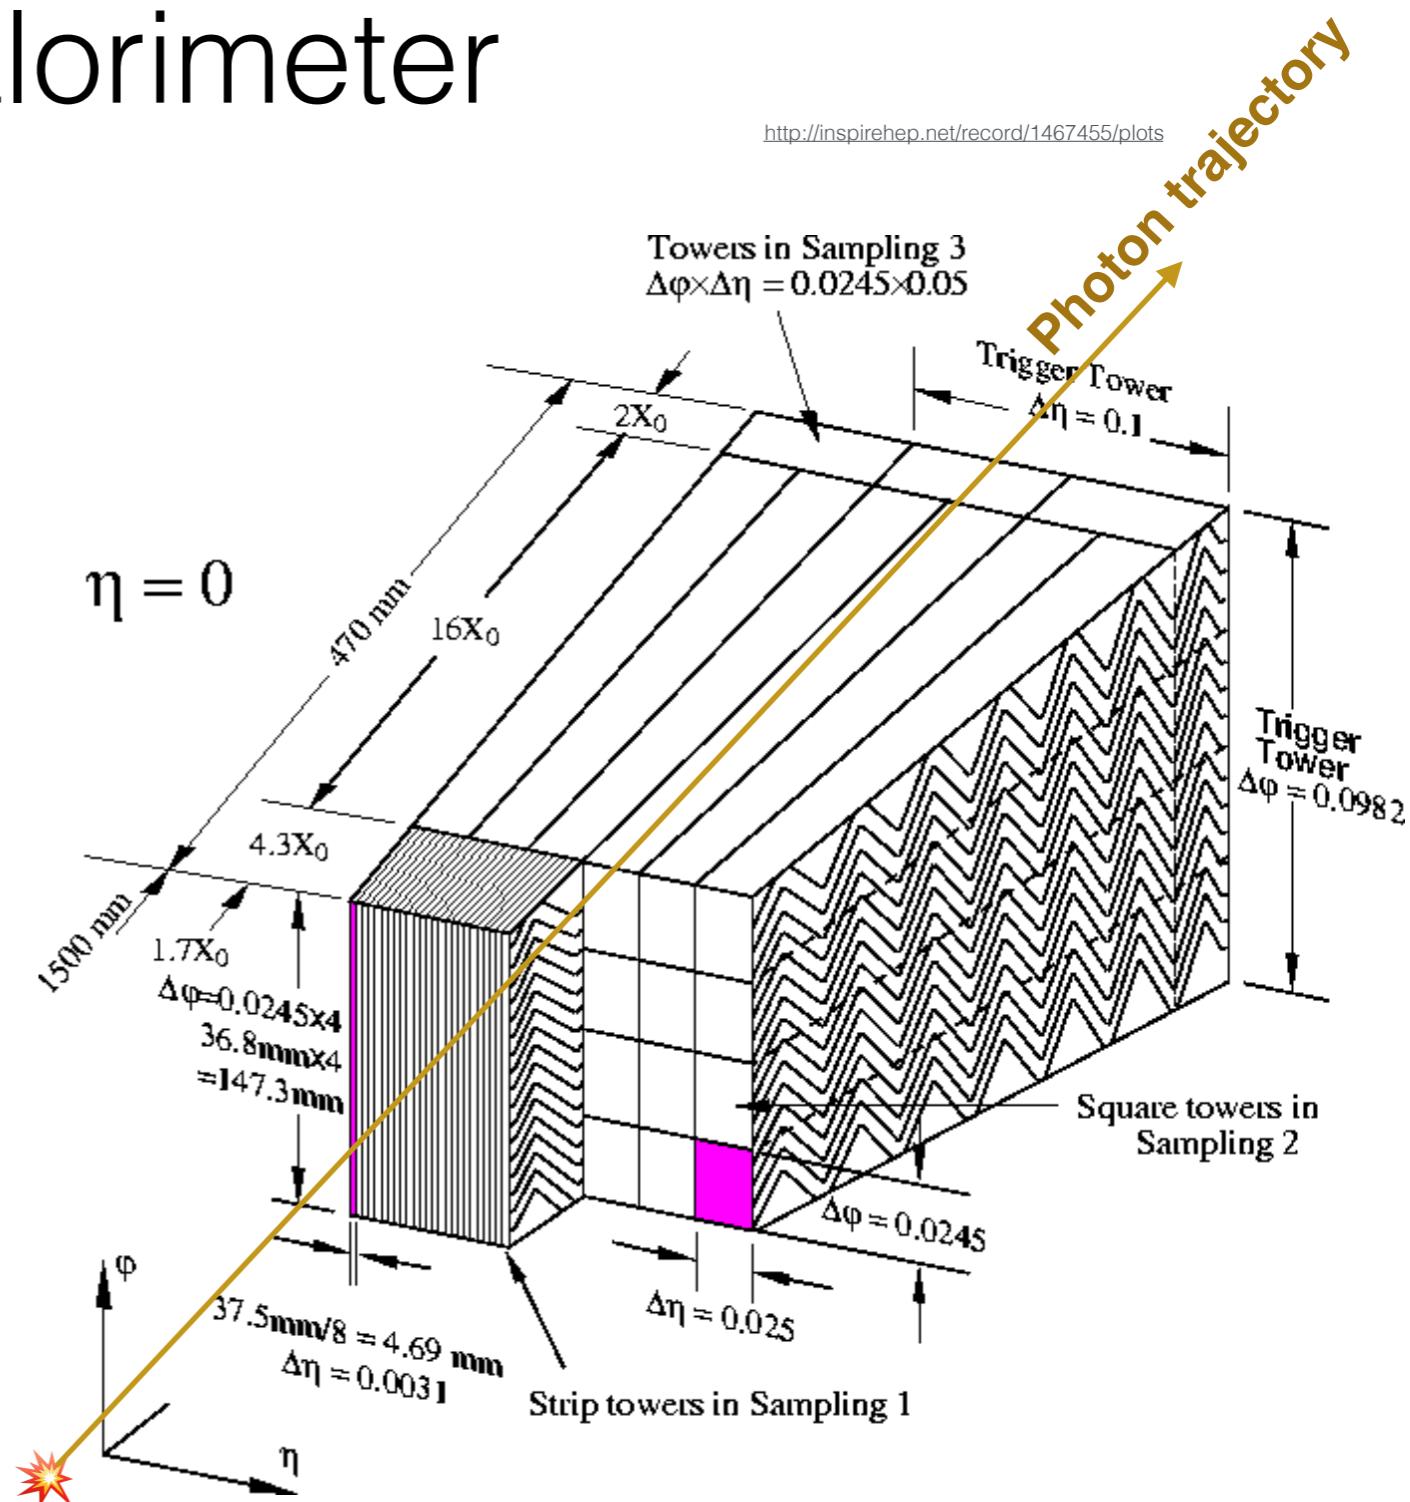
Particle goes through 4 layers in this order:

0. Pre-Sampler: Some energy deposit

1. Strips: Very granular in η ; more energy deposit

2. Middle: Thickest layer, maximum energy deposit

3. Back: Little Energy deposits



The Calorimeter

2-D Axis: ϕ vs η

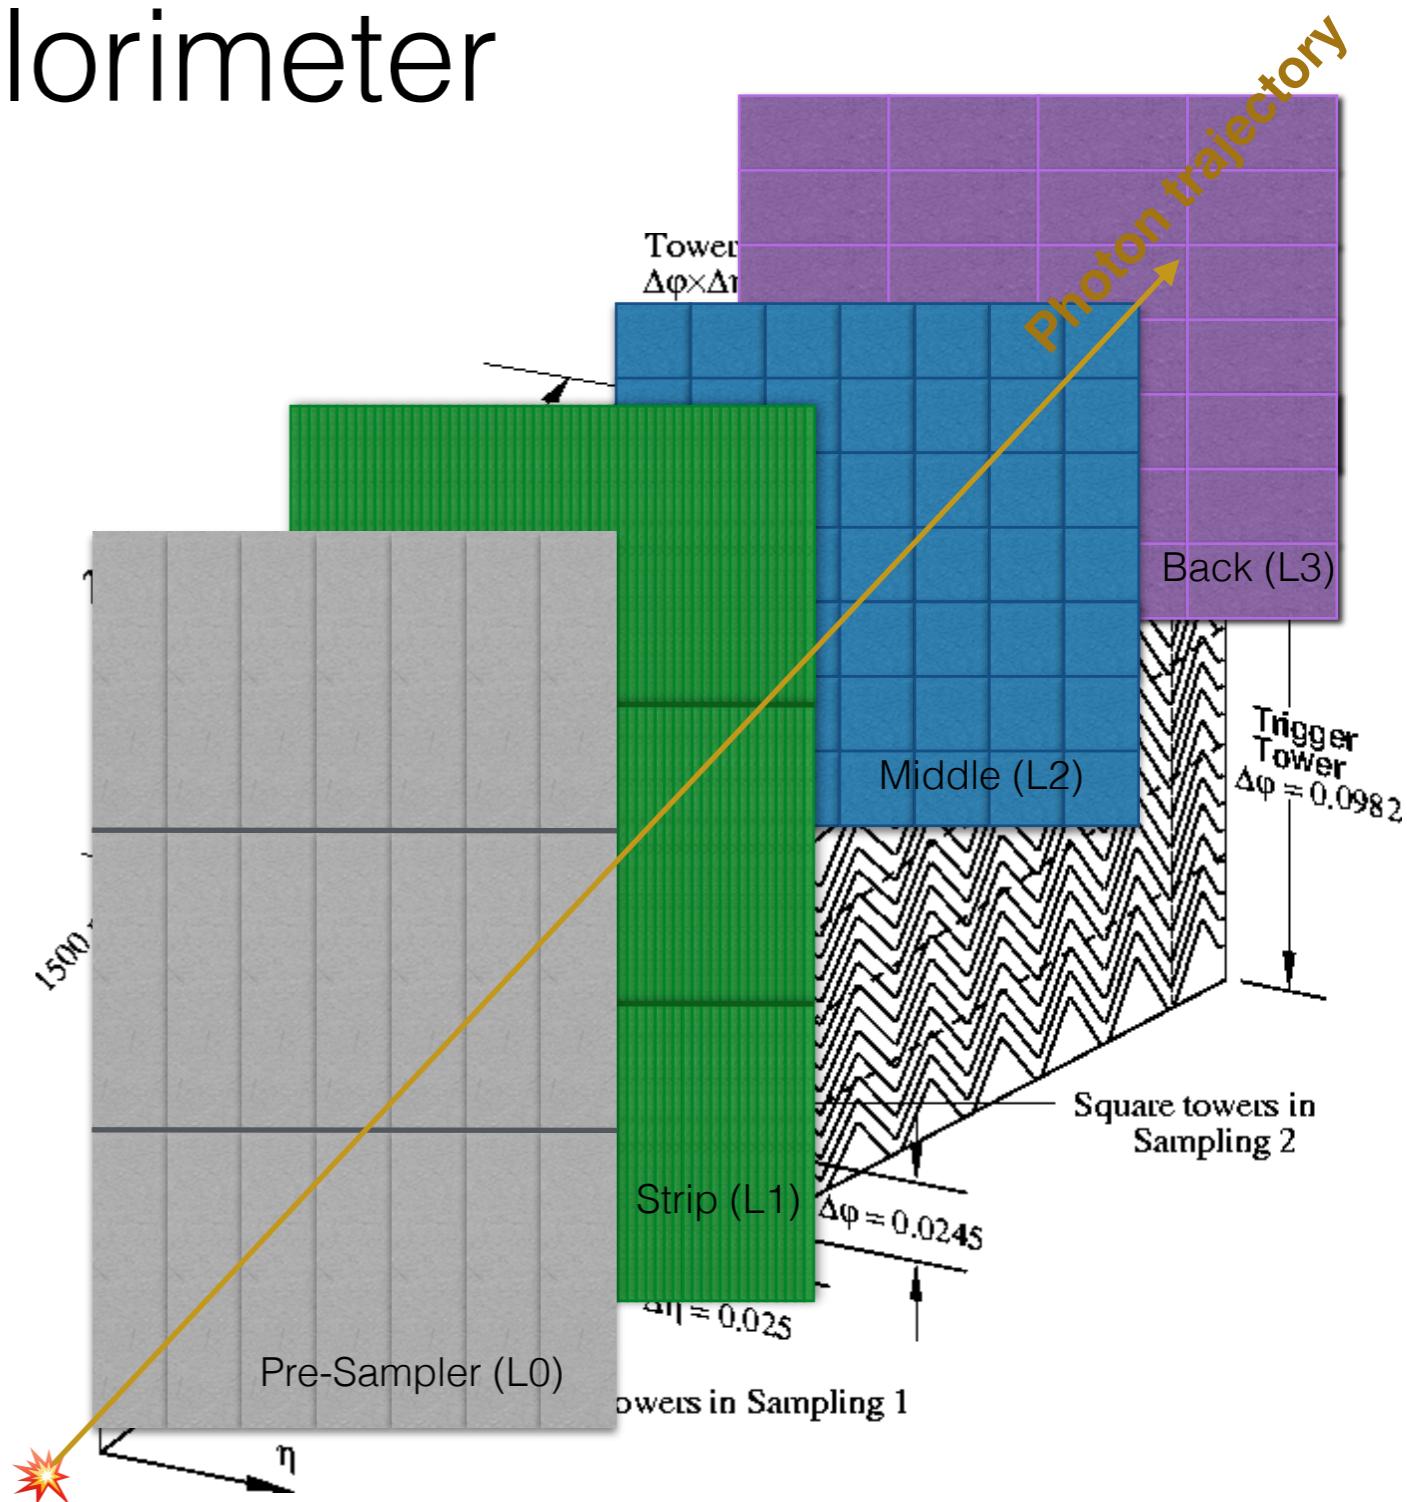
Particle goes through 4 layers in this order:

0. Pre-Sampler : Some energy deposit

1. Strips: Very granular in η ; more energy deposit

2. Middle: Thickest layer, maximum energy deposit

3. Back: Little Energy deposits



The Calorimeter

2-D Axis: ϕ vs η

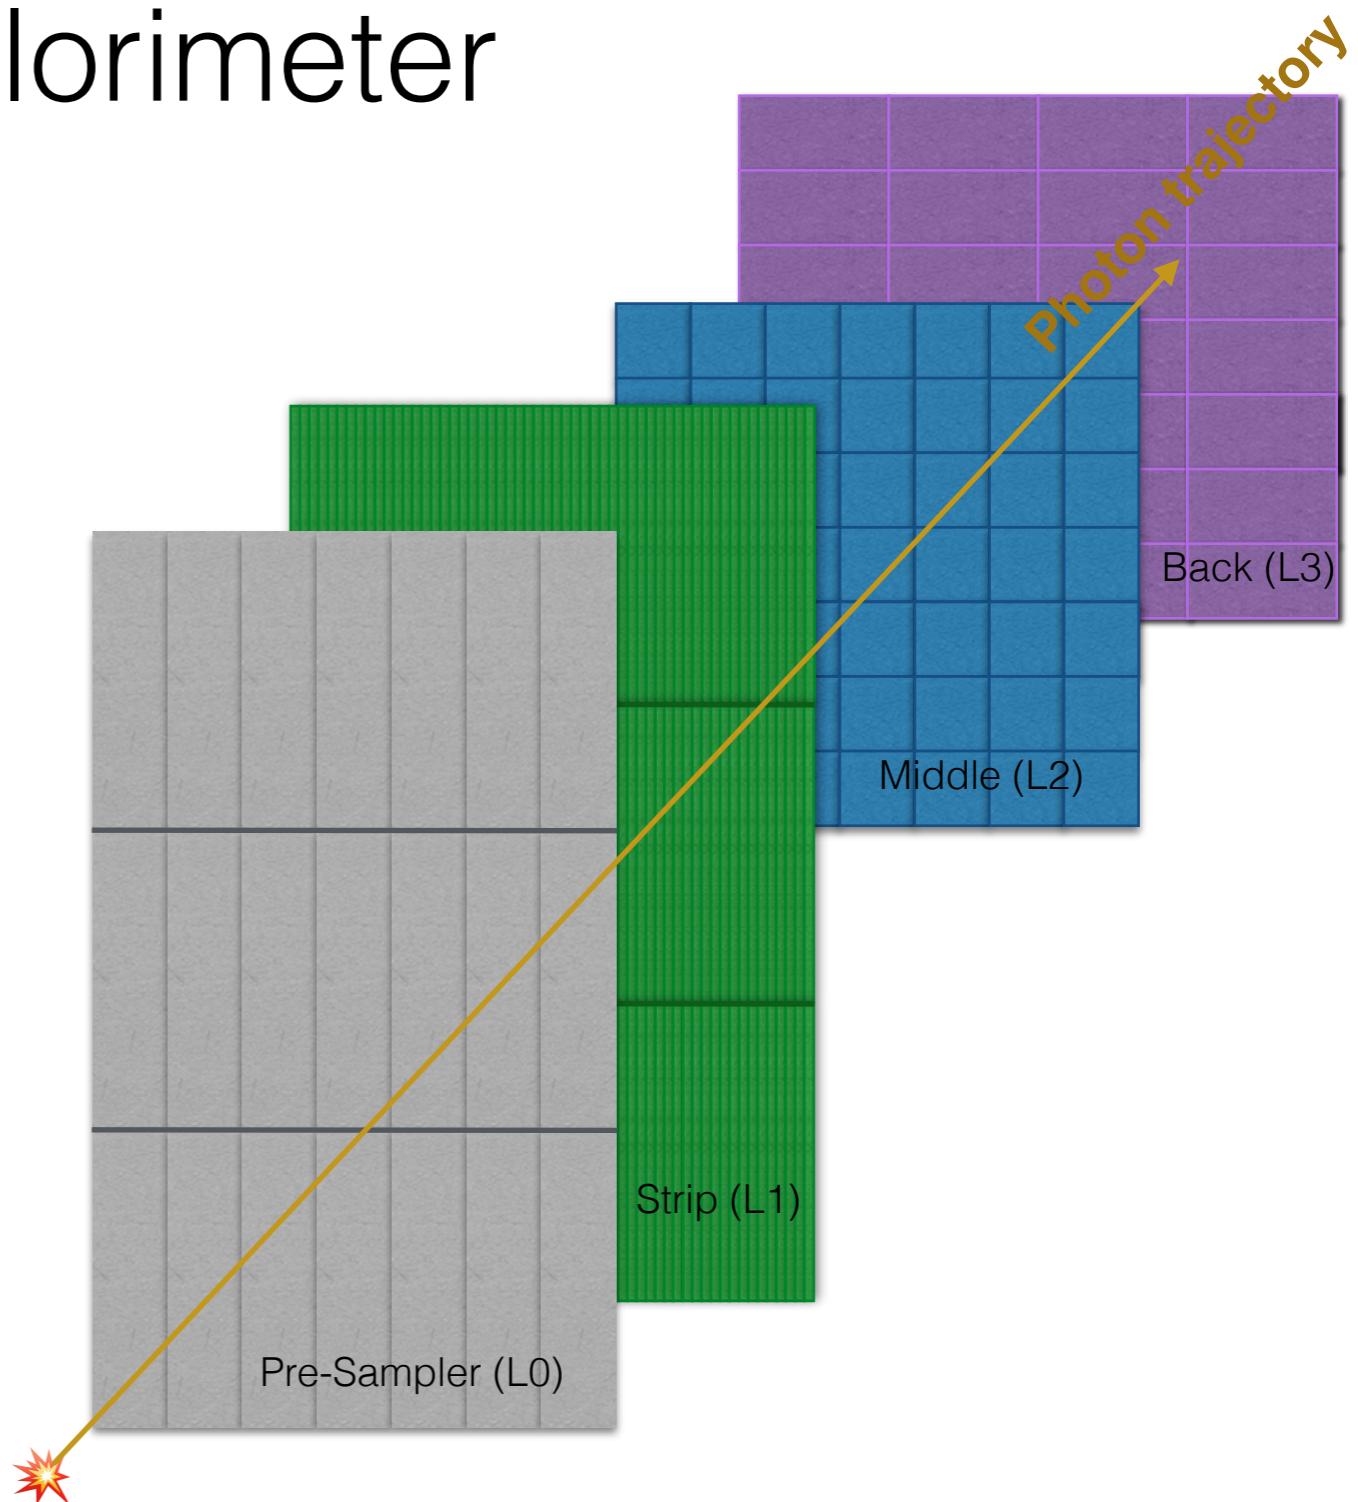
Particle goes through 4 layers in this order:

0. Pre-Sampler : Some energy deposit

1. Strips: Very granular in η ; more energy deposit

2. Middle: Thickest layer, maximum energy deposit

3. Back: Little Energy deposits



The Calorimeter

2-D Axis: ϕ vs η

Particle goes through 4 layers in this order:

0. Pre-Sampler: Some energy deposit

1. Strips: Very granular in η ; more energy deposit

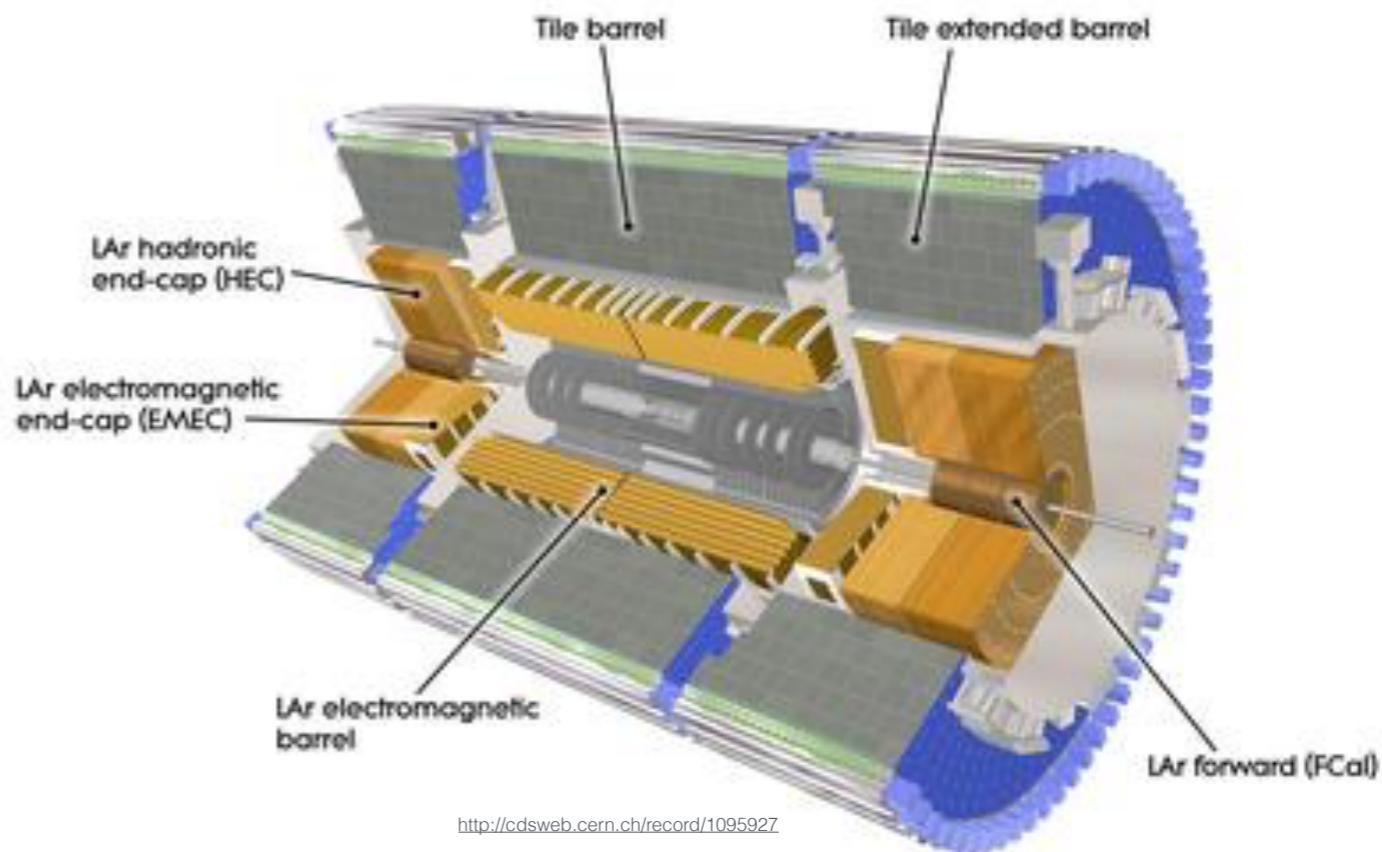
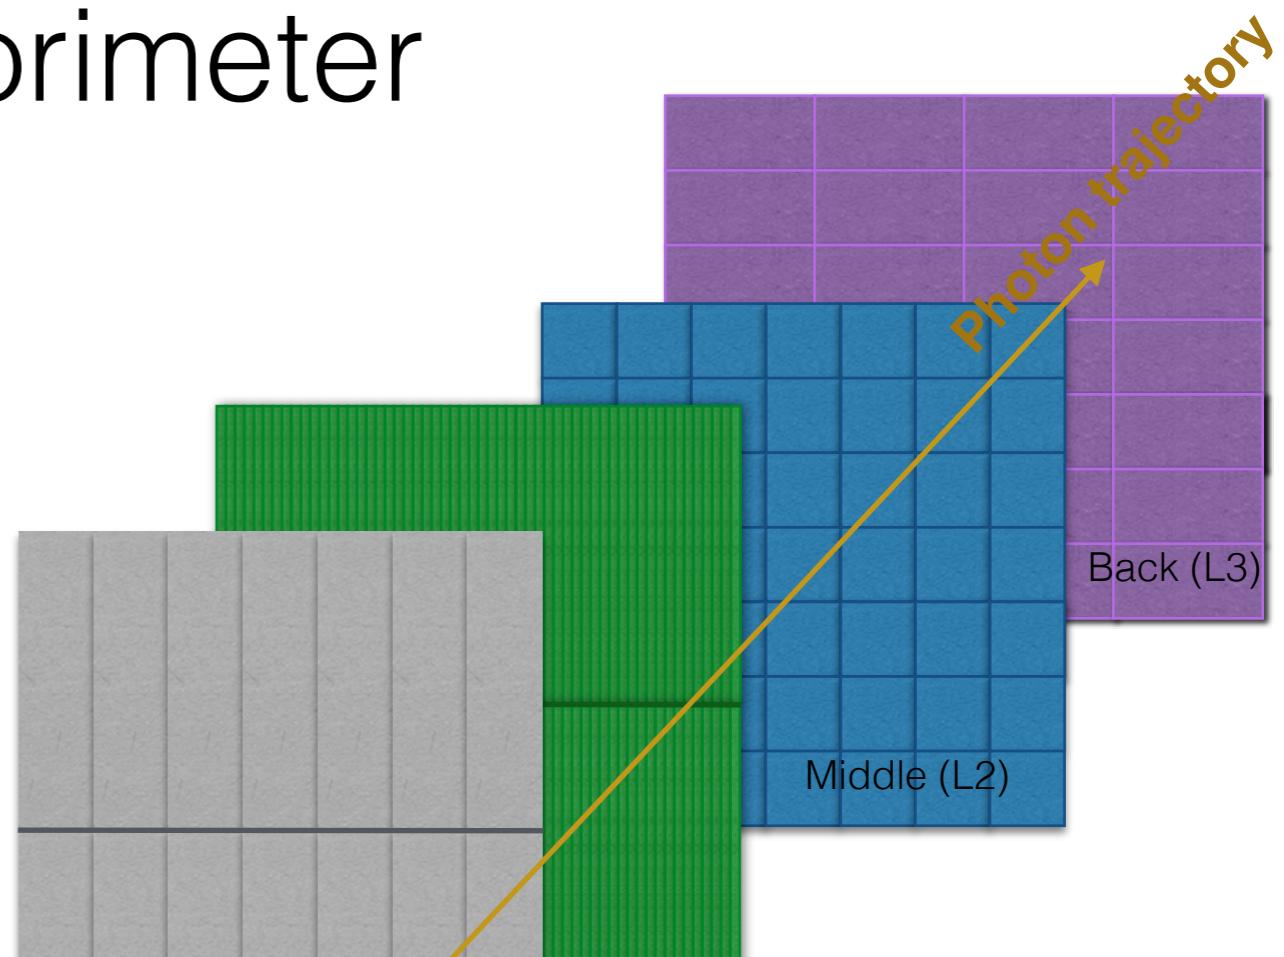
2. Middle: Thickest layer, maximum energy deposit

3. Back: Little Energy deposits

Due to misalignment of the two halves of the detector, cells are not perfectly well aligned.

Different widths of cells further complicate the alignment between cells of different layers

Cells not granular enough to see intricate details of shower pattern



Our 'Images'

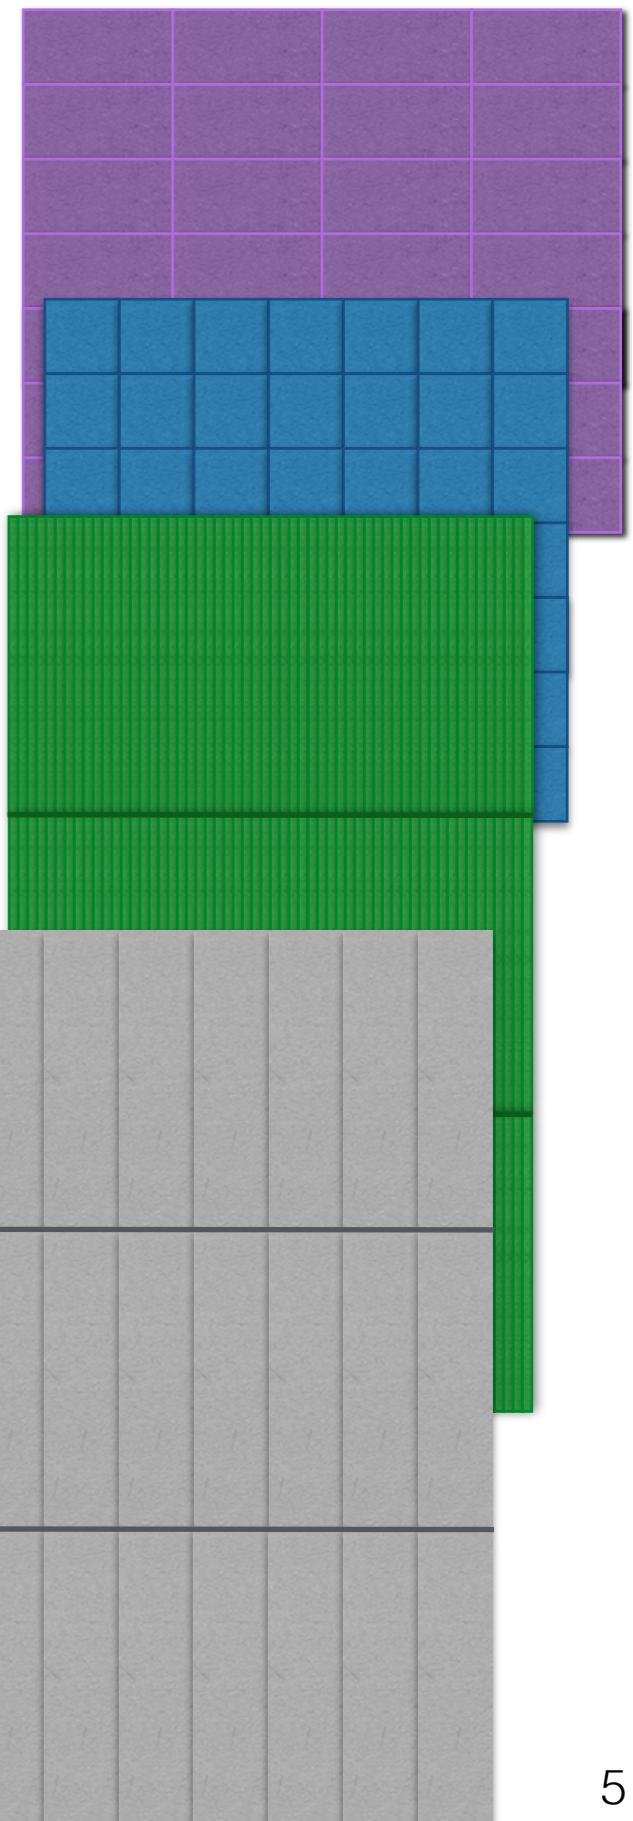
Calo Cell → Pixel

Energy deposited in Calo Cell → Pixel Intensity

Convert 266 Cells from 4 Calo Layers into 1-D Numpy Array Input for GAN

GAN should catch the correlations between cells (and correlations between calo layers)

Can it learn the 3D distribution without any spatial information?



Software

Train using **Keras** (high-level wrapper) with **Tensorflow** backend

Train on **1 GPU at CC Lyon**

Plan to apply using **Lightweight Trained Neural Network**:

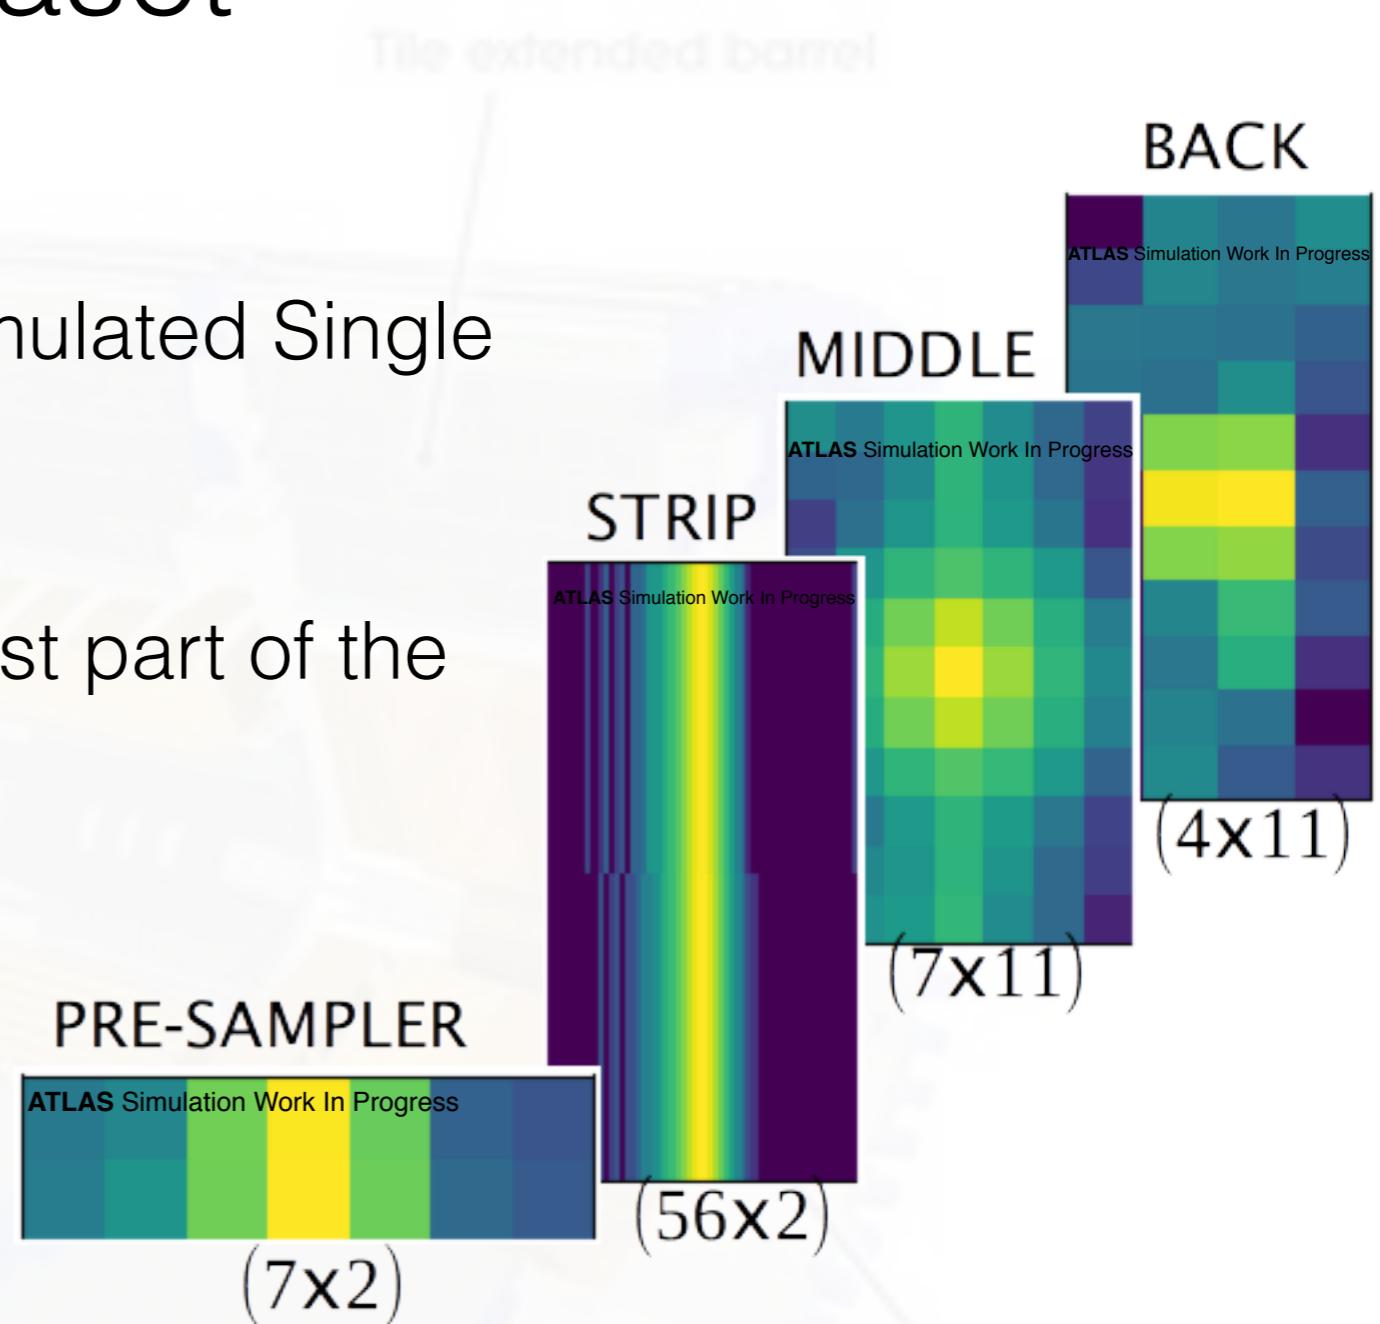
Converter: Keras → C++ framework

“If you have the flexibility to run any framework in your production environment, this package is *not* for you. **If you want to apply a network you've trained with Keras in a 6M line C++ production framework** that's only updated twice a year, you'll find this package very useful.”

-from [lwttn Github page](#)

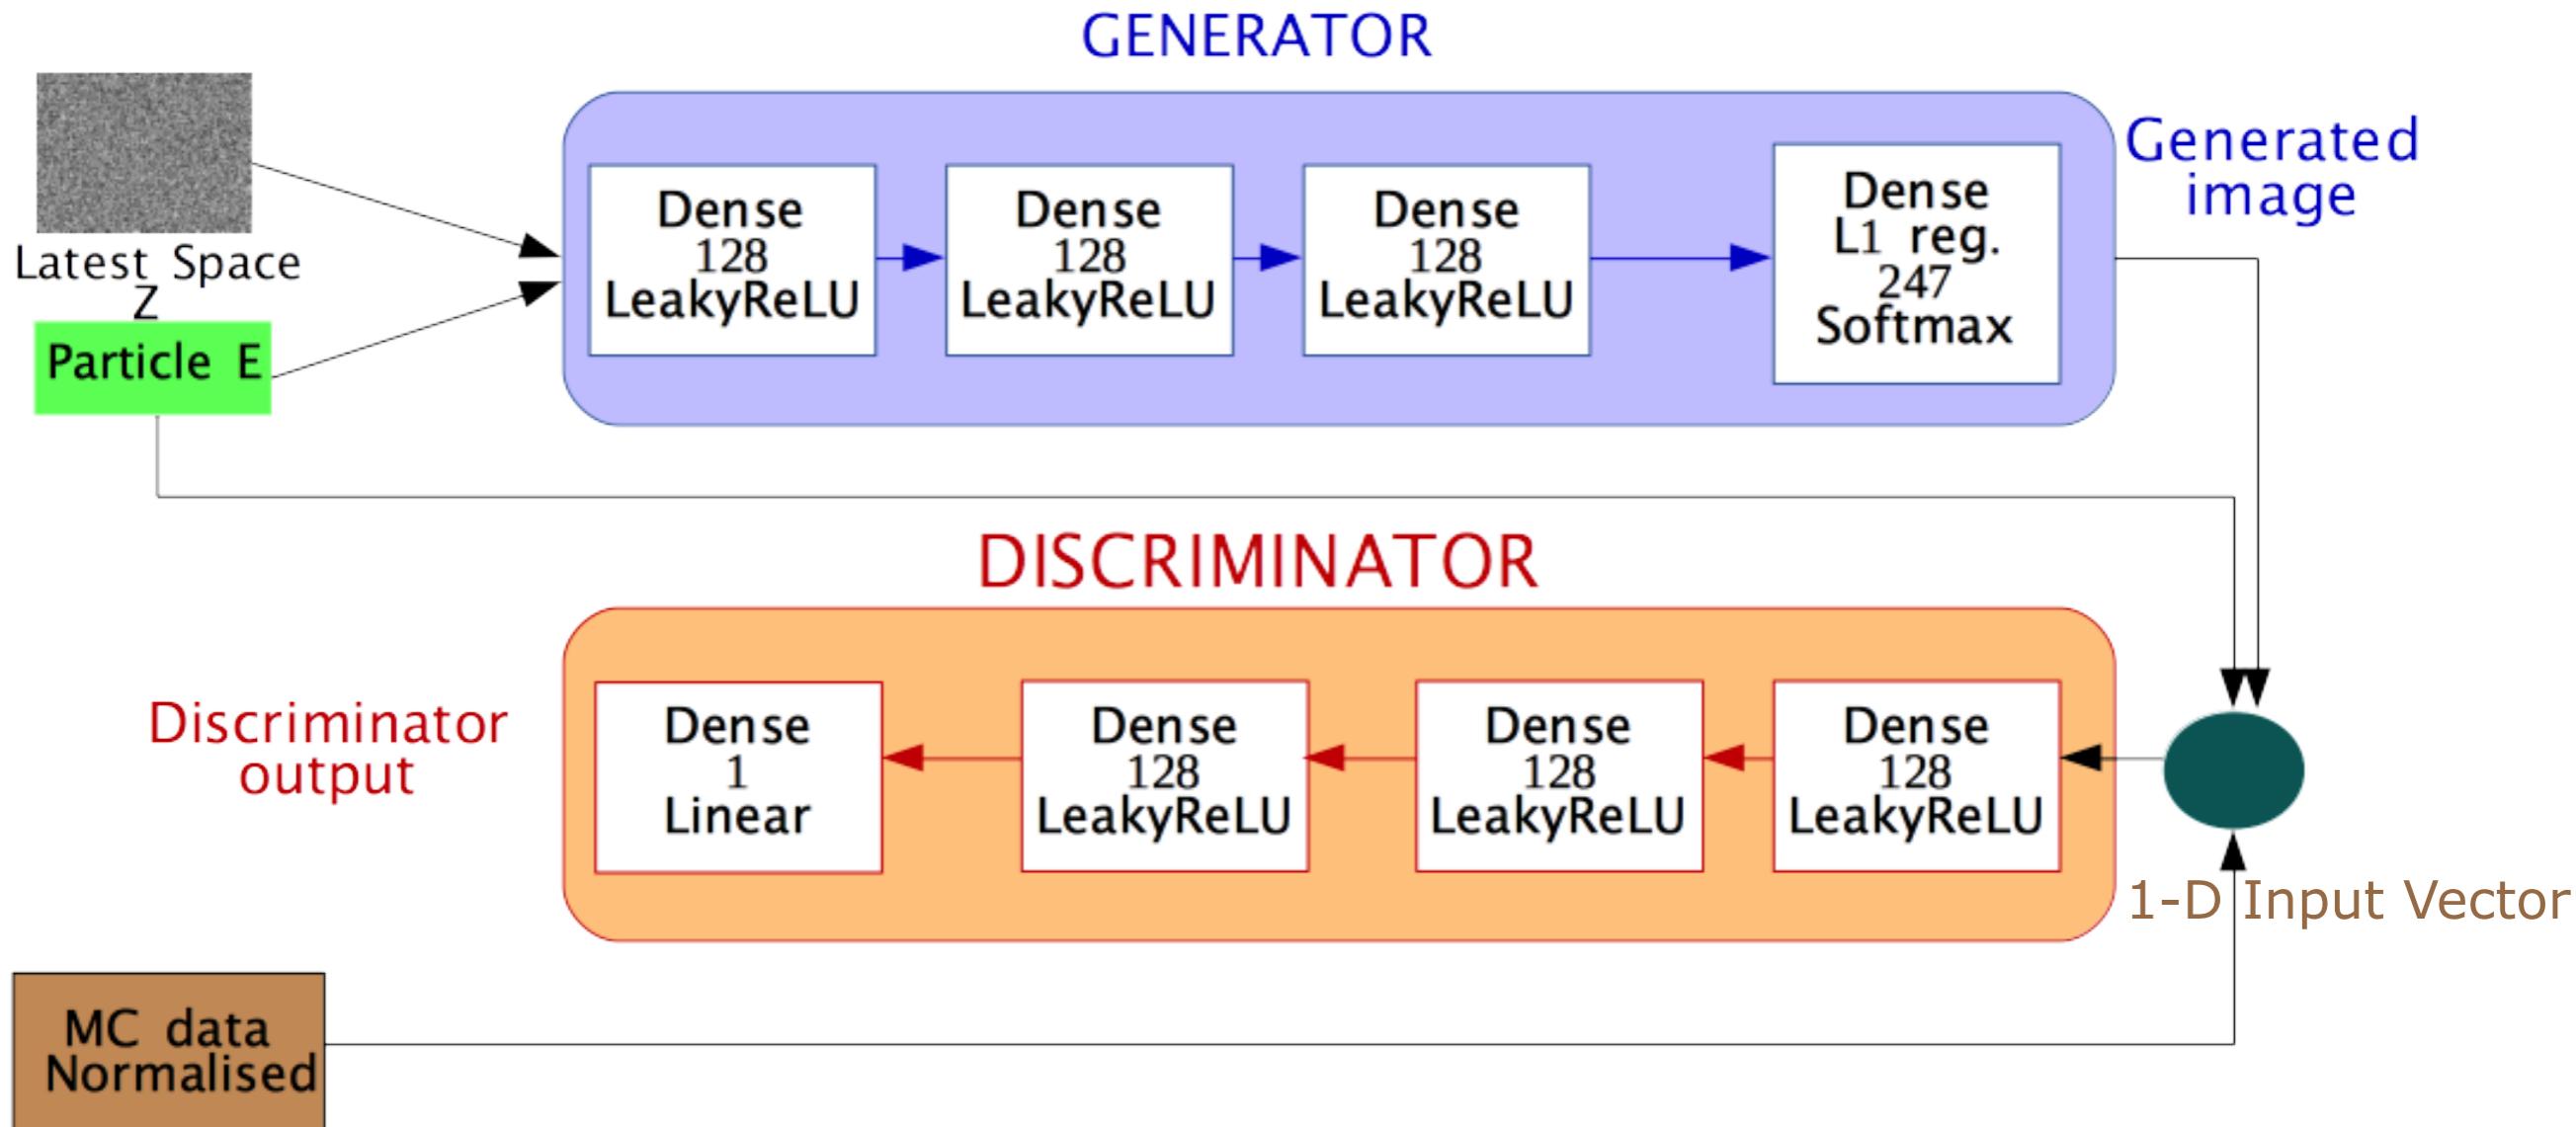
Dataset

- Train on **Geant4 Monte-Carlo** simulated Single Photon shower data
 - Use only the 'barrel' (simplest part of the calorimeter)
- GAN Training time: ~ 1h30m
- GPU speed: 2x over CPU
- GAN Generation Time: 1-2 seconds for 36k images



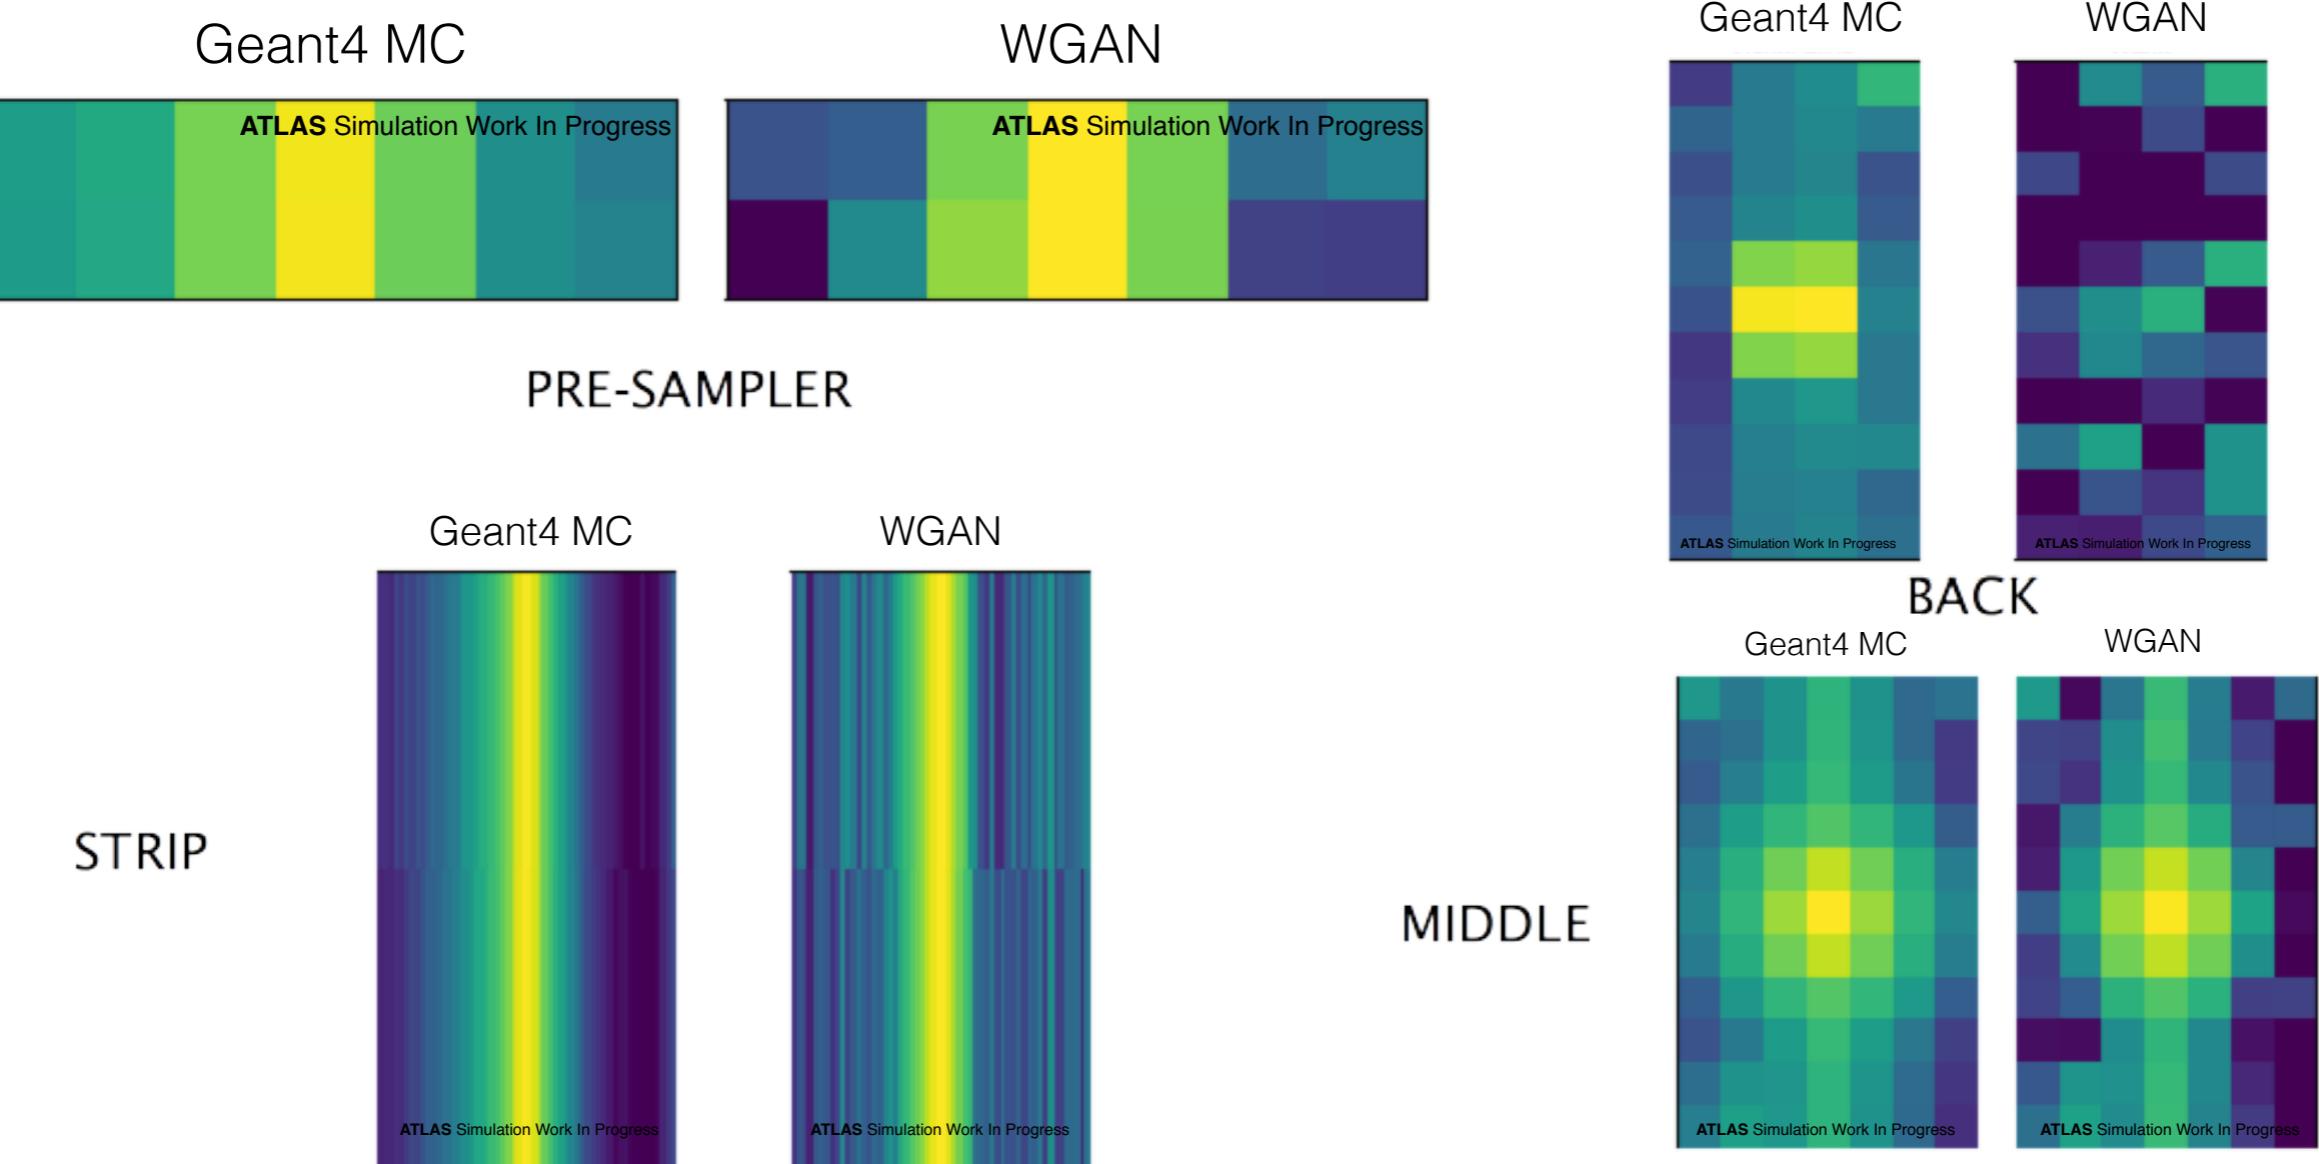
On the basis of Paul Klein's internship work with David Rousseau, Gilles Louppe, Kyle Cranmer at LAL & CERN, 2017

Our WGAN-GP Architecture



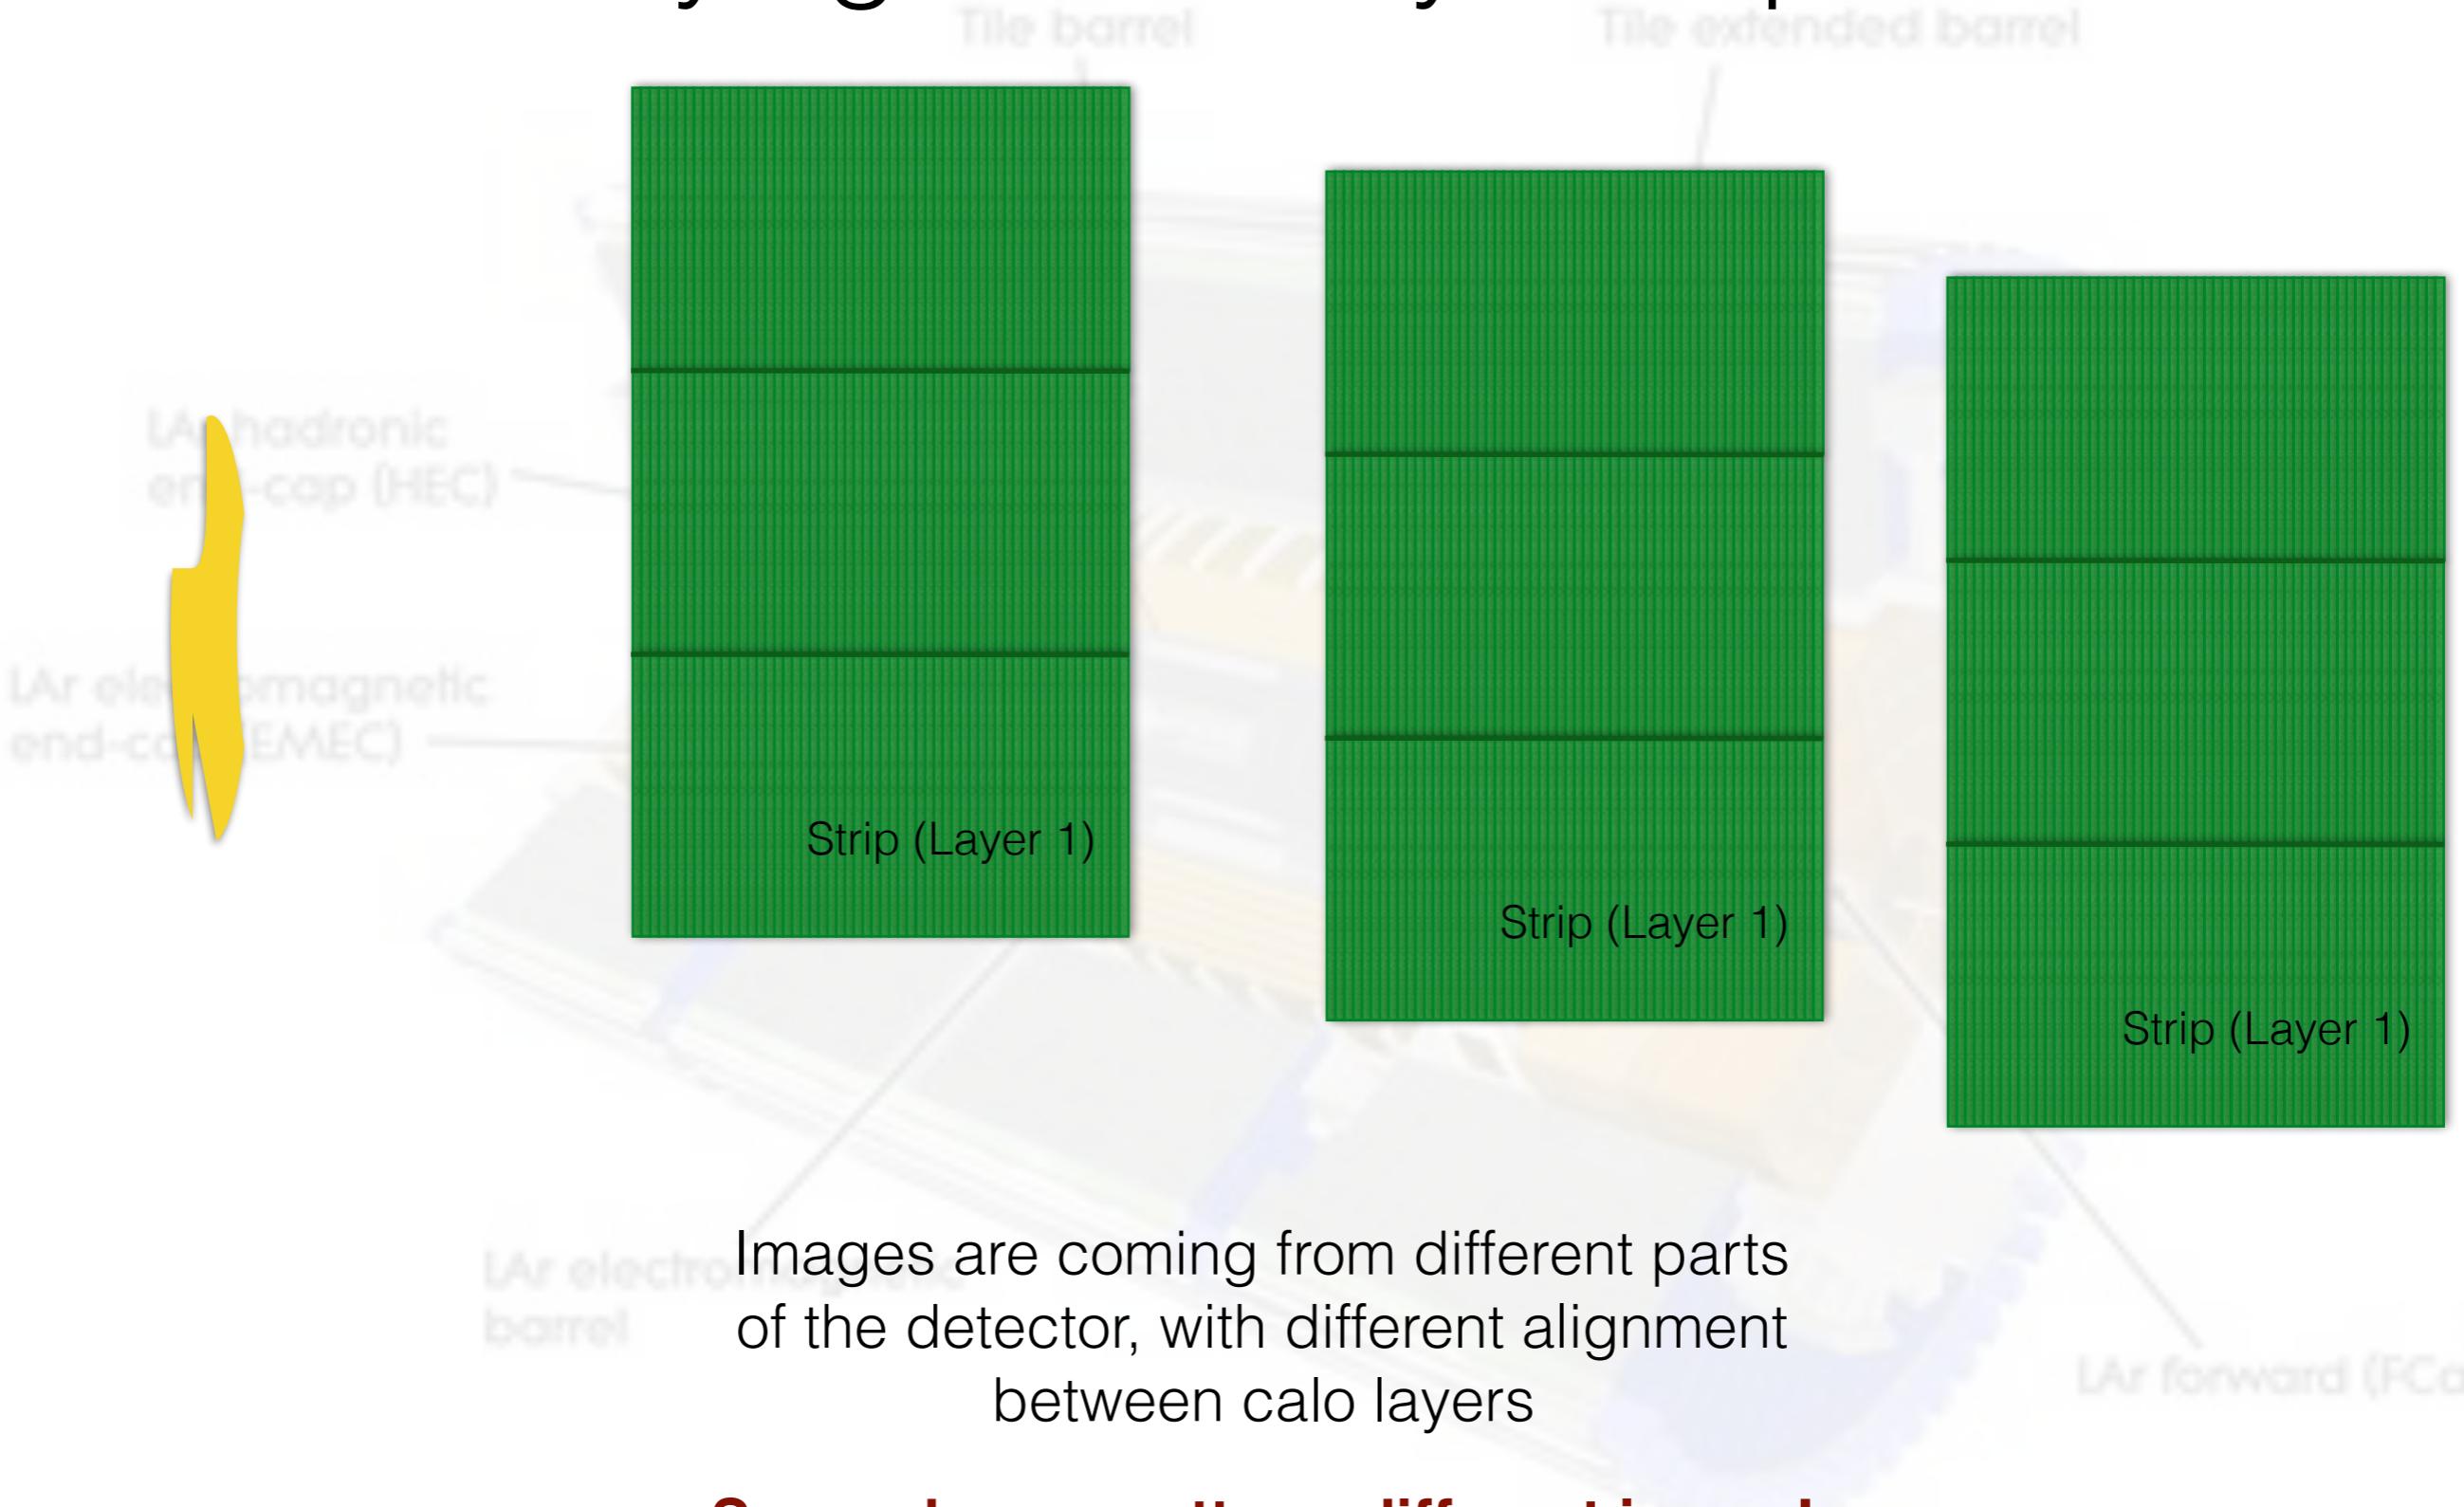
Total Trainable Parameters: G=103543, D=65025

Average Images

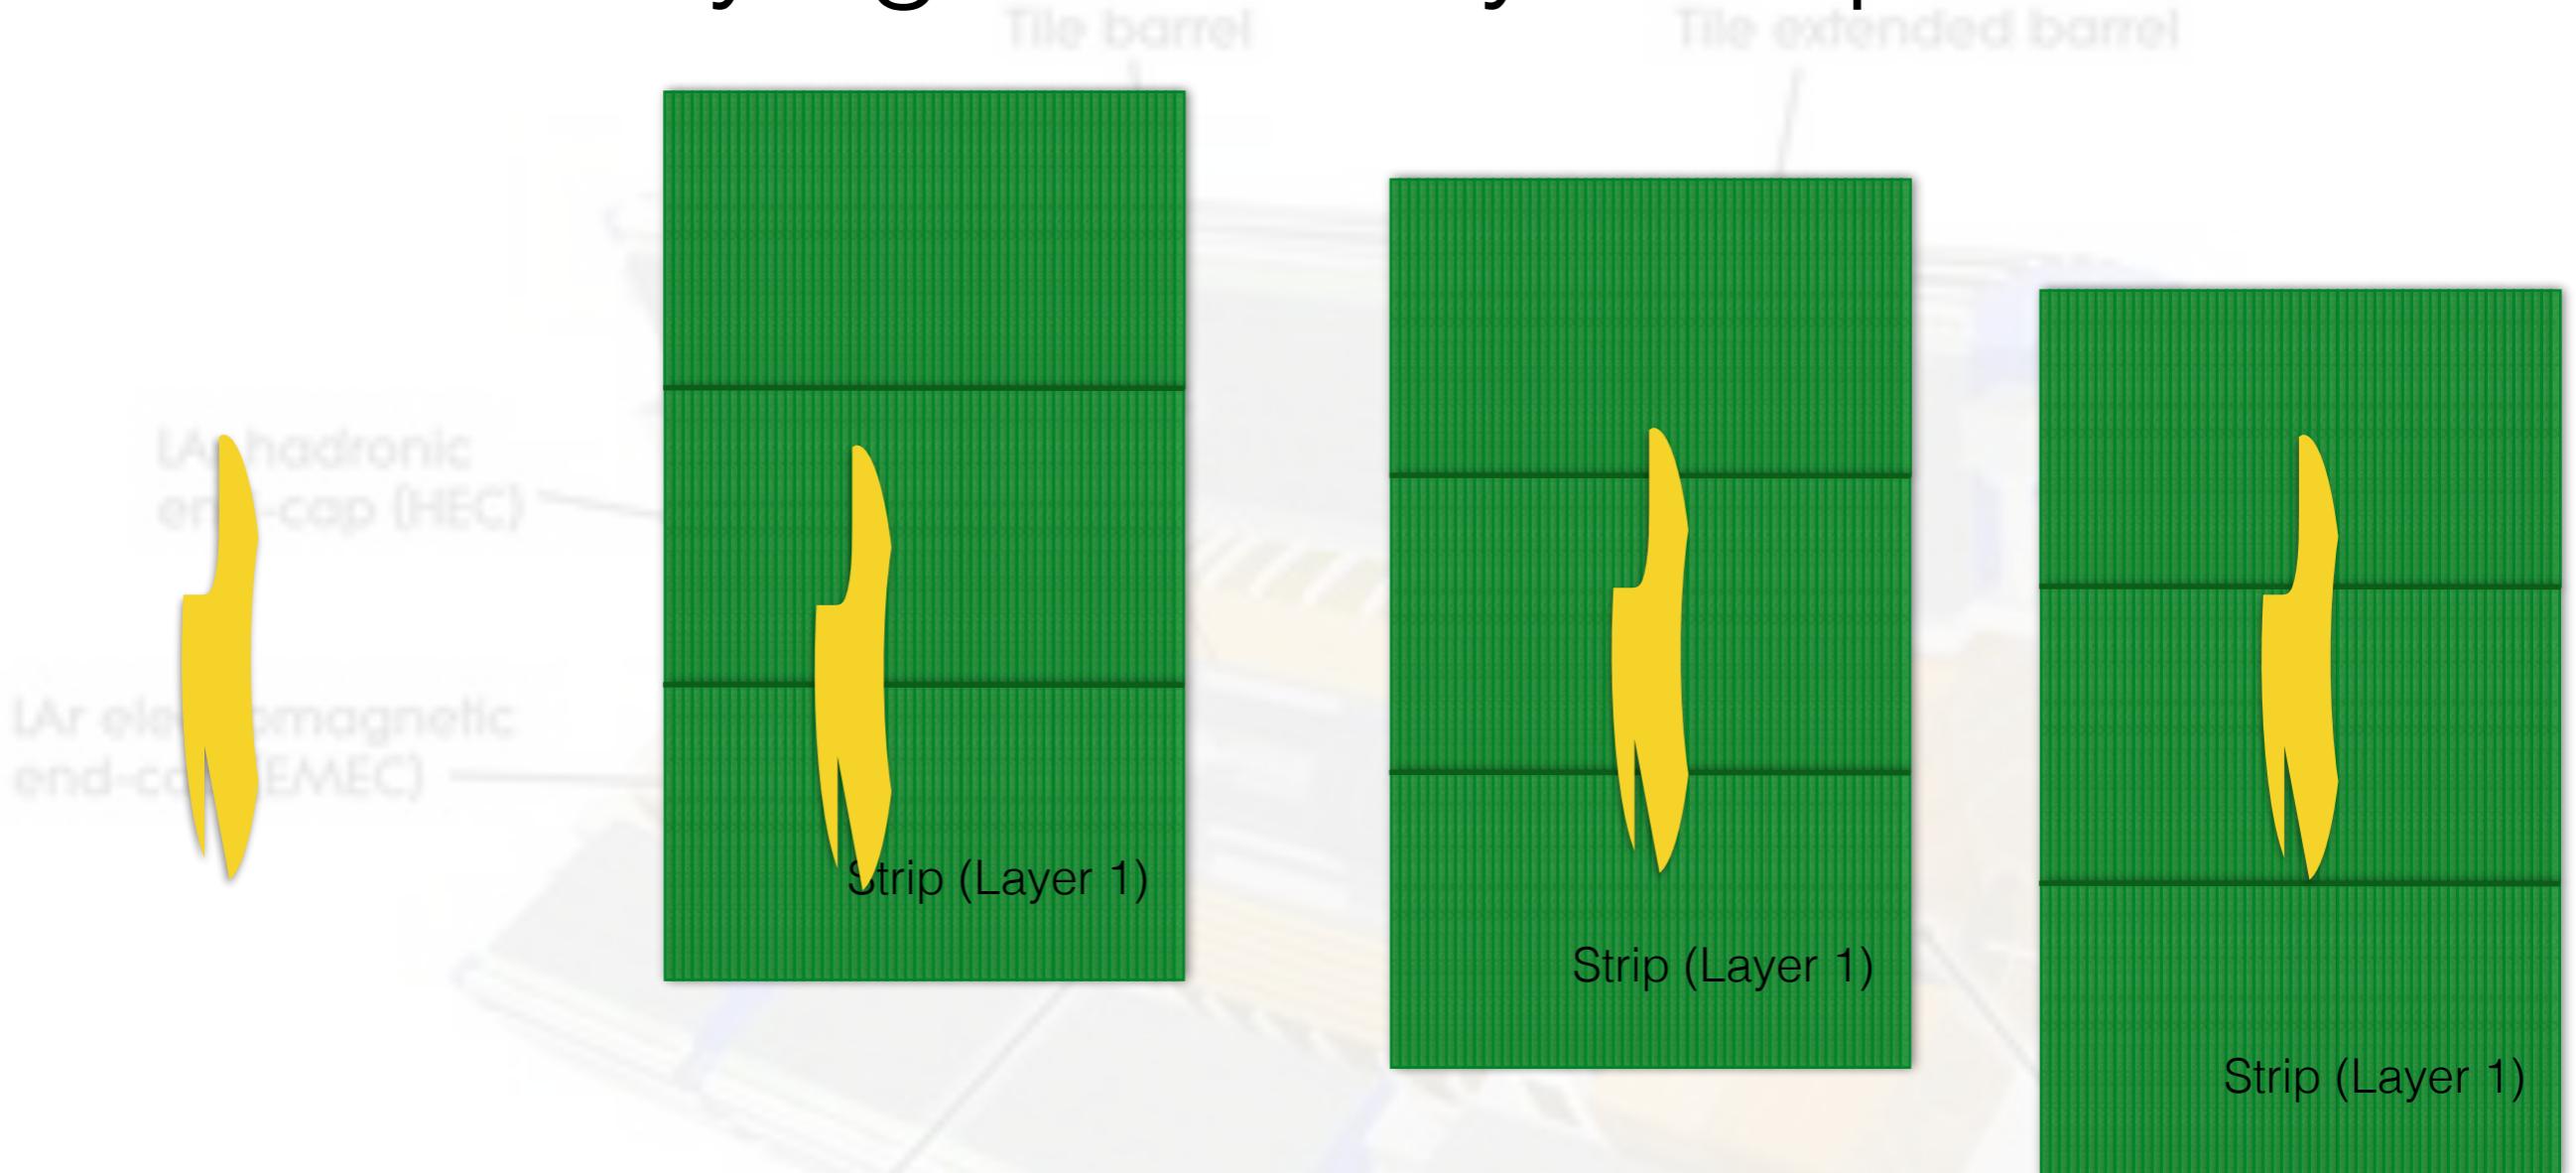


Warning: pixel sizes don't correspond to actual width of the cells

Underlying Geometry Complication



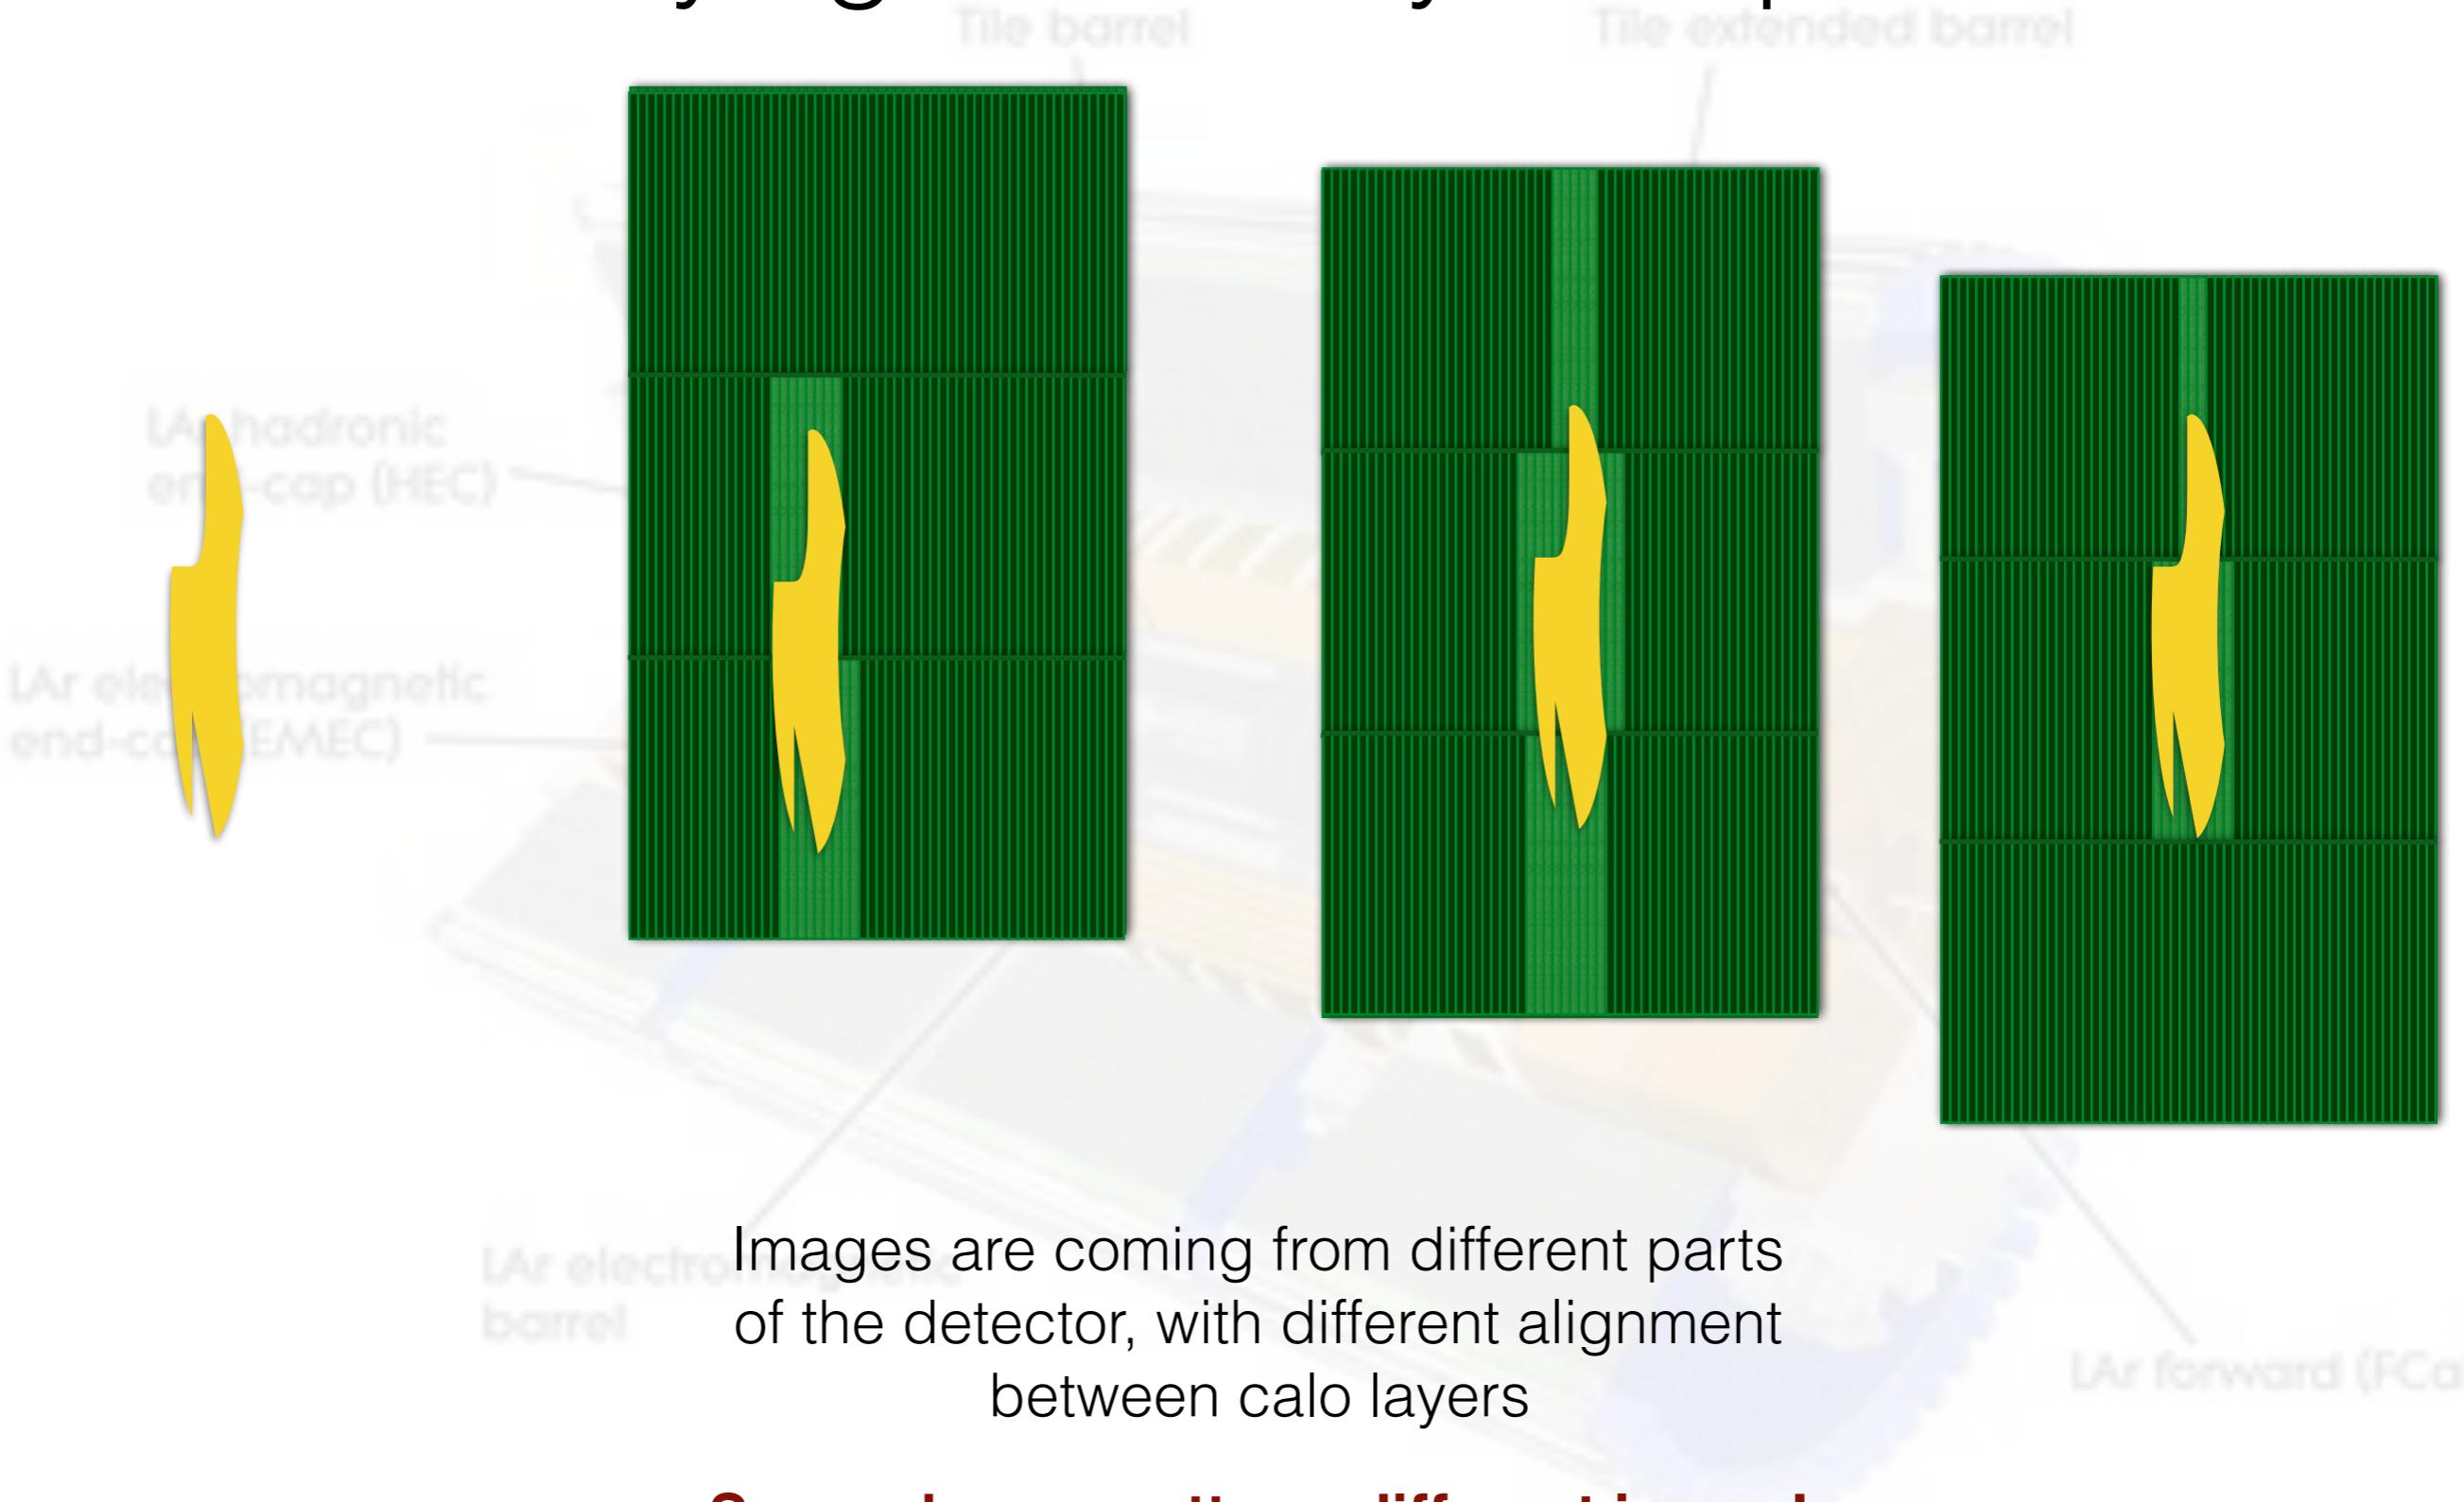
Underlying Geometry Complication



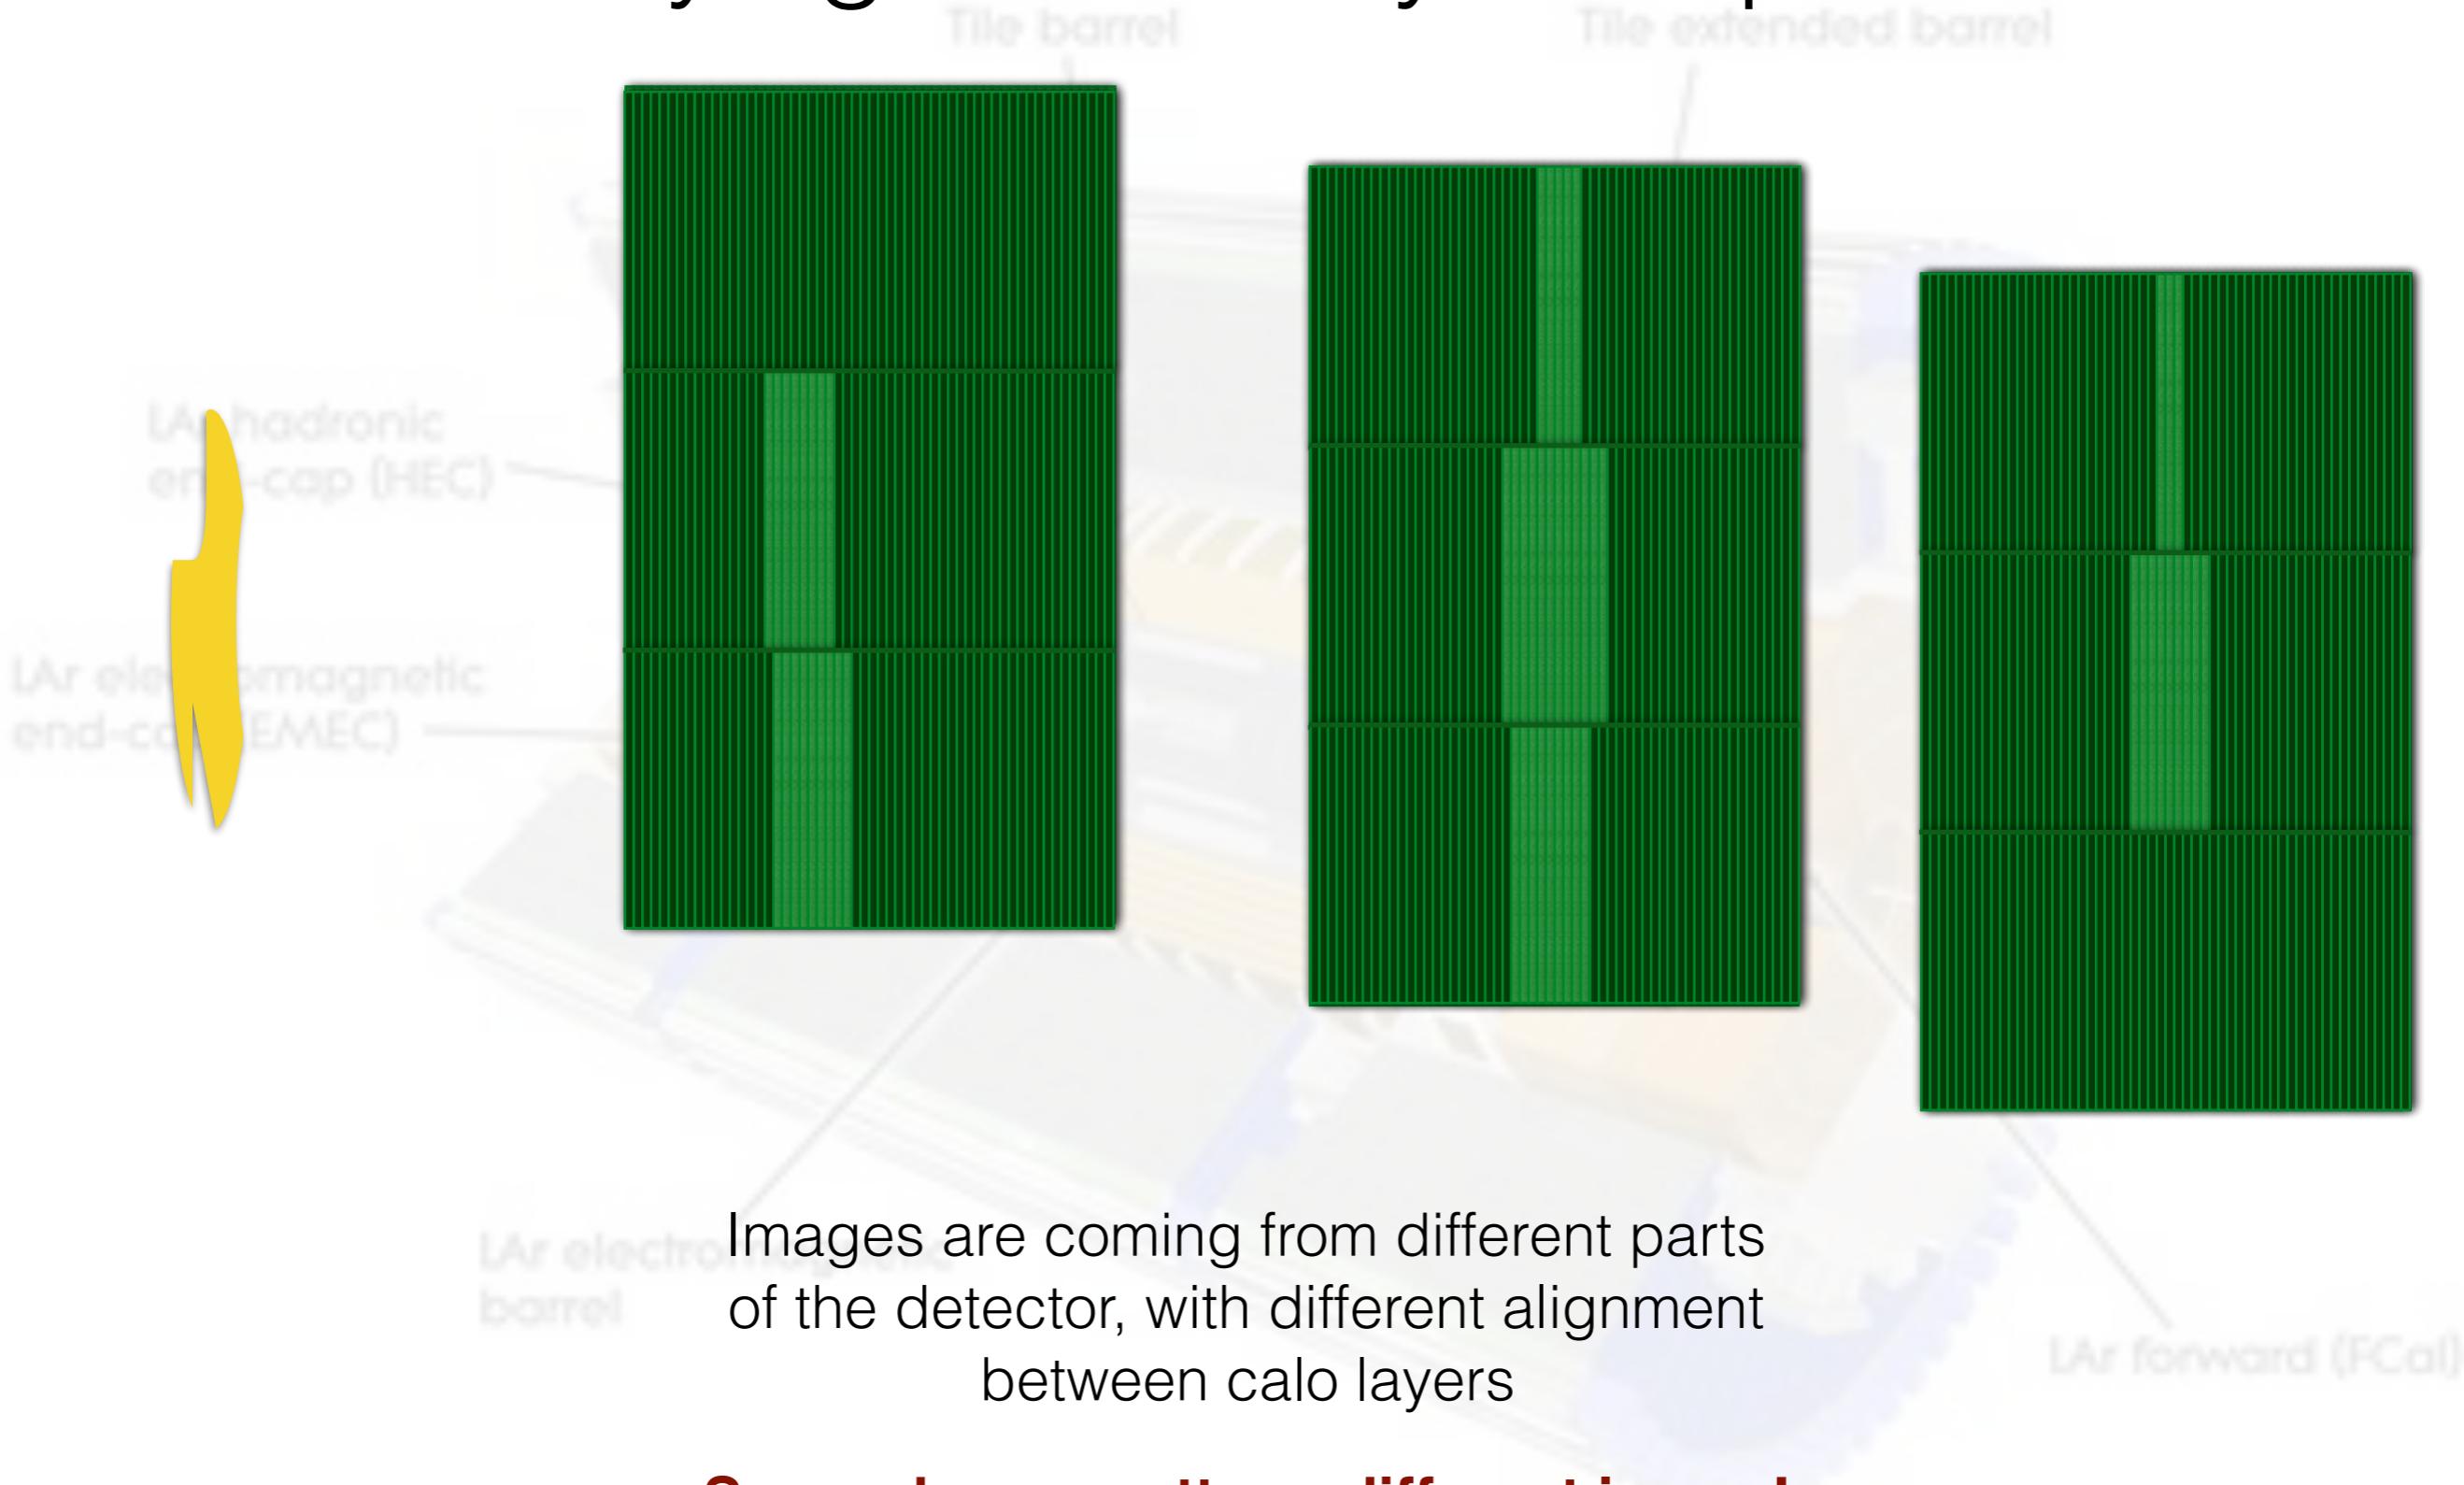
Images are coming from different parts
of the detector, with different alignment
between calo layers

Same shower pattern, different image!

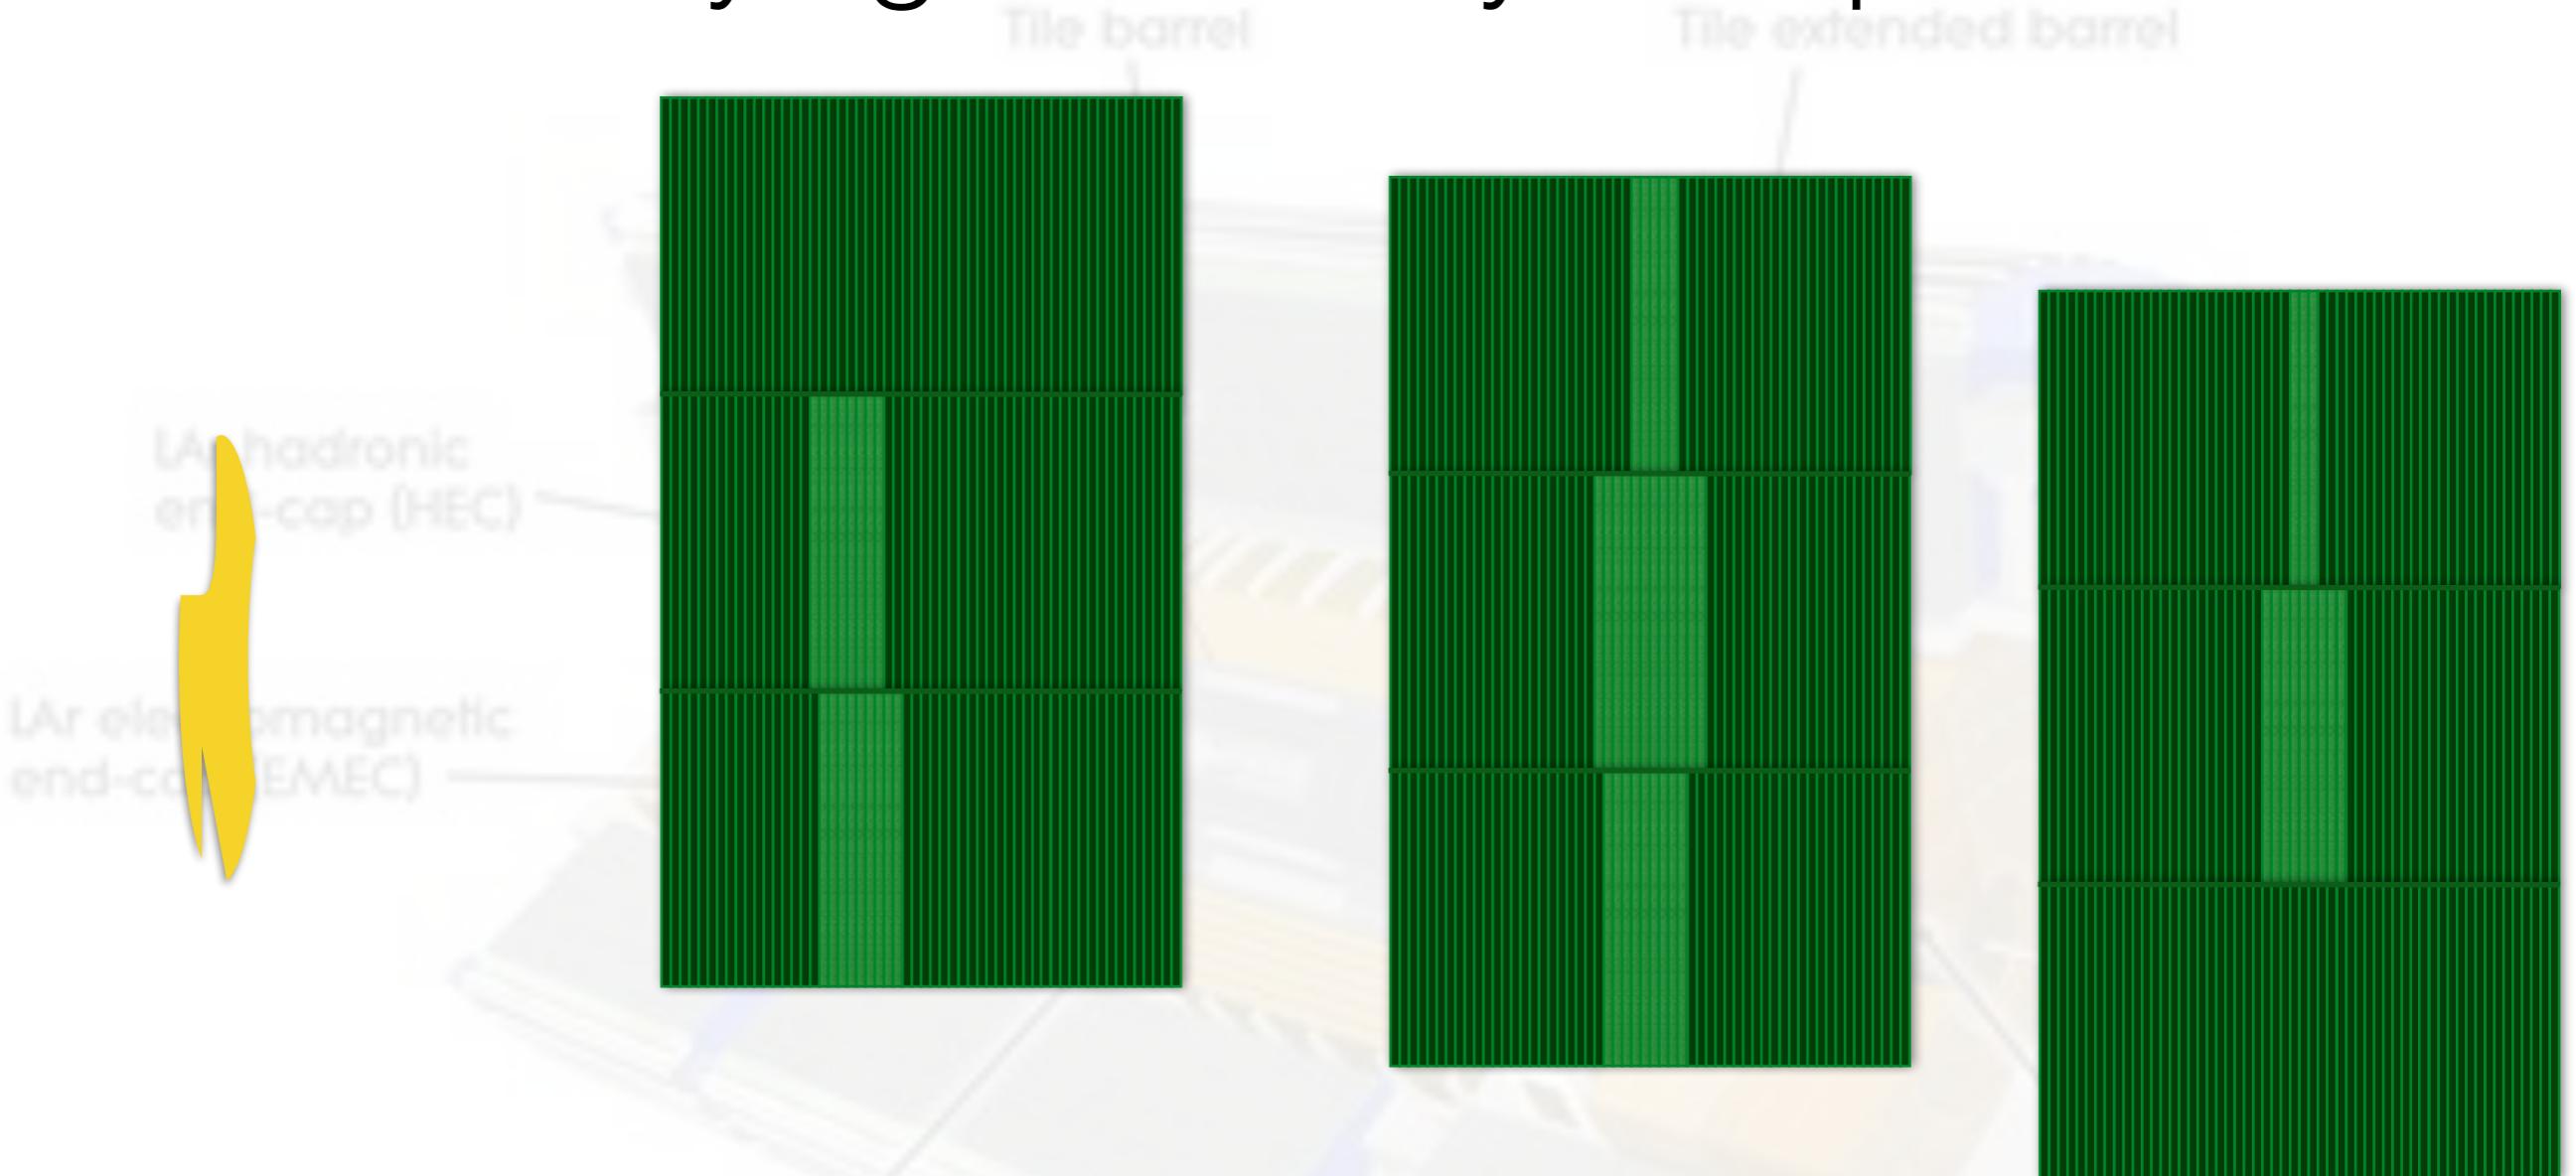
Underlying Geometry Complication



Underlying Geometry Complication



Underlying Geometry Complication



Train separately on images coming from
the same configuration of alignment
⇒ Smaller training dataset

Note: CaloGAN did not see this problem because particles shoot at single impact point in calorimeter

Variation in Training for Random Seed

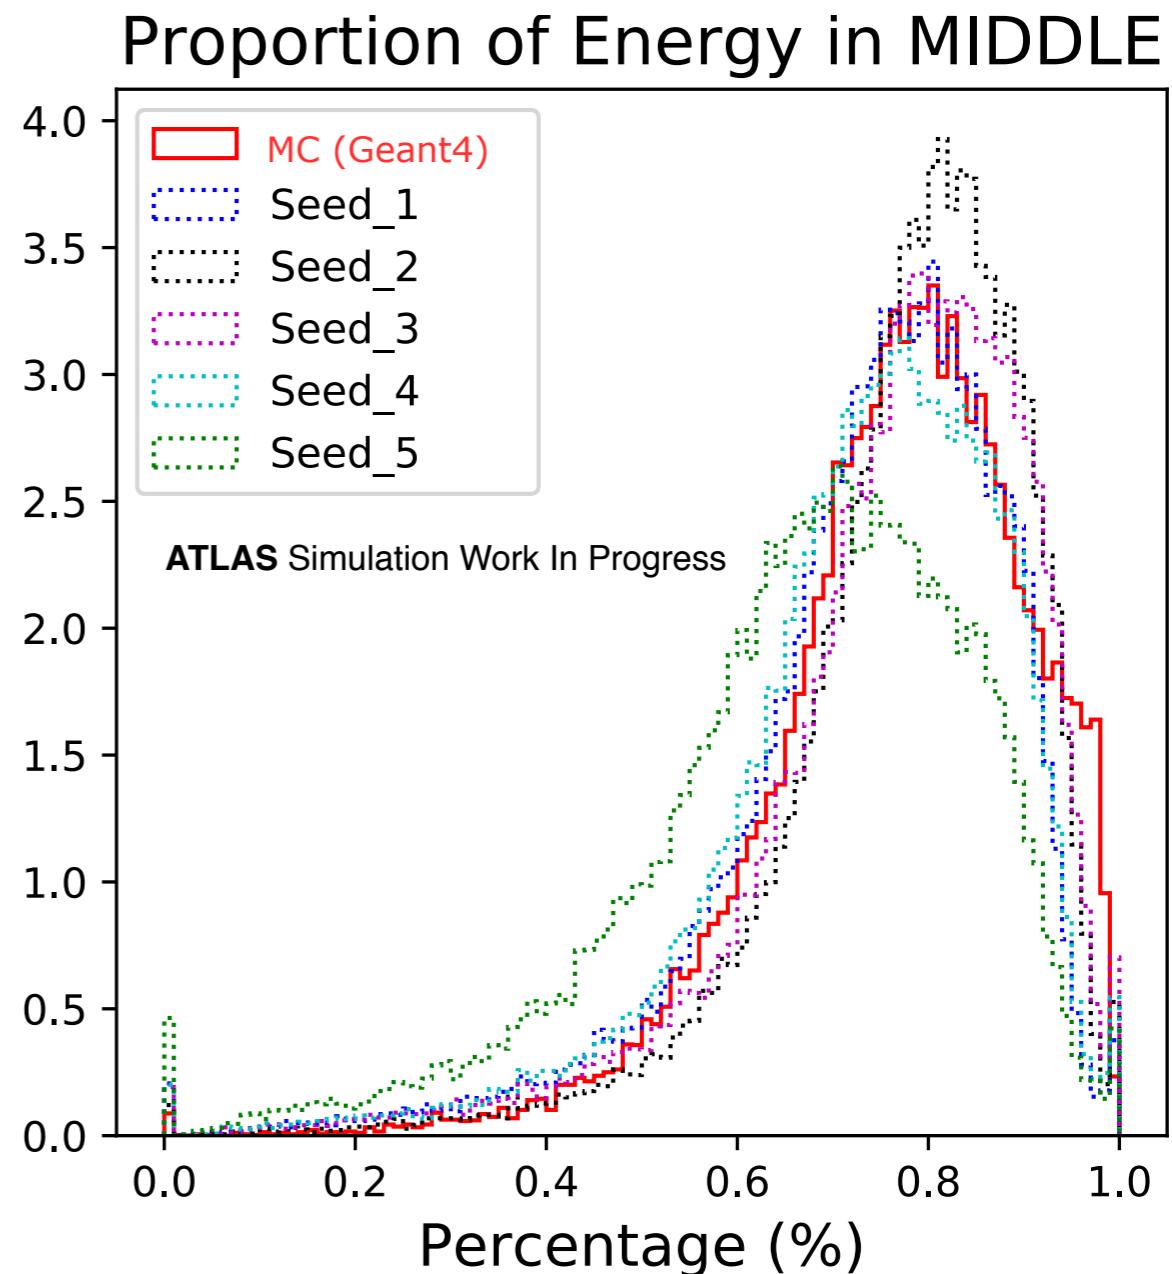
Training Seed Affects:

1. Initial random weights
2. Random noise input (Z) for generator
3. Batch Selection

Additional Randomness from GPU parallelisation

Do we pick the best one or do we average over 5?

Different training seeds



Initial training seed has significant impact on results. **Average over different seed** training to for Hyper-parameter optimisation

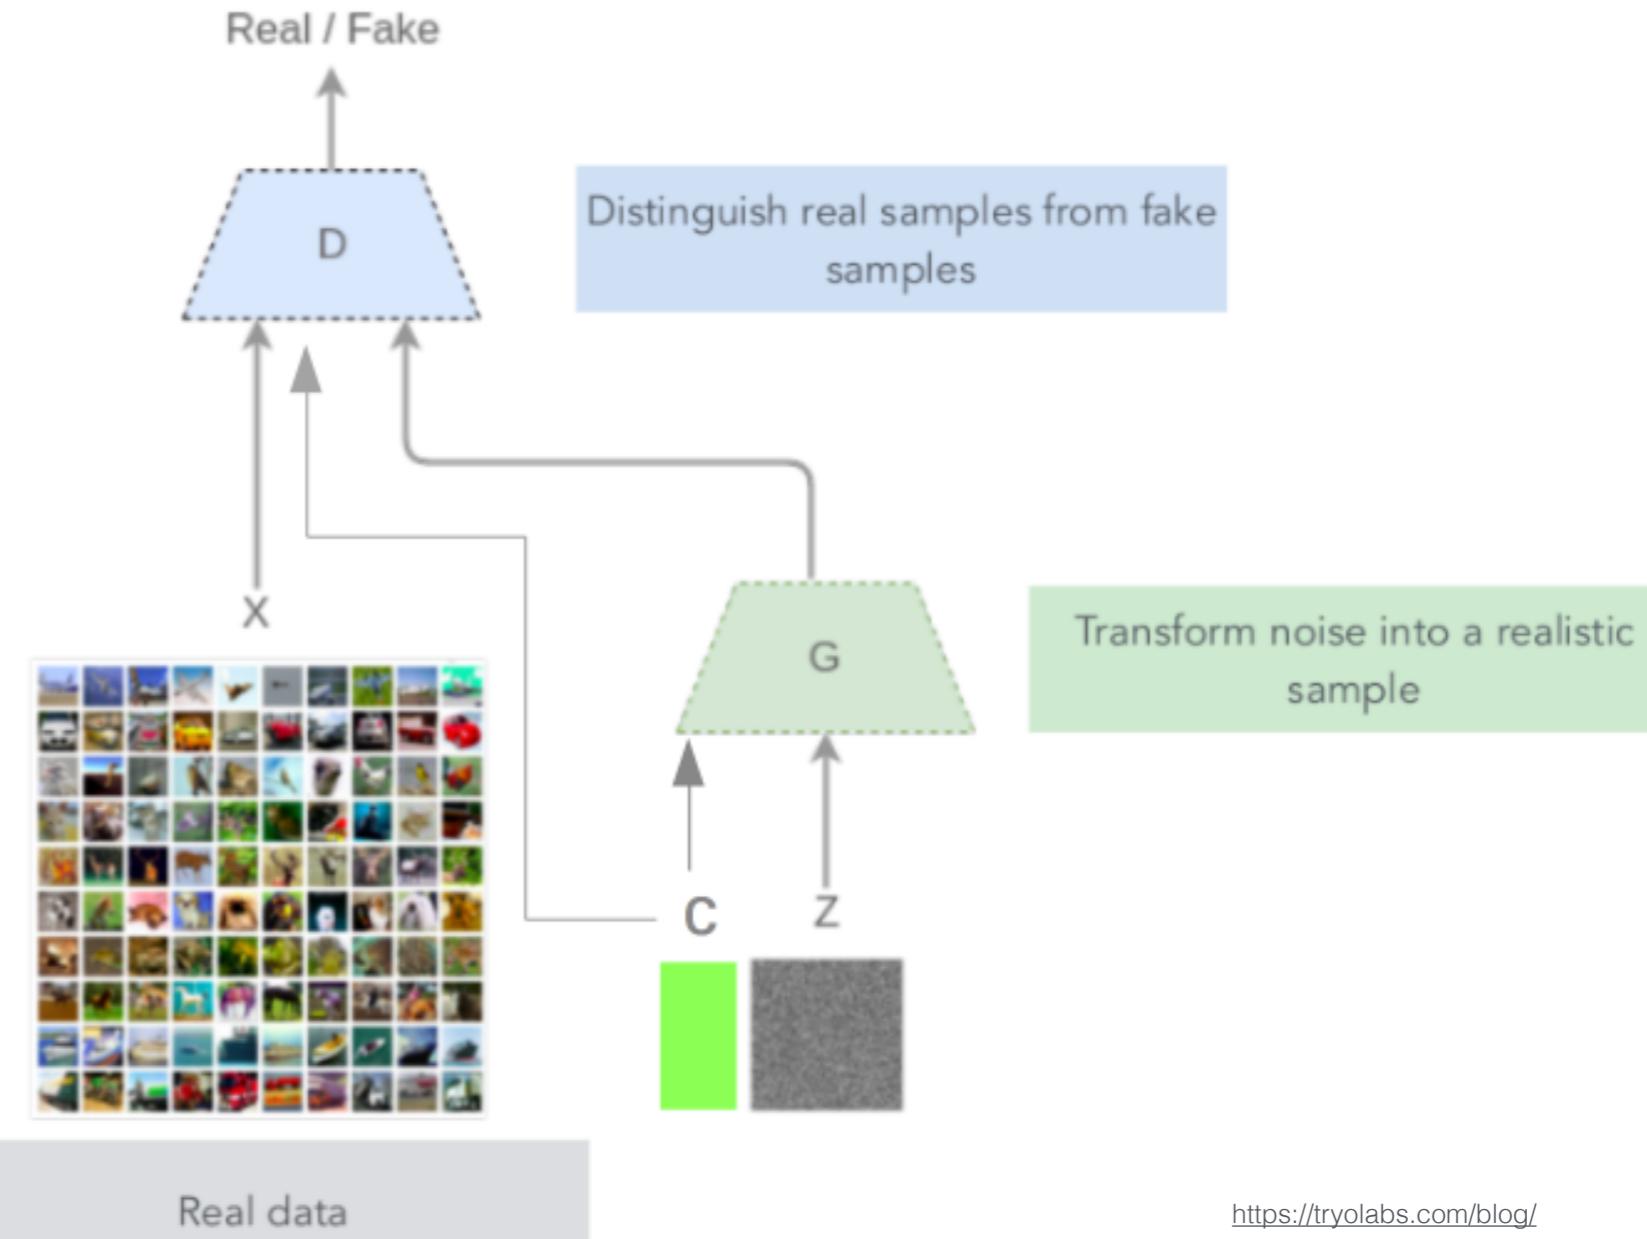
Remember the GAN game?

If D is trained too little, it will get fooled by G too easily, so G won't improve

Don't want to train D too much before every generator training

Training Ratio:

Number of discriminator training iterations for each generator iteration



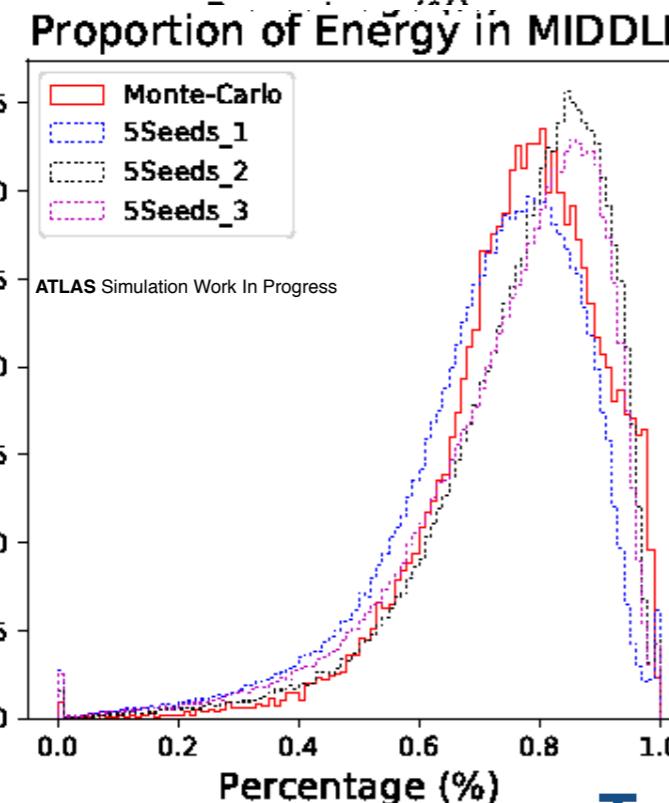
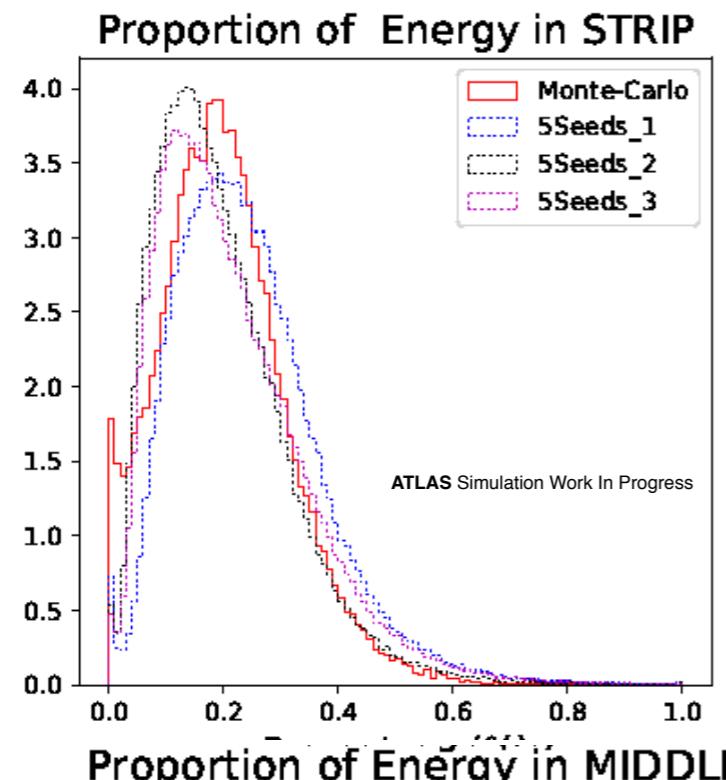
<https://tryolabs.com/blog/>

Hyperparameter Optimisation: Training Ratio

Same hyper-params,

3 sets of 5 separately trained GAN datasets

Random Variations in performance

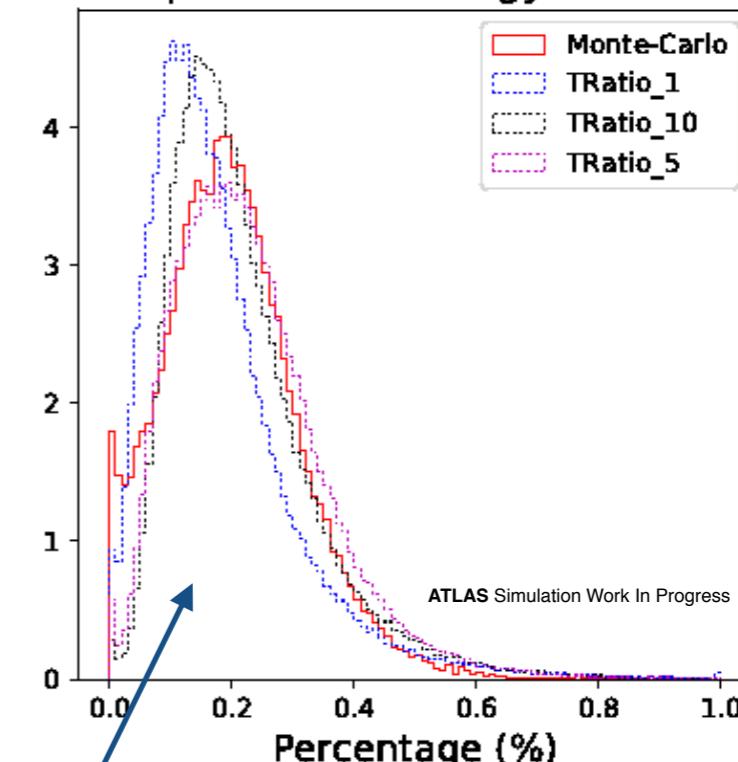


Monte-Carlo is from Geant4

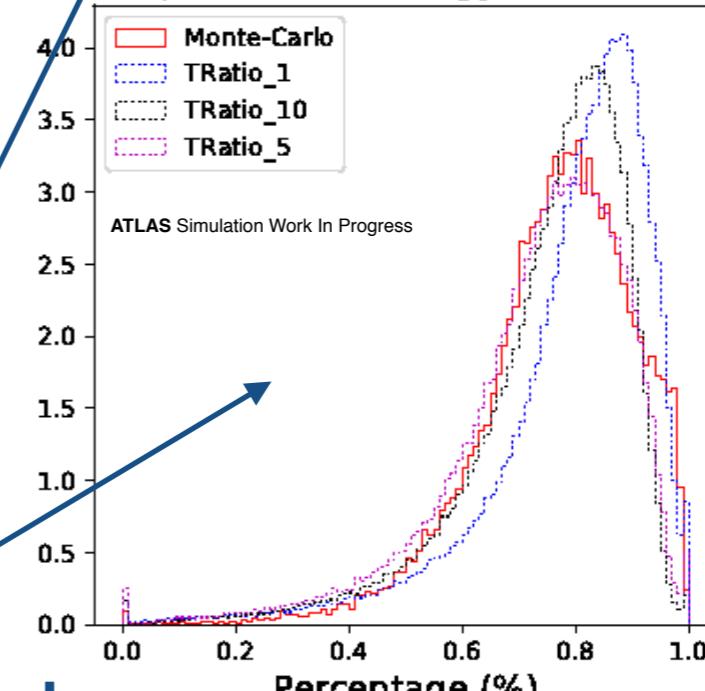
Training Ratio of 5 is good

1,5,10 Ratio; combine 5 trained GANs for each

Proportion of Energy in STRIP



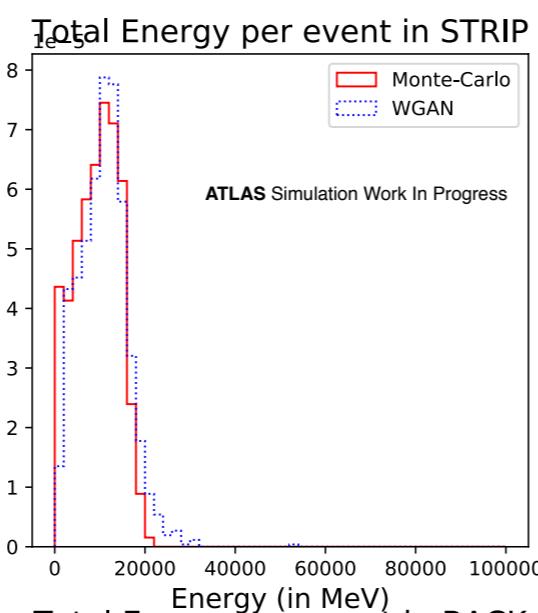
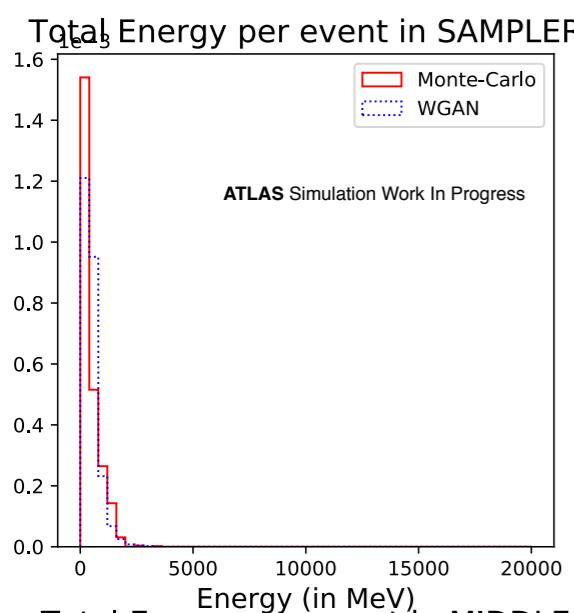
Proportion of Energy in MIDDLE



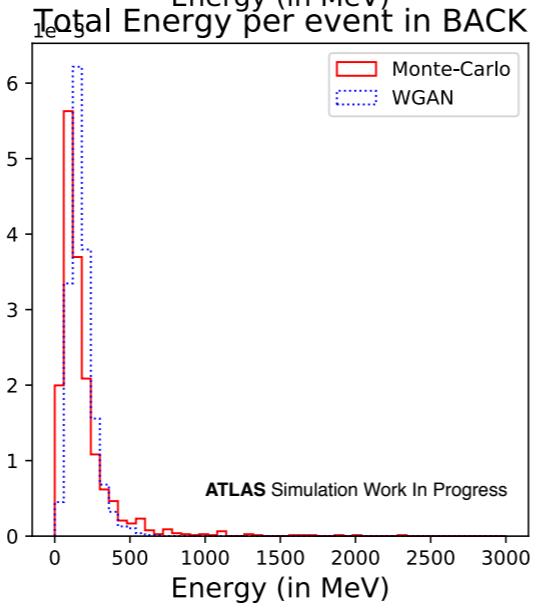
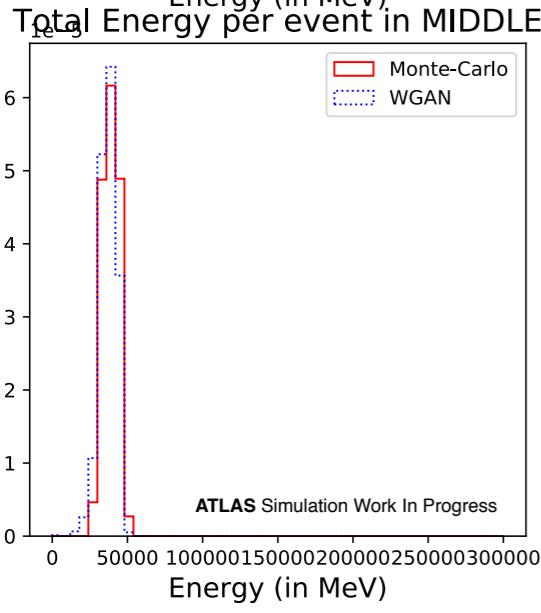
Non-random improvement in performance

Energy Distribution in Each Calo Layer

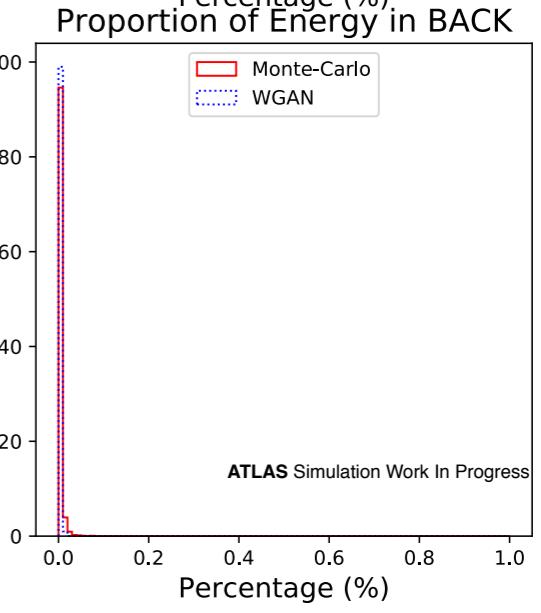
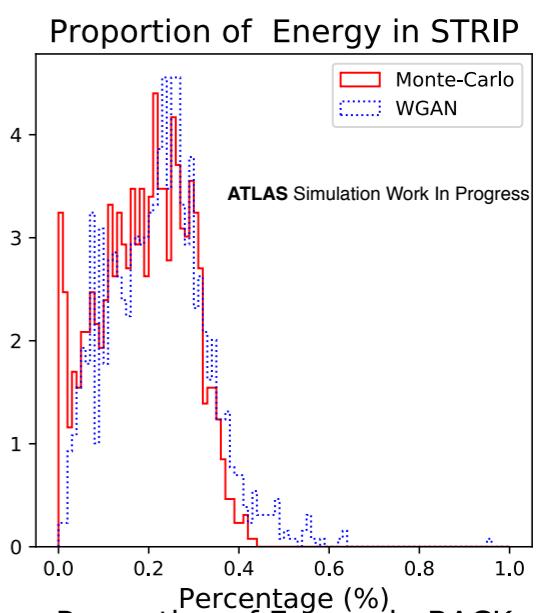
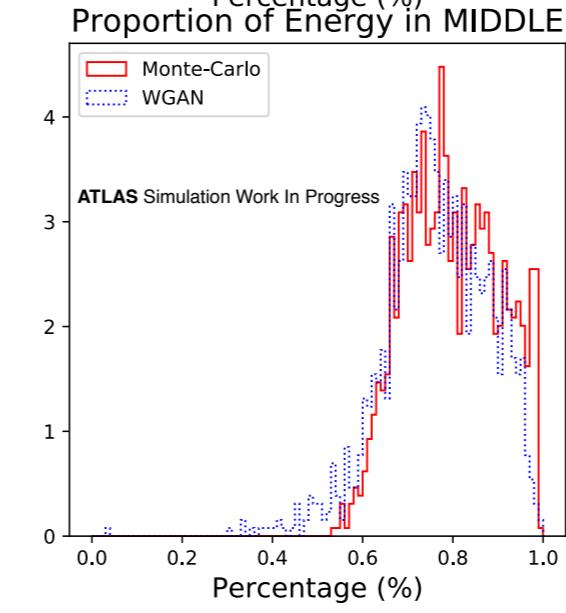
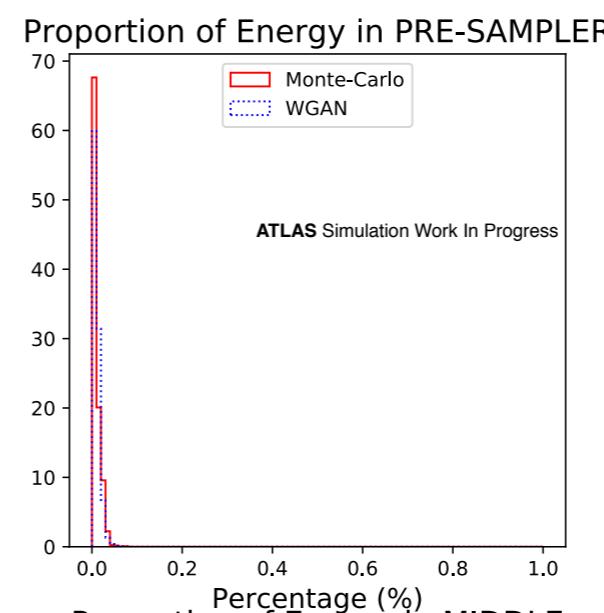
Total Energy



Monte-Carlo is from
Geant4



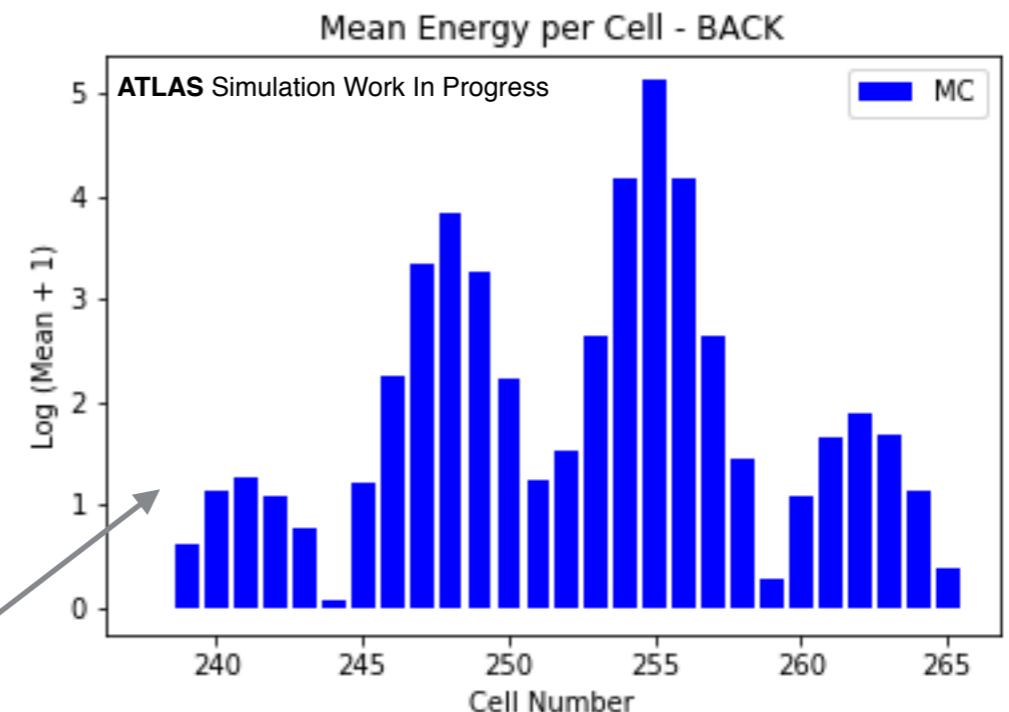
Fraction of Energy



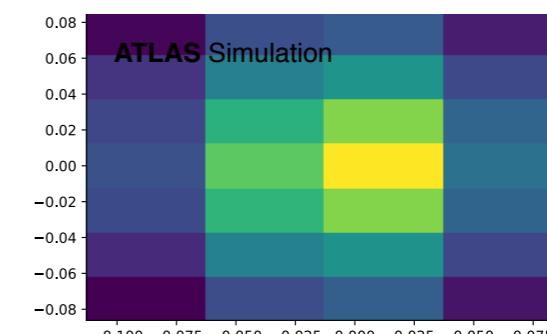
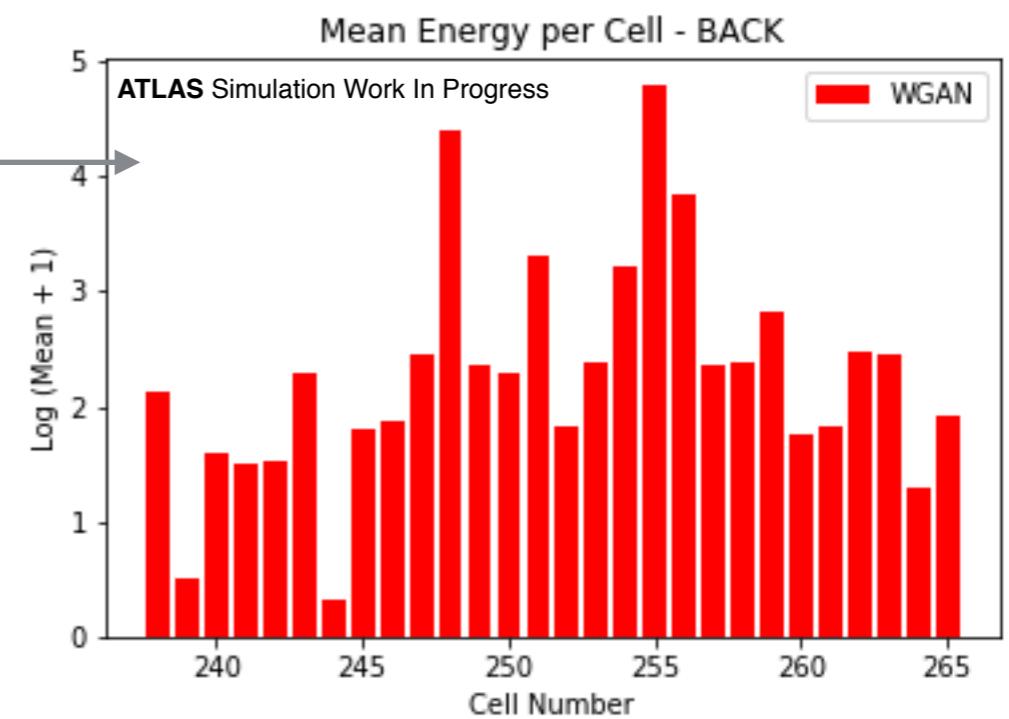
Energy Per Cell in BACK (L3)

The 4 peaks, periodic in 7 are now perfectly understandable

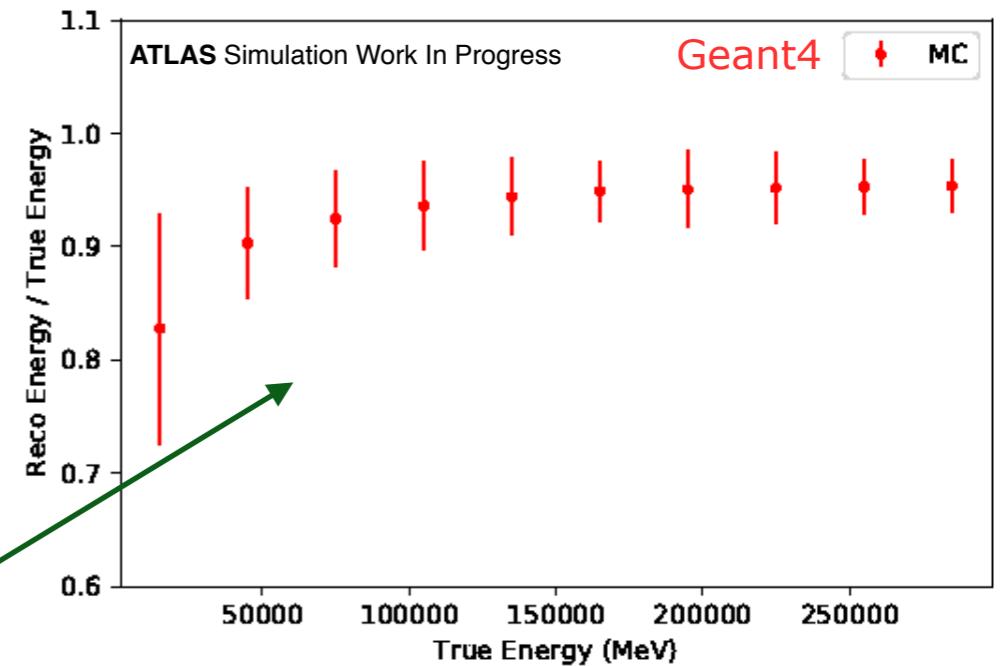
Monte-Carlo is from Geant4



GAN doesn't learn this structure well



Energy resolution of Detector



The energy resolution of the detector gets better at higher energies:

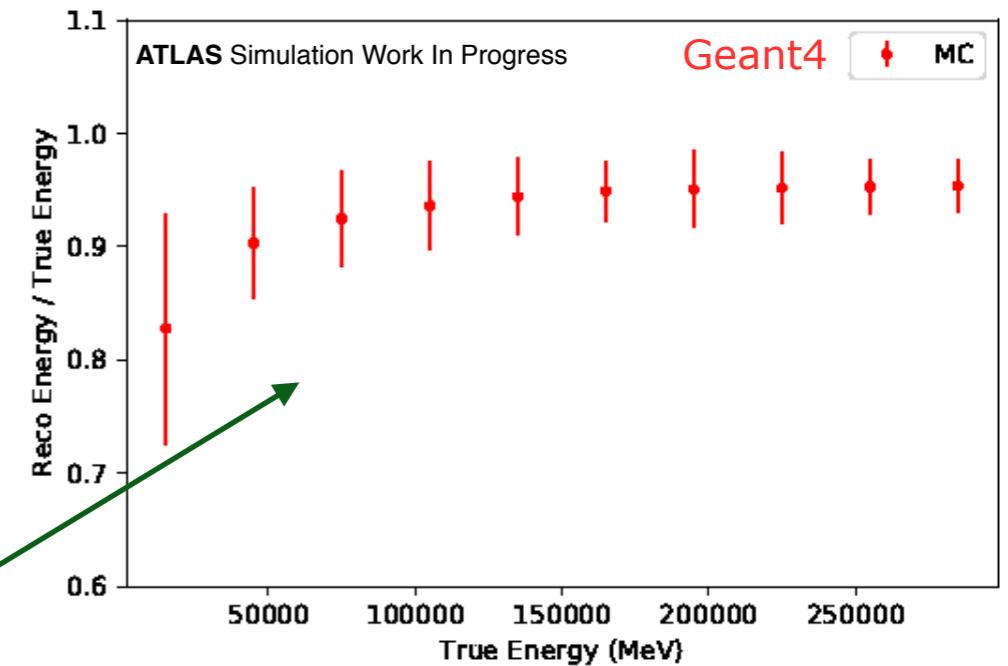
$$\sigma_E \sim 10\% \sqrt{E}$$

Can the GAN mimic this?

Energy resolution of Detector

To reproduce the energy resolution of the detector:

- Condition on True Energy of Particle
- ‘Normalise’ cell energies to by true energy of particle, **($\Sigma(\text{cell_energies}) \neq 1$ anymore)**
- Change activation function in the final layer from softmax to relu



$$\sigma_E \sim 10\% \sqrt{E}$$

The error bars are the **standard deviation in each bin**

Softmax:

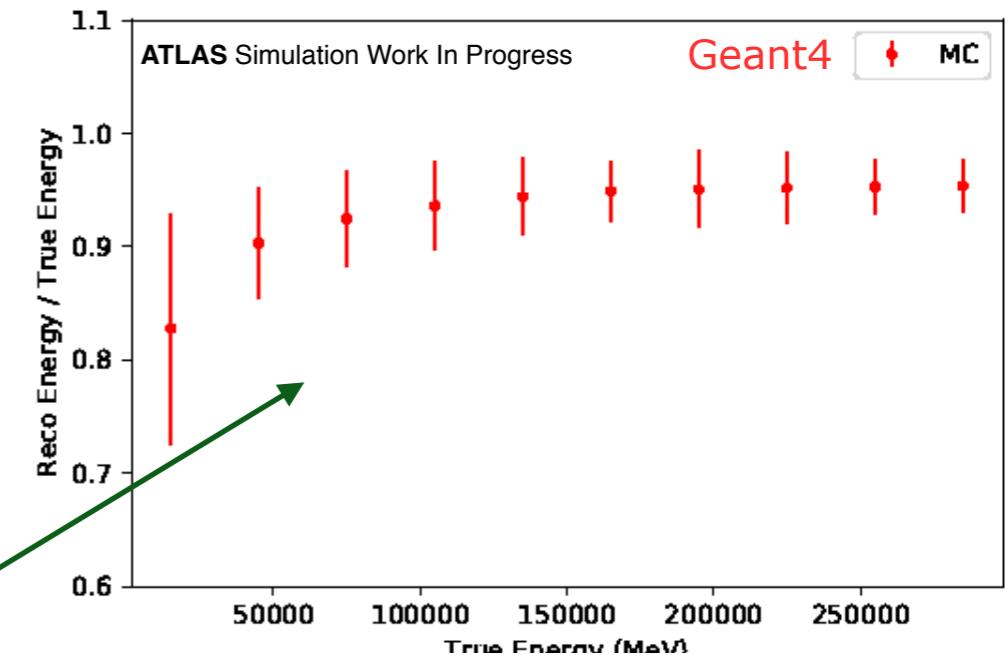
$$\sigma(z)_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

Enforces normalisation of output ($\Sigma = 1$)

Energy resolution of Detector

To reproduce the energy resolution of the detector:

- Condition on True Energy of Particle
- ‘Normalise’ cell energies to by true energy of particle, **($\Sigma(\text{cell_energies}) \neq 1$ anymore)**
- Change activation function in the final layer from softmax to relu



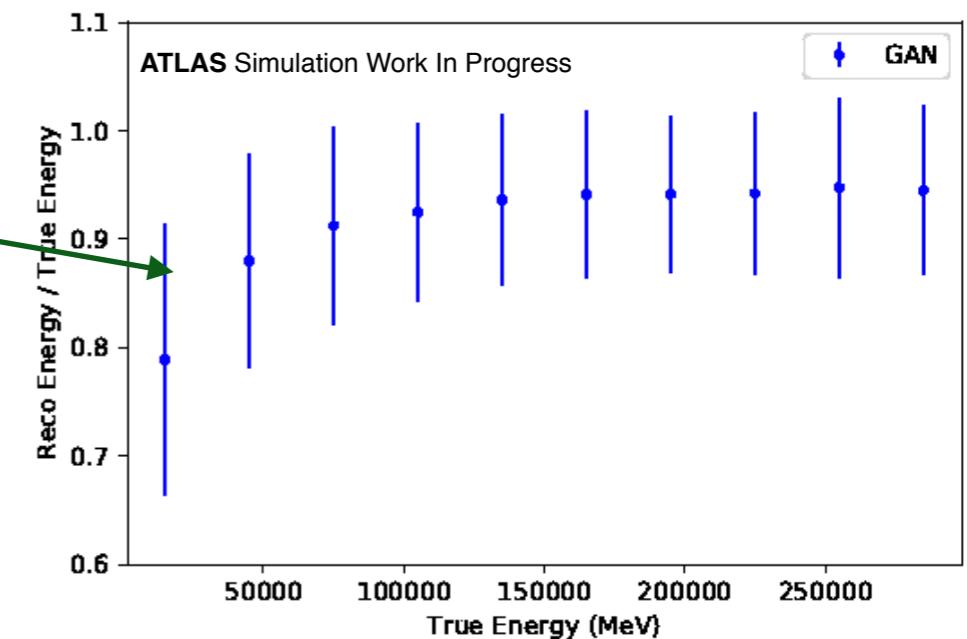
$$\sigma_E \sim 10\% \sqrt{E}$$

The error bars are the **standard deviation in each bin**

Softmax:

$$\sigma(z)_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

Enforces normalisation of output ($\Sigma = 1$)

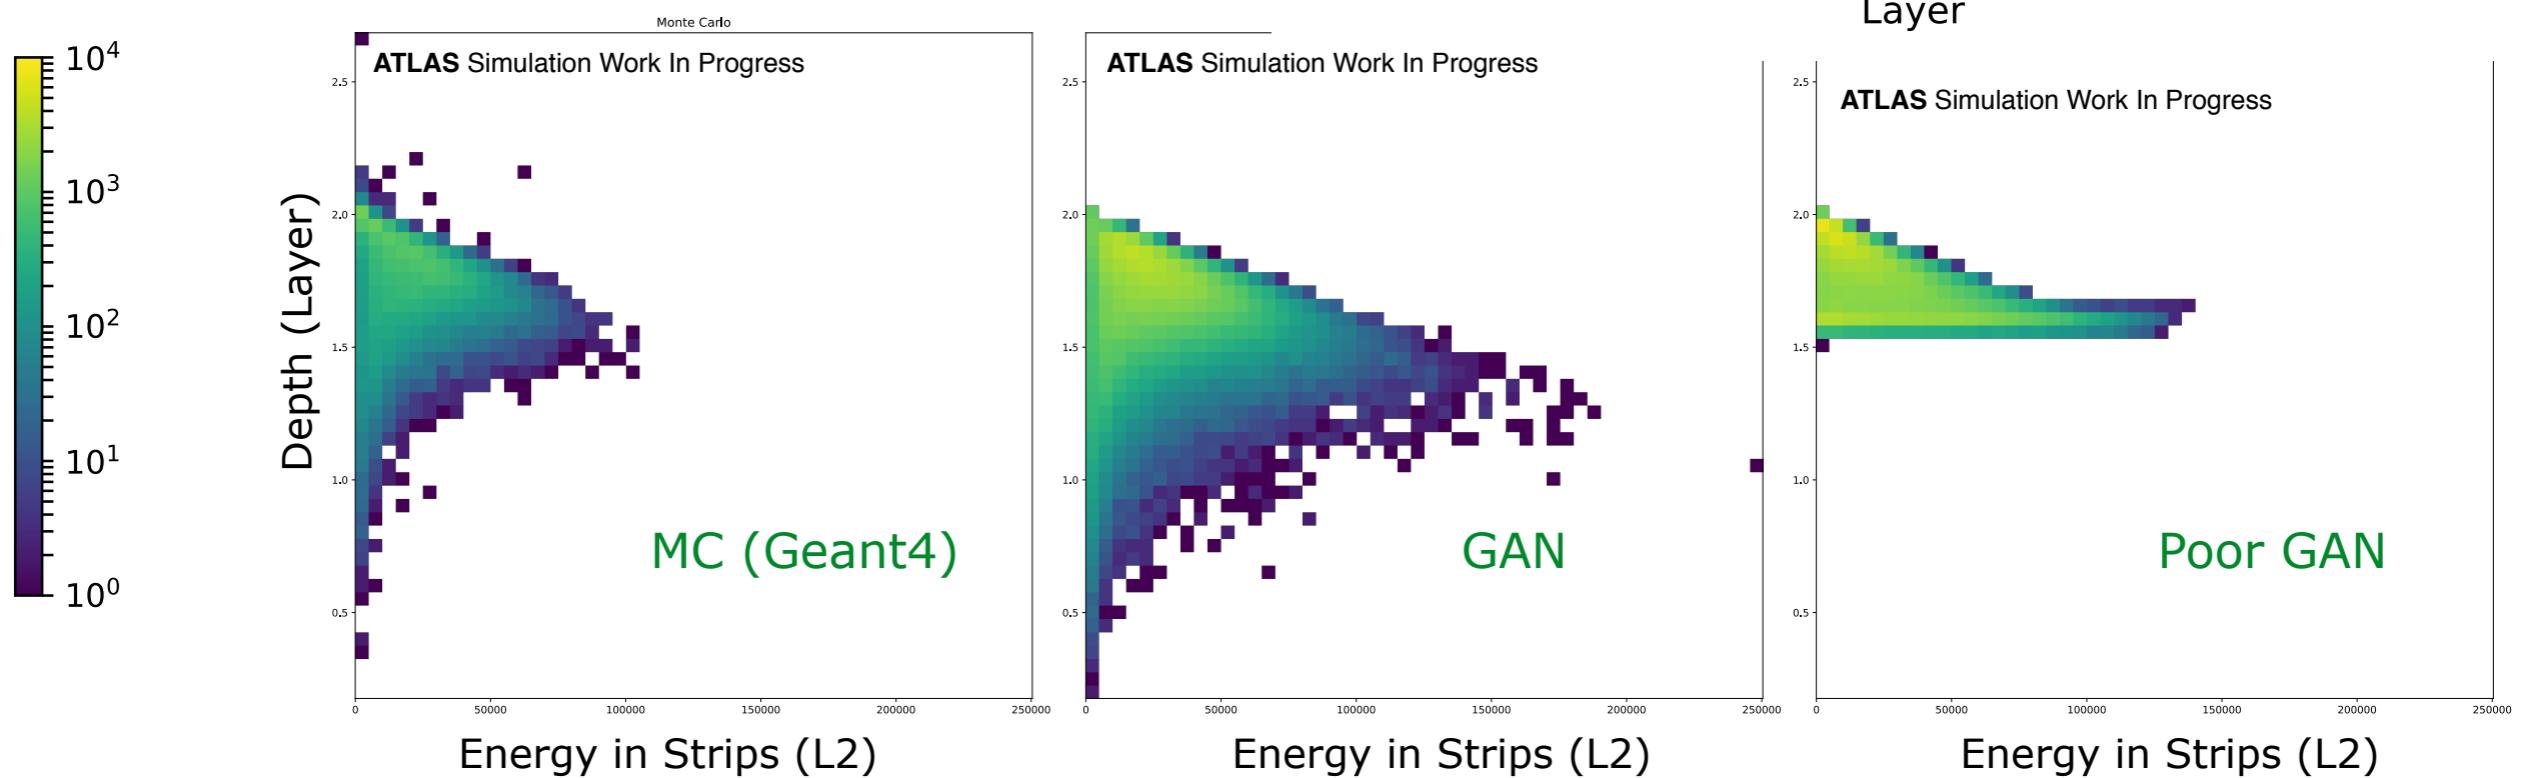
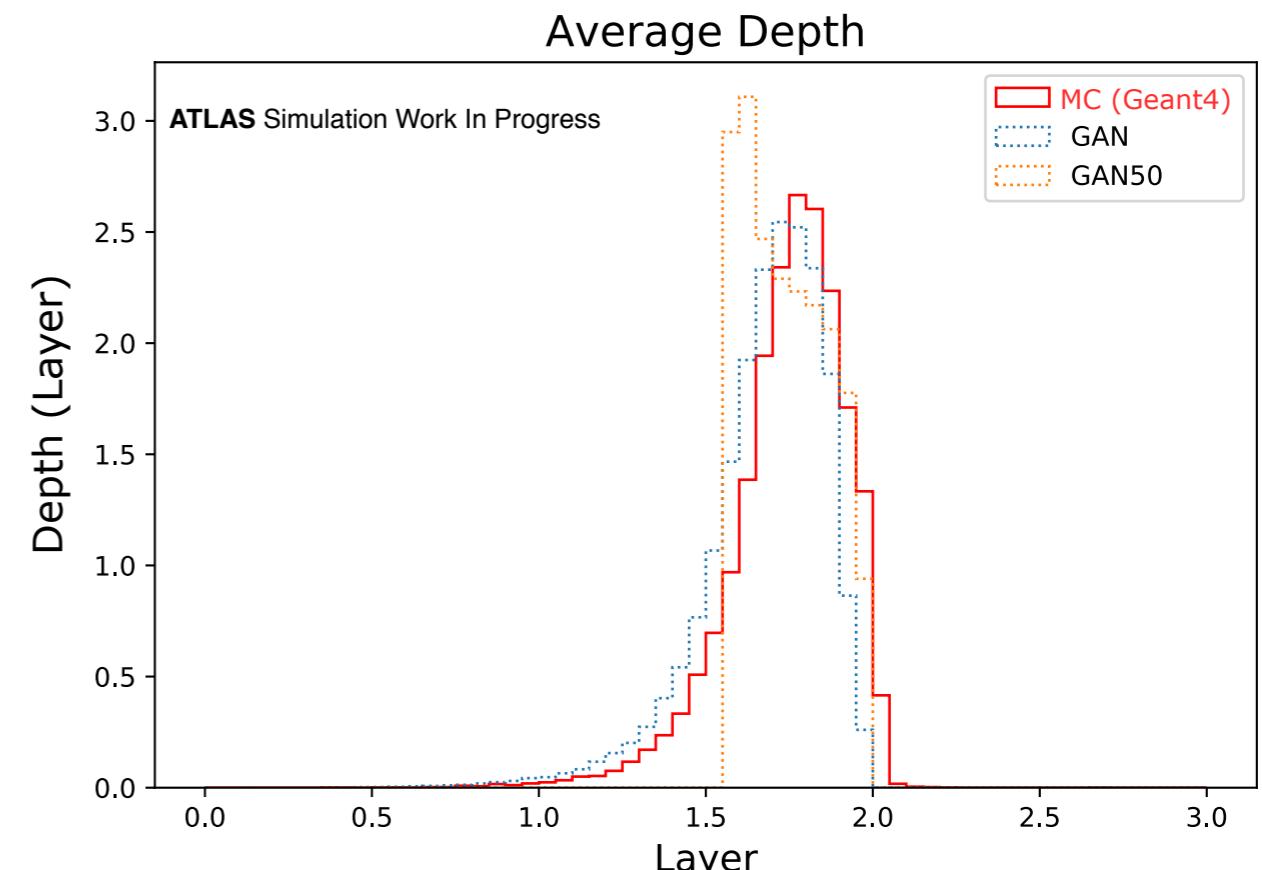


**GAN reproduces the mean
cannot yet reproduce $\sigma_E \sim 10\% \sqrt{E}$**

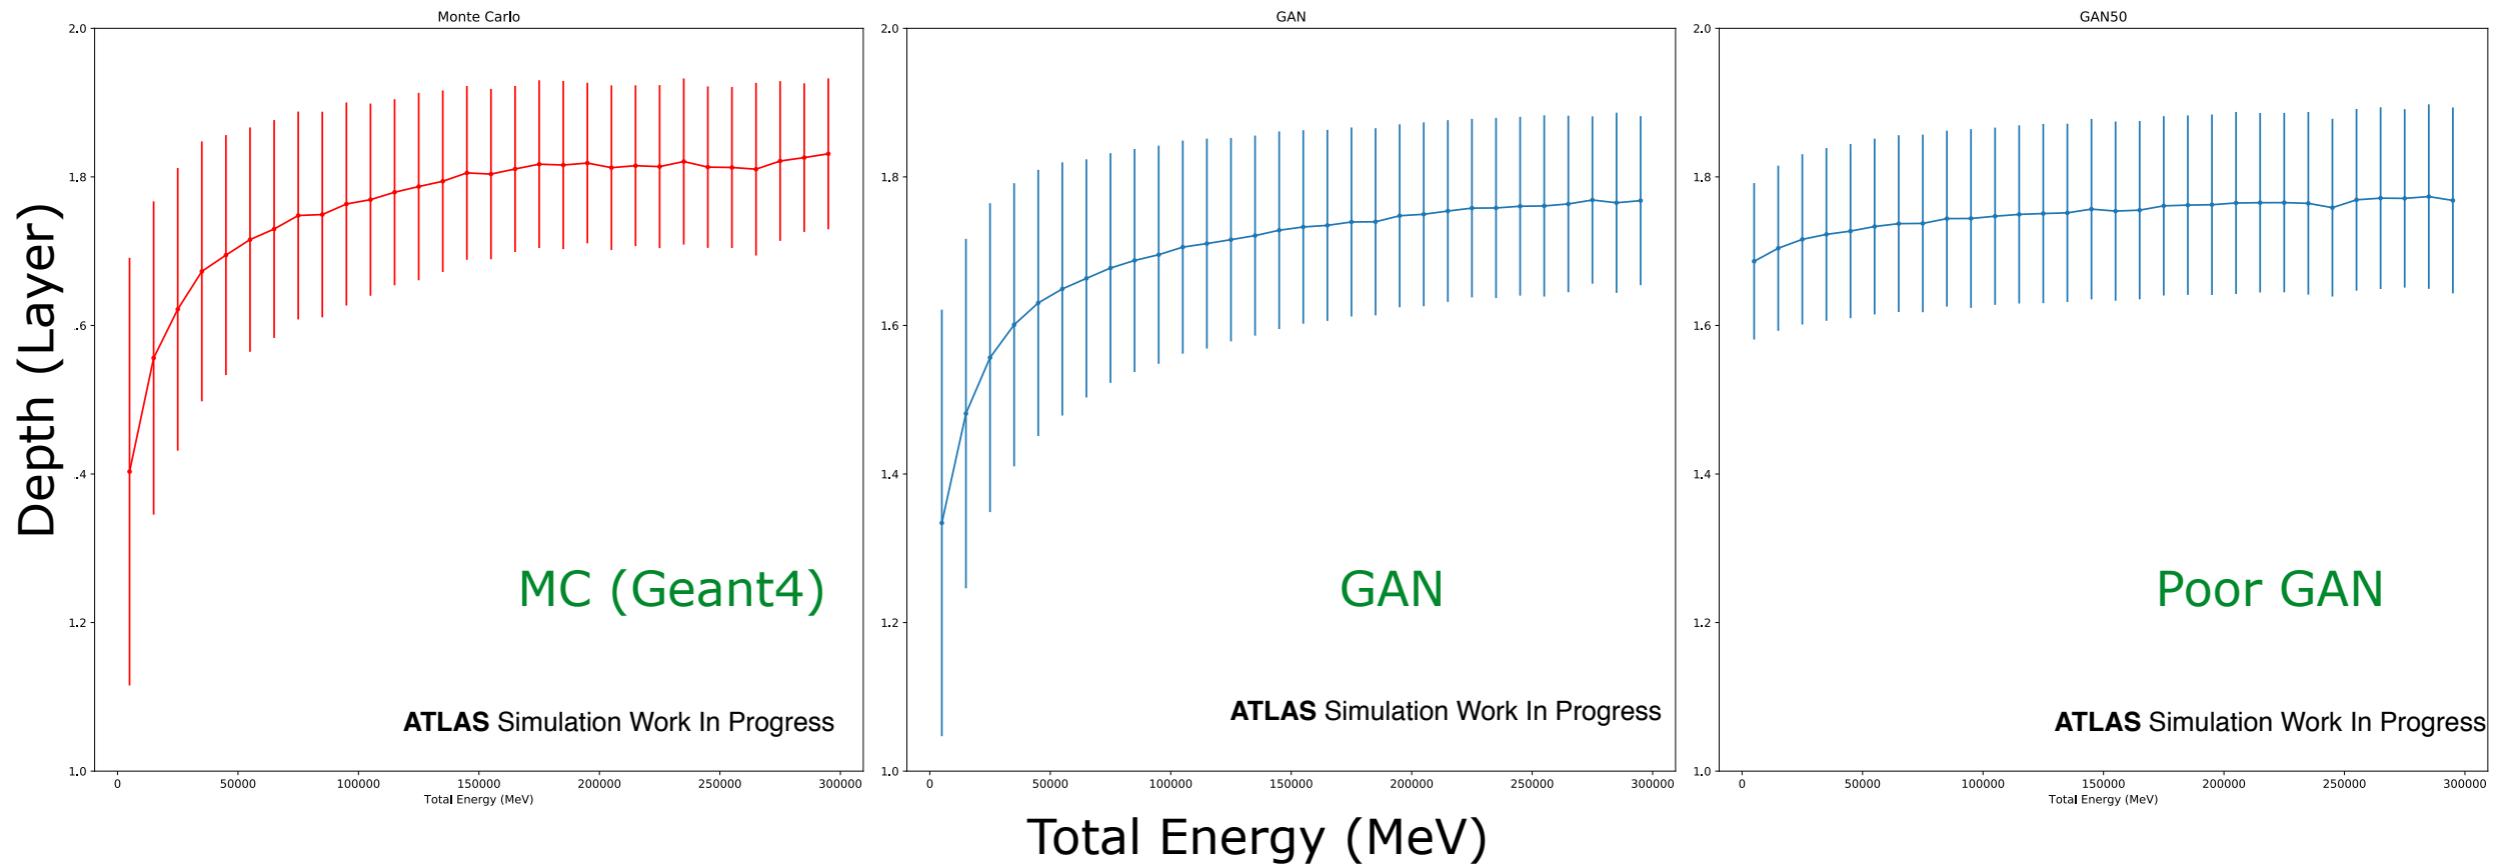
Investigation of Physics Variables: Depth

With our new Master 1 intern Jose-Antonio Mendez we are looking into new validation plots for depth

Comparing MC to GAN and ‘poor GAN’ trained for only 50 Epoch



Profile of Depth



The GAN **reproduces this non-trivial shape** of depth vs total energy very well, but is **shifted downwards**

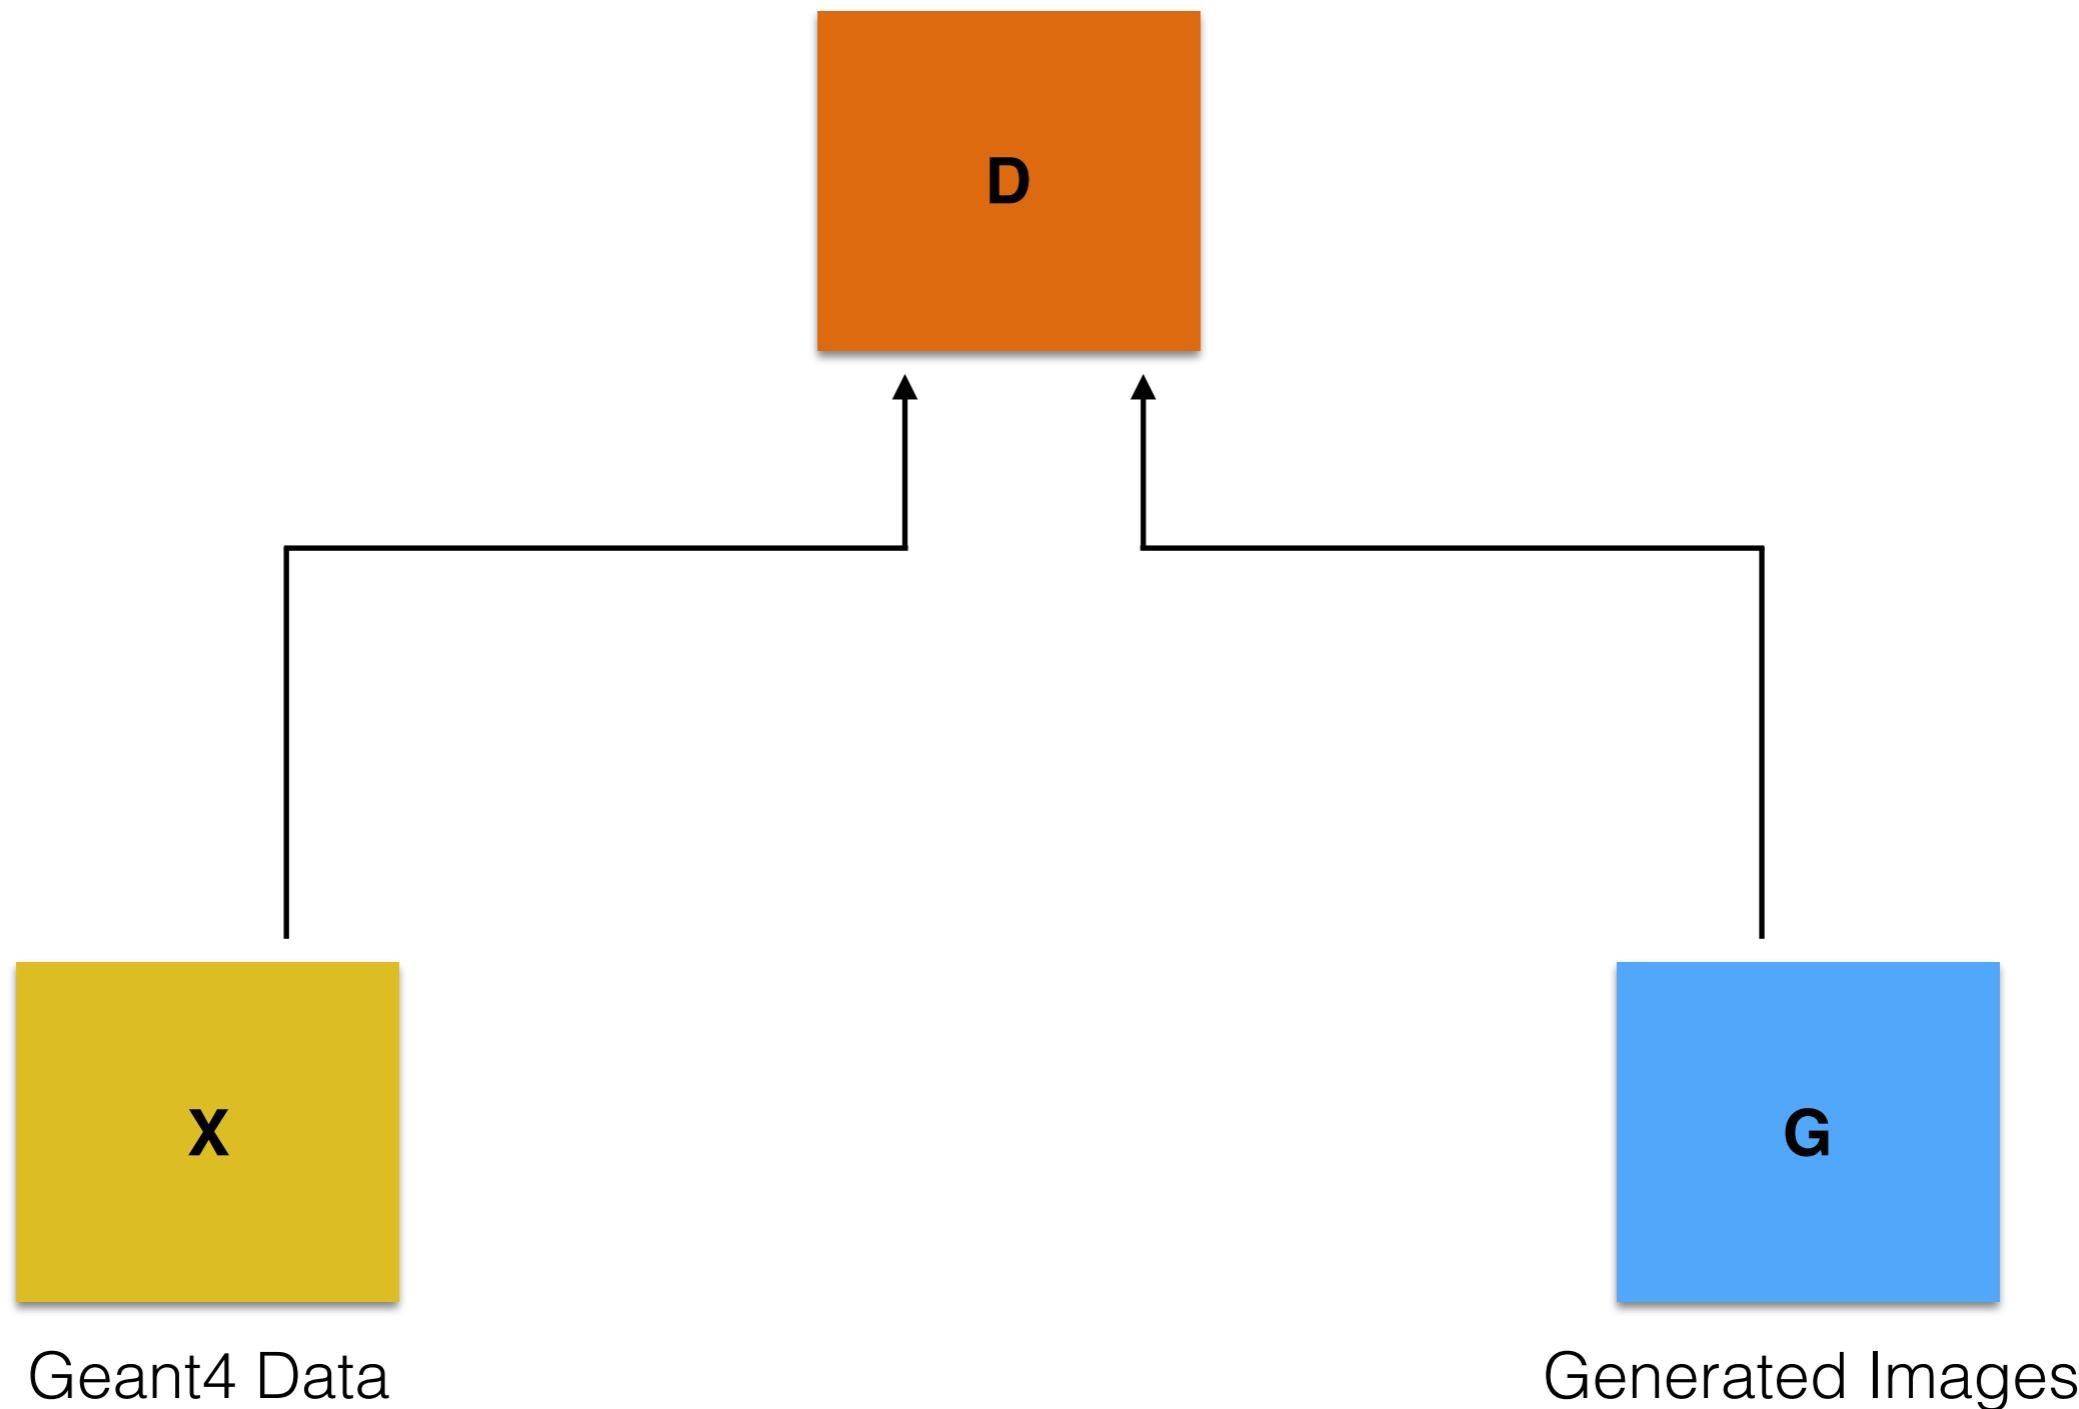
The error bars are the **Standard Deviation in each bin**

(Statistical Errors are negligible)

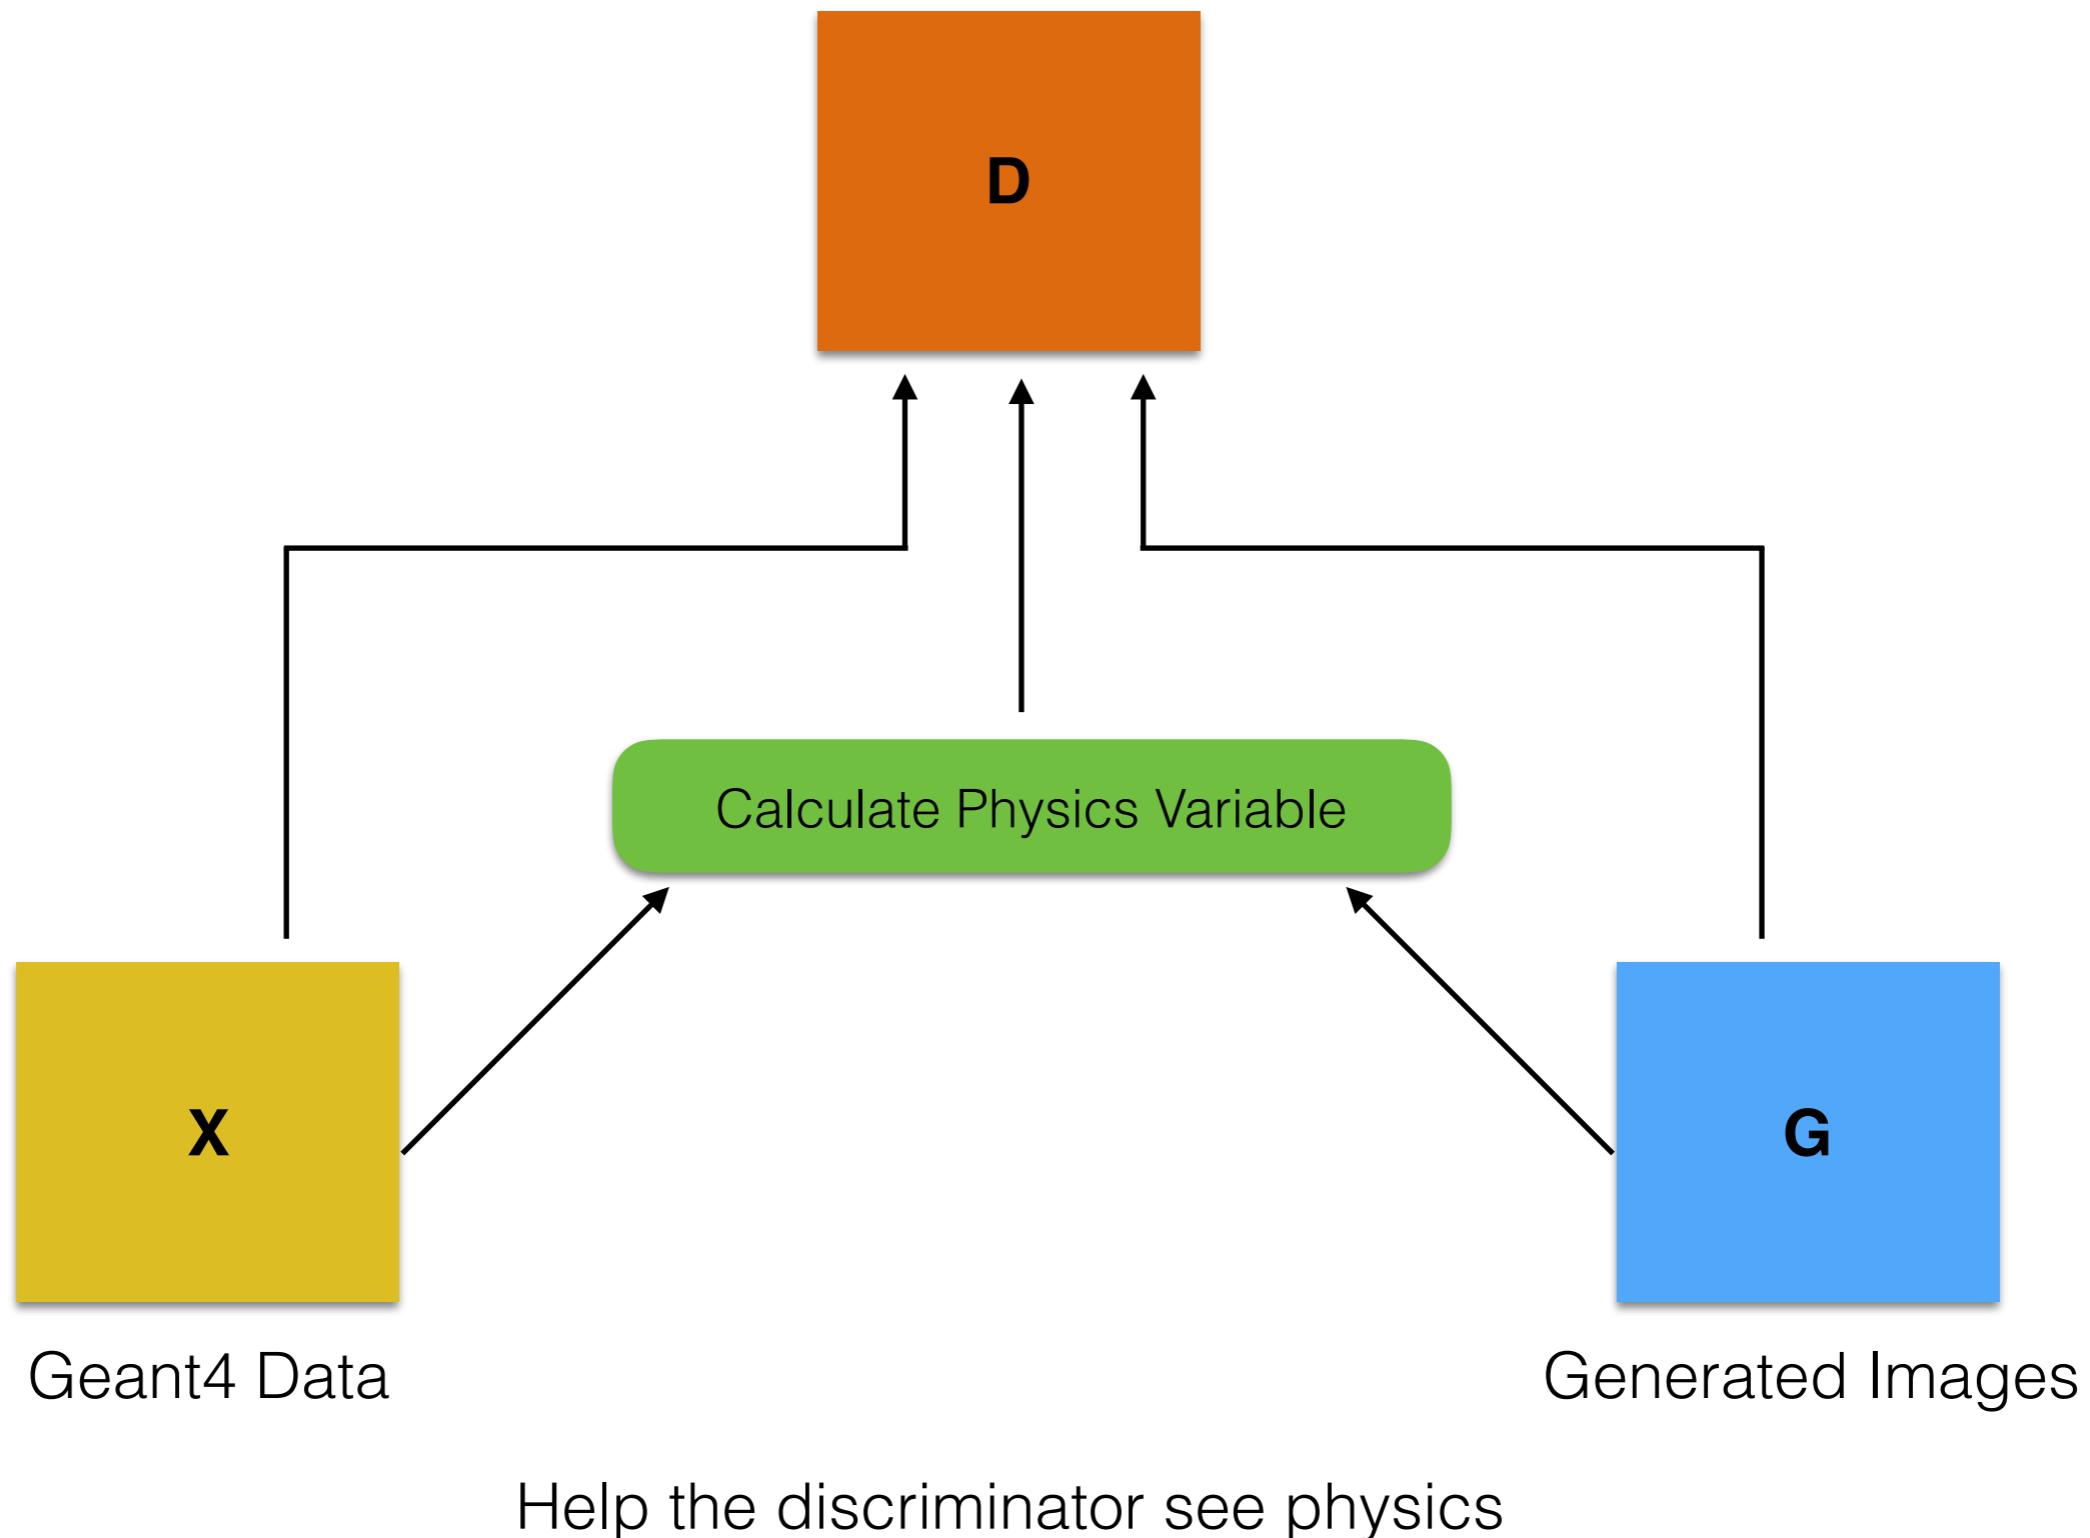
Next Plans

- Train GAN with knowledge about the alignment of calo layers (One-hot vector encoding) or use η/ϕ info
- Combine datasets to have more statistics, different Energy and η points
- Try new techniques:
 - Improve discriminator performance
 - Training Tricks: Train on $\log(\text{cell}_E)$, learning rate of G, Gradient penalty, non-RELU activations
 - Add physics variables in GAN training ([see slide](#))
- Try to train on more 'High Resolution', high granularity images to avoid the complications of calorimeter geometry
 - Harder to train on high resolution, but accuracy only needed at less granular 'cell level'
 - Use Progressive GAN with High Resolution images, other GAN flavours

Add Physics Variables in Training



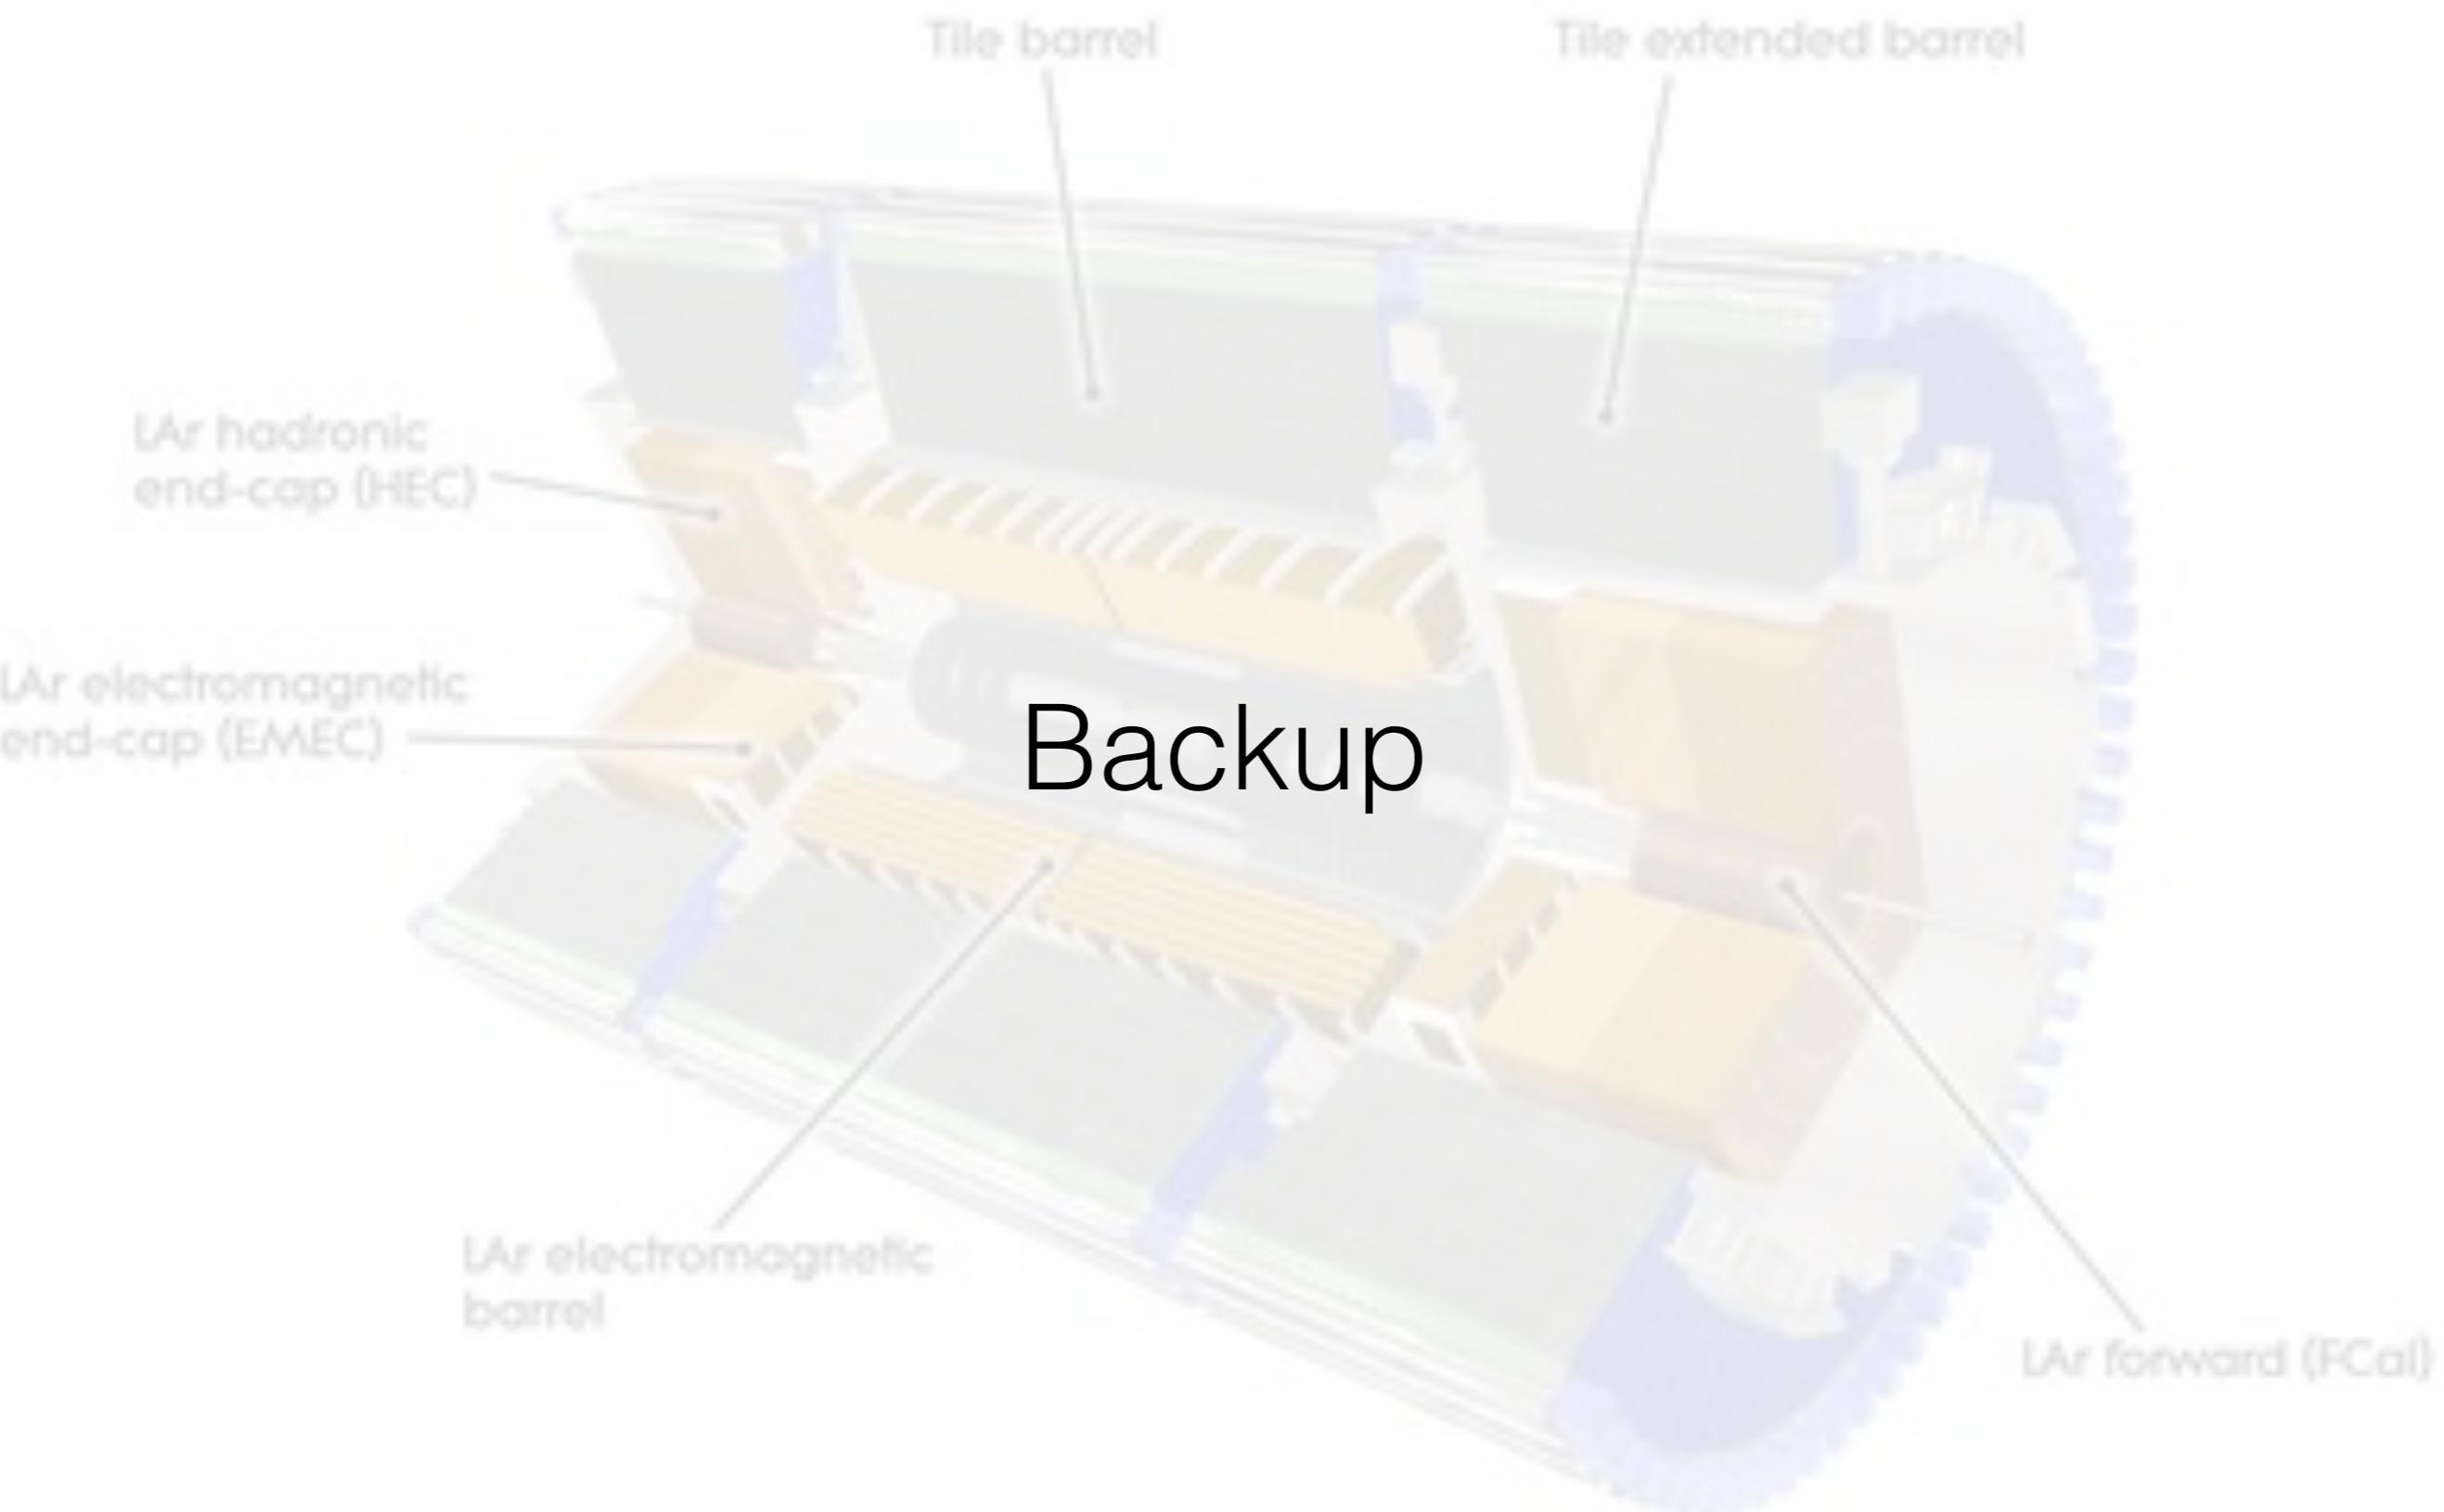
Add Physics Variables in Training



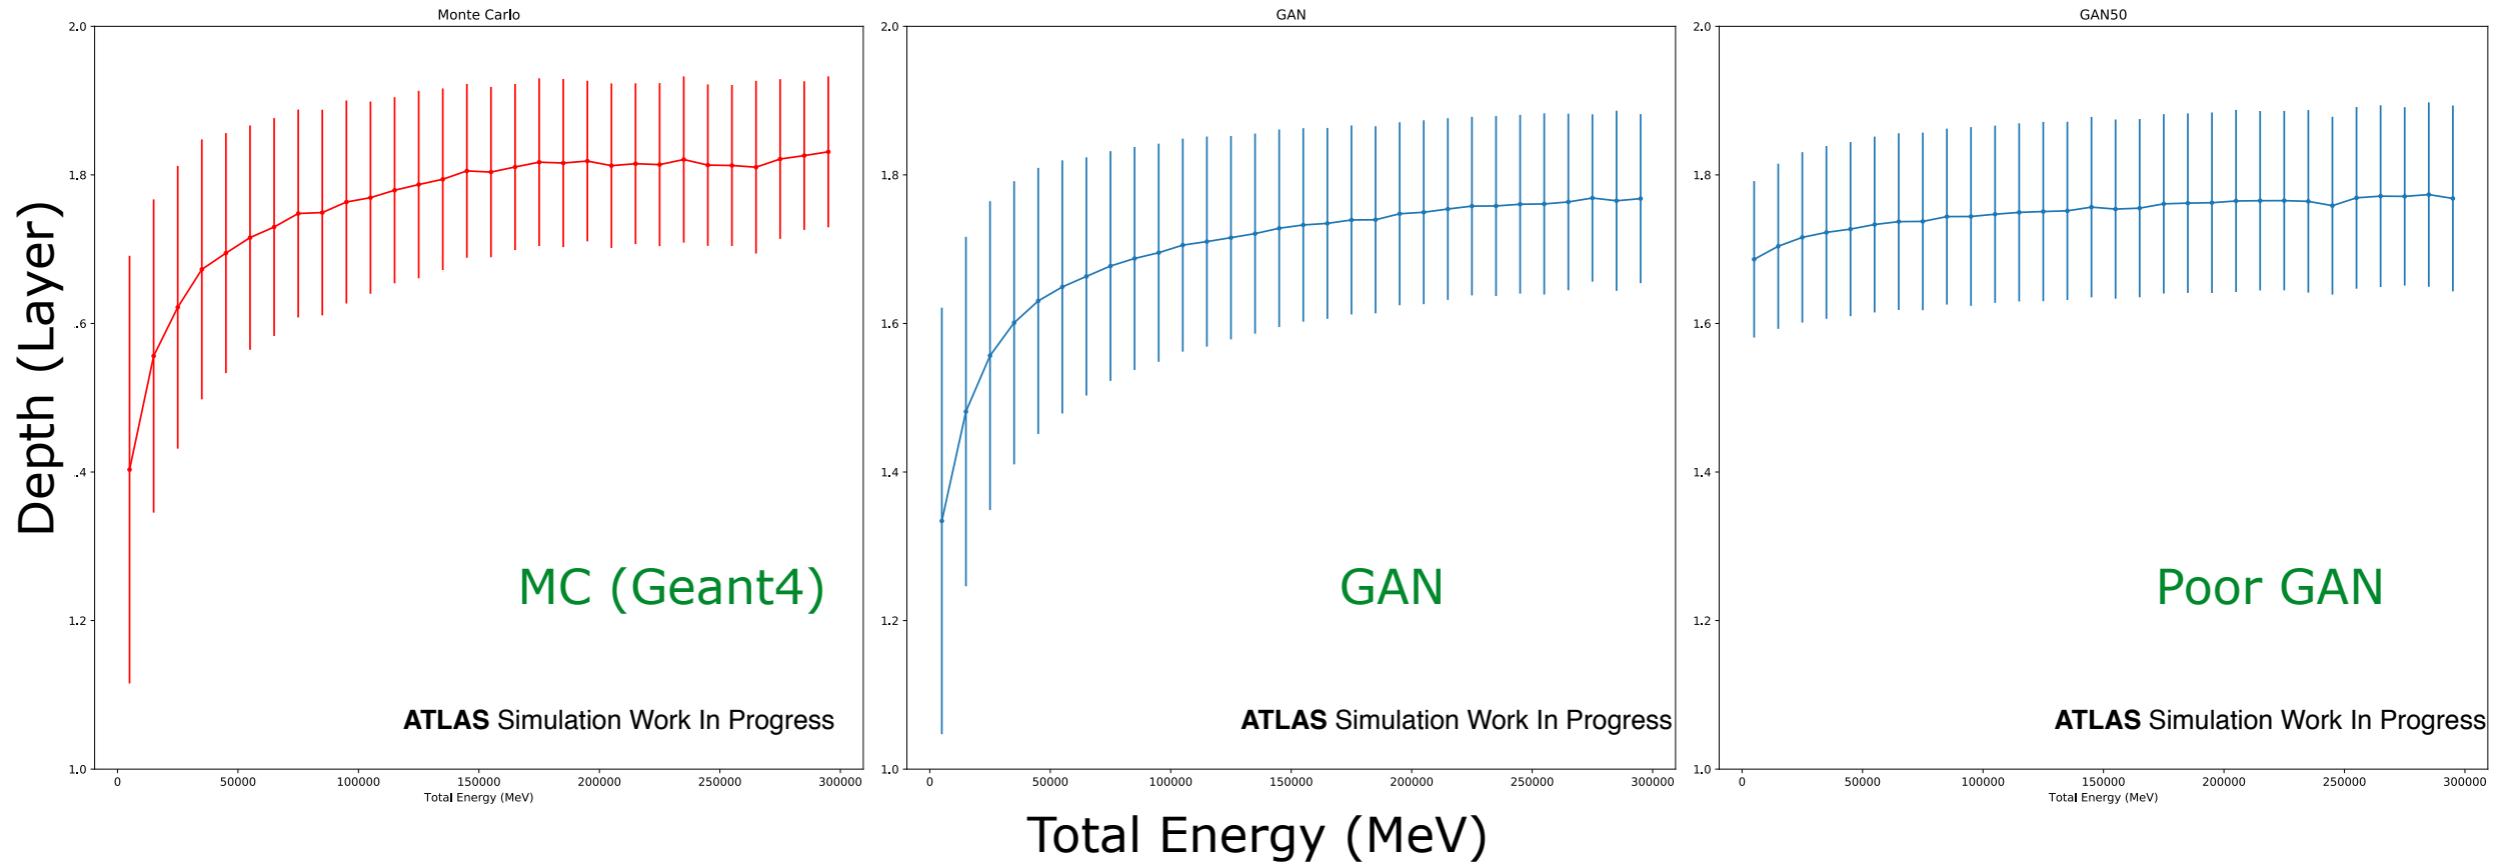
Conclusion

- Understanding the data structure remains essential for Machine Learning tasks
- Motivation to move to more granular, ‘High Resolution’ images
- Comparisons need to be made with FastCaloSim, VAE, which are other methods of fast simulation of EM showers
- GPU gives 2x speed up compared to CPU-only training
 - Might improve for larger Neural Nets, more high definition images

Backup



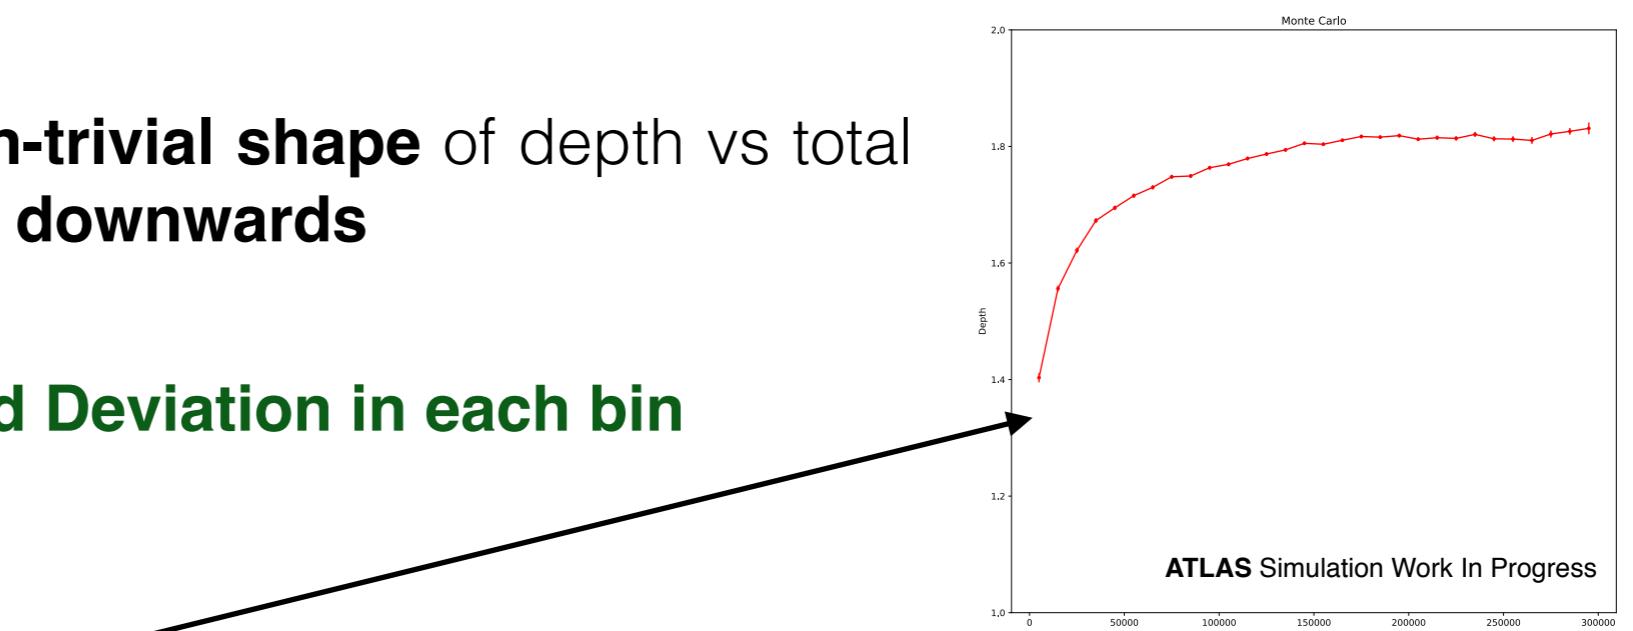
Profile of Depth



The GAN **reproduces this non-trivial shape** of depth vs total energy very well, but is **shifted downwards**

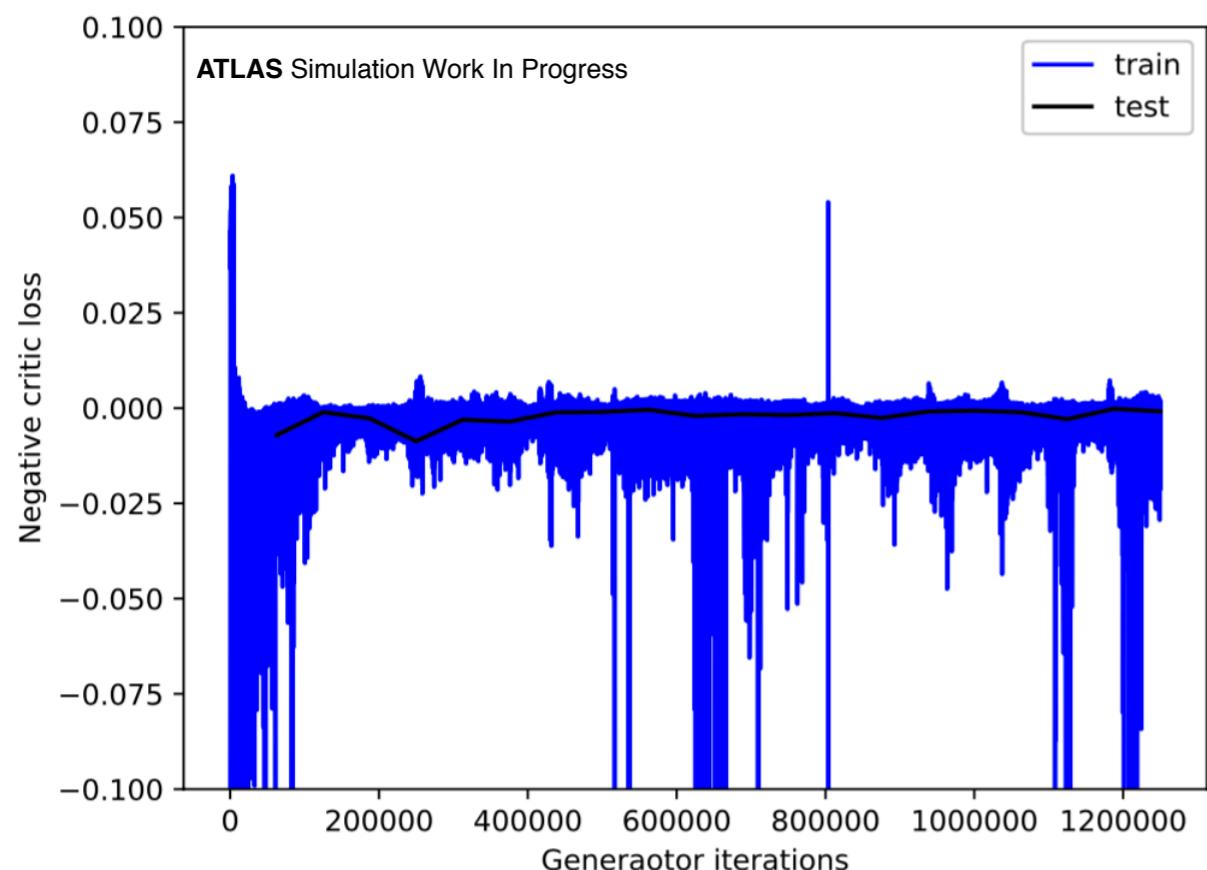
The error bars are the **Standard Deviation in each bin**

Statistical Errors are negligible

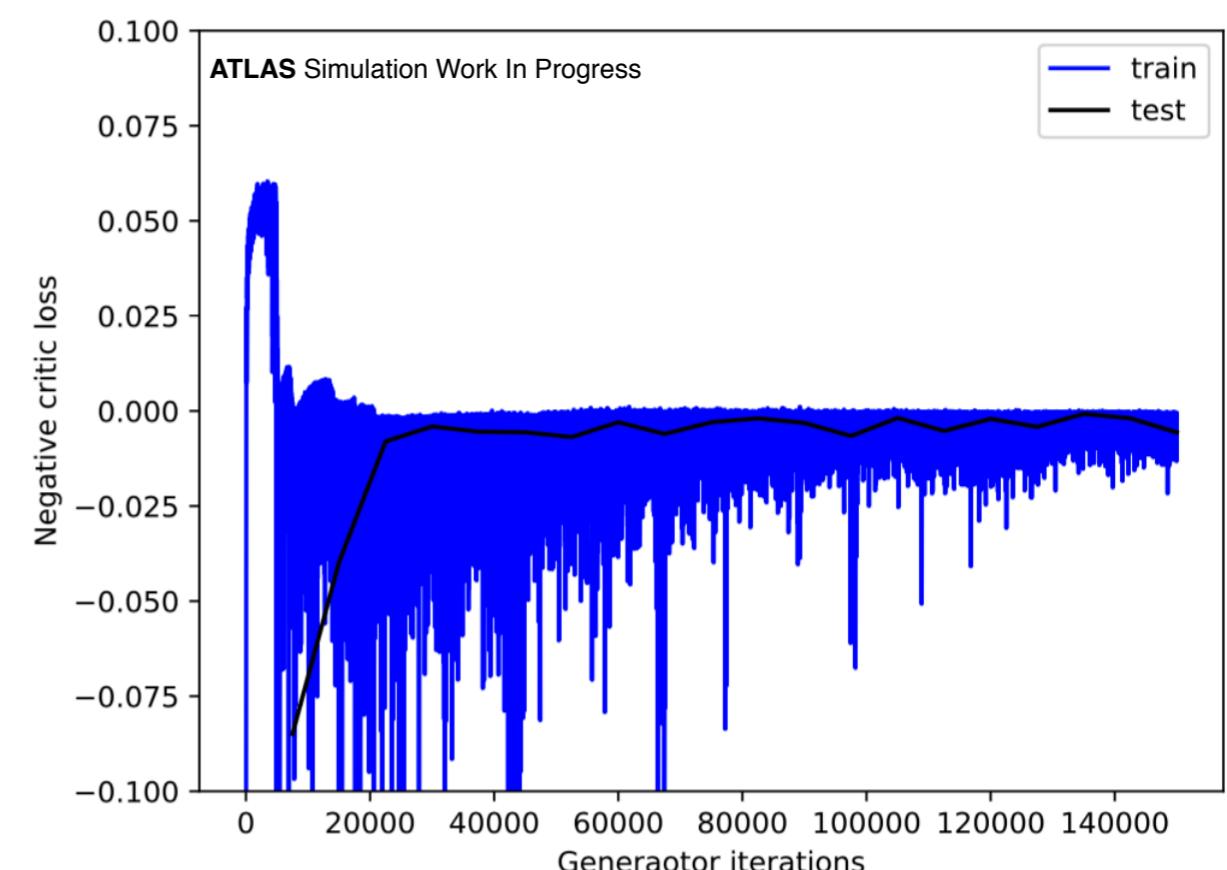


Mixed Config: Critic Loss

10k Events



Only 10k/8 Events

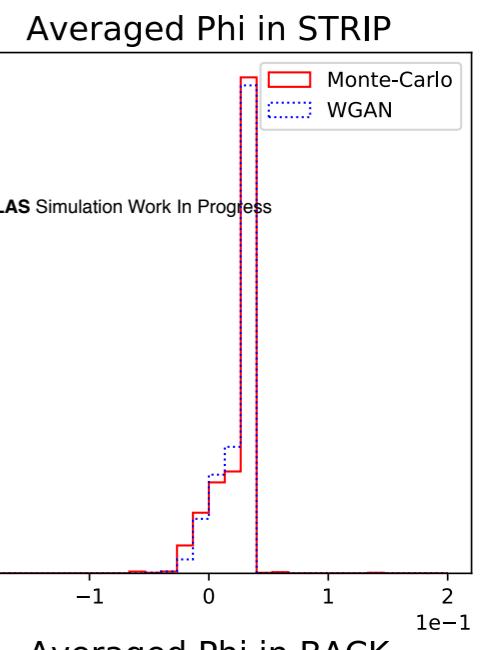
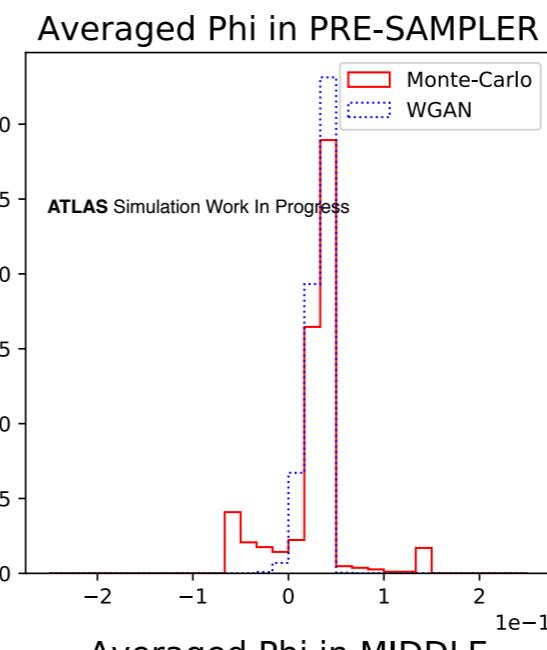
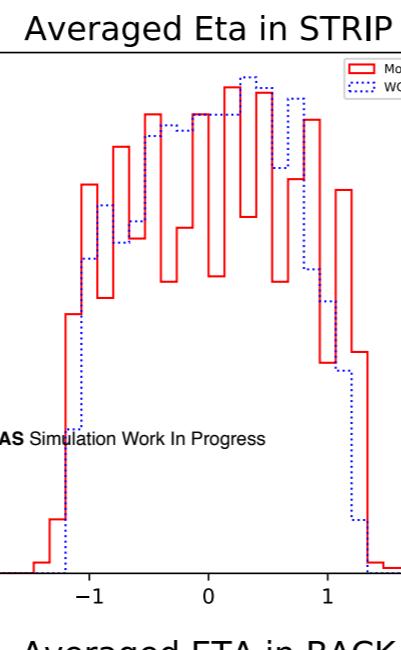
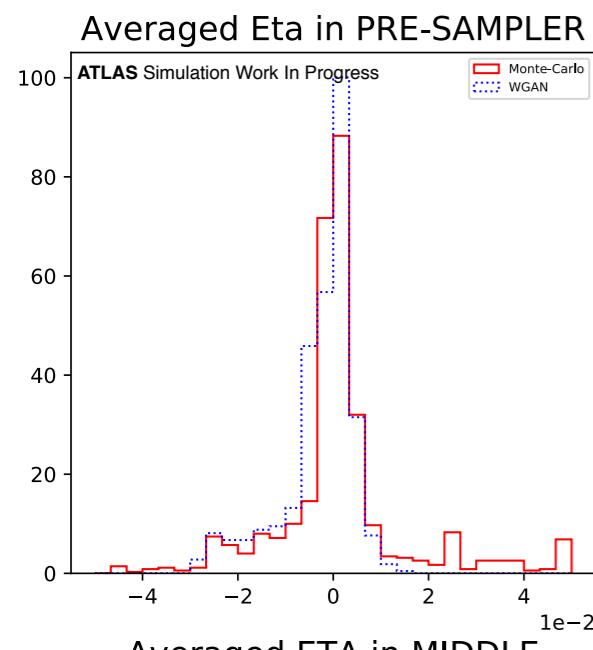


‘Configurations’

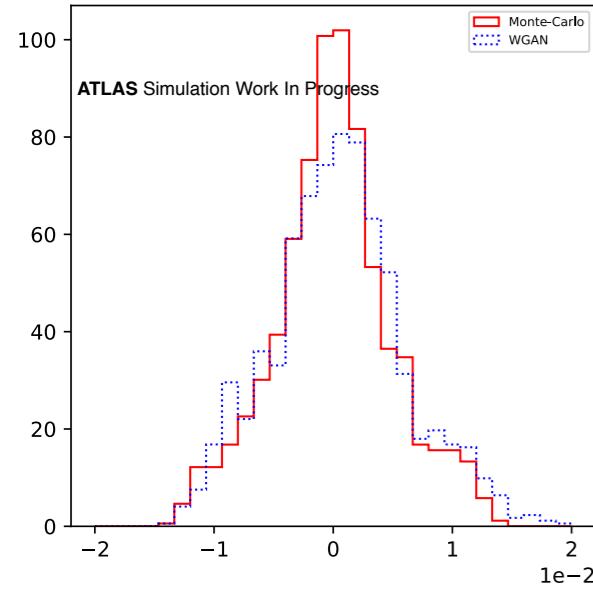
- The calo layers are one behind the other : PreSampler, Strip, Middle, Back. But they are not aligned in the same way for every event. Since they have different widths, the exact same EM shower would look different at different parts of the calorimeter.
- If we select a 7x7 cluster in the middle layer, then in front of it, the strip could have 4 different ‘configurations’. [See](#). For a fixed middle 7x7, there could be 4 different ways the strip cells are aligned w.r.t the middle cells. In config 0 and 7, 2 cells of the strip can cover all the 7 cells of the middle layer in phi (y-axis). Whereas for the other 2 configurations 3 cells are required.
- A similar case for the back layer’s alignment w.r.t. the middle layer means that there are 2 distinct ways the back might be aligned.
- $4 \text{ (strips in phi)} * 2 \text{ (back in eta)} = 8$ different arrangements of the cells in the calorimeter.
- The Pre-Sampler behaves the same as Strip in phi, and behaves the same as Middle in eta. The Back behaves the same as Middle in phi. 8 strip cells have the same width as 1 Middle cell in eta, so there is no problem of misalignment.
- Configurations 0,1,2,3 are the same as 4,5,6,7 respectively for Strips/Pre-Sampler. All configurations in the set {0,1,2,3} are equivalent for the Back, all configurations in the set {4,5,6,7} are equivalent, but configurations from different sets are not equivalent to each other for the Back (e.g. config 0 not equivalent to config 4).
- The configurations in the Strip are periodic every 4 cells in phi, and the back are periodic every 2 cells in eta in the EM calorimeter.

Eta/Phi for Config 0

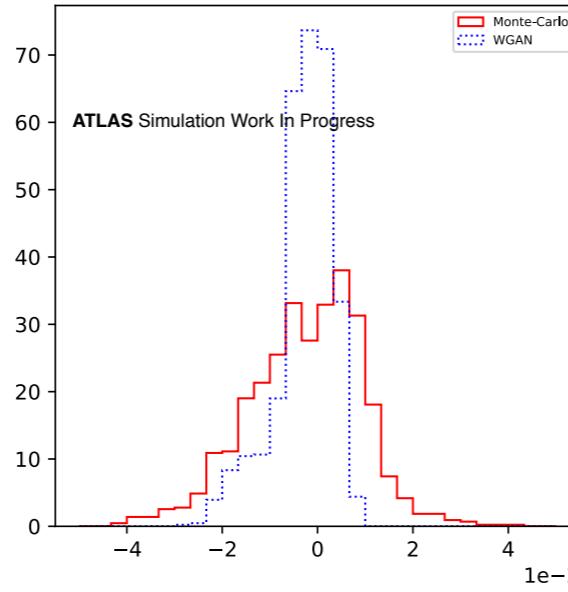
Average Eta



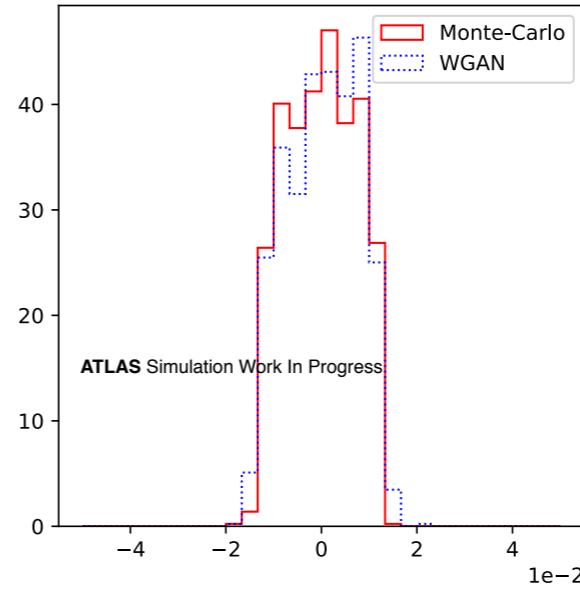
Averaged ETA in MIDDLE



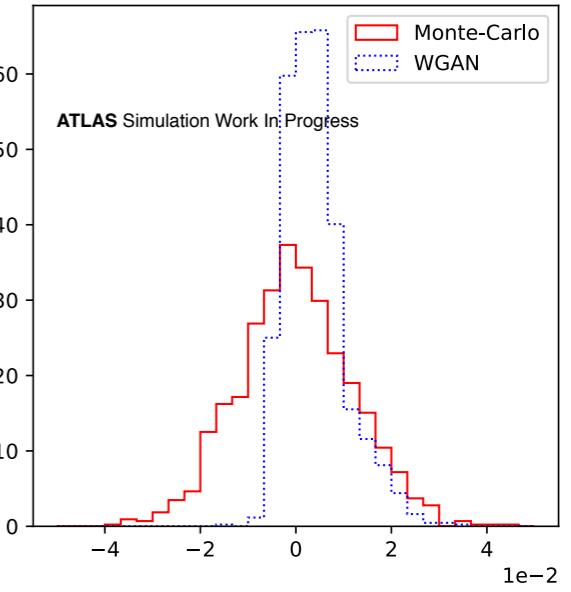
Averaged ETA in BACK



Averaged Phi in MIDDLE

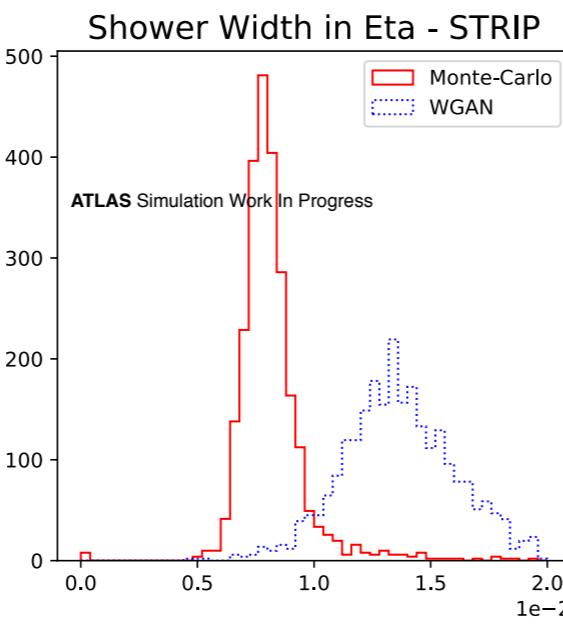
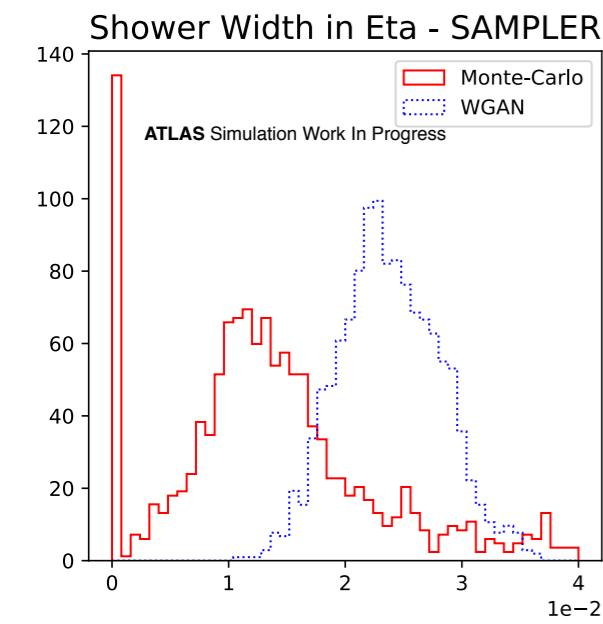


Averaged Phi in BACK

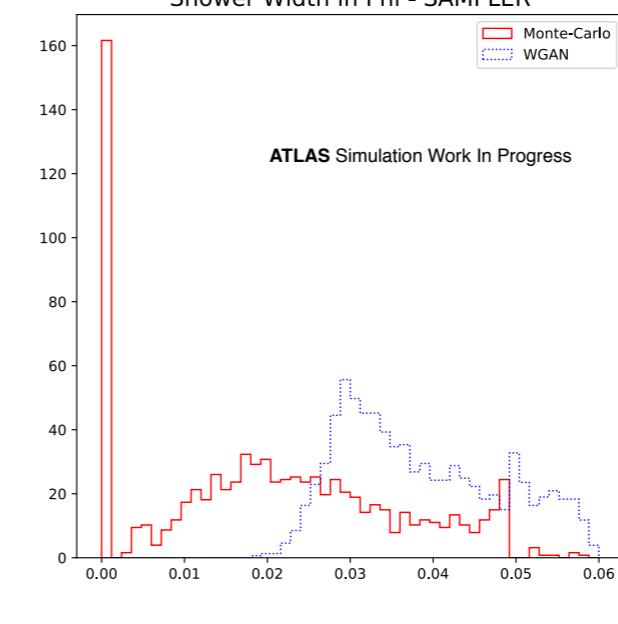


Width for Config 0

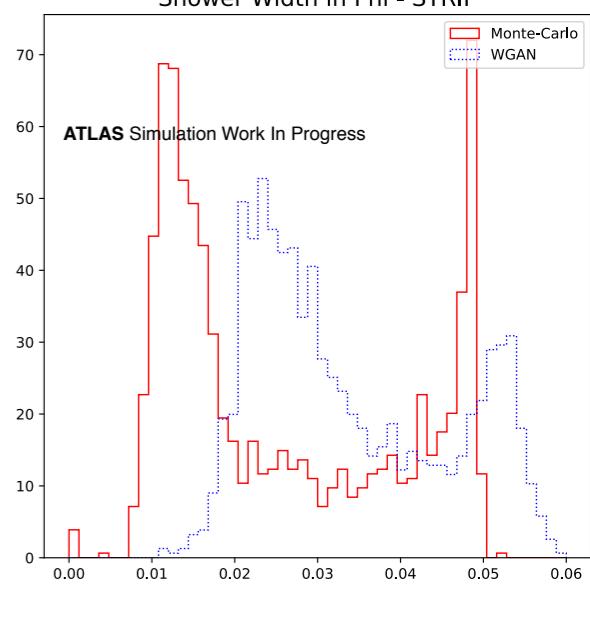
Eta Width



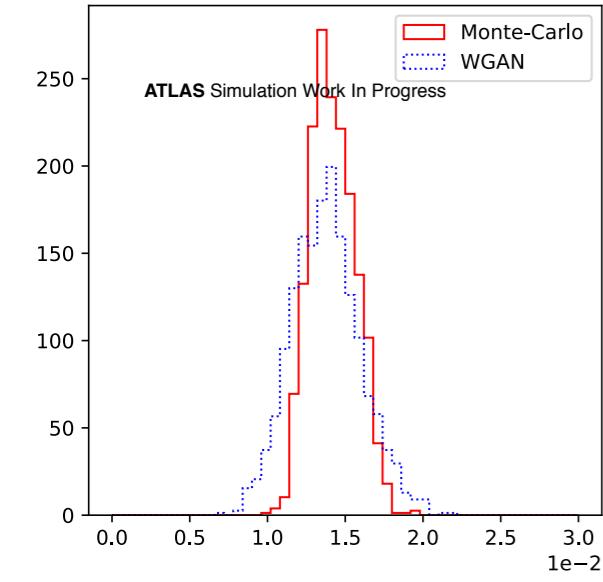
Shower Width in Phi - SAMPLER



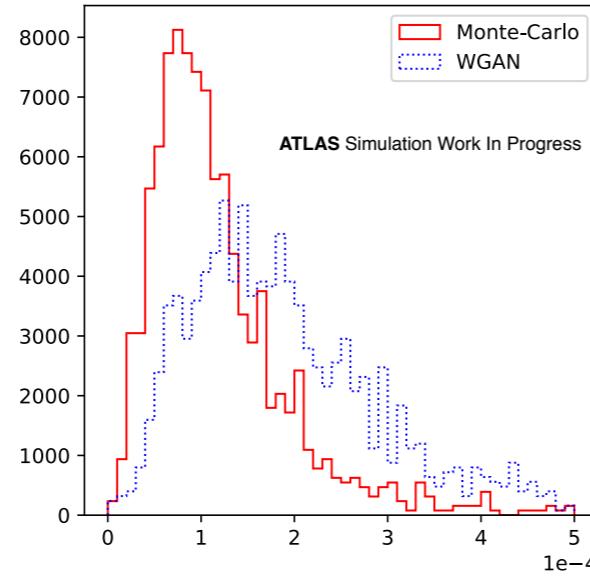
Shower Width in Phi - STRIP



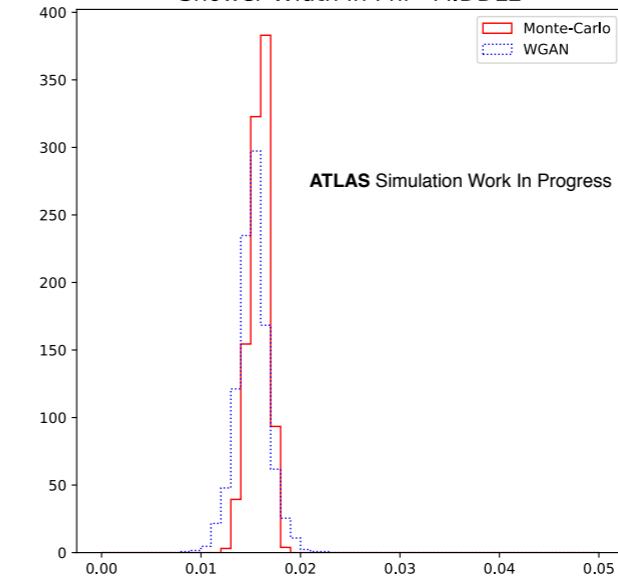
Shower Width in Eta - MIDDLE



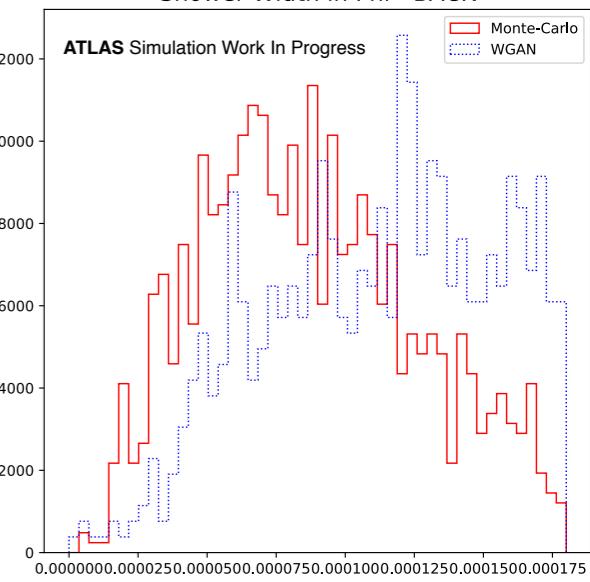
Shower Width in Eta - BACK



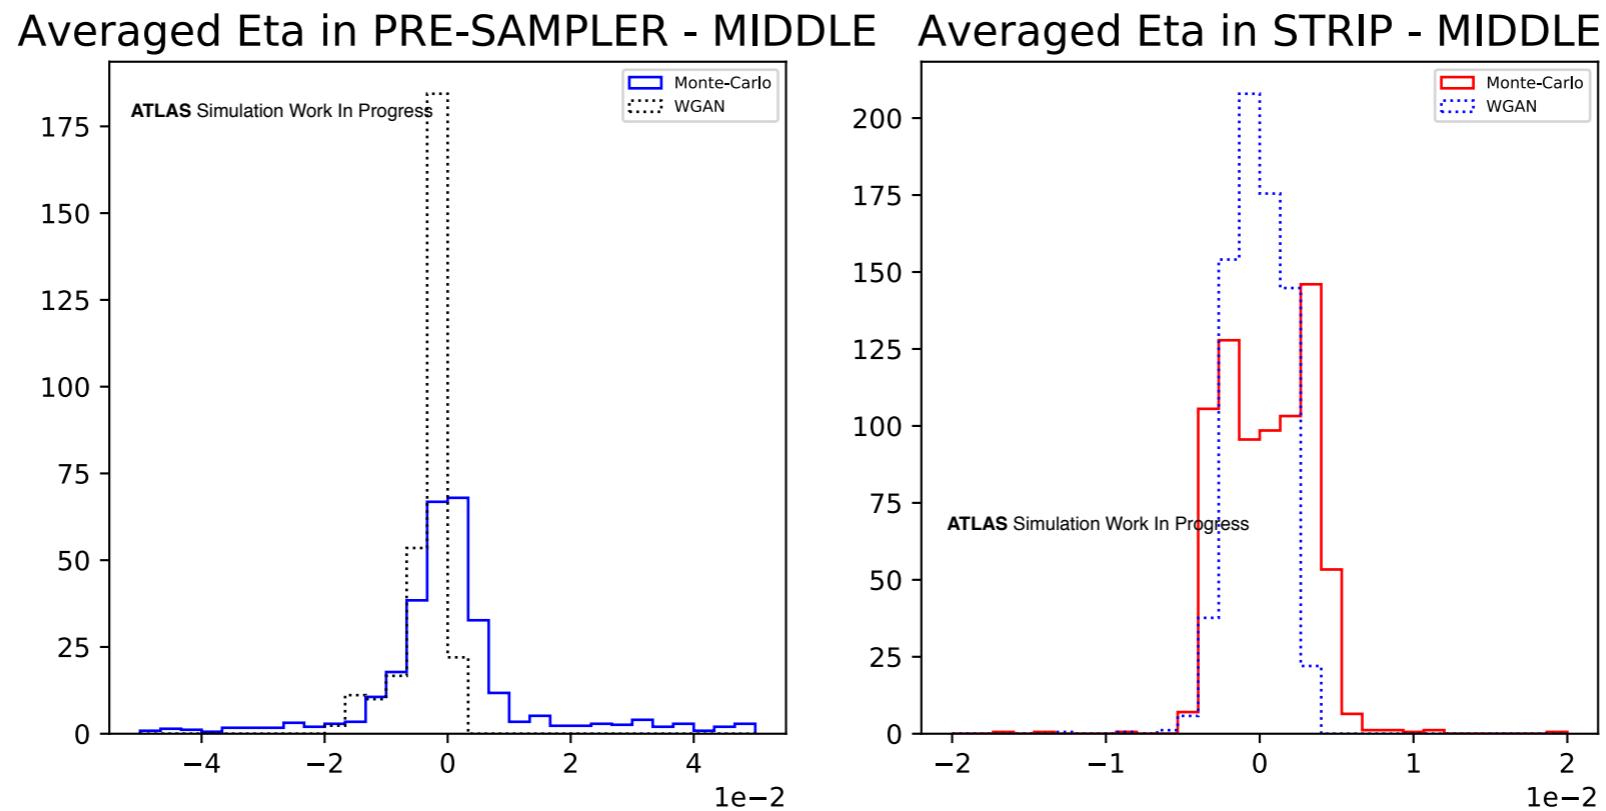
Shower Width in Phi - MIDDLE



Shower Width in Phi - BACK

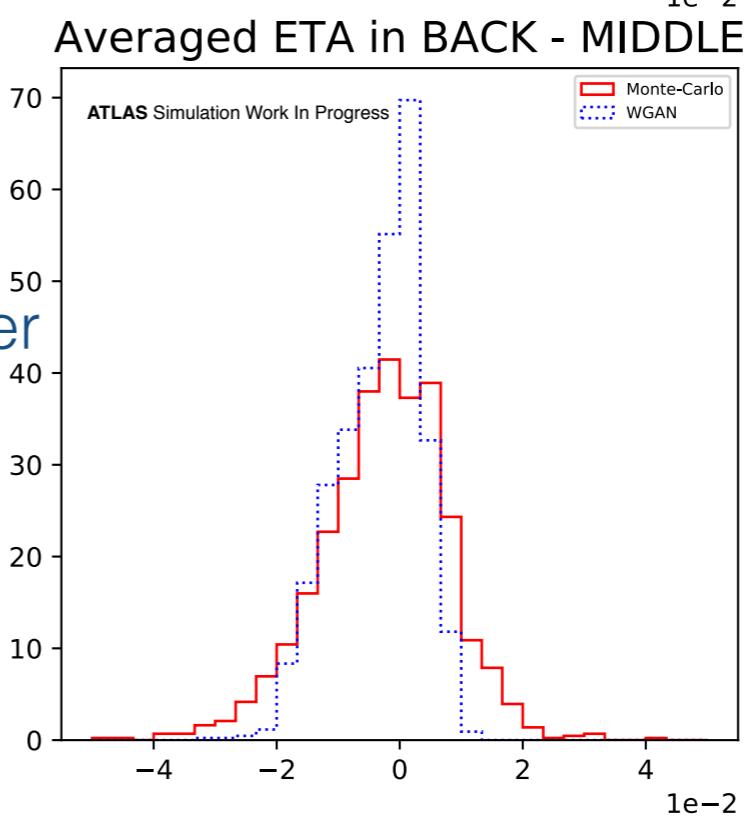


Layer-to-Layer Correlation (Config 0)



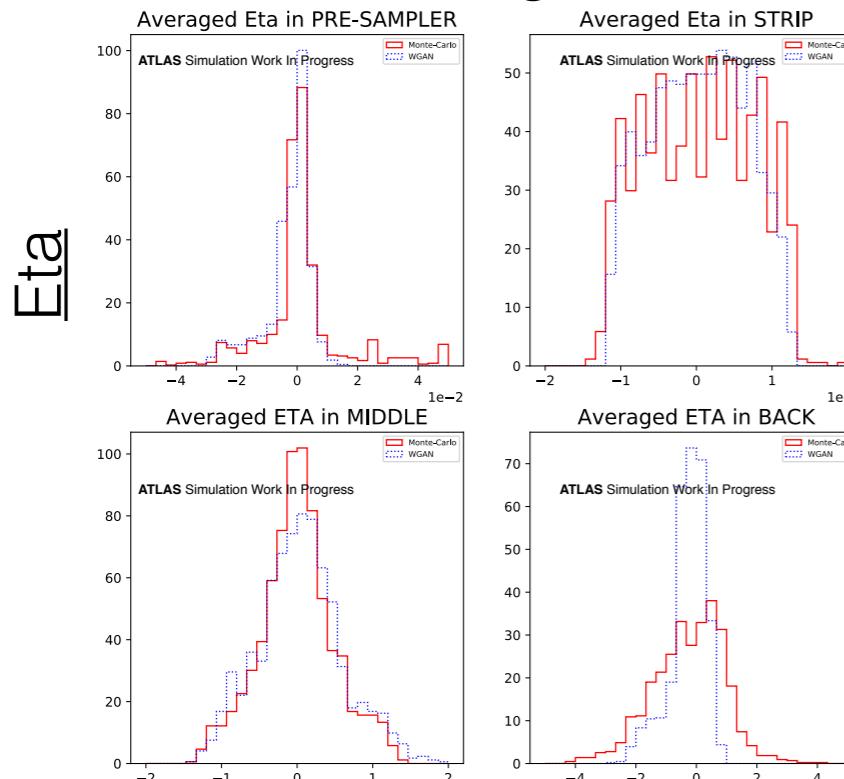
Compute:

Average Eta in Layer - Average Eta in Middle Layer

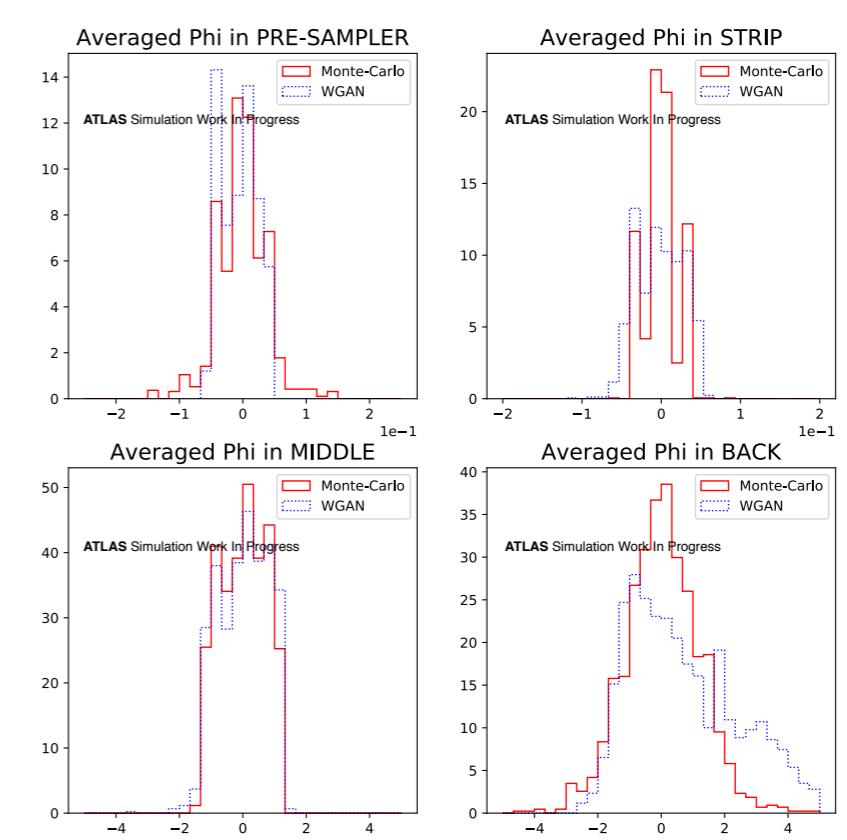
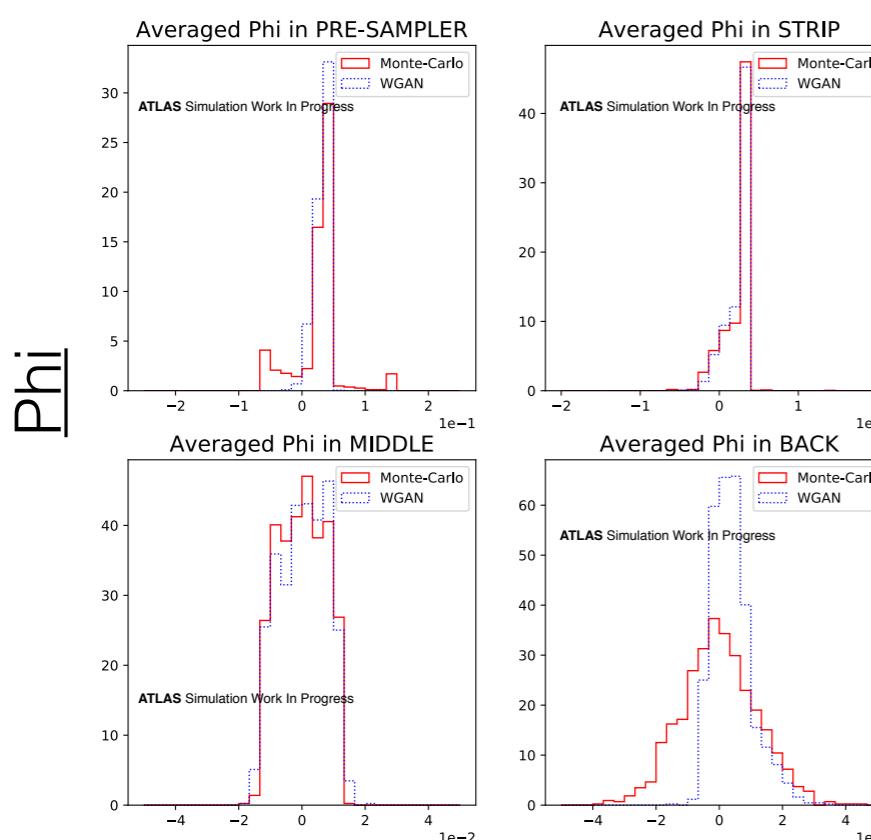
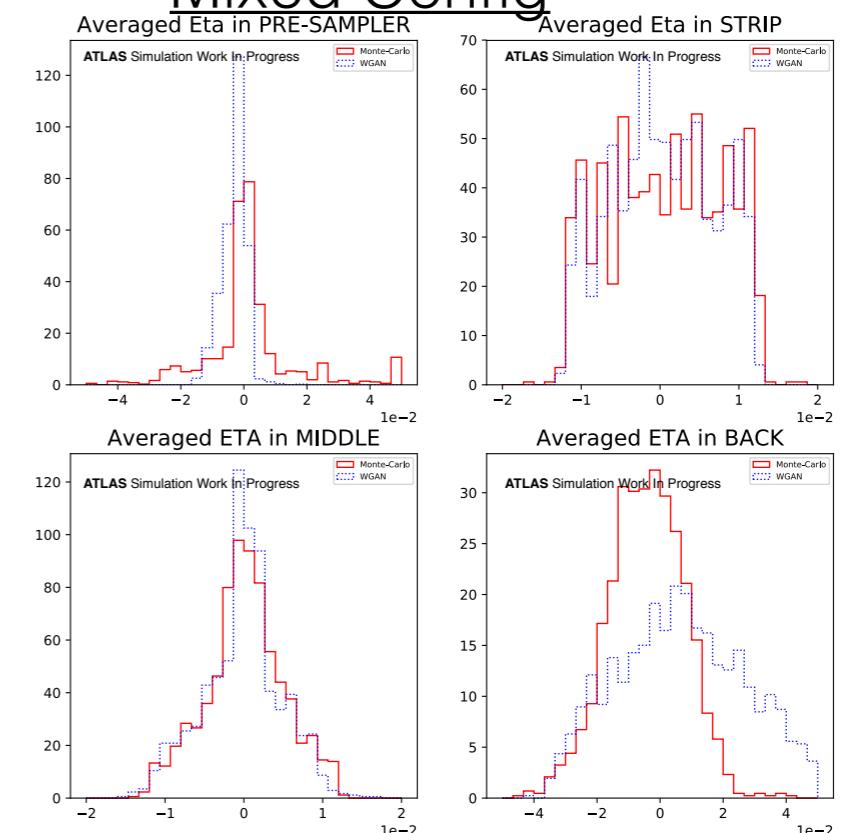


Compare with Mix of 8 Configs (same number of events)

Config 0

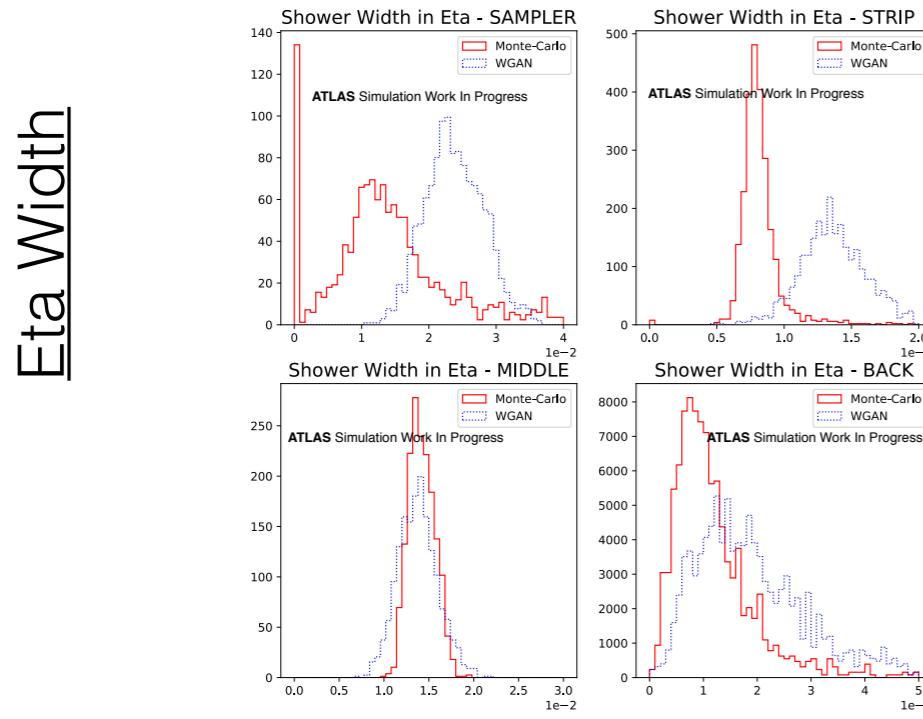


Mixed Config

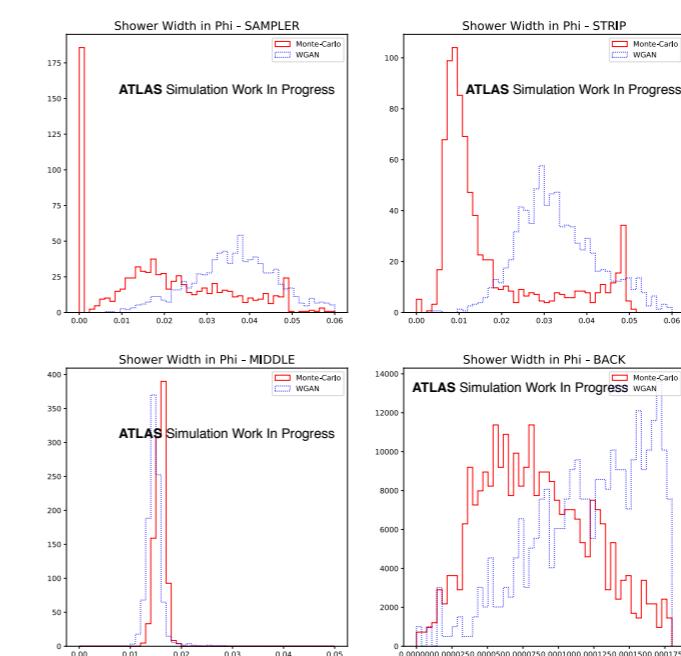
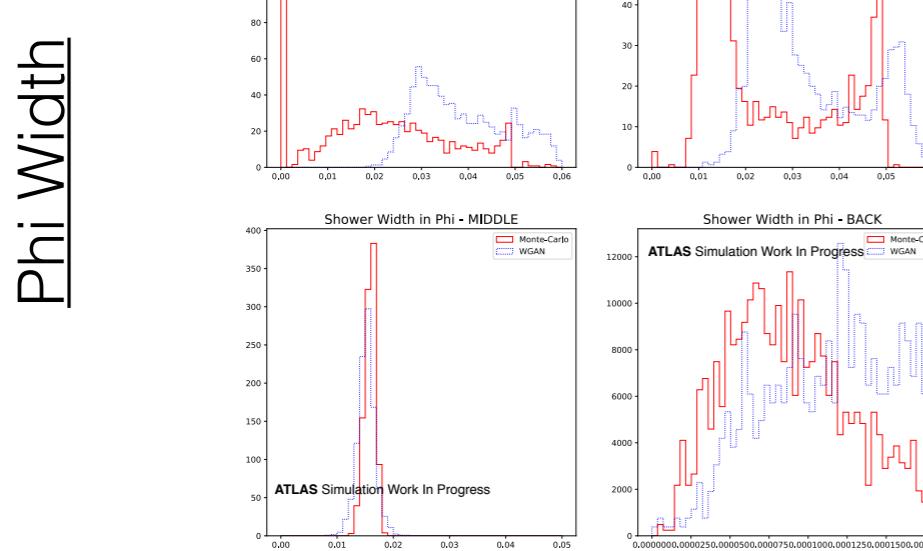
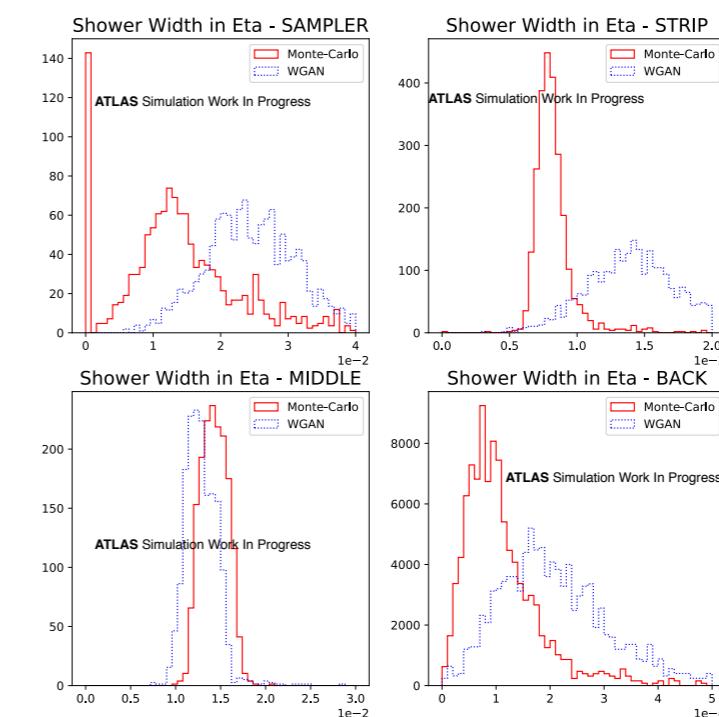


Compare with Mixed Configs (same number of events)

Config 0

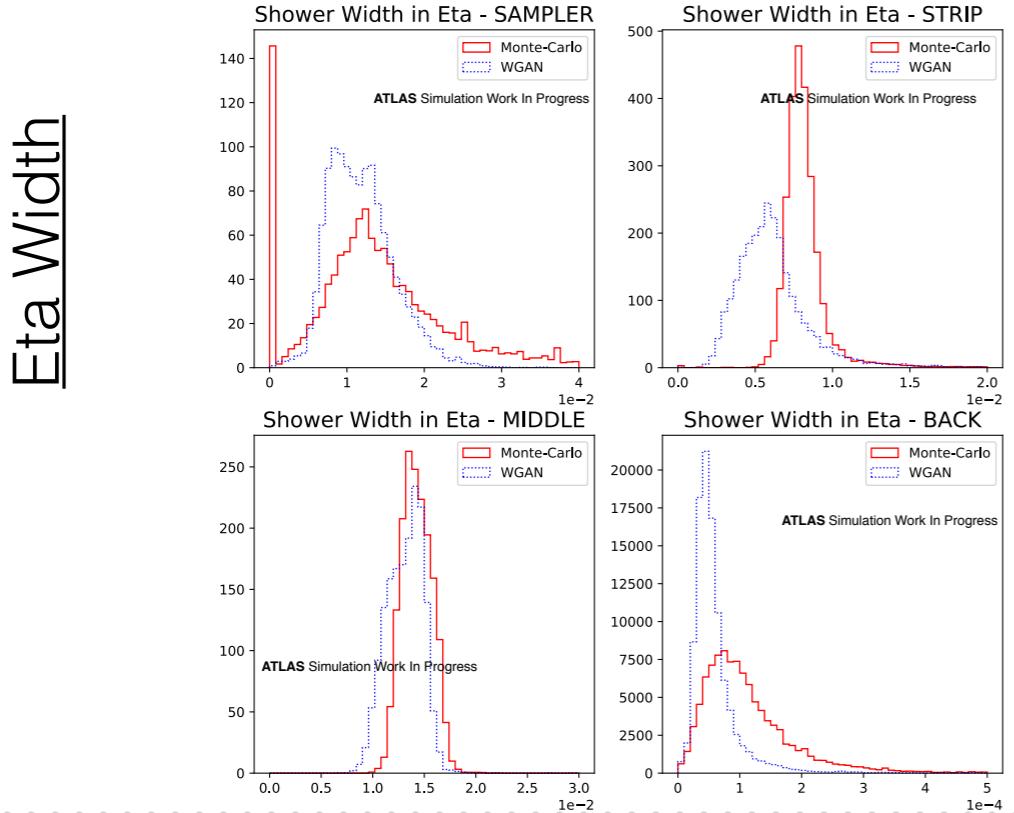


Mixed Config

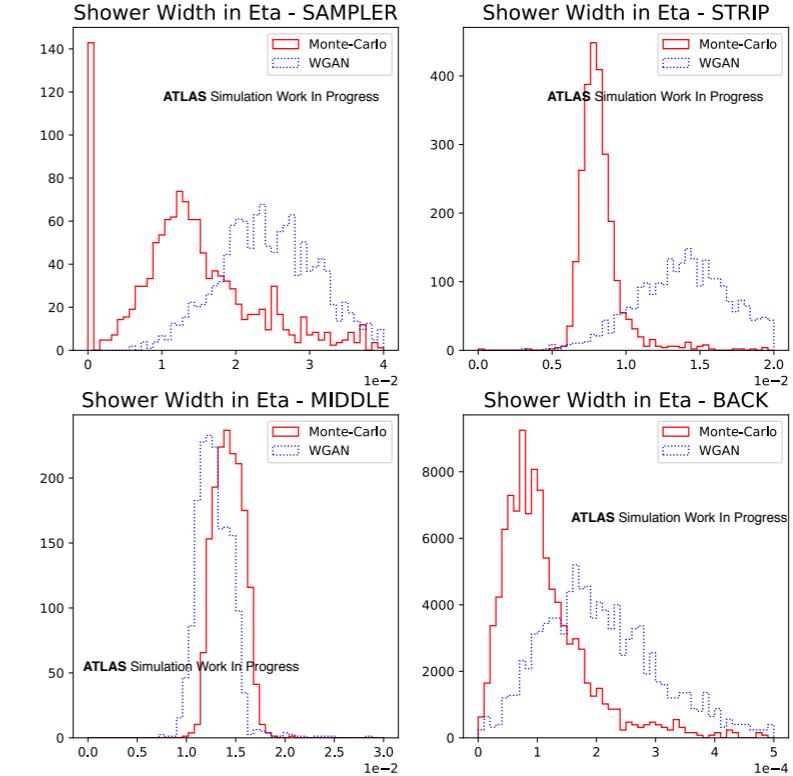


Mixed Config: All 10k Events vs 10k/8 Events

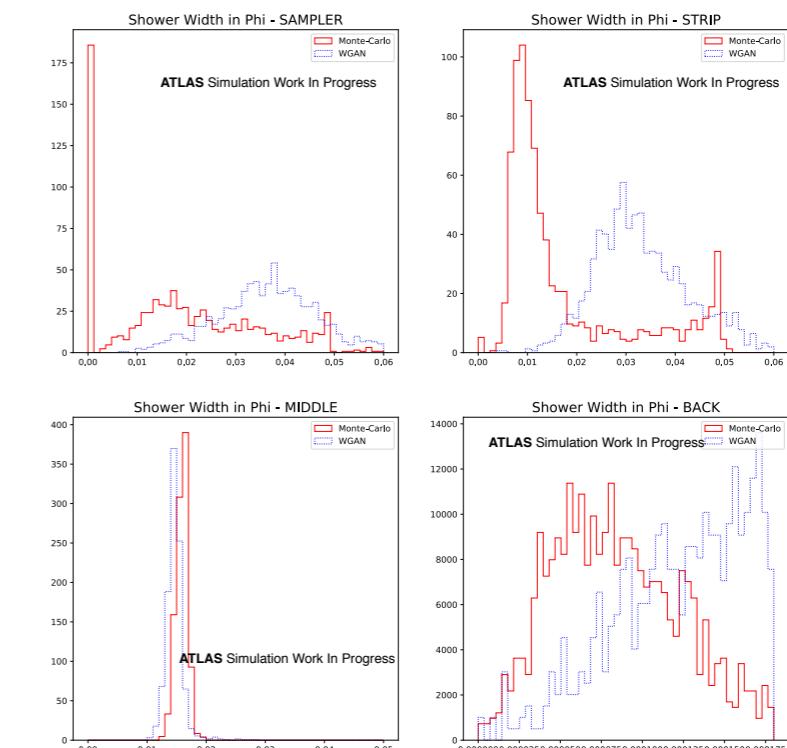
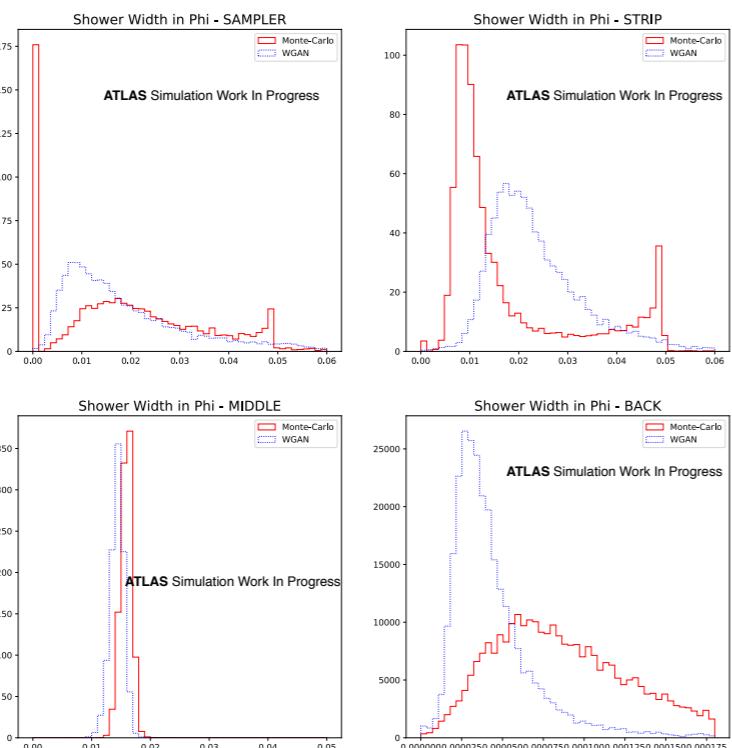
10k Events



Only 10k/8 Events



More statistics helps GAN learn widths (but no other distributions)

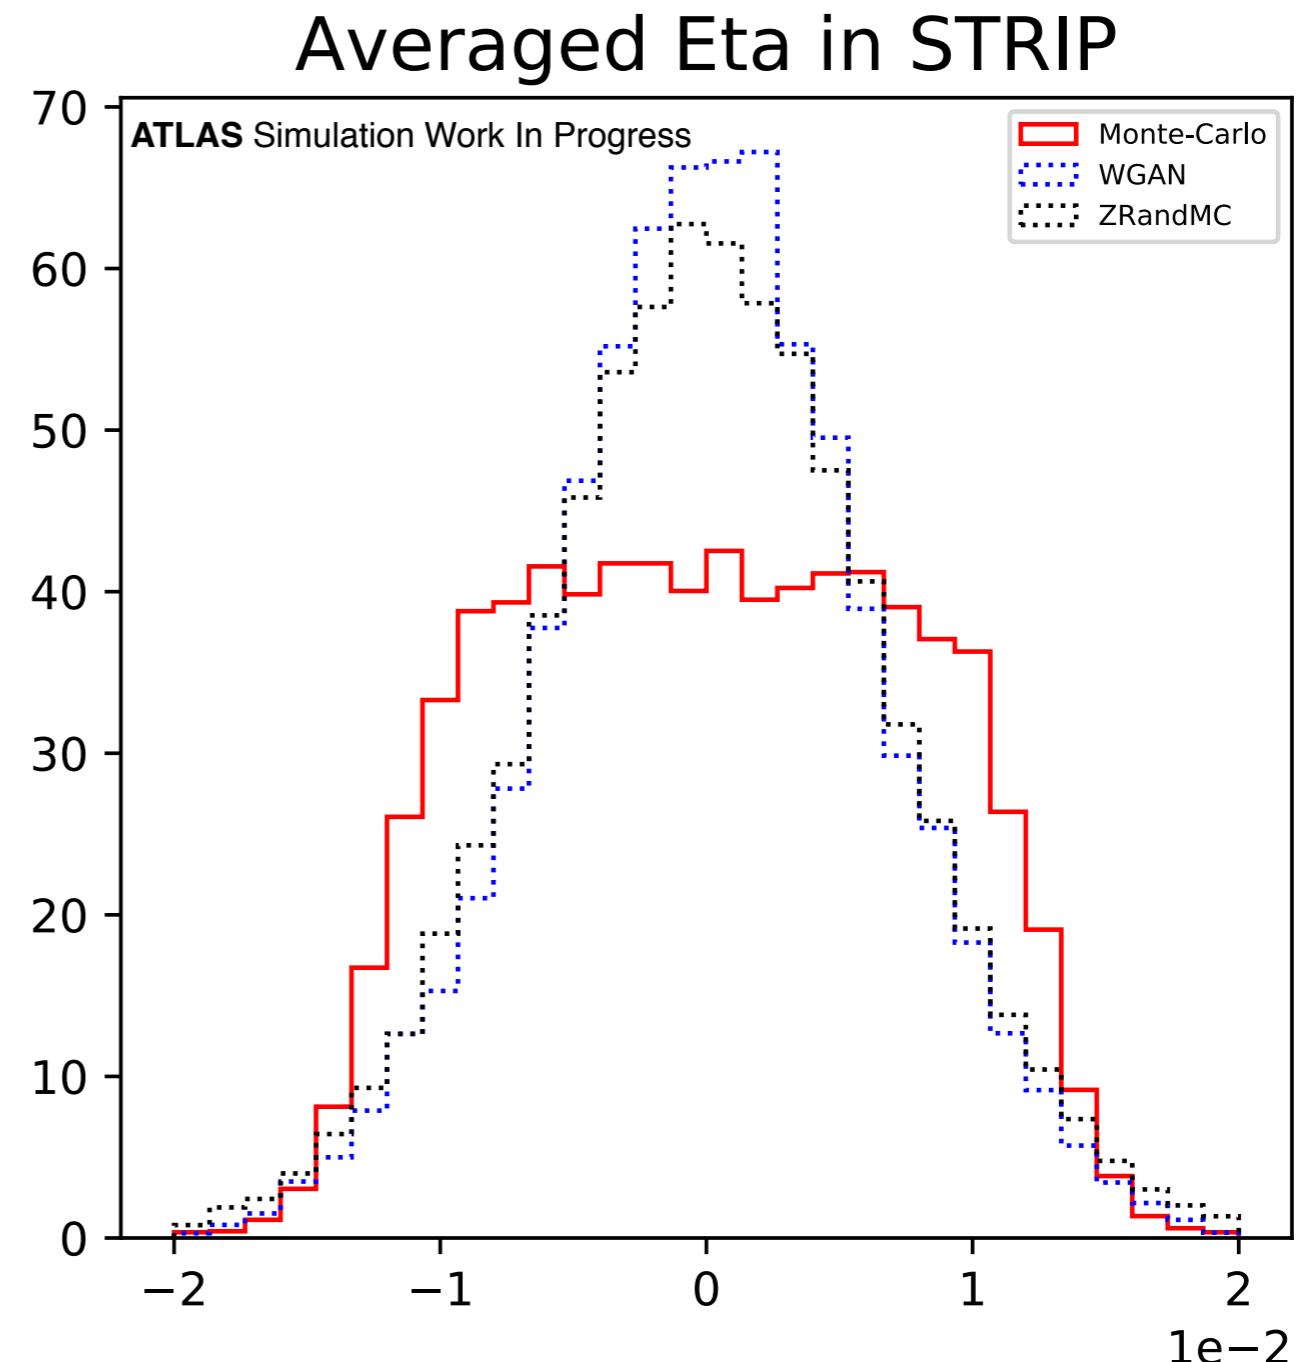


Unfair comparison for GAN in xAOD dataset

	Red	Blue	Black
Cell Energy (E)	MC	GAN	MC
Cell η	From same event	From an MC event	From different MC event
E- η Correspondence	Correct	N/A	Randomised

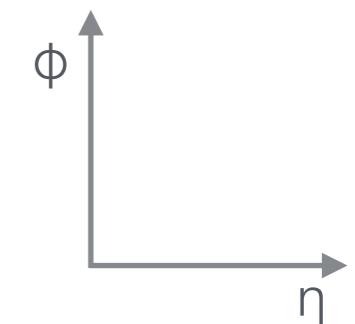
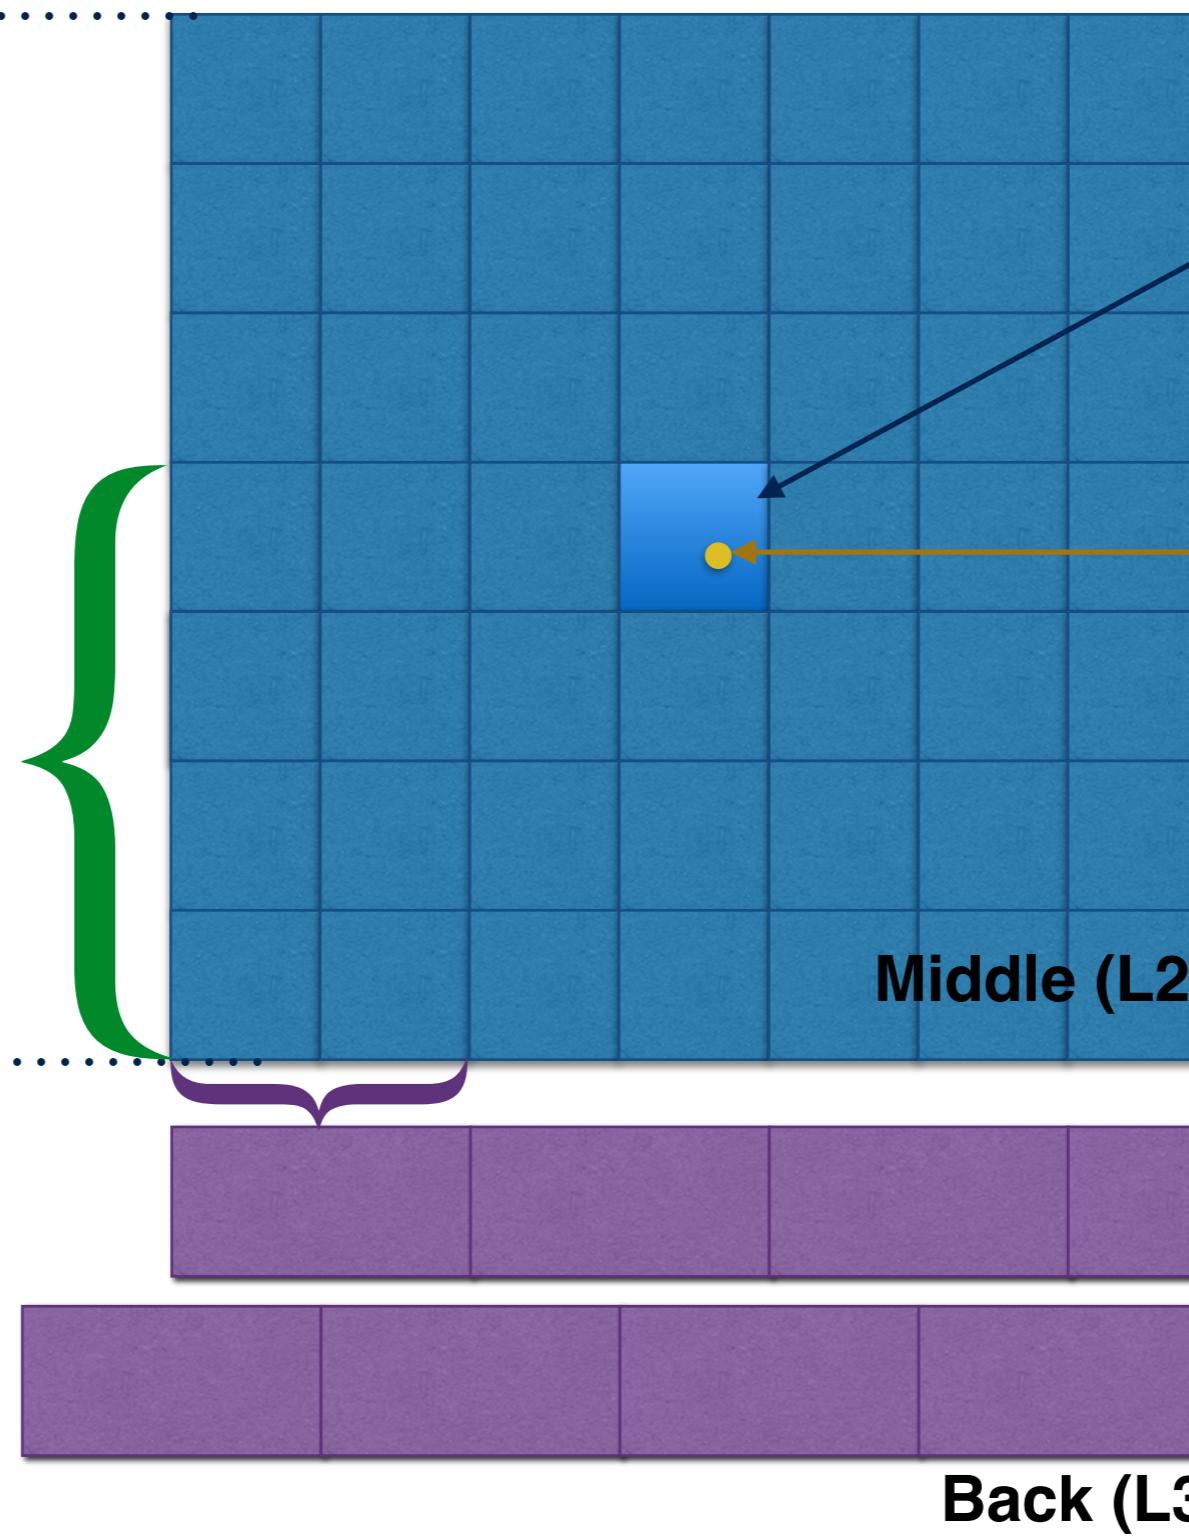
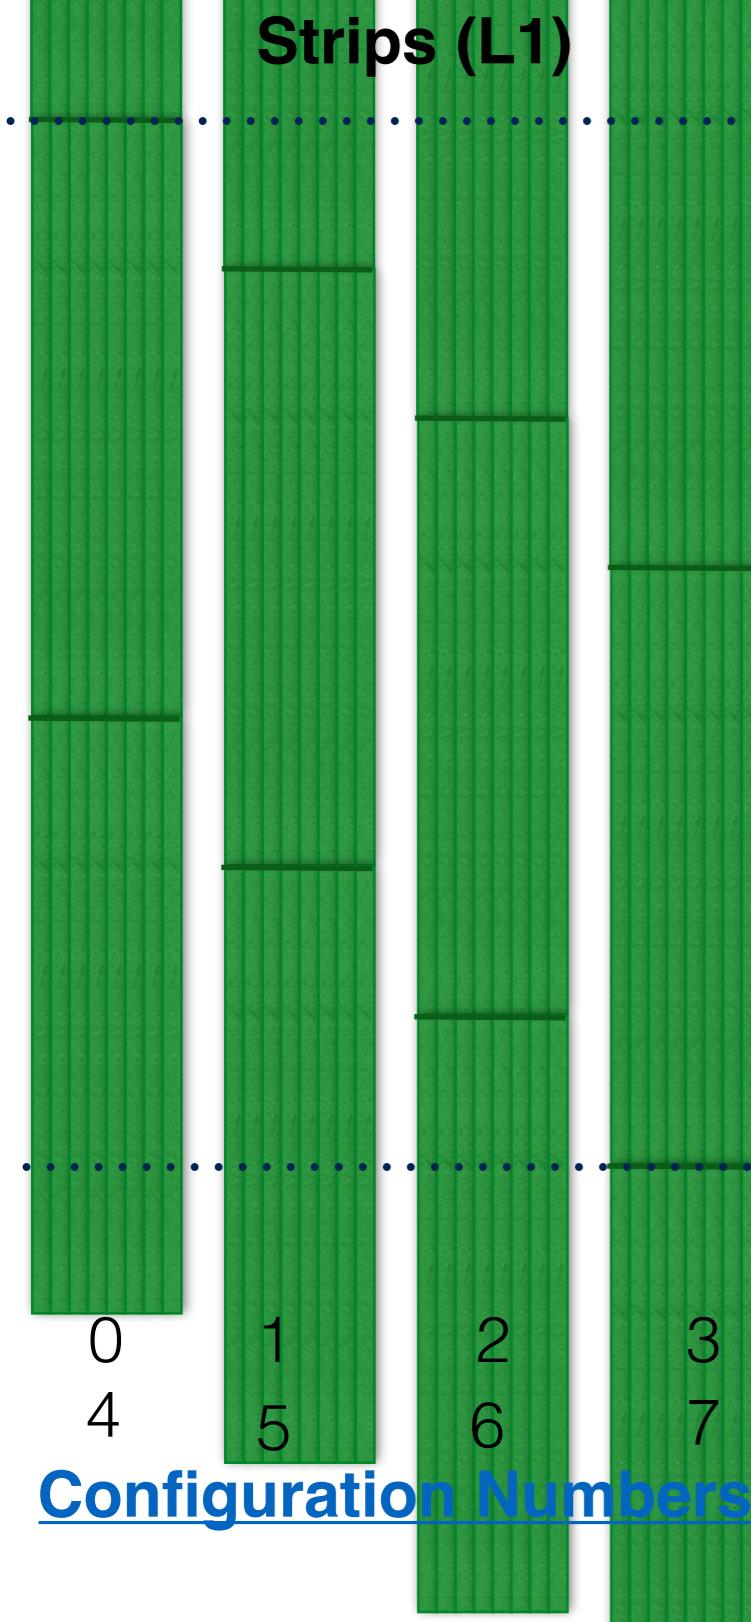
The previous analysis
(see [1](#), [2](#)) did not account
for this inherent advantage
of MC over GAN

GAN only knows Energy, no η/ϕ info.
For validation, η/ϕ came in from MC
data in previous analysis



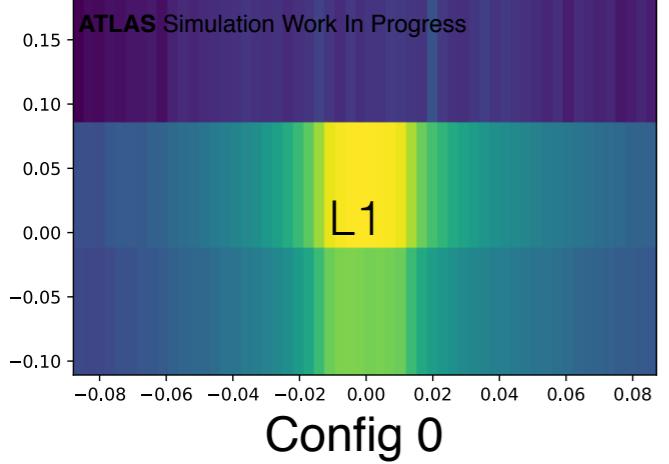
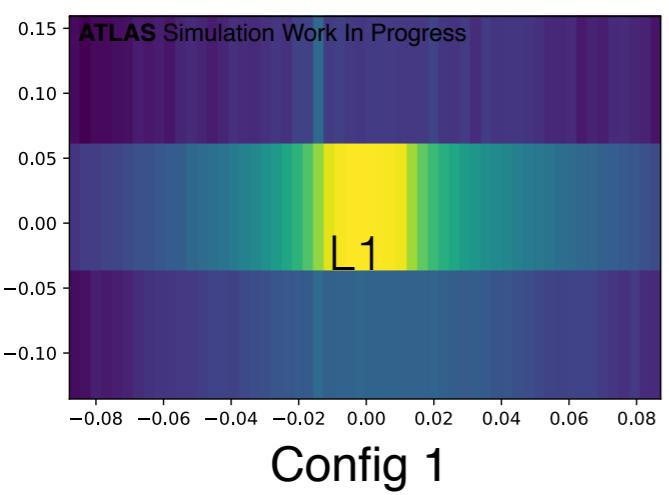
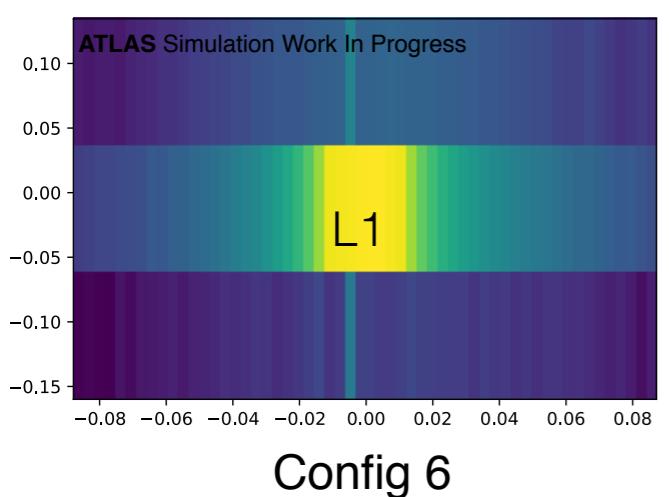
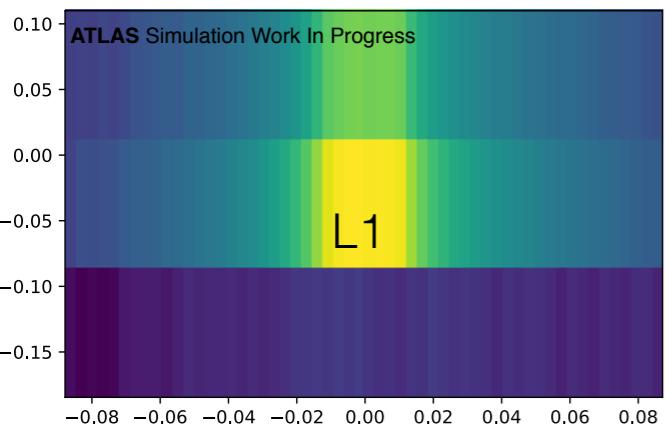
Eight Different Possible Configurations

Pre-Sampler (L0) ~ Strip (L1) for DelPhi



Configuration
Numbers

ATLAS Simulation



Strip (L1)

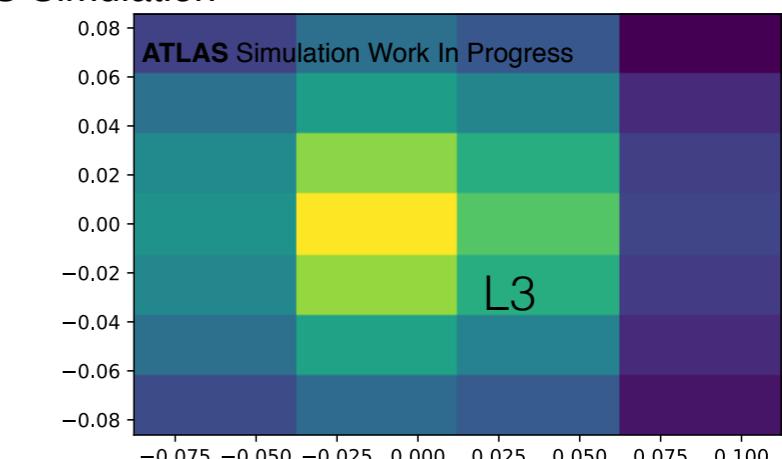
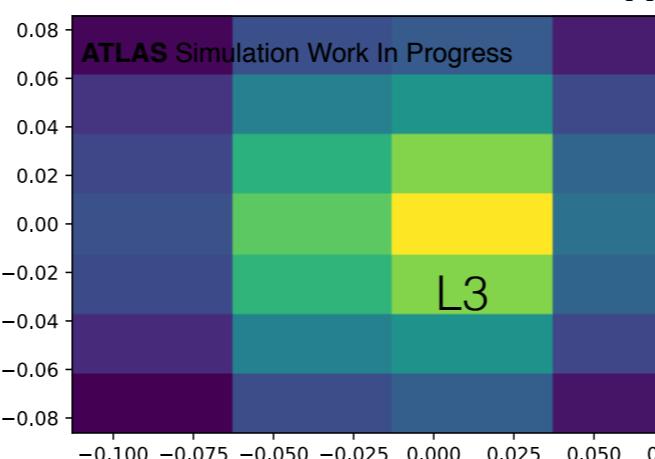
Strip (L1)

0 1 2 3 4 5 6 7

Back (L3)

Back (L3)

ATLAS Simulation



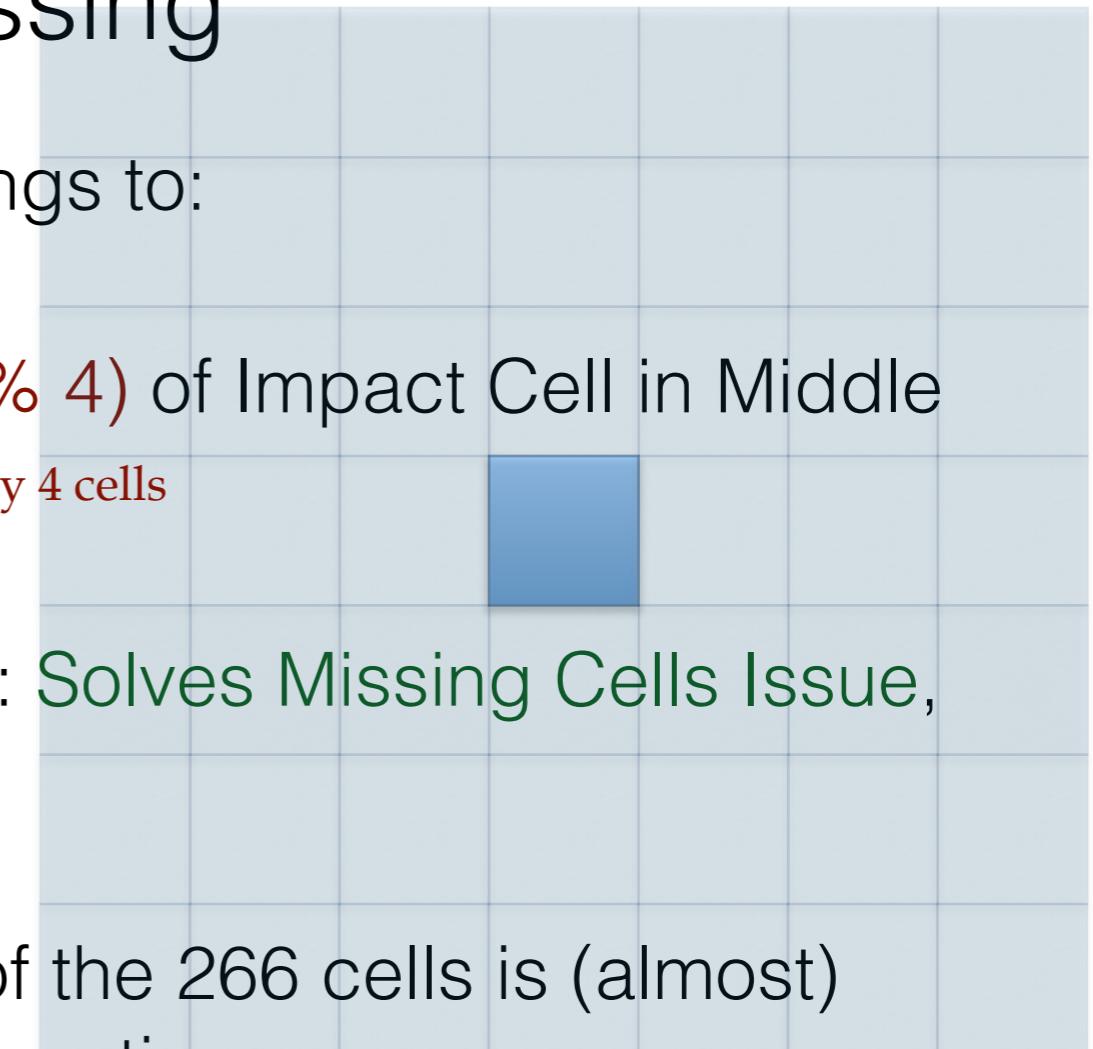
Config 0

Config 7

So How Can We Train On These Images? :

Pre-Processing

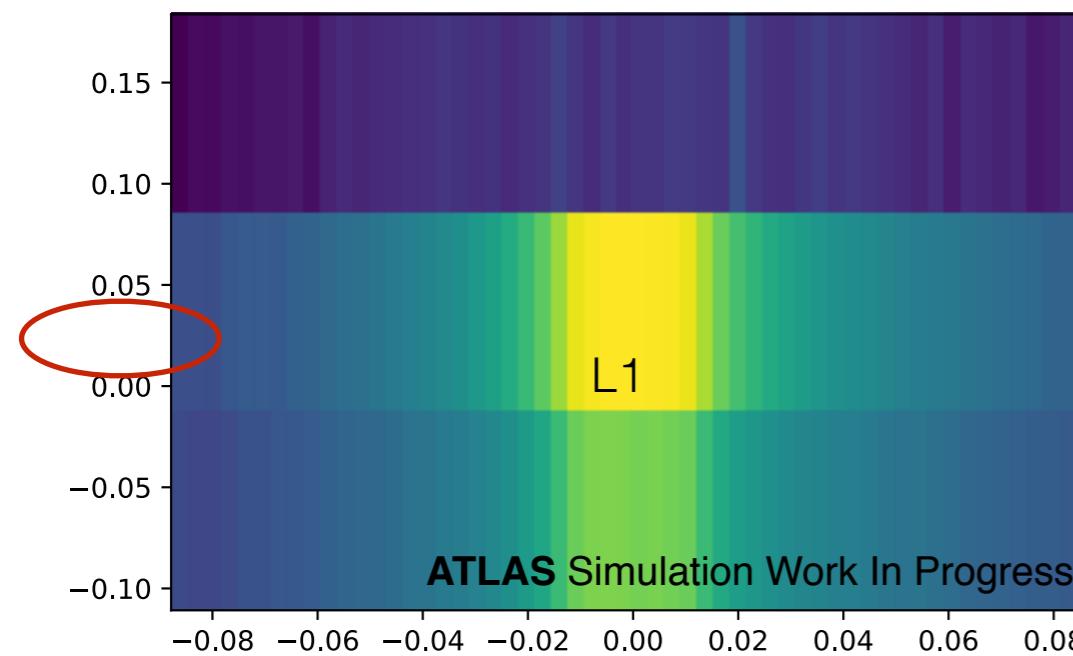
- Identify which config each event belongs to:
 - Using $(\text{eta_index \% 2}), (\text{phi_index \% 4})$ of Impact Cell in Middle Layer
 - Create 8 grids of η, ϕ for 8 configs : **Solves Missing Cells Issue**, brings back $\eta=0, \phi=0$
 - Verify that DelEta, DelPhi for each of the 266 cells is (almost) identical for all events of one configuration
- **Train separately** for each of the 8 configuration (but low stats: 10k/8 events)



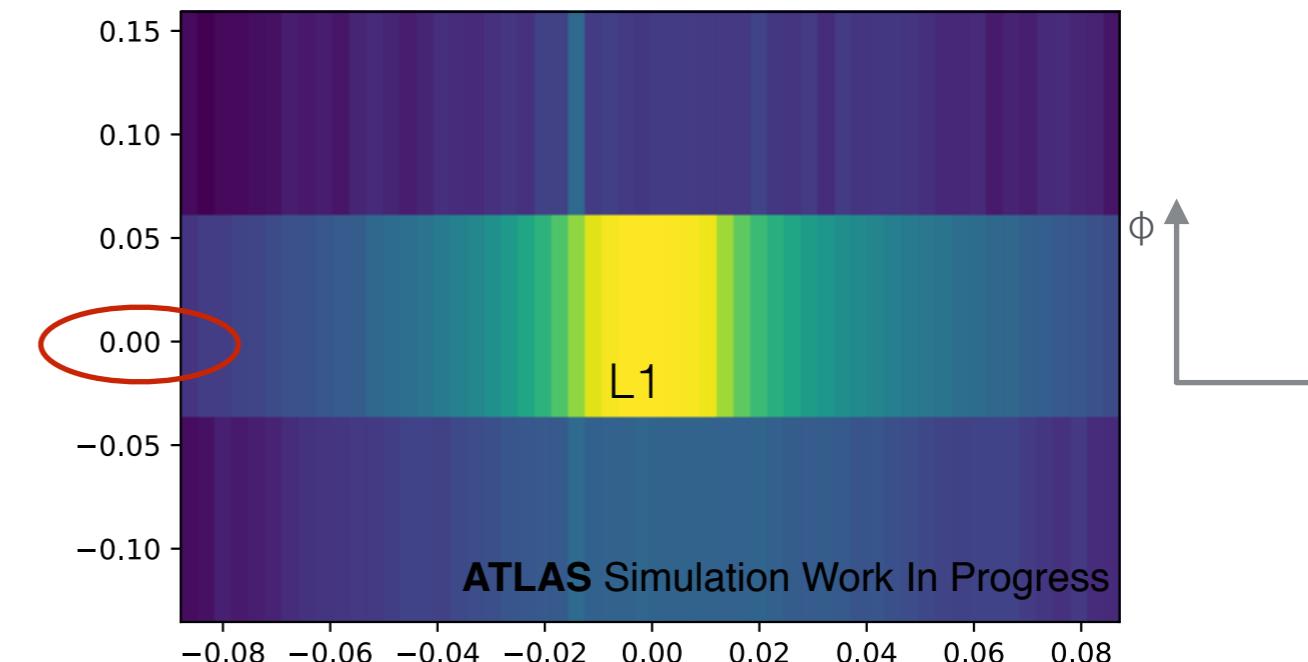
To Do:

- Tell the GAN about 8 configurations: one-hot encoding

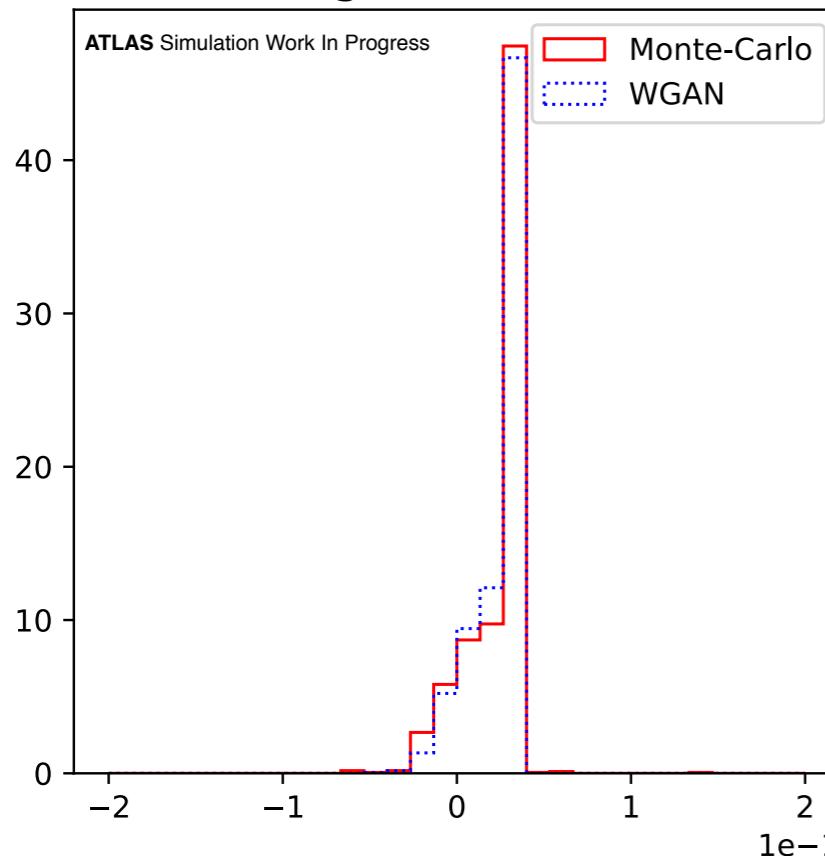
Strips Config 0



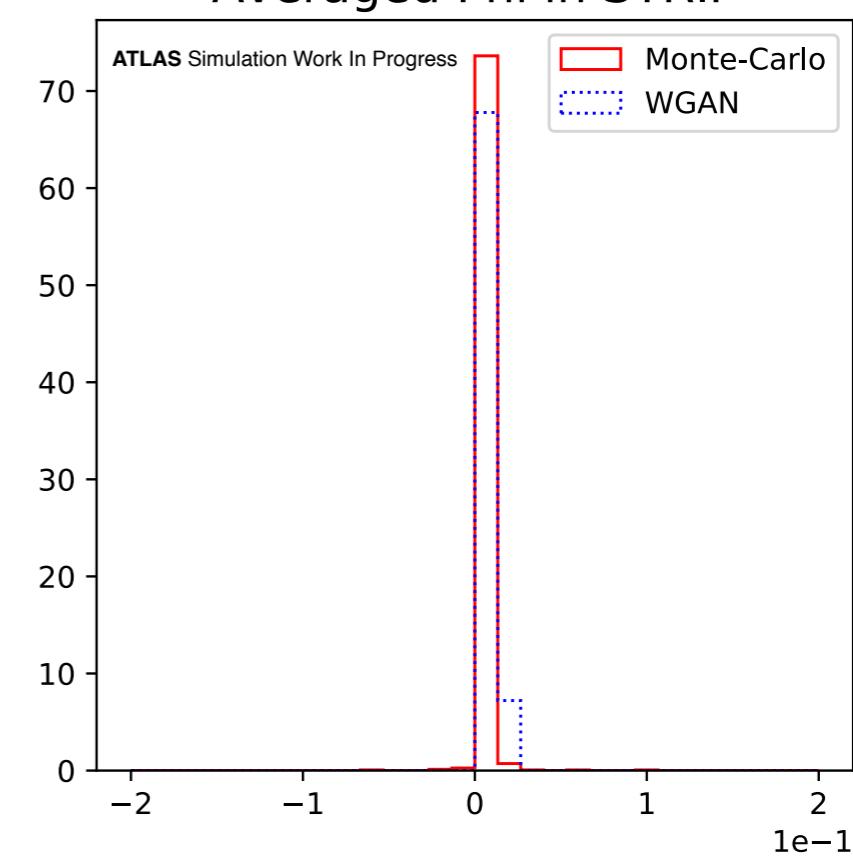
Strips Config 1



Averaged Phi in STRIP



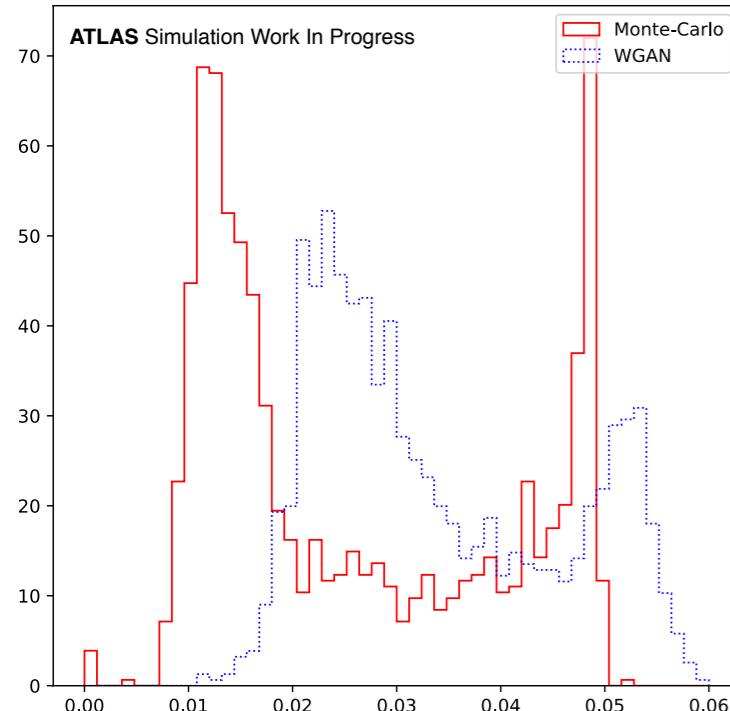
Averaged Phi in STRIP



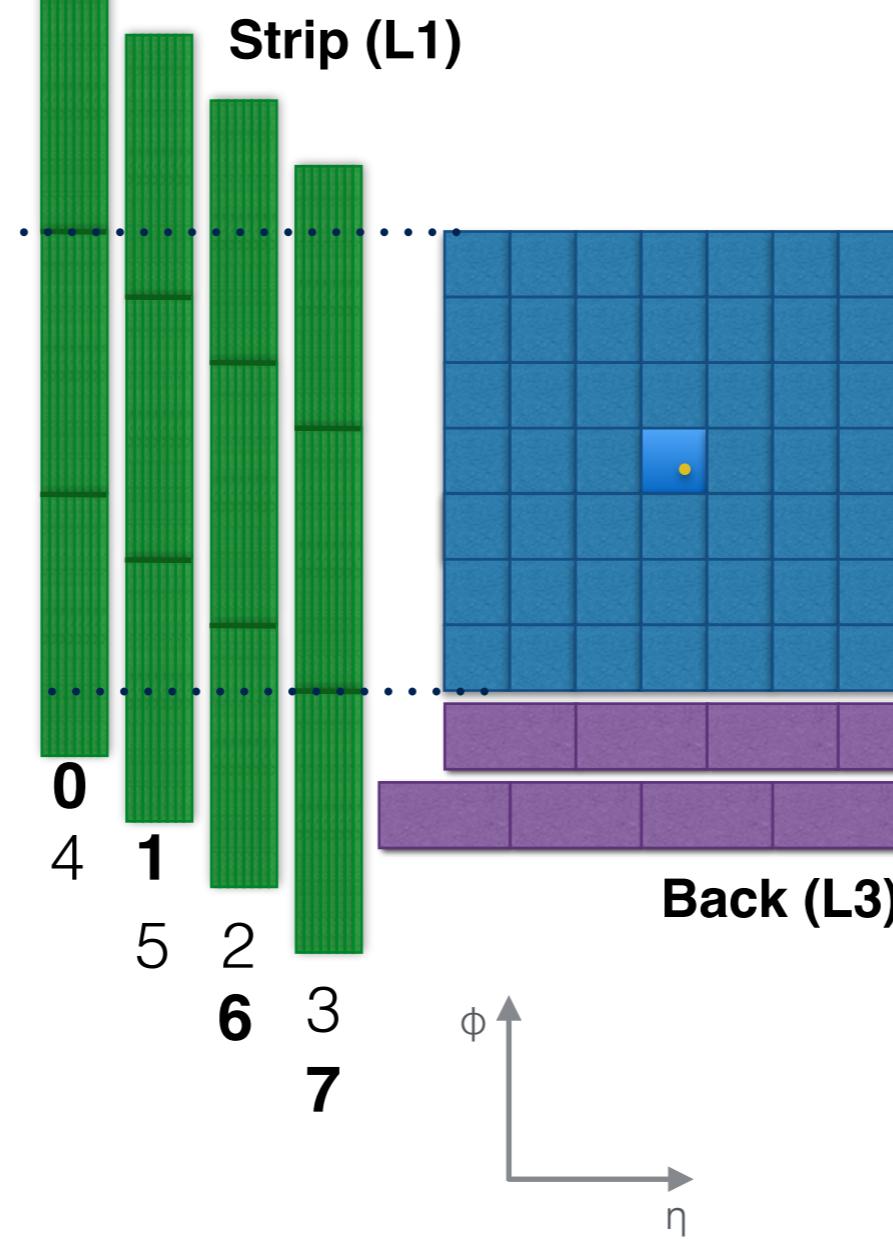
The two GANs learn their respective distributions well

Double Peak in ϕ Width in Strips

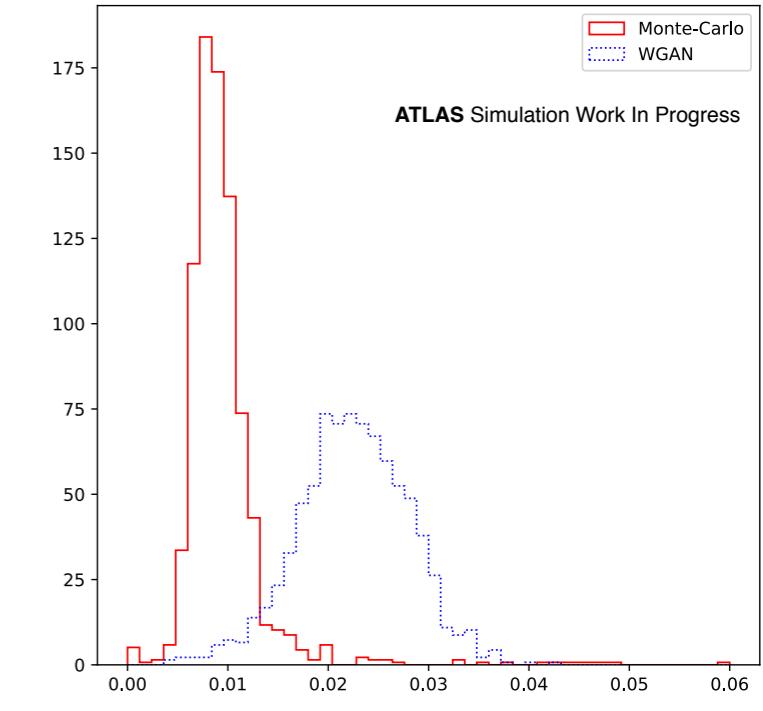
Shower Width in Phi - STRIP



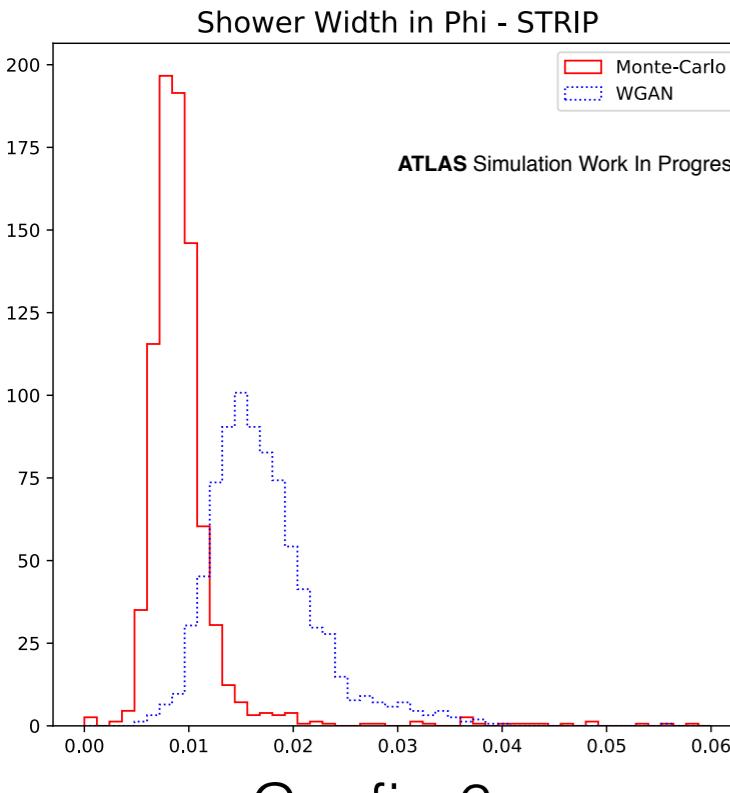
Config 0



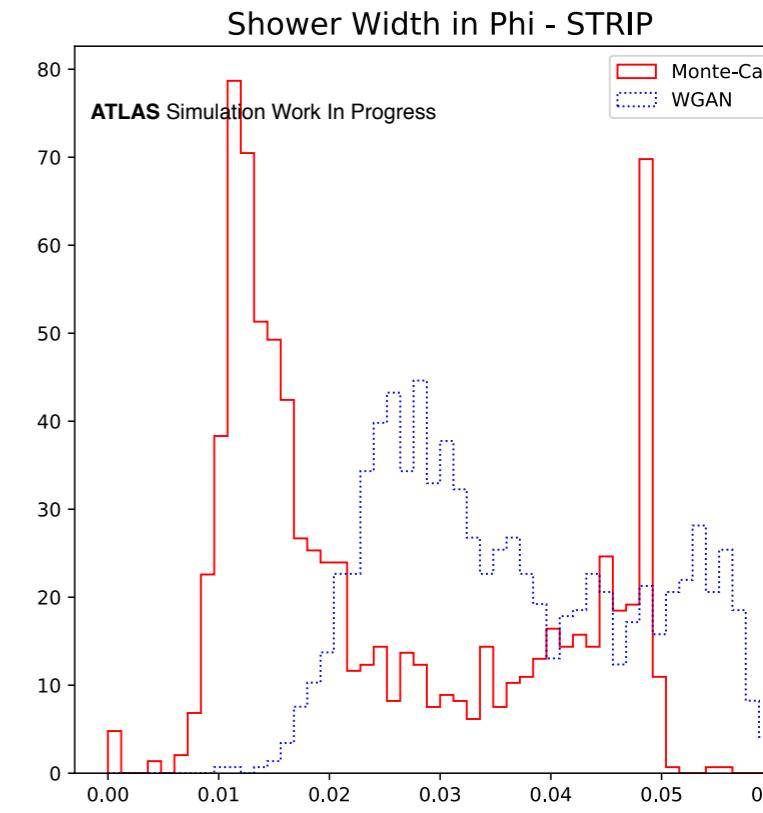
Shower Width in Phi - STRIP



Config 1



Config 6

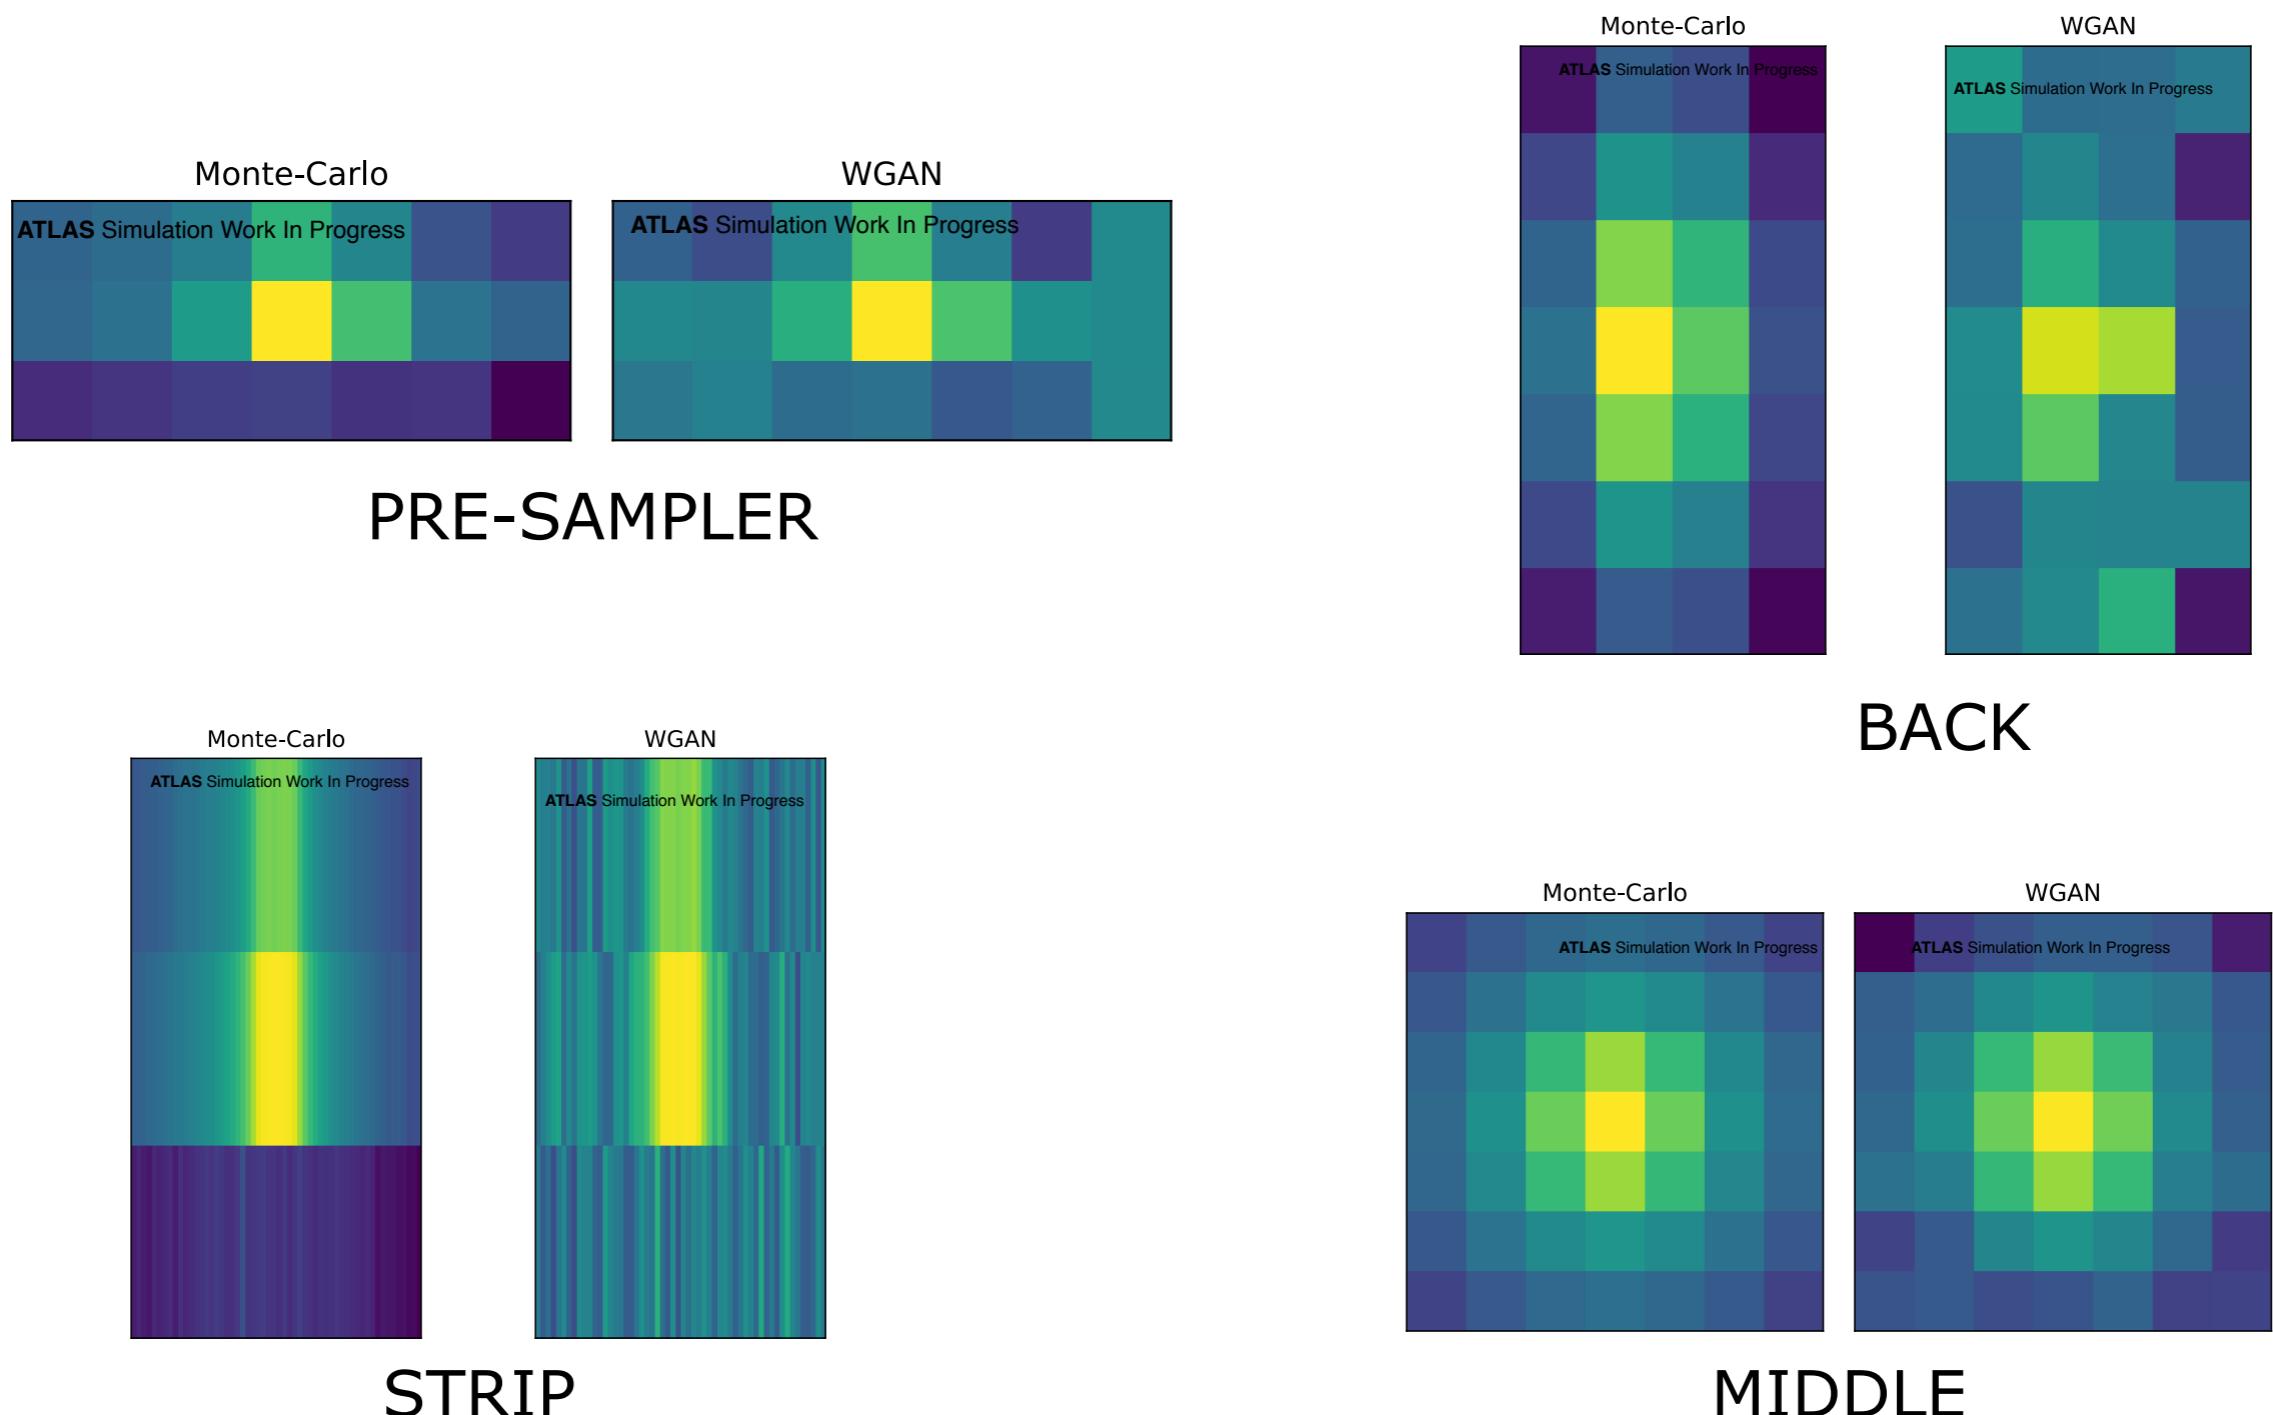


Config 7

Double peaks in config {0,7}, single peak in {1,6}

Training on Config 0

Average Images



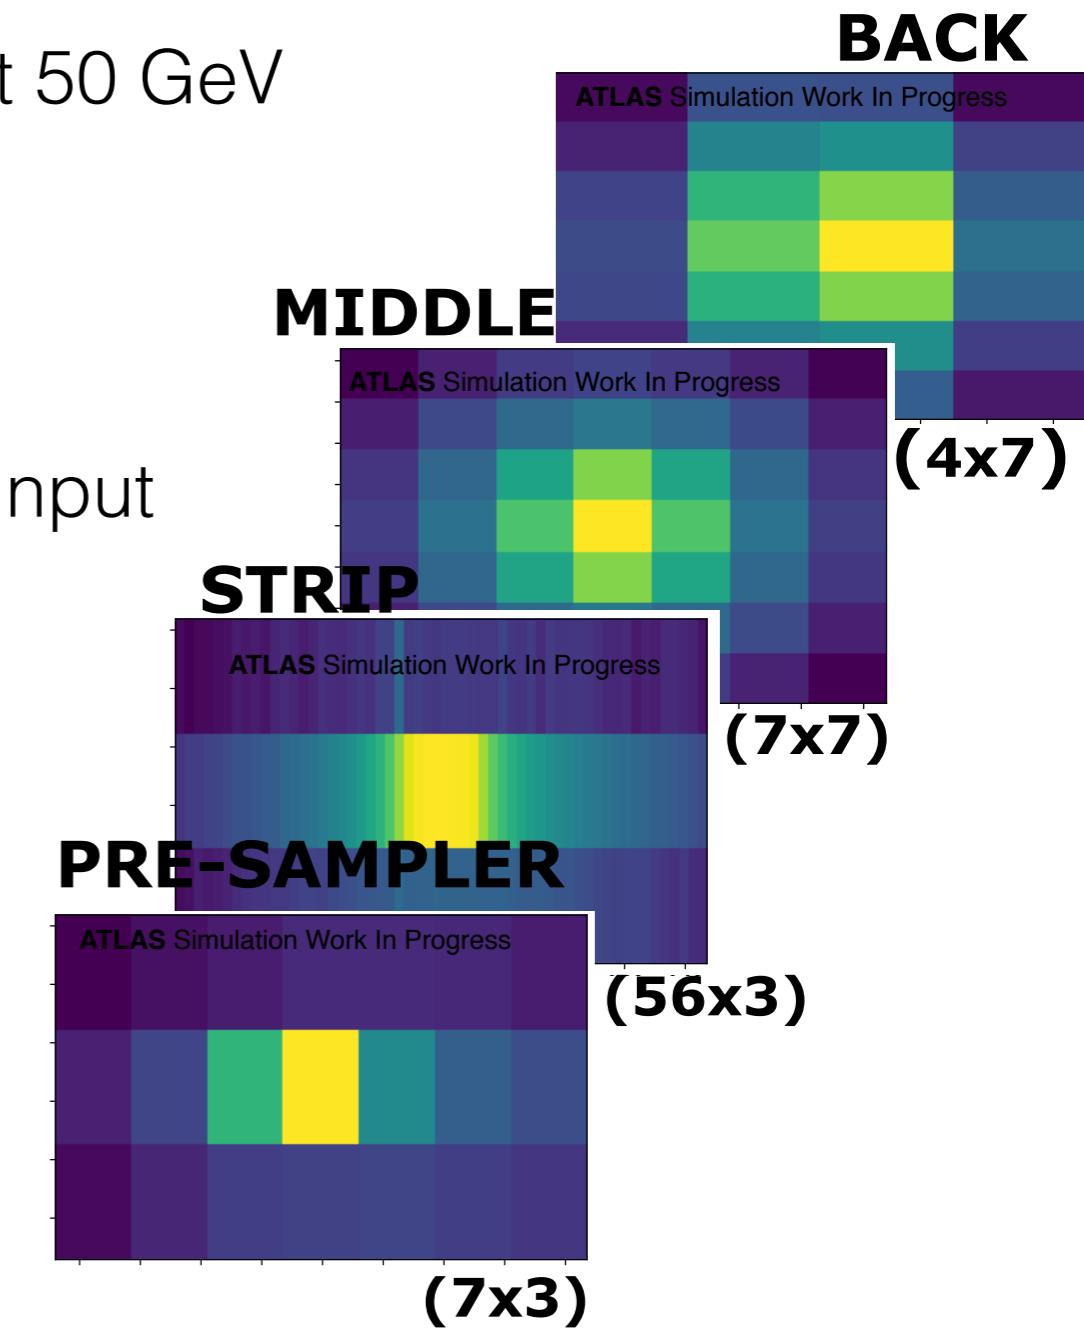
Warning: pixel sizes don't correspond to actual width of the cells

Improvement over mixed configurations training

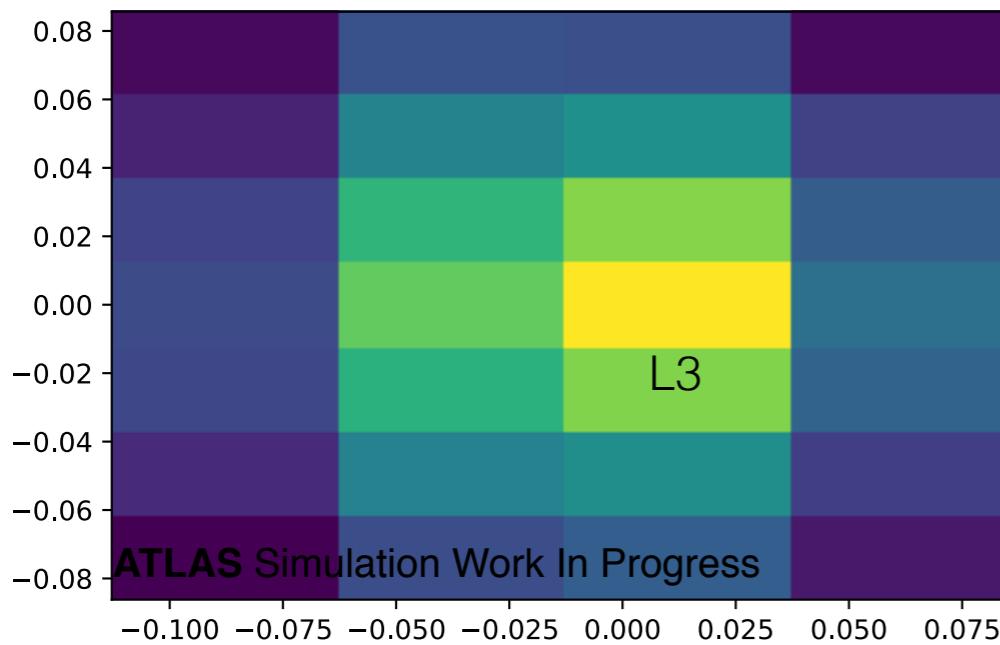
Dataset & Software

Taking over from Paul Klein, 2017

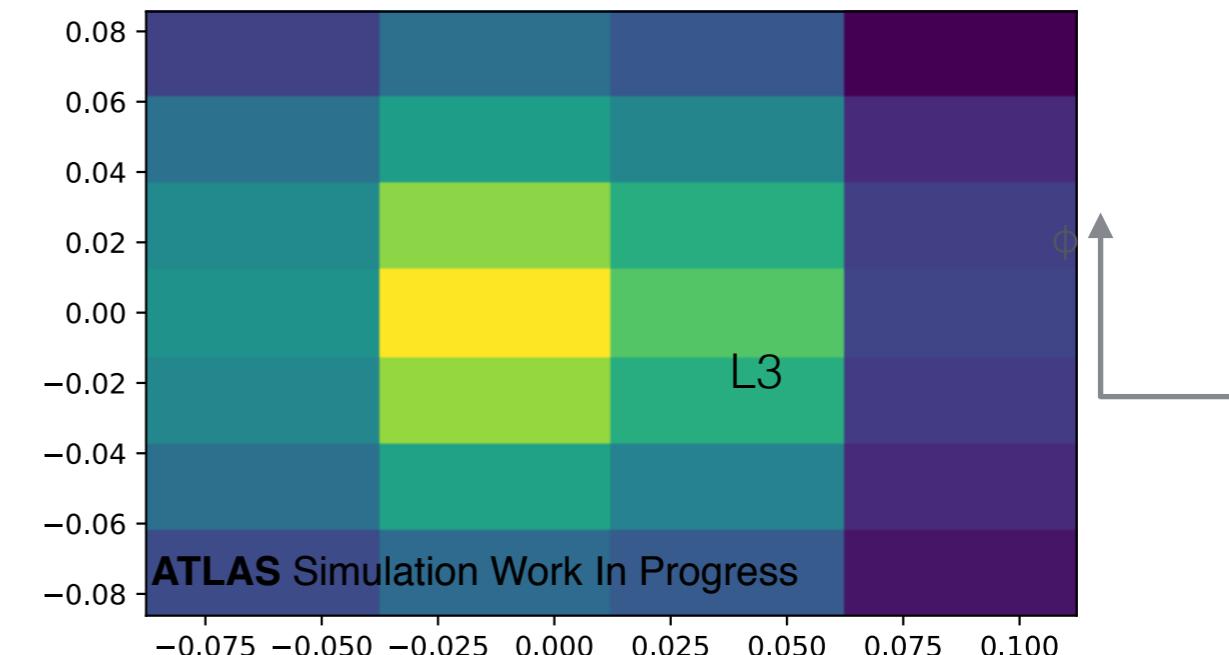
- Geant4 simulated single photon showers at 50 GeV
- $|\eta| \in [0.20, 0.25]$
- 266 Cells (Covering 4 layers) as 1-D array input
- 10k Events
- Train using: **Keras with Tensorflow**
- GAN Training time: ~ 2h50m
- Training Platform: **1 GPU at CC Lyon**
- GAN Generation Time: 1-2 seconds for 10k events



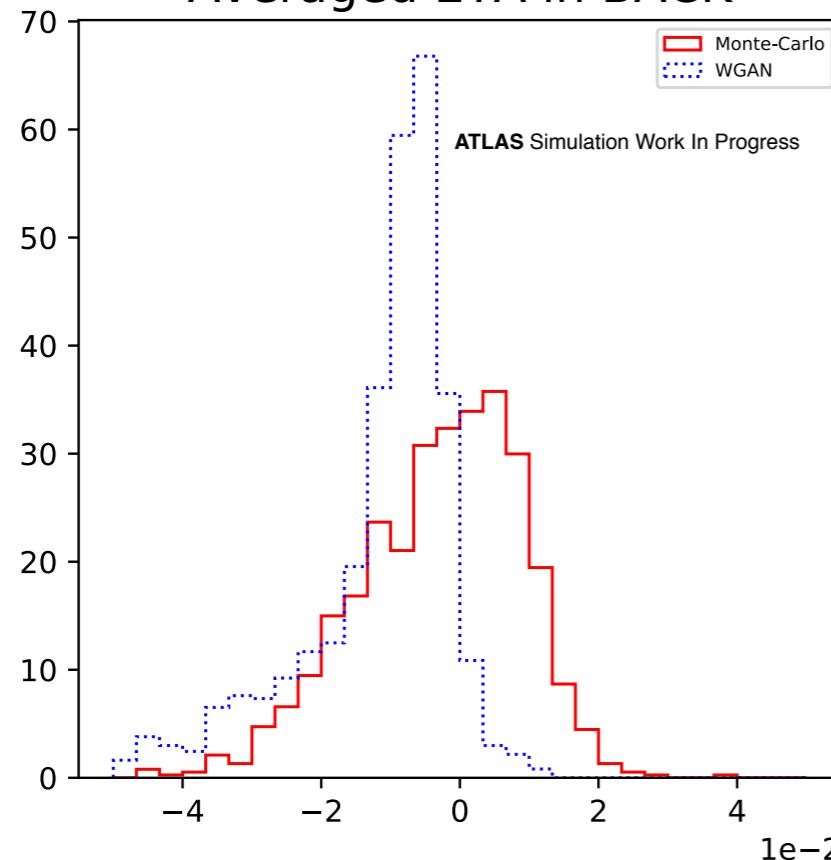
Back Config 1



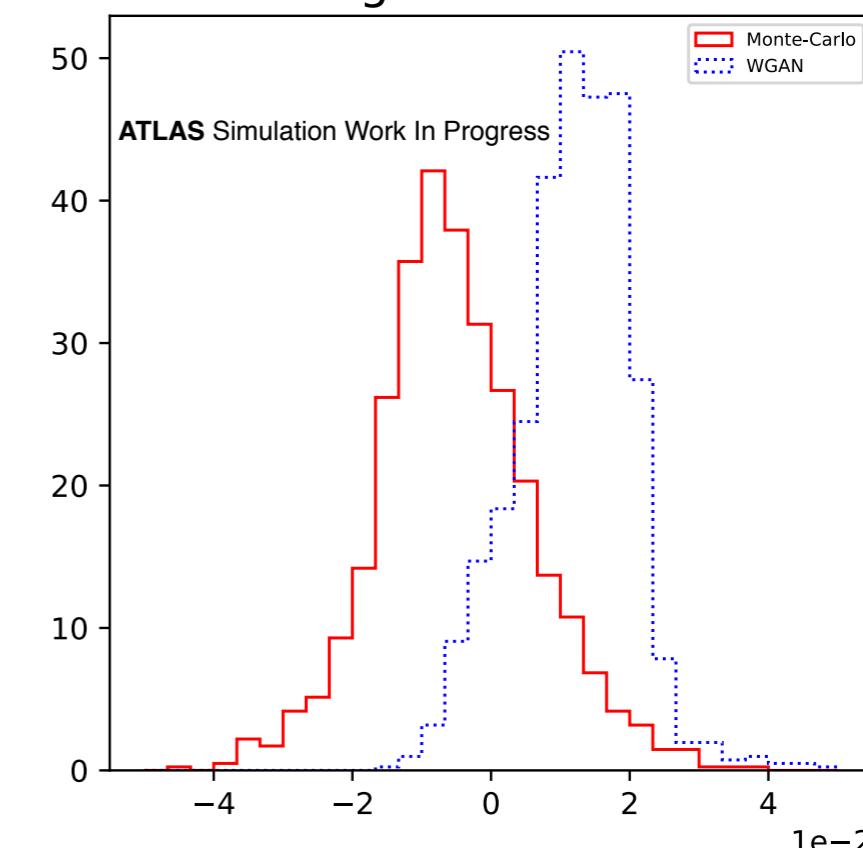
Back Config 7



Averaged ETA in BACK

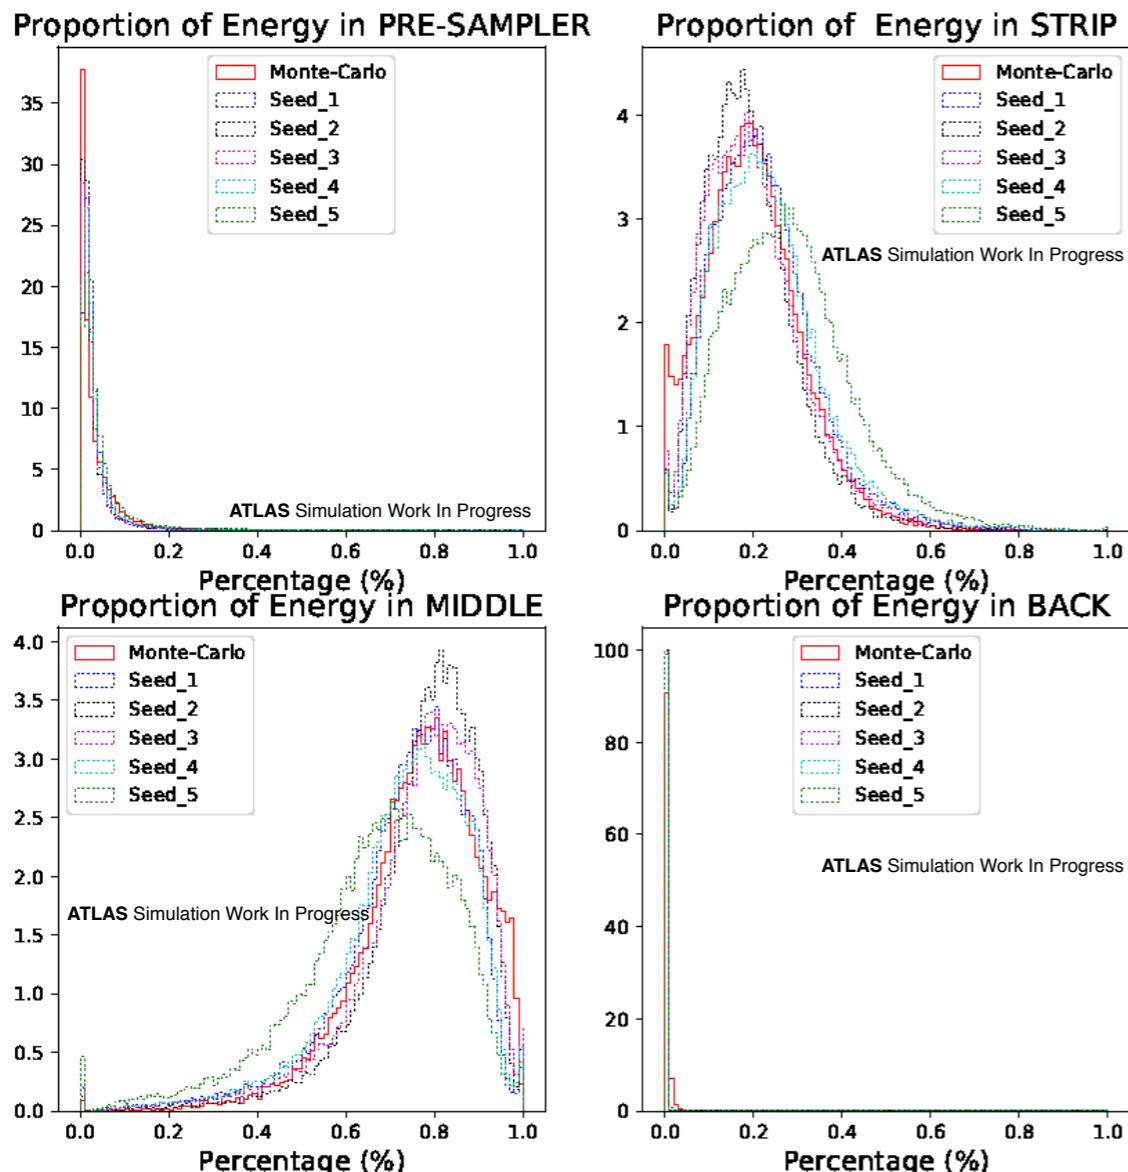


Averaged ETA in BACK

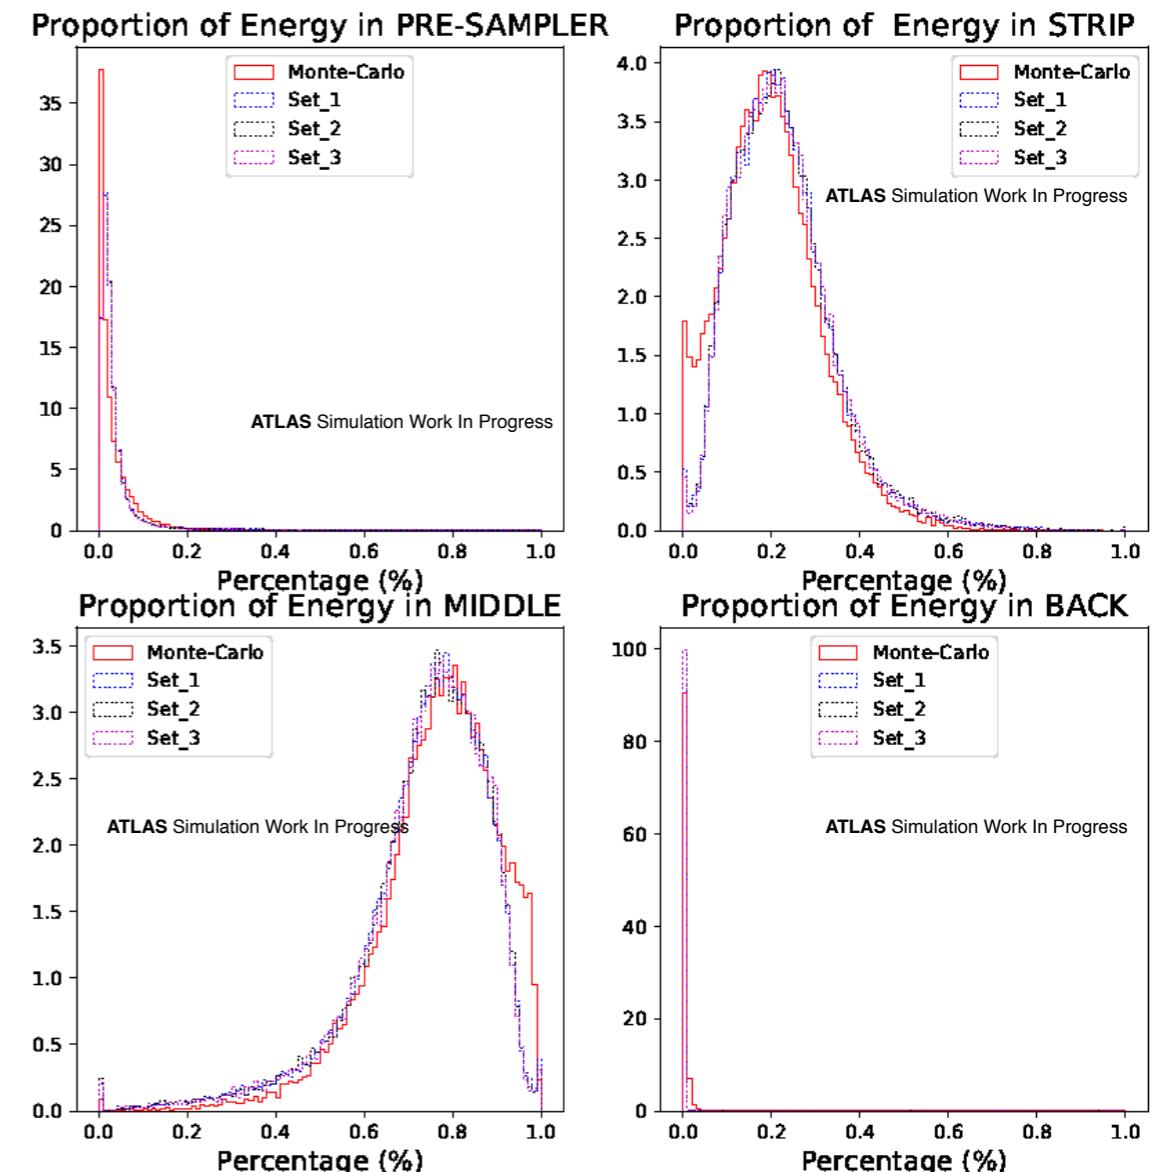


Variation in Training for Random Seed

Different training seeds



Same training seed;
different generation seeds



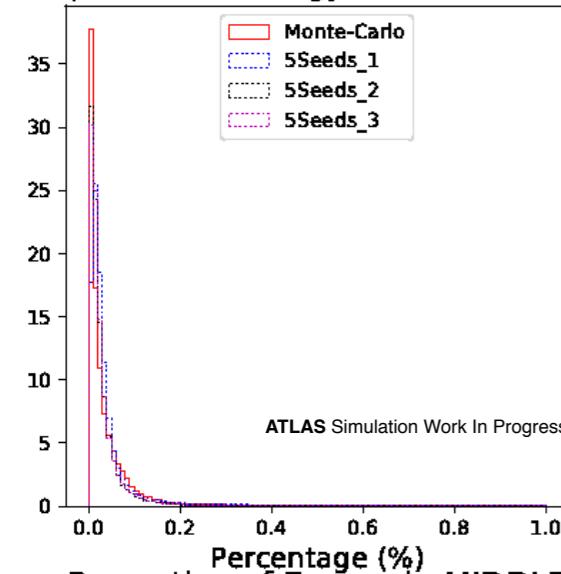
Initial training seed has significant impact on results. Need to average over different seed training to assess the GAN

Hyperparameter Optimisation: Training Ratio

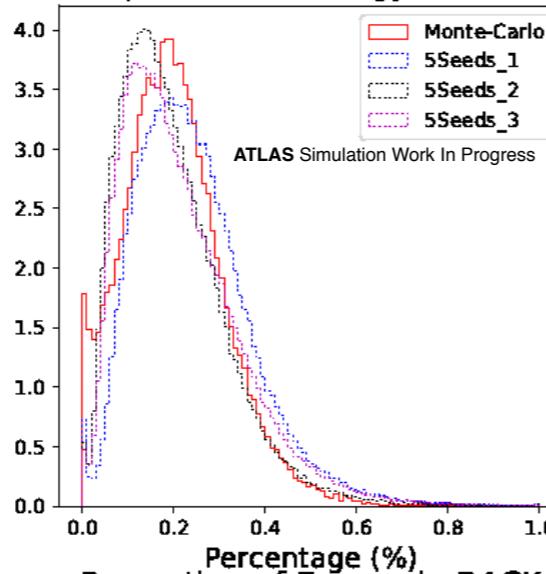
Same hyper-params,
3 sets of 5 separately trained GAN datasets

1,5,10 Ratio; combine 5 trained GANs for each

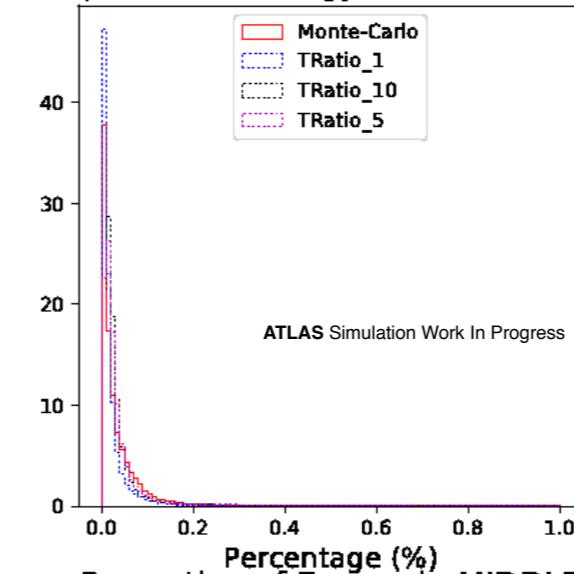
Proportion of Energy in PRE-SAMPLER



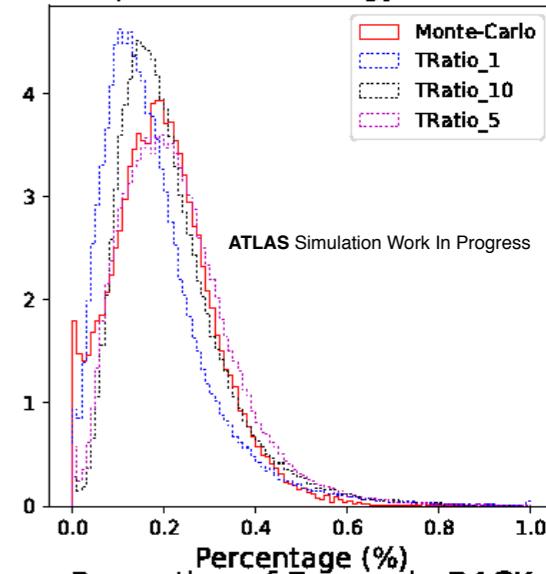
Proportion of Energy in STRIP



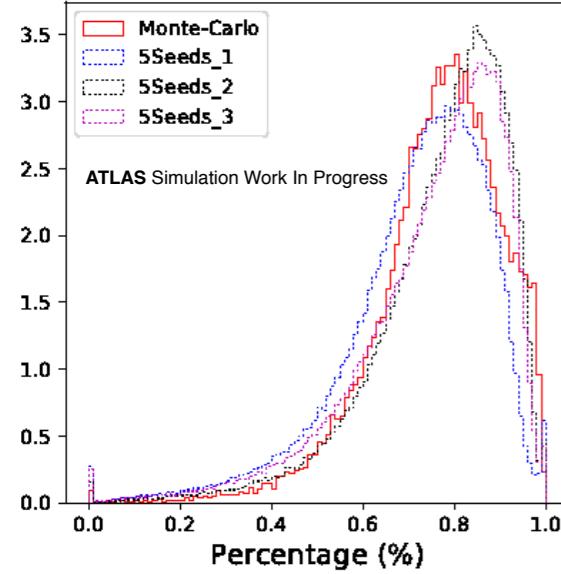
Proportion of Energy in PRE-SAMPLER



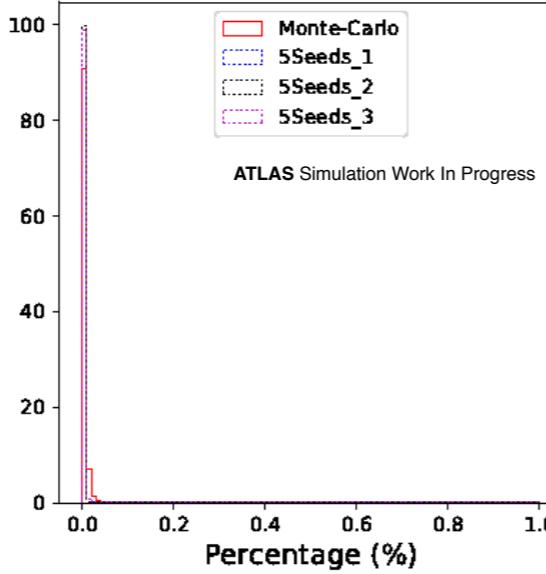
Proportion of Energy in STRIP



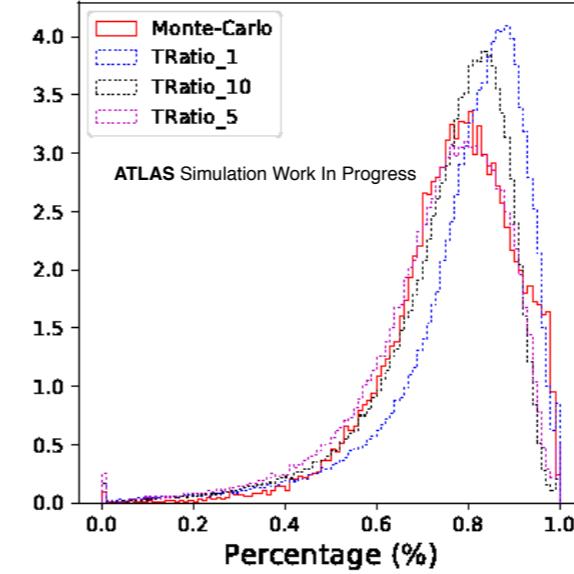
Proportion of Energy in MIDDLE



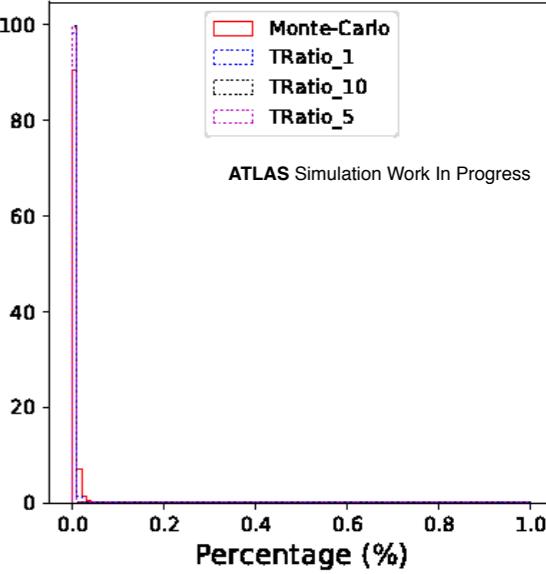
Proportion of Energy in BACK



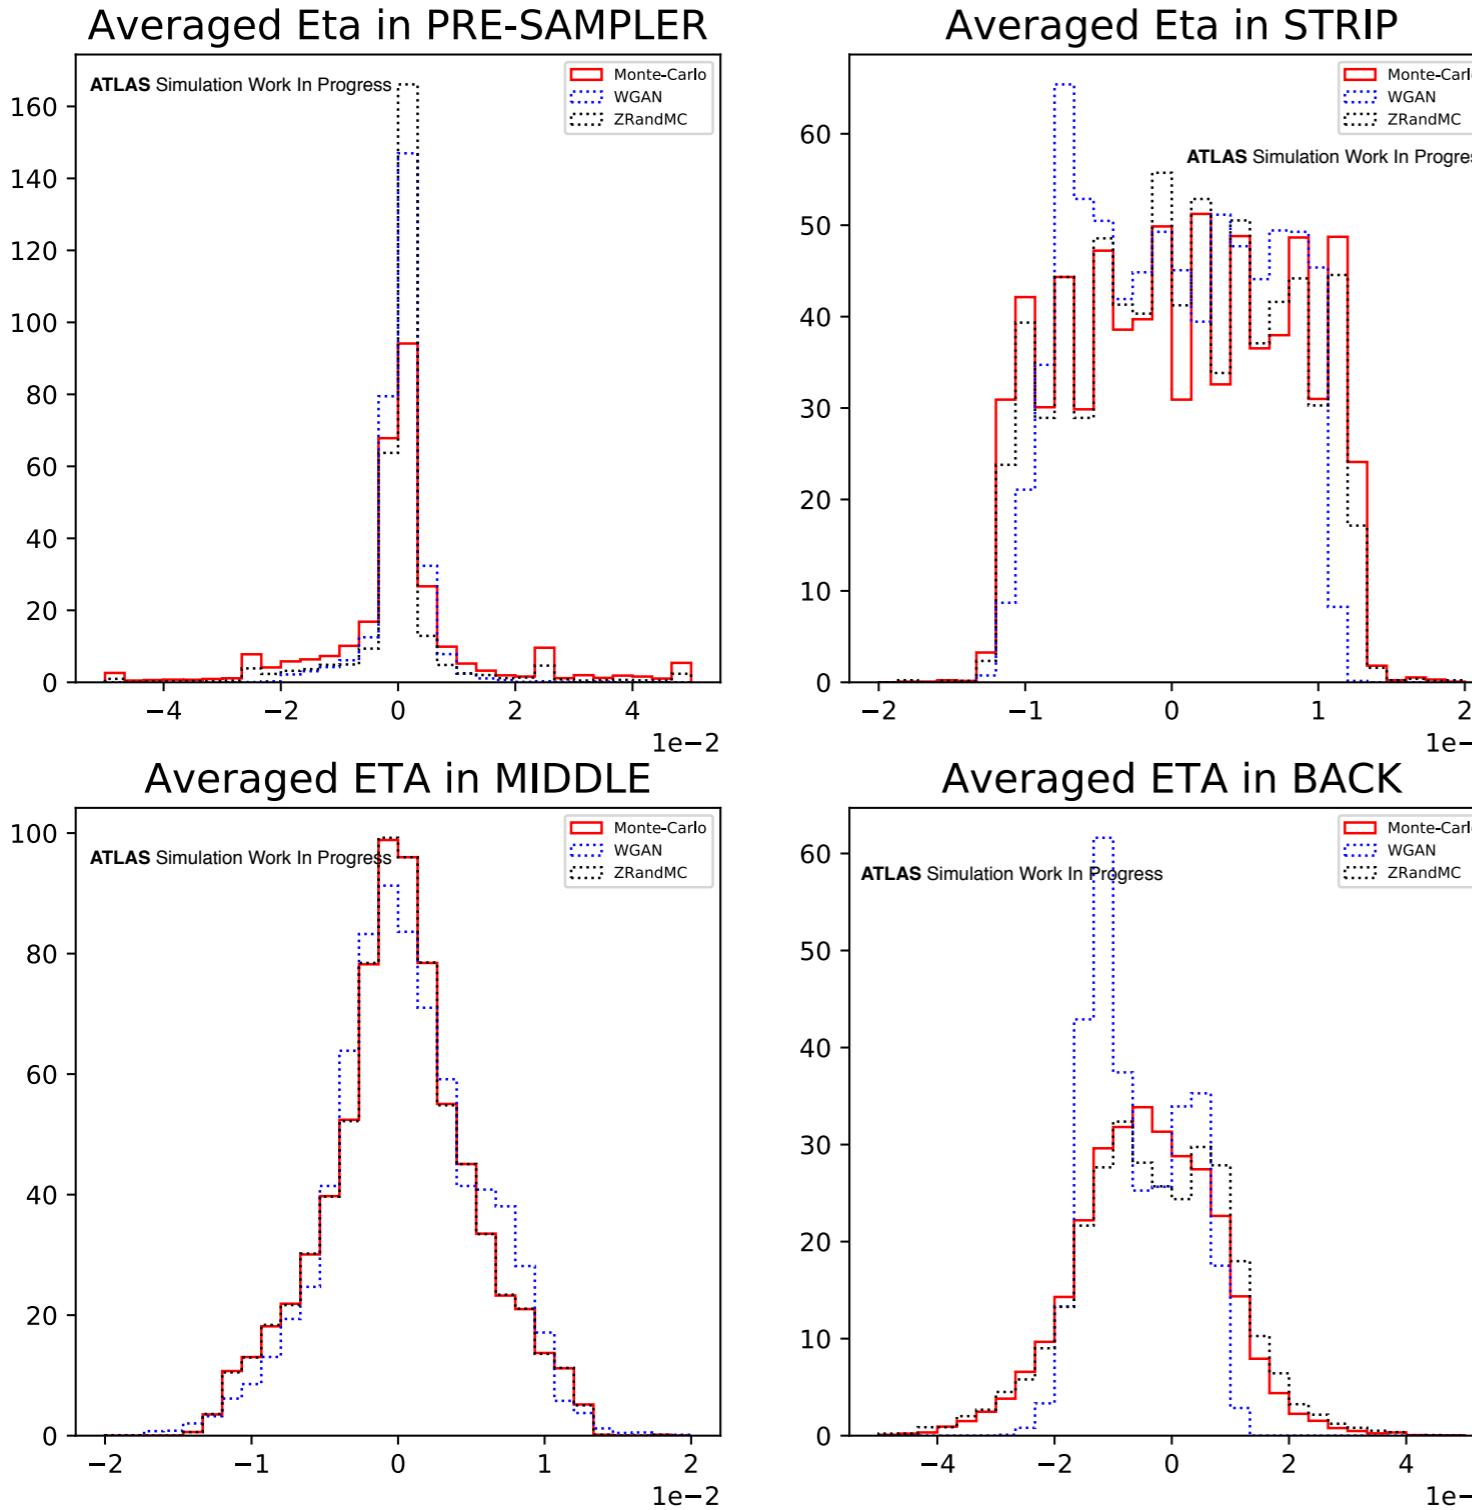
Proportion of Energy in MIDDLE



Proportion of Energy in BACK



Unfair comparison for GAN

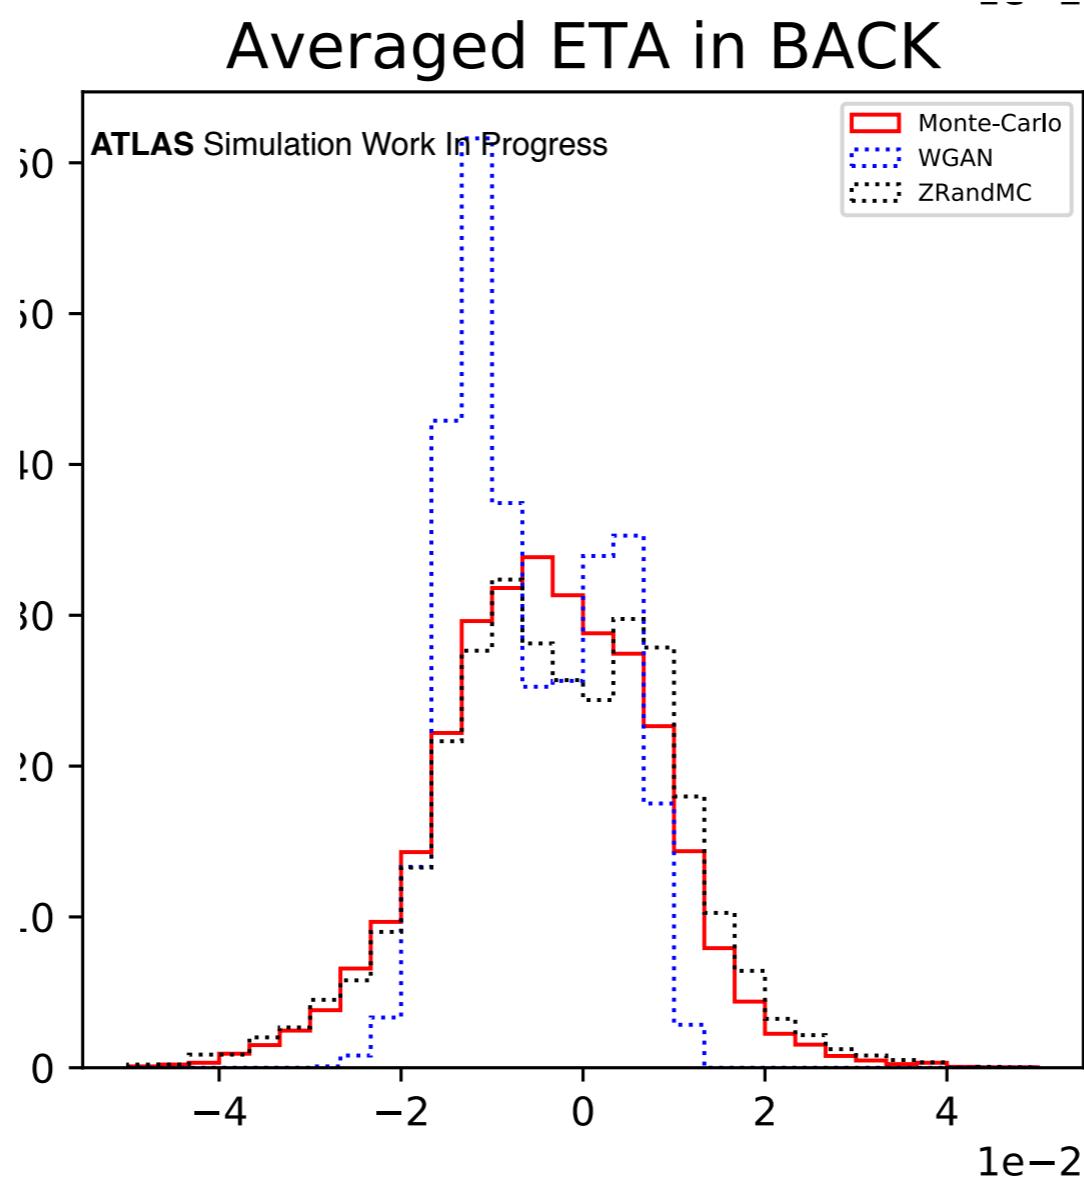


Very clearly that we're mixing these 8 different configurations.

MC Events Plotted with Eta of the Corresponding Events, WGAN plotted with Eta of MC Events, MC Events Plotted with Eta of Random Event

MC vs GAN vs MC

	Red	Blue	Black
Cell Energy (E)	MC	GAN	MC
Cell η	From same event	From an MC event	From different MC event
E- η Correspondence	Correct	N/A	Randomised

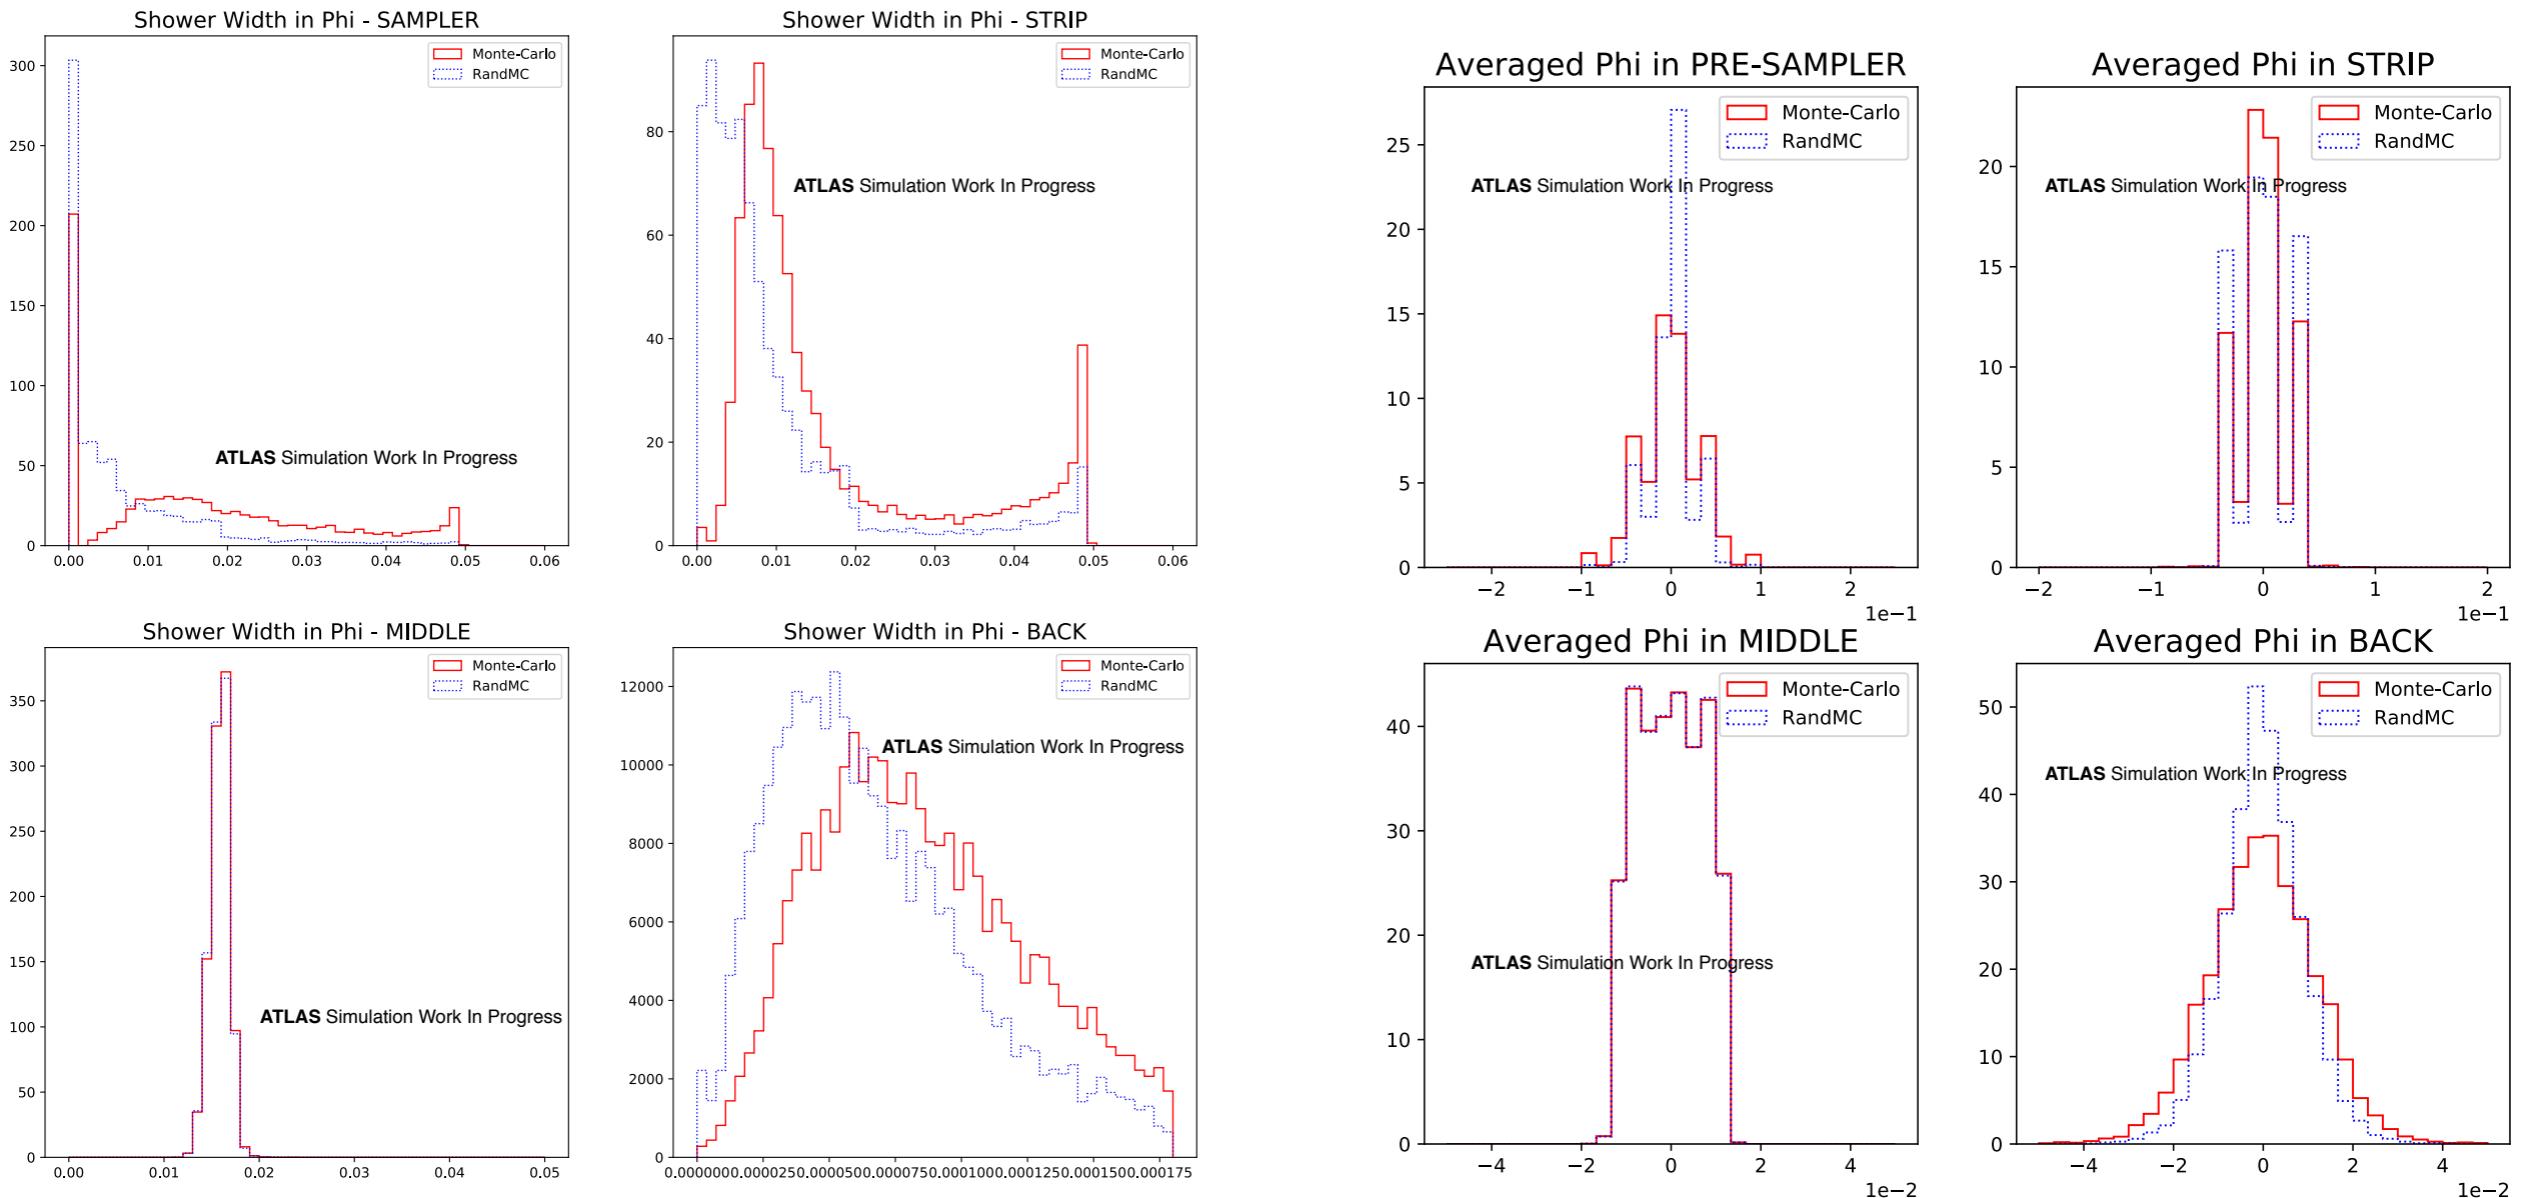


Two peaks for the **blue** and **black**

MC vs MC

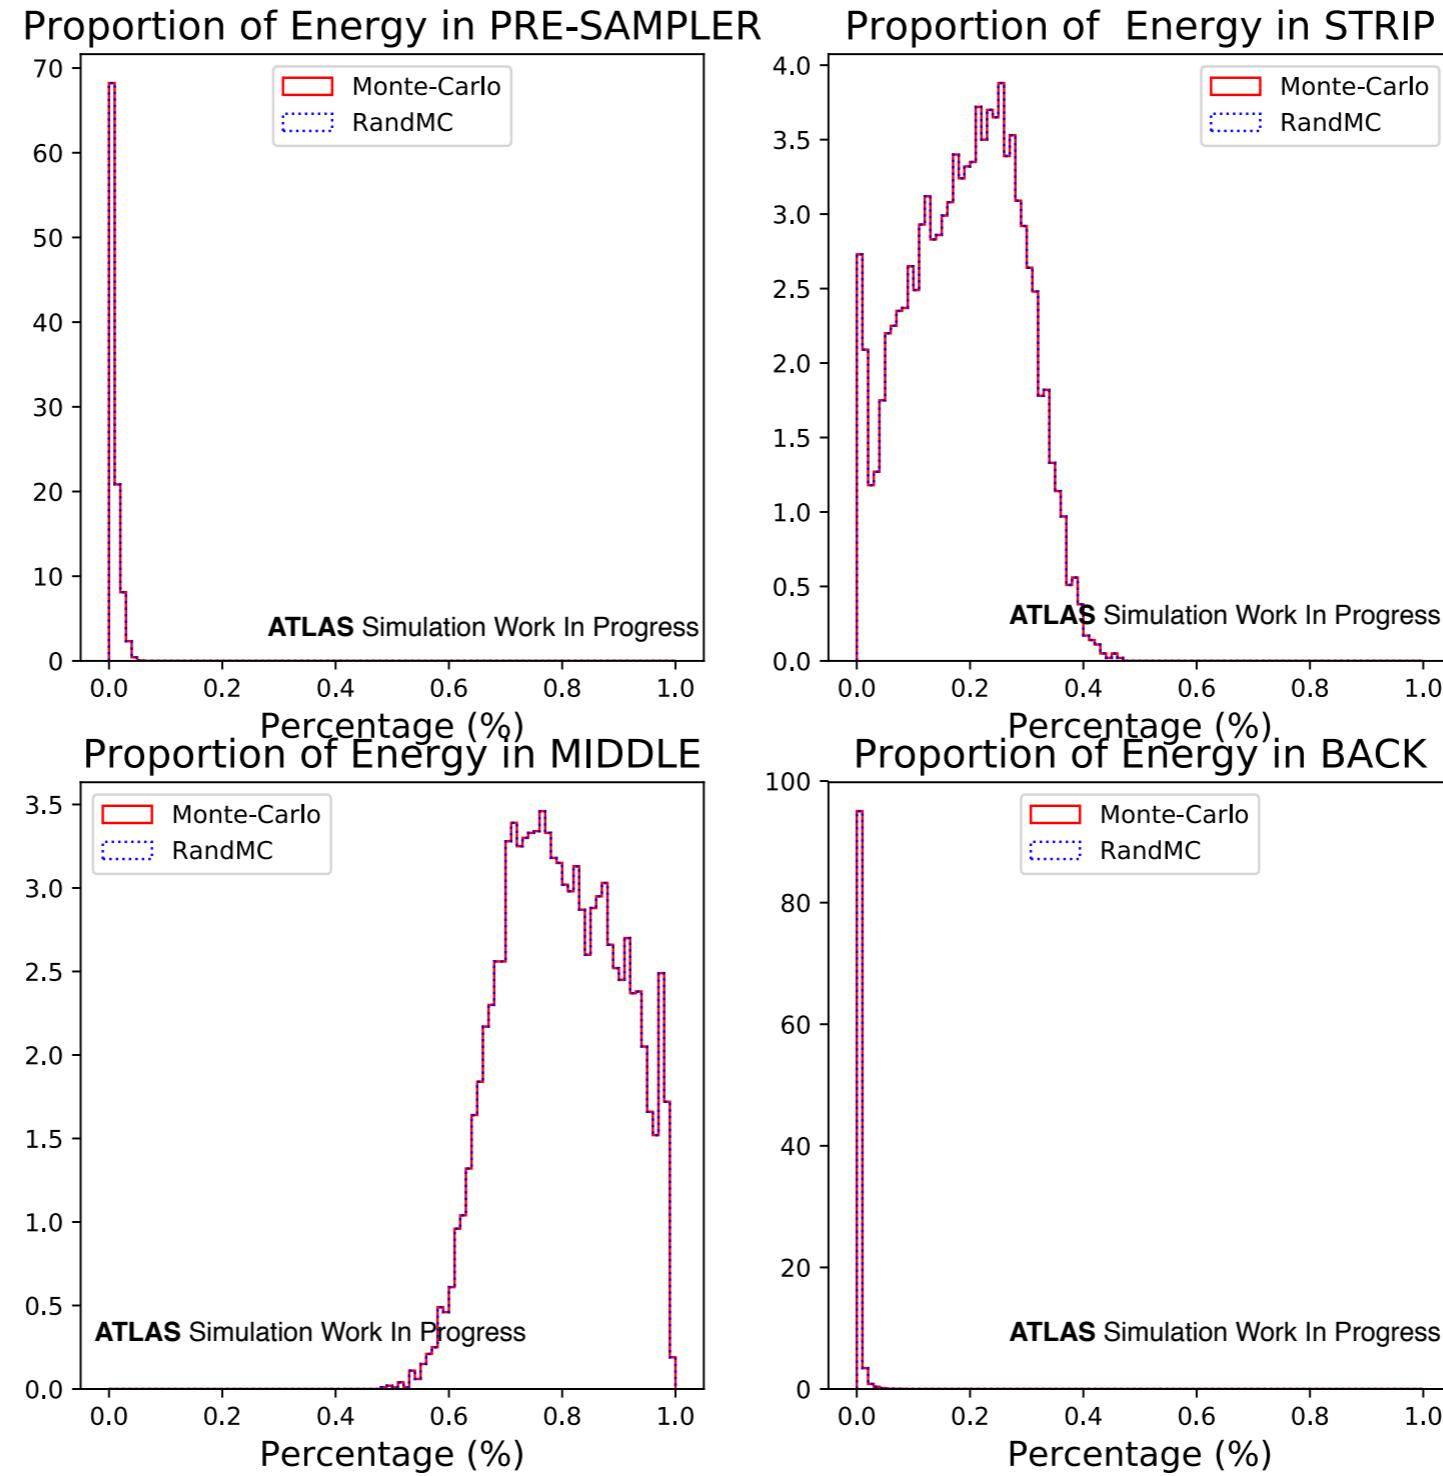
Phi from Corresponding Event vs Phi from another random event

BLUE HERE IS ALSO GEANT4 DATA (NOT GAN!)

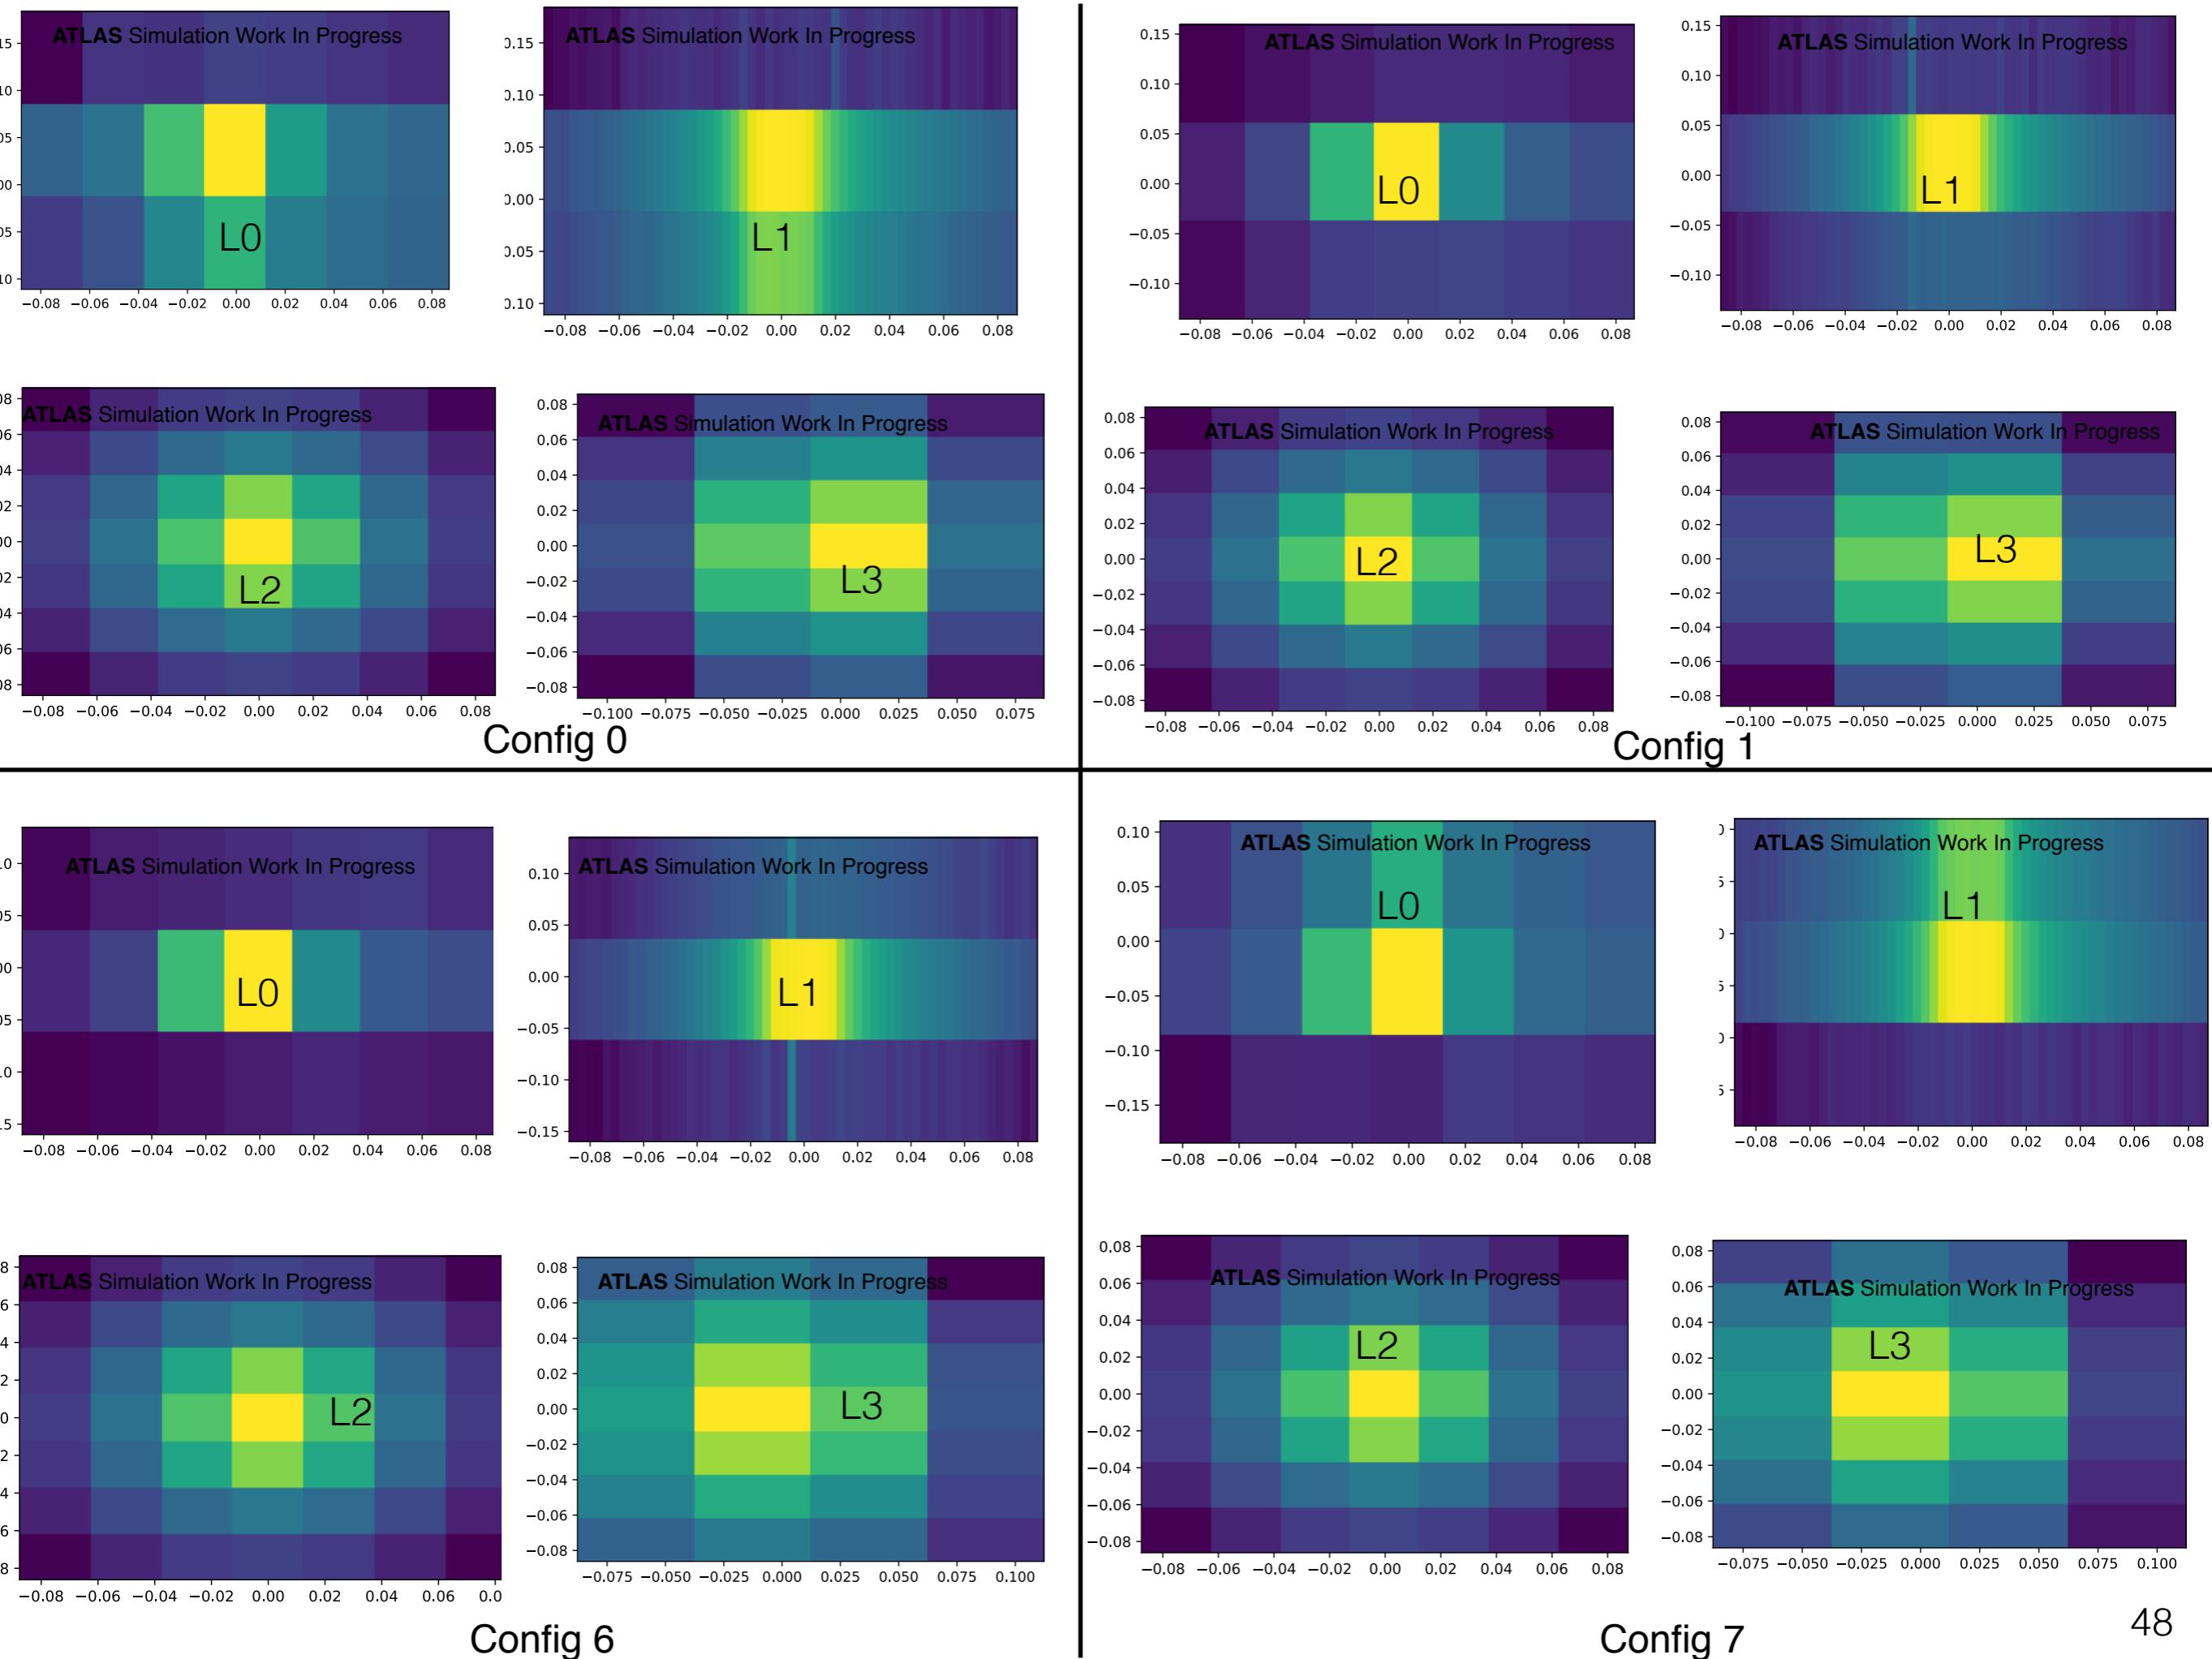


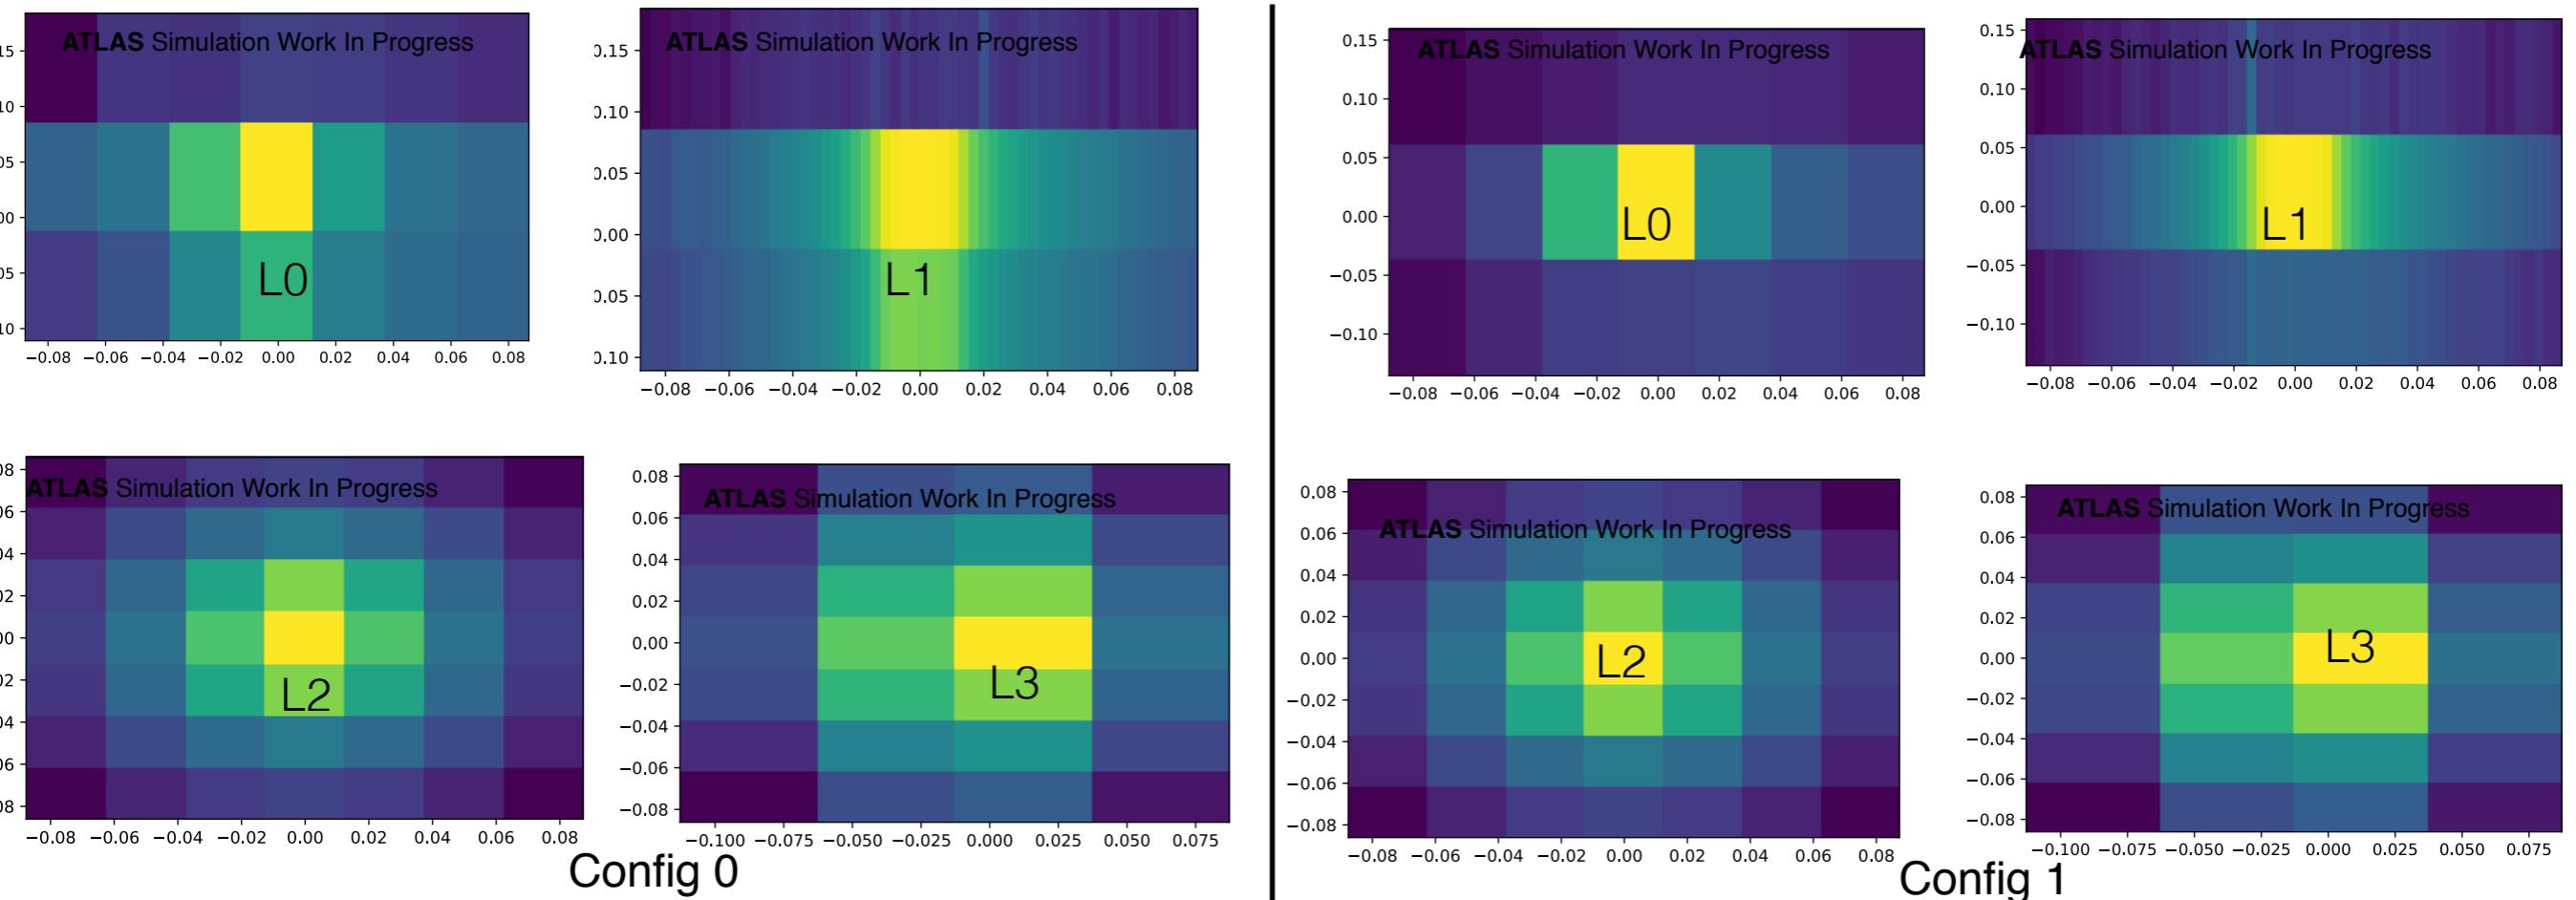
Distributions for Middle Layer are almost perfect

MC vs MC

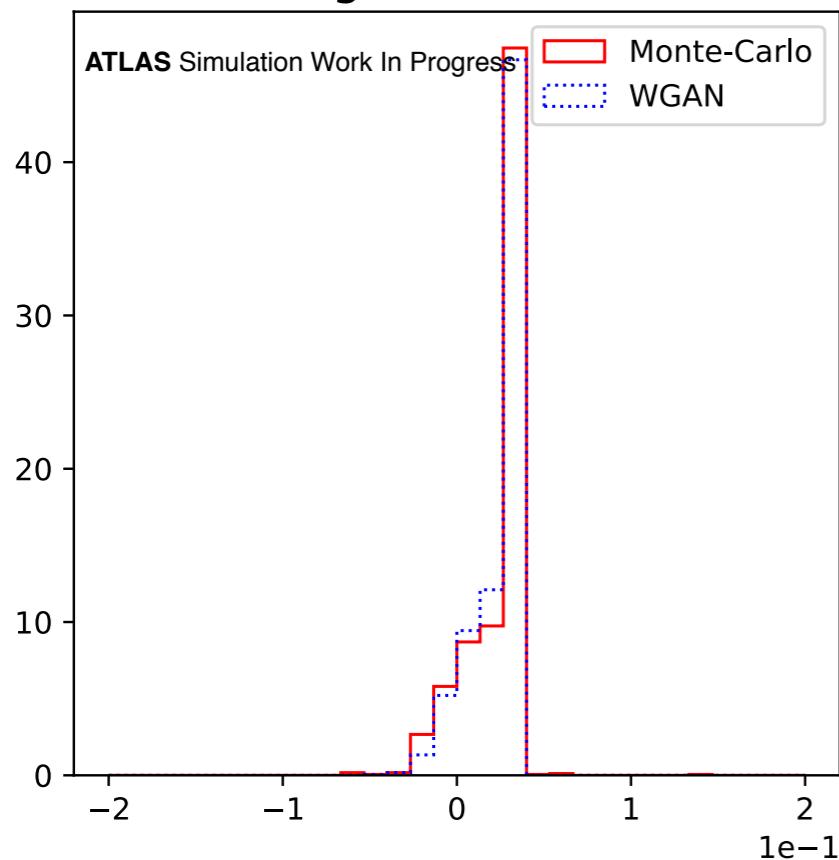


Energy Distributions Match Perfectly

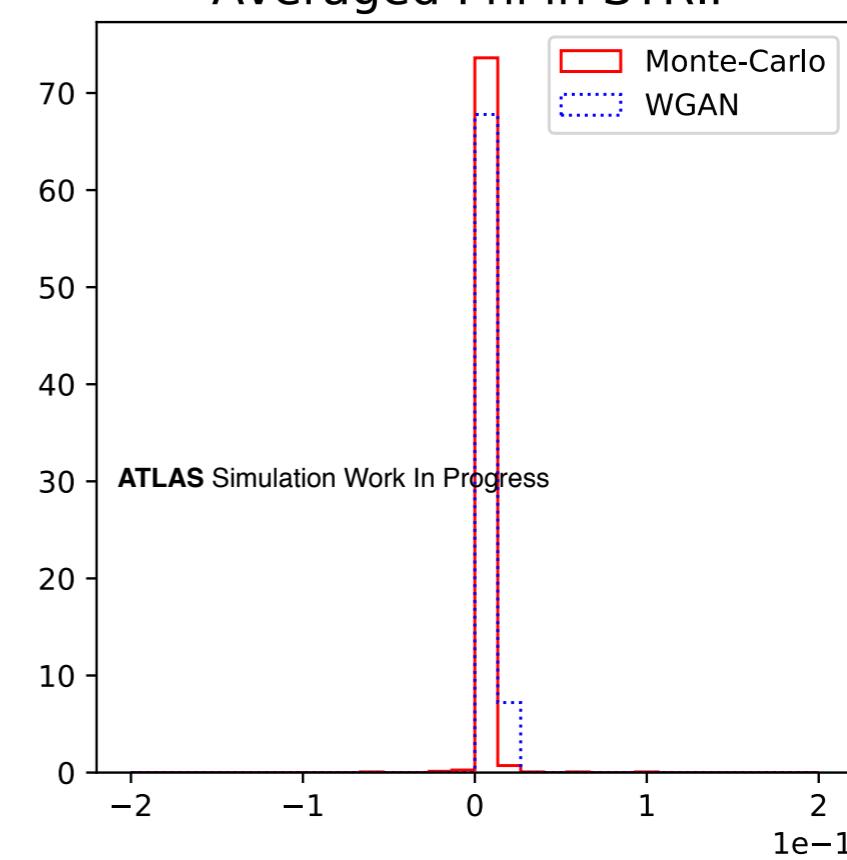


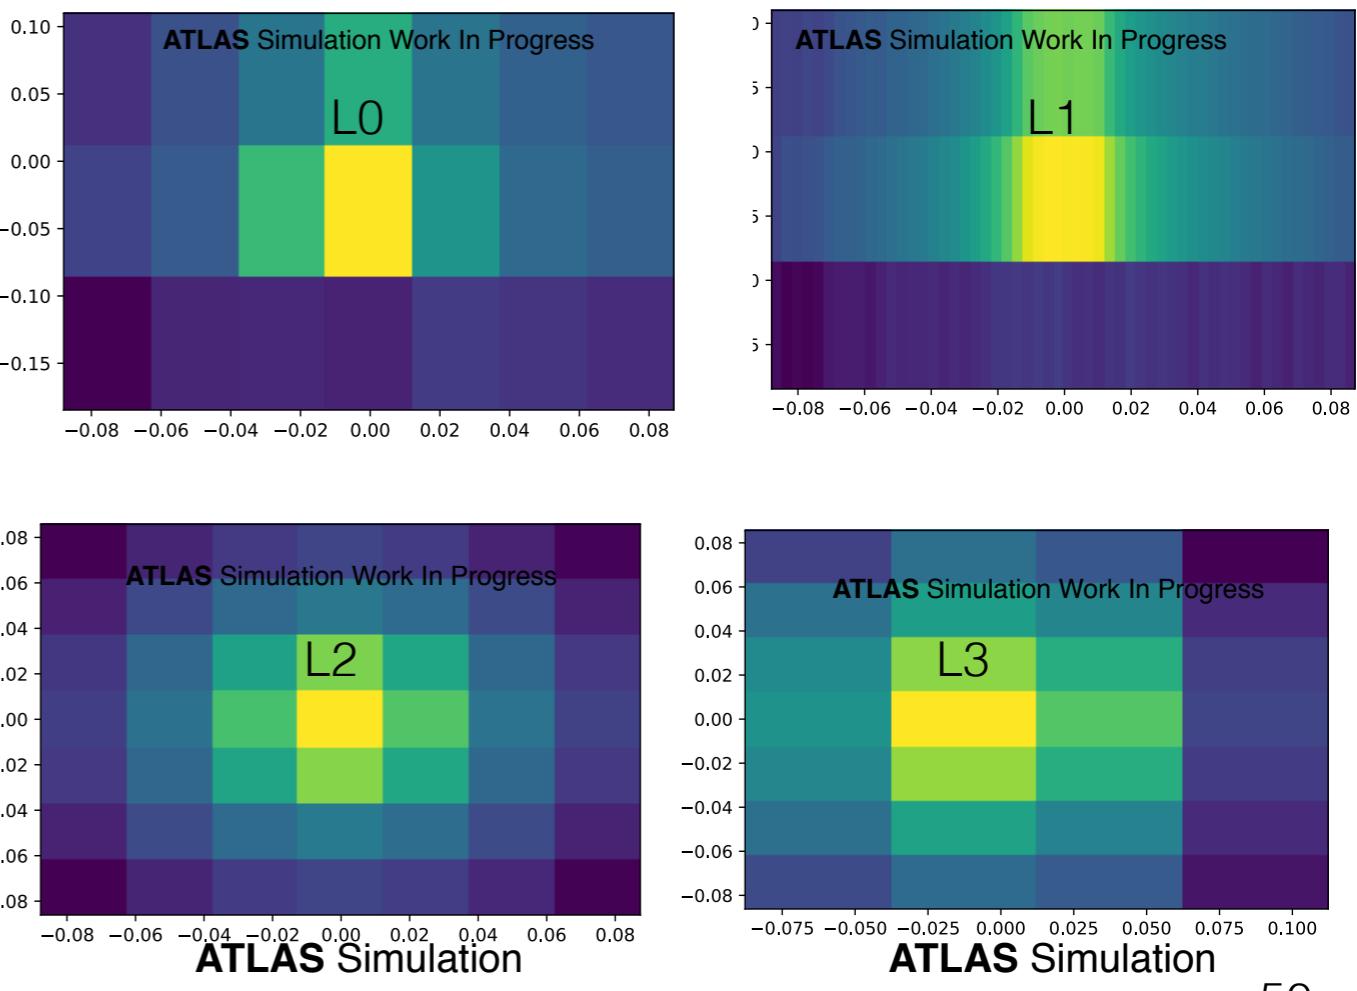
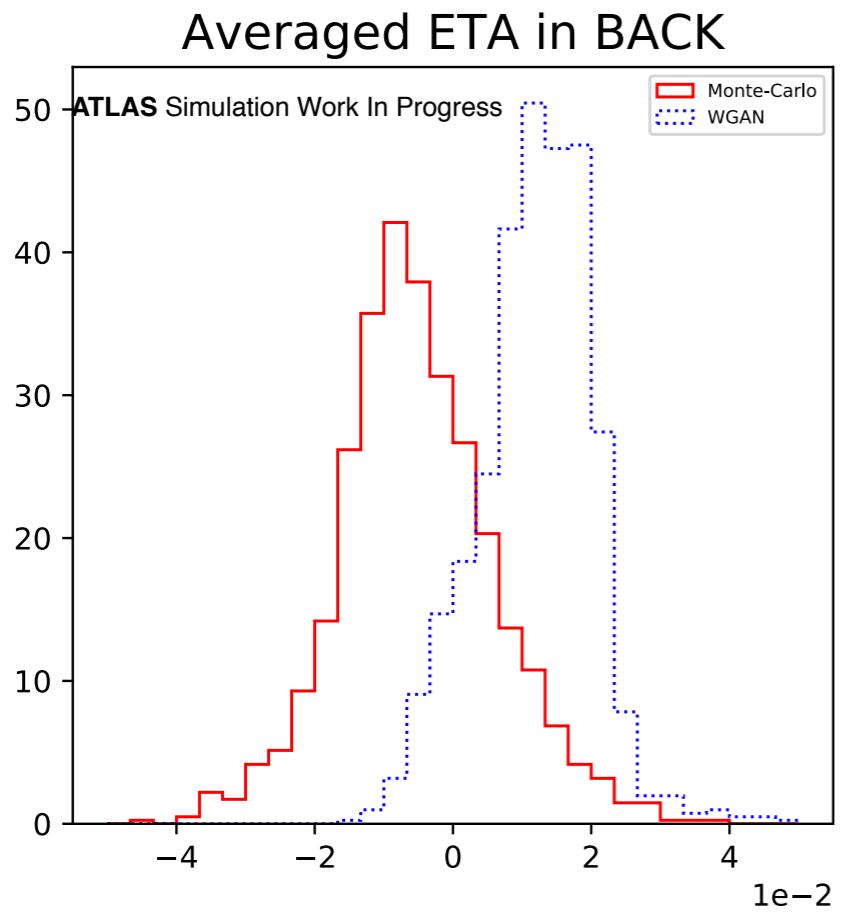
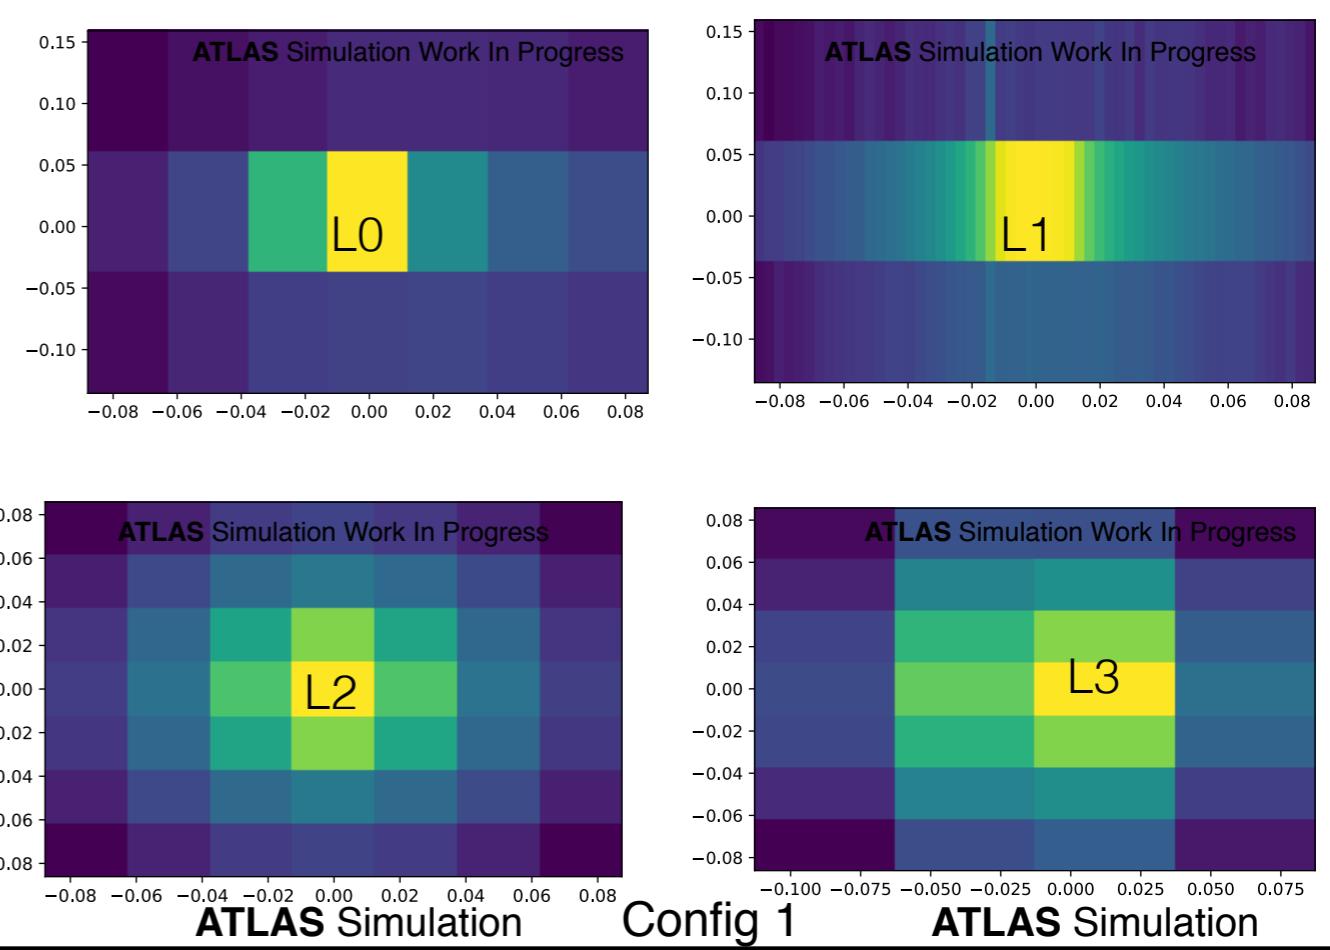
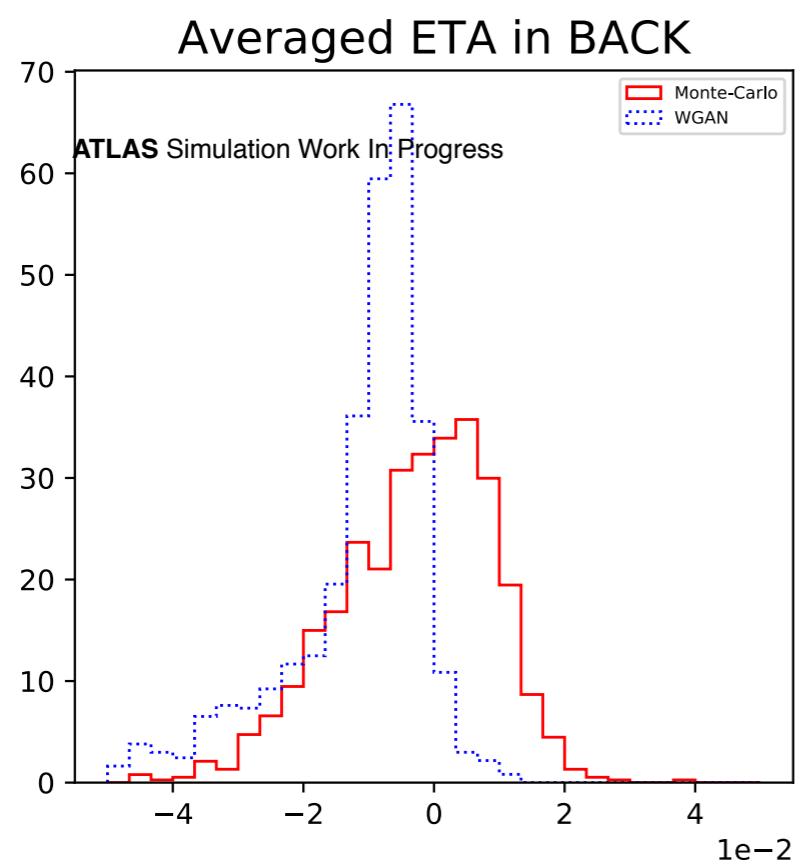


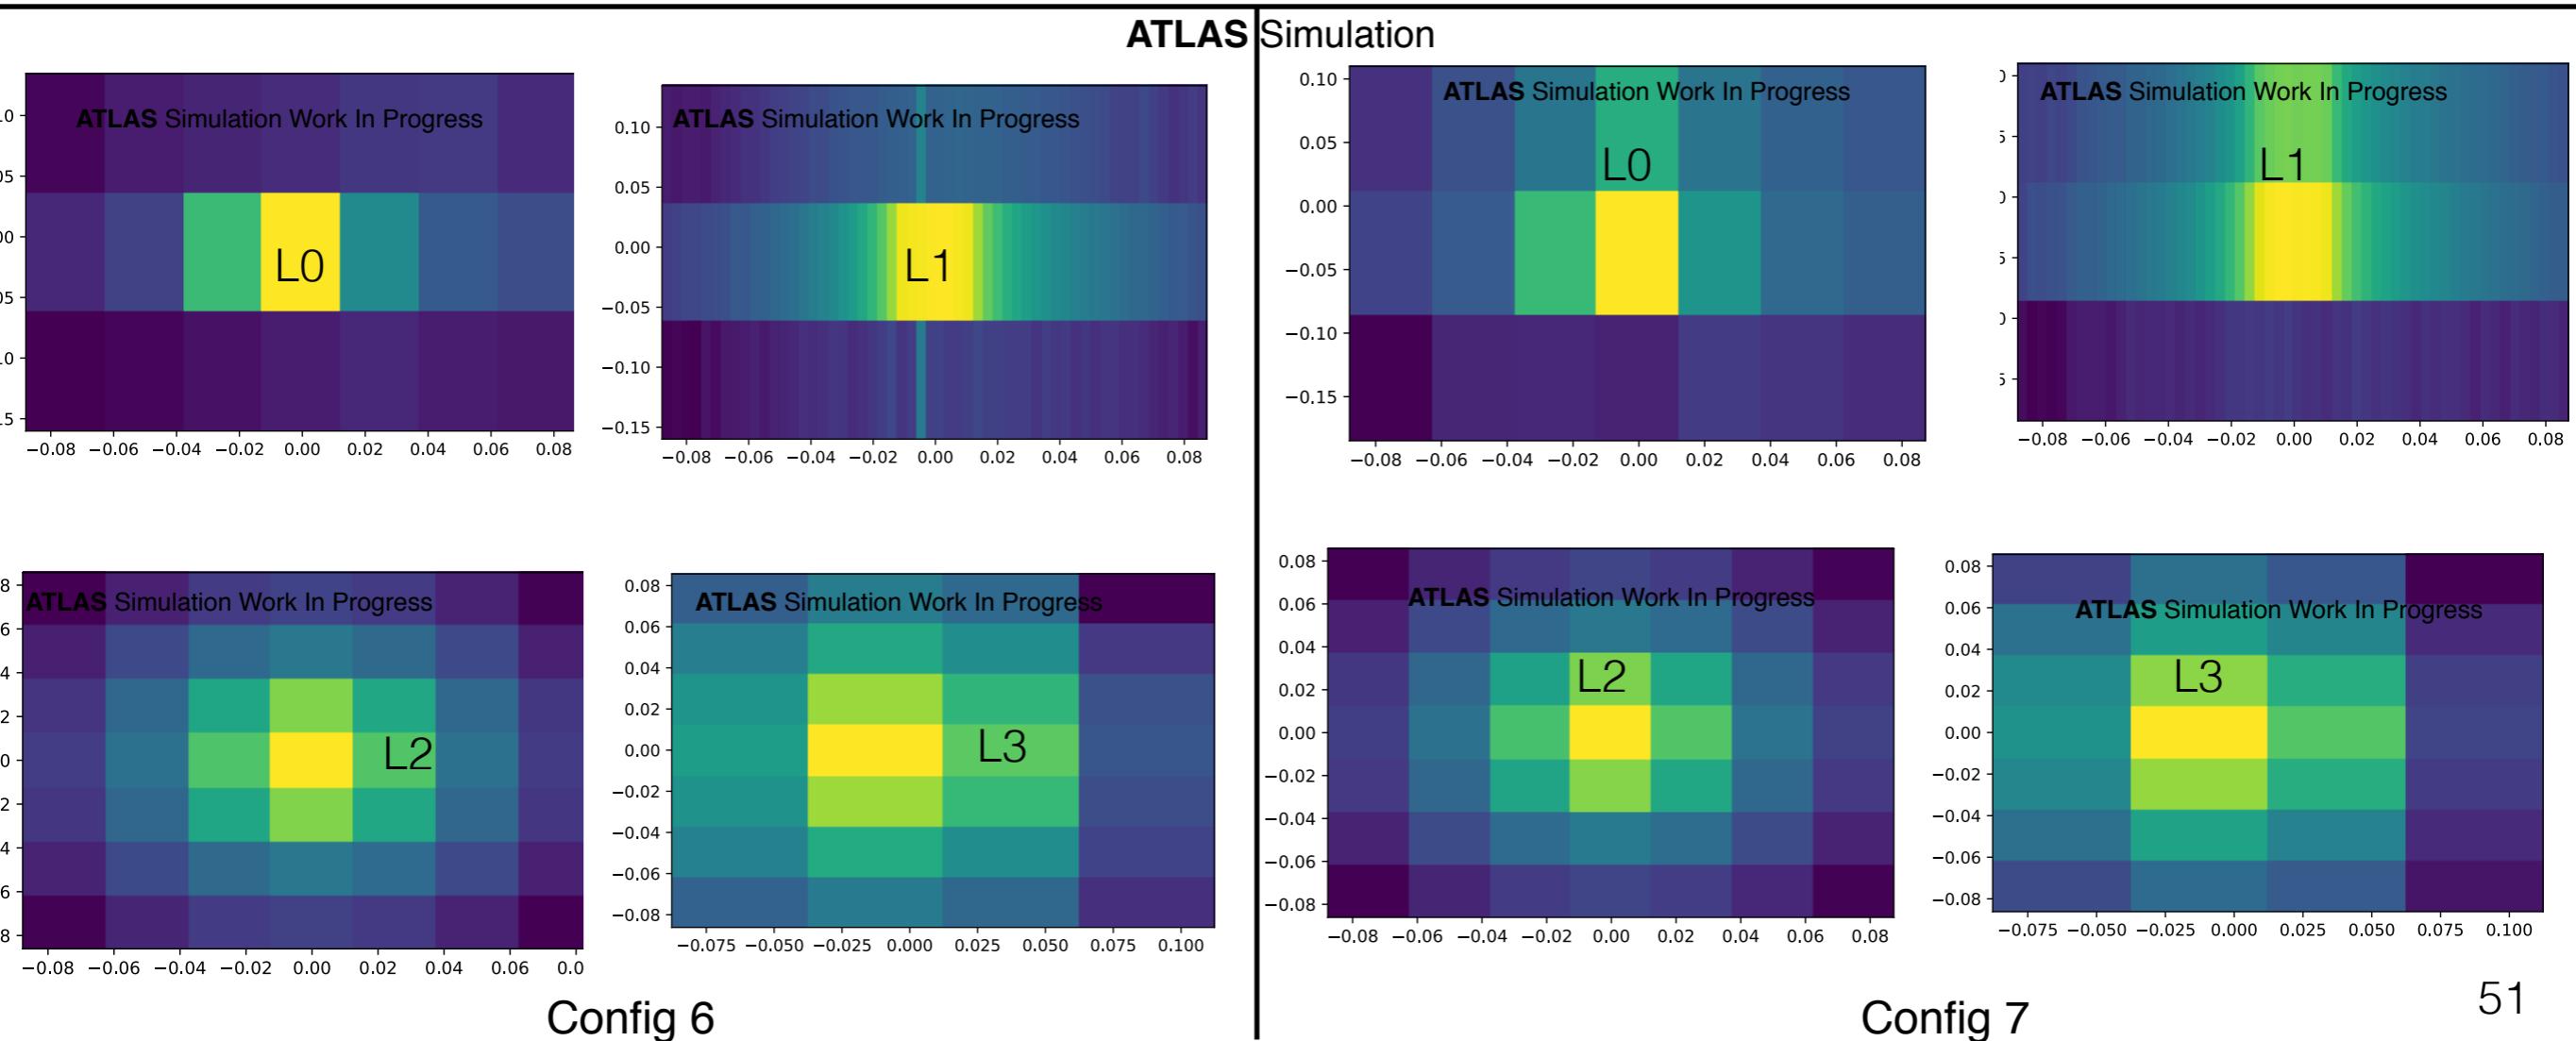
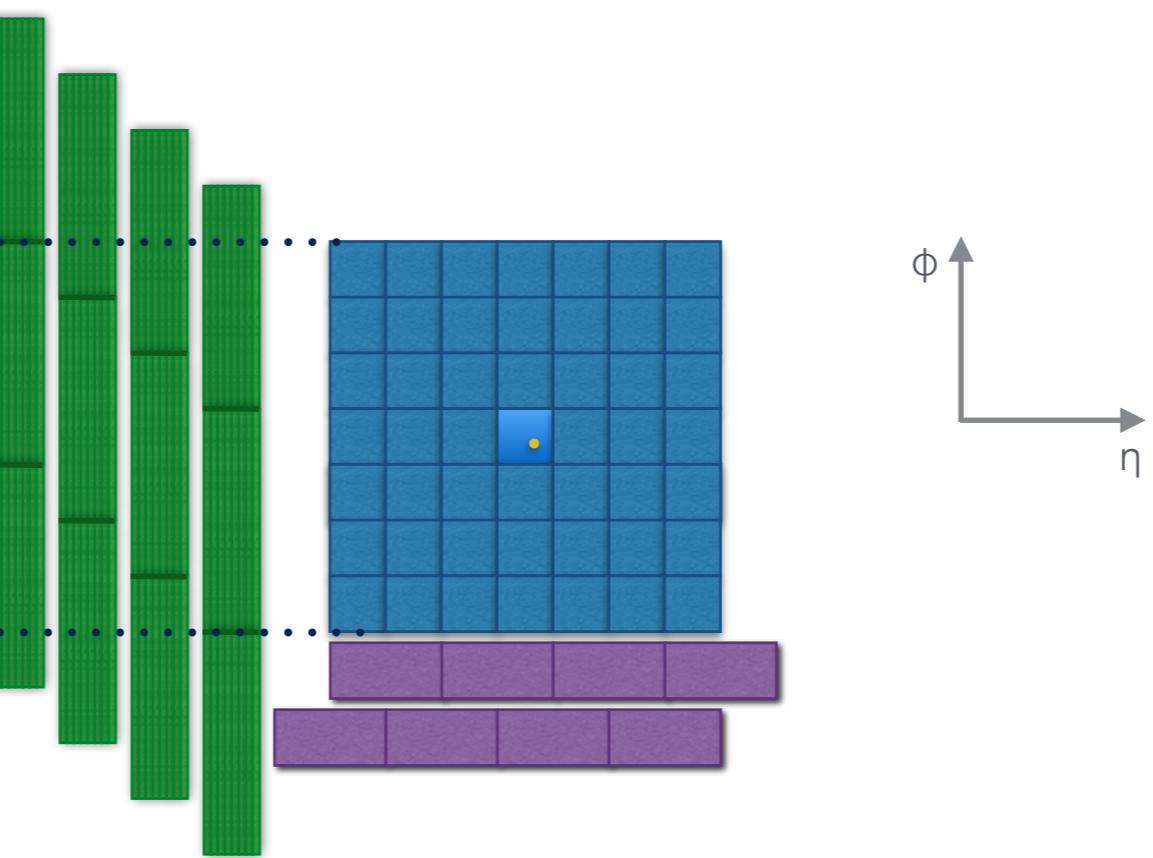
Averaged Phi in STRIP



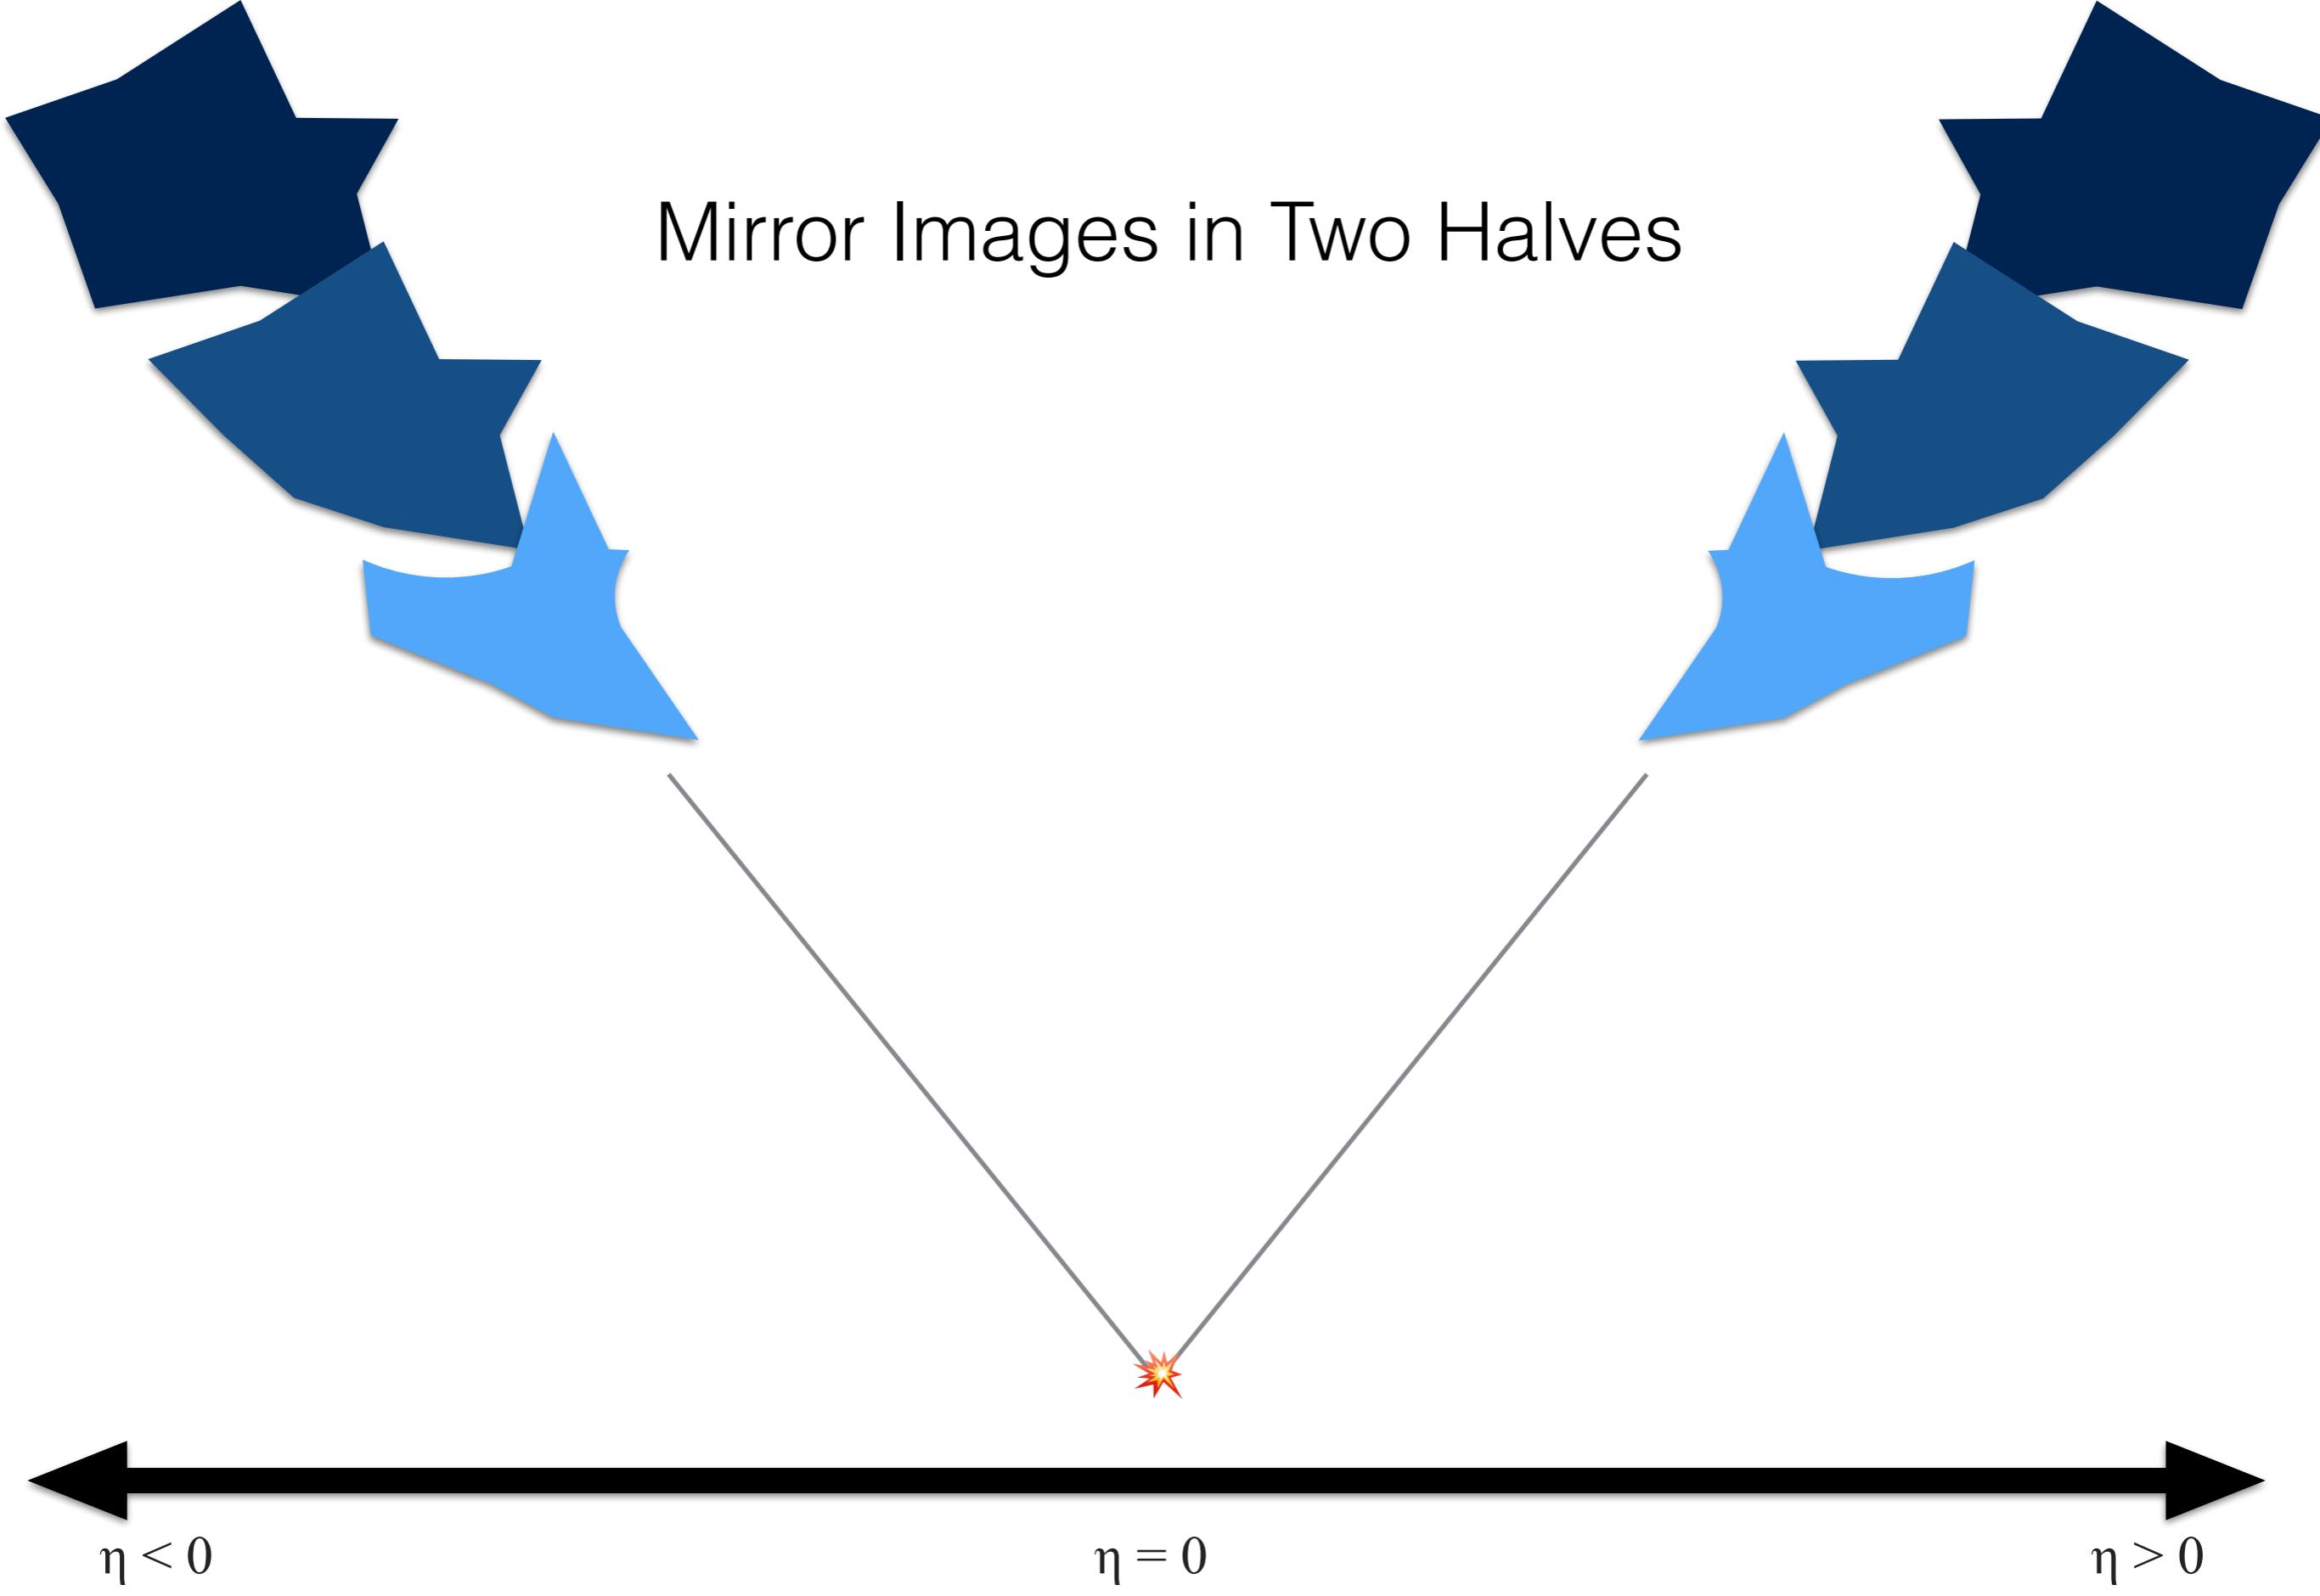
Averaged Phi in STRIP





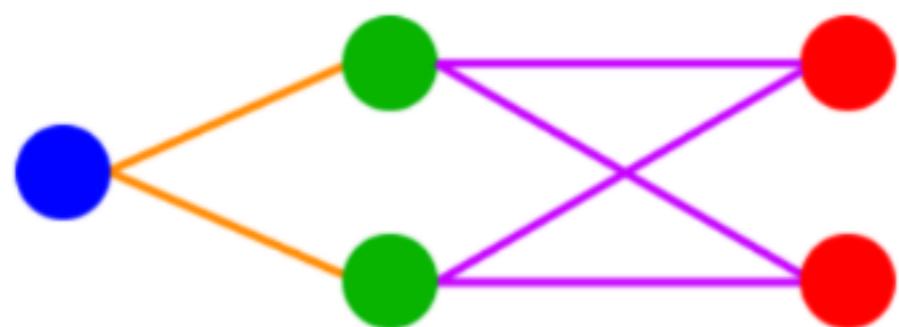


Mirror Images in Two Halves



What are Neural Networks?

- Just a lot of linear combinations followed by non-linear activation:



$$f_{NN} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

$$\text{Input : } x = \begin{bmatrix} x_{1,1} \\ x_{2,1} \end{bmatrix}$$

$$\text{Weight : } W_1 = \begin{bmatrix} W_{1,1} & W_{1,2} \\ W_{2,1} & W_{2,2} \end{bmatrix}$$

$$\text{Bias : } b_1 = \begin{bmatrix} b_{1,1} \\ b_{2,1} \end{bmatrix}$$

Activation : $\sigma(x) = \tanh(x)$ (as example)

Activation is applied elementwise!

From Stefan Wunsch

- We can have multiple outputs:

